

LuaT_EX

Reference

snapshot 2007-07-11



LuaT_EX

Reference

Manual

copyright: LuaT_EX development team
more info: www.luatex.org
version: July 11, 2007(snapshot 2007-07-11)

Contents

1	Introduction	5
2	Basic T _E X enhancements	7
2.1	Version information	7
2.2	UNICODE support	7
2.3	Wide math characters	8
2.4	Extended register tables	8
2.5	Attribute registers	9
2.6	LUA related primitives	9
2.6.1	<code>\directlua</code>	9
2.6.2	<code>\latelua</code>	9
2.6.3	<code>\luaescapestring</code>	9
2.6.4	<code>\closeslua</code>	10
2.7	New ε -T _E X primitives	10
2.7.1	<code>\clearmarks</code>	10
2.7.2	<code>\noligs</code> and <code>\nokerns</code>	10
2.7.3	<code>\formatname</code>	10
2.7.4	<code>\scantextokens</code>	10
2.7.5	Catcode tables	11
2.7.6	Font syntax	12
3	LUA general	13
3.1	Initialization	13
3.1.1	LUAT _E X as a LUA interpreter	13
3.1.2	Other commandline processing	13
3.2	LUA changes	14
3.3	LUA Modules	15
4	LUAT _E X LUA Libraries	17
4.1	The <code>tex</code> library	17
4.1.1	Integer parameters	17
4.1.2	Dimension parameters	19
4.1.3	Direction parameters	19
4.1.4	Glue parameters	19
4.1.5	Muglue parameters	20
4.1.6	Tokenlist parameters	20
4.1.7	Convert commands	20
4.1.8	attribute, count, dimension and token registers	20
4.1.9	Box registers	21
4.1.10	Print functions	22
4.1.11	Helper functions	23
4.2	The <code>token</code> library	24



4.2.1	<code>token.get_next</code>	24
4.2.2	<code>token.is_expandable</code>	24
4.2.3	<code>token.expand</code>	24
4.2.4	<code>token.is_activechar</code>	24
4.2.5	<code>token.create</code>	25
4.2.6	<code>token.command_name</code>	25
4.2.7	<code>token.command_id</code>	25
4.2.8	<code>token.csname_name</code>	25
4.2.9	<code>token.csname_id</code>	26
4.3	The <code>node</code> library	26
4.3.1	Node handling functions	26
4.3.2	Attribute handling	30
4.4	The <code>texio</code> library	31
4.4.1	Printing functions	31
4.5	The <code>pdf</code> library	32
4.6	The <code>callback</code> library	32
4.6.1	File discovery callbacks	33
4.6.2	File reading callbacks	35
4.6.3	Data processing callbacks	38
4.6.4	Node list processing callbacks	39
4.6.5	Information reporting callbacks	41
4.6.6	Font-related callbacks	42
4.7	The <code>lua</code> library	43
4.7.1	Variables	43
4.7.2	LUA bytecode registers	43
4.8	The <code>kpse</code> library	43
4.8.1	<code>kpse.find_file</code>	43
4.8.2	<code>kpse.set_program_name</code>	45
4.8.3	<code>kpse.init_prog</code>	45
4.8.4	<code>kpse.readable_file</code>	45
4.8.5	<code>kpse.expand_path</code>	45
4.8.6	<code>kpse.expand_var</code>	46
4.8.7	<code>kpse.expand_braces</code>	46
4.8.8	<code>kpse.var_value</code>	46
4.9	The <code>status</code> library	46
4.10	The <code>texconfig</code> table	48
4.11	The <code>font</code> library	49
4.11.1	Loading a TFM file	49
4.11.2	Loading a VF file	49
4.11.3	The fonts array	49
4.11.4	Checking a font's status	50
4.11.5	Defining a font directly	50
4.11.6	Currently active font	50
4.11.7	Maximum font id	50



4.11.8	Iterating over all fonts	50
4.12	The fontforge library	51
4.12.1	Getting quick information on a font	51
4.12.2	Loading an OPENTYPE or TRUETYPE file	51
5	Font structure	63
5.1	Real fonts	66
5.2	Virtual fonts	68
6	Nodes	71
6.1	LUA node representation	71
6.1.1	Auxiliary items	71
6.1.2	Main text nodes	72
6.1.3	whatsit nodes	75
7	Modifications	83
7.1	Changes from T _E X 3.141592	83
7.2	Changes from ϵ -T _E X 2.2	83
7.3	Changes from PDFT _E X 1.40	83
7.4	Changes from ALEPH RC4	84
7.5	Changes from standard WEB2C	84
8	Implementation notes	87
8.1	Primitives overlap	87
8.2	Memory allocation	87
8.3	Sparse arrays	87
8.4	Simple single-character csnames	88
8.5	Compressed format	88
8.6	Binary file reading	88
9	Known bugs and limitations	89
10	TODO	91





1 Introduction

This book will eventually become the reference manual of L^AT_EX. At the moment, it simply reports the behaviour of the executable matching the snapshot or beta release date in the title page.

Features may come and go. The current version of L^AT_EX is not meant for production and users cannot depend on functionality staying the same.

Nothing in the API is considered stable just yet. This manual therefore simply reflects the current state of the executable. ***Absolutely nothing*** on the following pages is set in stone. When the need arises, anything can (and will) be changed without prior notice.

If you are not willing to deal with this situation, you should wait for the stable version. The current estimate expects that to become available sometime in the summer of 2008.

L^AT_EX consists of a number of interrelated but (still) distinguishable parts:

- PDF_T_EX version 1.40.3
- ALEPH RC4 (from the T_EX_{LIVE} repository)
- Functionality of ϵ -T_EX 2.2
- LUA 5.1.2
- Dedicated LUA libraries
- Various T_EX extensions
- Parts of the Font Parser from FONTFORGE 2007.06.07
- Compiled source code to glue it all together

Neither I/O translation processes, nor tcx files, nor ENC_T_EX can be used. All these encoding-related functions are superseded by a LUA-based solution (reader callbacks).





2 Basic T_EX enhancements

2.1 Version information

There are three new primitives to test the version of L^AT_EX:

primitive	explanation
<code>\luatexversion</code>	A combination of major and minor number, as in pdfT _E X. Current value: 10
<code>\luatexrevision</code>	The revision, also as in pdfT _E X. Current value: 0
<code>\luatexdatestamp</code>	A combination of the local date and hour when the current executable was compiled, the syntax is identical to <code>\luatexrevision</code> . Value for the executable that generated this document: 2007071111

Note that the `\luatexdatestamp` depends on both the compilation time and place of the current executable, it is defined in terms of the local time. The purpose of this primitive is solely to be an aid in the development process, do not use it for anything besides debugging.

2.2 UNICODE support

Text input and output is now considered to be UNICODE text, so characters can use the full range of UNICODE ($2^{20} + 2^{16} = 10FFFF = 1114111$).

For now, it only makes sense to use values above the base plane ("**0xFFFF**") for `\mathcode` and `\catcode` assignments, since the hyphenation patterns are still limited to at the most 16-bit values, so the other commands will not know what to do with those high values.

Many primitives are affected by this. For instance, `\char` now accepts values between 0 and 1114111. This should not be a problem for well-behaved input files, but it could create incompatibilities for input that would have generated an error when processed by older T_EX-based engines. The maximum number of allocations is "**10FFFF**" or $2^{20} + 2^{16}$ (21 bits). The maximum value that can be assigned are:

primitive	bits	hex	numeric
<code>\char</code>	21	10FFFF	$2^{20} + 2^{16}$
<code>\chardef</code>	21	10FFFF	$2^{20} + 2^{16}$
<code>\lccode</code>	21	10FFFF	$2^{20} + 2^{16}$
<code>\uccode</code>	21	10FFFF	$2^{20} + 2^{16}$
<code>\sfcode</code>	15	7FFF	2^{15}
<code>\catcode</code>	4	F	2^4
<code>\mathchardef</code>	15	8000	$2^3 \times 2^8 \times 2^4$
<code>\mathcode</code>	15	8000	$2^3 \times 2^8 \times 2^4$
<code>\delcode</code>	27	7FFFFFFF	$2^3 \times 2^4 \times 2^8 \times 2^4 \times 2^8$



As far as the core engine is concerned, all input and output to text files is UTF-8 encoded. Input files can be pre-processed using the `reader` callback. This will be explained in a later chapter.

Output in byte-sized chunks can be achieved by using characters just outside of the valid unicode range, starting at the value 1.114.112 (0x110000). When the times comes to print a character $c \geq 1.114.112$, L^AT_EX will actually print the single byte corresponding to $c - 1.114.112$.

Output to the terminal uses `^^` notation for the lower control range ($c < 32$), with the exception of `^^I`, `^^J` and `^^M`. These are considered safe and therefore printed as-is.

Normalization of the UNICOD^E input can be handled by a macro package during callback processing (will be explained below).

2.3 Wide math characters

Text is now extended up to the full UNICOD^E range, but math mode deals mostly with glyphs in fonts directly, and fonts tend to be 16-bit at maximum.

Therefore, the math primitives from ALEPH are kept mostly as-is, except for the ones that convert from input to math commands. The extended commands (with the `o` prefix) accept 16-bit glyph indices in one of 256 possible families. The traditional T_EX primitives are unchanged, their arguments are upscaled internally.

primitive	max index/bits	hex	numeric
<code>\mathchar</code>	15	7FFF	$2^3 * 2^8 * 2^4$
<code>\delimiter</code>	27	7FFFFFFF	$2^3 * 2^4 * 2^8 * 2^4 * 2^8$
<code>\omathchar</code>	27	7FFFFFFF	$2^3 * 2^{16} * 2^8$
<code>\odelimiter</code>	27+24	7FFFFFFF + FFFFFFFF	$2^3 * 2^8 * 2^{16} + 2^8 * 2^{16}$
<code>\omathchardef</code>	21=27	10FFFF = 8000000	$2^{20} + 2^{16} = 2^3 * 2^{16} * 2^8$
<code>\omathcode</code>	21=27	10FFFF = 8000000	$2^{20} + 2^{16} = 2^3 * 2^{16} * 2^8$
<code>\odelcode</code>	21=27+24	10FFFF = 7FFFFFFF + FFFFFFFF	$2^{20} + 2^{16} = 2^3 * 2^8 * 2^{16}$ $+ 2^8 * 2^{16}$

2.4 Extended register tables

All registers can be 16 bit numbers as in ALEPH. The affected commands are:

<code>\count</code>	<code>\countdef</code>	<code>\unhbox</code>	<code>\ht</code>
<code>\dimen</code>	<code>\dimendef</code>	<code>\unvbox</code>	<code>\dp</code>
<code>\skip</code>	<code>\skipdef</code>	<code>\copy</code>	<code>\setbox</code>
<code>\muskip</code>	<code>\muskipdef</code>	<code>\unhcopy</code>	<code>\vsplit</code>
<code>\marks</code>	<code>\toksdef</code>	<code>\unvcopy</code>	
<code>\toks</code>	<code>\box</code>	<code>\wd</code>	



2.5 Attribute registers

Attributes are a completely new concept to LuaTeX . Syntactically, they behave a lot like counters: attributes obey TeX 's nesting stack and can be used after `\the` etc. just like the normal `\count` registers.

```
\attribute <16-bit number> <optional equals> <31-bit number>\crlf \attributedef  
<cname> <optional equals> <16-bit number>
```

However, there are some differences as well. Conceptually, an attribute is either set or unset. Set attributes can only have values of 0 or more, otherwise they are considered unset and automatically remapped to -1 . All attributes start out in the unset state (in iniTeX).

Attributes can be used as extra counter values, but their usefulness comes from the fact that the numbers and values of all set attributes are attached to all nodes created in their scope. These can then be queried from any Lua code that deals with node processing. Further information is given in the chapter on 6.

2.6 LUA related primitives

In order to merge Lua code with TeX input, a few new primitives are needed. LuaTeX has support for 65536 separate Lua interpreter states. States are automatically created based on the integer argument to the primitives `\directlua` and `\latelua`.

2.6.1 `\directlua`

The primitive `\directlua` is used to execute Lua code immediately. The syntax is

```
\directlua <16-bit number> <general text>
```

The `<general text>` is fed into the Lua interpreter state indicated by the `<16-bit number>`. If the state does not exist yet, then it will be initialized automatically.

This command is expandable.

2.6.2 `\latelua`

`\latelua` stores Lua code in a whatsit that will be processed inside the output routine. It's intended use is very similar to `\pdfliteral`.

Within the Lua code, you should use `pdf.print` to print stuff directly to the pdf file.

```
\latelua <16-bit number> <general text>
```

2.6.3 `\luaescapestring`

This primitive converts a TeX token string so that it can be safely used as the contents of a Lua string: embedded backslashes, double quotes and single quotes are escaped by prepending an extra token consisting of a backslash with catcode 12.



`\luaescapestring` \langle general text \rangle

2.6.4 `\closelua`

This primitive allows you to close a LUA state, freeing all of its used memory.

`\closelua` \langle 16-bit number \rangle

You cannot close LUA state zero (0), any attempt to do so will be silently ignored.

States are only closed automatically when a fatal (out of memory) error occurs, but at that point `LUATEX` will exit anyway.

States are not closed immediately, but only when the output routine comes into play next (because there may be pending `\latelua` calls)

2.7 New ϵ -`TEX` primitives

2.7.1 `\clearmarks`

This primitive clears a marks class completely, resetting all three connected mark texts to empty.

`\clearmarks` \langle 16-bit number \rangle

2.7.2 `\noligs` and `\nokerns`

These primitives prohibit ligature and kerning insertion at the time when the initial node list is built by `LUATEX`'s main control loop (It is a temporary trick that will be removed soon).

`\noligs` \langle integer \rangle `\nokerns` \langle integer \rangle

2.7.3 `\formatname`

`\formatname`'s syntax is identical to `\jobname`.

In `INITTEX`, the expansion is empty. Otherwise, the expansion is the value that `\jobname` had during the `INITTEX` run that dumped the currently loaded format.

2.7.4 `\scantextokens`

The syntax of `\scantextokens` is identical to `\scantokens`.

This is a slightly adapted version of ϵ -`TEX`'s `\scantokens`. The differences are:



- The last (and usually only) line does not have a `\endlinechar` appended
- `\scantextokens` never raises an EOF error, and it does not execute `\everyeof` tokens.
- The while end of file tests are not executed, allowing the expansion to end on a different grouping level or while a conditional is still incomplete

2.7.5 Catcode tables

Catcode tables are a new feature that allows you to switch to a predefined catcode regime in a single statement. You can have a practically unlimited number of different tables (at this moment up to 268.435.456. The limit depends on an array allocation).

The subsystem is backward compatible: if you never use the following commands, your document will not notice any difference in behavior compared to traditional T_EX.

The contents of each catcode table is independent of any other catcode tables, and their contents is stored and retrieved from the format file.

2.7.5.1 `\catcodetable`

`\catcodetable` *<28-bit number>*

The `\catcodetable` switches to a different catcode table. Such a table has to be previously created using one of the two primitives below, or it has to be zero (table zero is initialized by `INITEX`).

2.7.5.2 `\initcatcodetable`

`\initcatcodetable` *<28-bit number>*

The `\initcatcodetable` creates a new table with catcodes identical to those defined by `INITEX`:

0	<code>\</code>		escape
5	<code>^^M</code>	return	<code>car_ret</code>
9	<code>^^@</code>	null	ignore
10	<code><space></code>	space	<code>spacer</code>
11	<code>a – z</code>		letter
11	<code>A – Z</code>		letter
12	everything else		other
14	<code>%</code>		comment
15	<code>^^?</code>	delete	<code>invalid_char</code>

The new catcode table is allocated globally: it will not go away after the current group has ended. If the supplied number is the currently active table, an error is raised.

2.7.5.3 `\savecatcodetable`

`\savecatcodetable` *<28-bit number>*



`\savecatcodetable` copies the current set of catcodes to a new table with the requested number. The definitions in this new table are all treated as if they were made in the outermost level.

The new table is allocated globally: it will not go away after the current group has ended. If the supplied number is the currently active table, an error is raised.

2.7.6 Font syntax

LuAT_EX will accept a braced argument as a font name:

```
\font\myfont = {cmr10}
```

This allows for embedded spaces, without the need for double quotes. Macro expansion takes place in the argument.



3 LUA general

3.1 Initialization

3.1.1 L^AT_EX as a LUA interpreter

There are two situations that make L^AT_EX behaves like it is a LUA interpreter only:

- If a `--luaonly` option is given
- If the executable is named `texlua` (or `luatexlua`)

In this mode, it will set LUA's `arg[0]` to the found script name, pushing preceding options in negative values and the rest of the commandline in the positive values, just like the LUA interpreter.

L^AT_EX will exit immediately after executing the specified LUA script and is, in effect, a somewhat bulky standalone LUA interpreter.

3.1.2 Other commandline processing

Whenever the L^AT_EX executable starts, it looks for a `--lua` commandline option. If such an option is present, it will enter an alternative mode of commandline parsing.

In this mode, it will only interpret a very small subset of the commandline directly:

```
-luaonly  execute a LUA script, then exit
-lua=s    load and execute a LUA init script
-safer    disable easily exploitable LUA commands
-help     display help and exit
-version  display version and exit
```

If a requested LUA script can not be found using the actual name given on the commandline, a second attempt is made by prepending the value of the environment variable `LUATEXDIR`, if that variable is defined.

Then the script is loaded and executed. It will find the entire commandline in the table `arg`, beginning with `arg[0]`, that is the name of the executable.

L^AT_EX will fetch some of the other commandline options from the `texconfig` table at the end of script execution (see the description of the `texconfig` table later on in this document).

Commandline processing happens very early on. So early, in fact, that none of T_EX's initializations have taken place yet. For that reason, the `tex`, `token`, `node` and `pdf` tables are off-limits during the execution of the startup file (they are nilled). Special care is taken that `texio.write` and `texio.write_nl`



function properly, so that you can at least report your actions to the log file when (and if) it eventually becomes opened (note that T_EX does not even know it's `\jobname` yet at this point).

The file is loaded into LUA state 0, and everything you do will remain visible during the rest of the run, with the exception of the `tex`, `token`, `node` and `pdf` tables: those will be restored to their normal meaning right after the execution of the script.

We recommend you use the startup file only for your own T_EX-independant initializations (if you need any), to parse the commandline, set values in the `texconfig` table, and register the callbacks you need.

You can use the `--safer` switch to disable some commands that can easily be abused by a malicious document. At the moment, this switch `nils` the following functions:

library functions

os	execute exec setenv rename remove
io	popen output tmpfile
lfs	rmdir mkdir chdir lock touch

And it makes `io.open()` fail on files that are opened for anything besides reading.

Unless the `texconfig` table tells it not to start KPATHSEA at all (set `texconfig.kpse_init` to `false` for that), it also acts on three other commandline options:

flag	meaning
<code>--fmt=s</code>	set the format name
<code>--progname=s</code>	set the progname (only for KPATHSEA)
<code>--ini</code>	enable INIT _E X mode

In order to initialize the built-in KPATHSEA library properly, L^UA_T_EX needs to know the correct progname to use, and for that it needs to check `-progname` (and `-ini` and `-fmt`, if `-progname` is missing).

If there is no `--lua` option, the commandline is interpreted in a similar fashion as in traditional PDF_T_EX and ALEPH.

3.2 LUA changes

The `read("*line")` function from the `io` library has been adjusted so that it is line-ending neutral: any of `LF`, `CR` or `CR+LF` are accepted.

The `tostring()` printer for numbers has been changed so that it return 0 instead of something like `2e-5` (which confused T_EX enormously) when the value is so small that T_EX cannot distinguish it from zero.

Dynamic loading of `.so` and `.dll` files is disabled on all platforms.

`luafilesystem` has been extended with two extra boolean functions (`isdir(filename)` and `isfile(filename)`) and one extra string field in the attributes table (`permissions`).

The `string` library has an extra function: `string.explode(s[,m])`. This function returns an array containing the string argument `s` split into substrings based on the value of the string argument `m`. The second argument is a string that is either empty (this splits the string into characters), a single character (this splits on each occurrence of that character, possibly introducing empty strings), or a single character



followed by the plus sign `+` (this special version does not create empty substrings). The default value for `m` is `' +'` (multiple spaces).

The `string` library also has six extra iterators that return strings piecemeal:

- `string.utfvalues(s)` (returns an integer value in the UNICODE range)
- `string.utfcharacters(s)` (returns a string with a single UTF-8 token in it)
- `string.characters(s)` (a string of length one)
- `string.characterpairs(s)` (two strings of length one) will produce an empty second string in the string length was odd.
- `string.bytes(s)` (a single byte value)
- `string.bytepairs(s)` (two byte values) Will produce nil instead of a number as its second return value if the string length was odd.

The `os` library has a few extra functions and variables:

- `os.exec('command')` is a non-returning version of `os.execute`. The advantage of this command is that it cleans out the current process before starting the new one, making it especially useful for use in T_EX_{LUA}.
- `os.setenv('key', 'value')` This sets a variable in the environment. Passing `nil` instead of a value string will remove the variable.
- `os.environ` This is a read-only hash table containing all of the variables and values in the process environment.

3.3 LUA Modules

Some modules that are normally external to Lua are statically linked in with L^AT_EX, because they offer useful functionality:

- `slnunicode`, from the Selene libraries, <http://luaforge.net/projects/sln>. (version 1.1)
- `luazip`, from the kepler project, <http://www.keplerproject.org/luazip/>. (version 1.2.1, but patched for compilation with lua 5.1)
- `luafilesystem`, also from the kepler project, <http://www.keplerproject.org/luafilesystem/>. (version 1.2, but patched for compilation with lua 5.1)
- `lpeg`, by Roberto Ierusalimschy, <http://www.inf.puc-rio.br/~roberto/lpeg.html>. (version 0.6)
- `lzlib`, by Tiago Dionizio, <http://mega.ist.utl.pt/~tngd/lua/>. (version 0.2)
- `md5`, by Roberto Ierusalimschy <http://www.inf.puc-rio.br/~roberto/md5/md5-5/md5.html>.





4 L^AT_EX LUA Libraries

The interfacing between T_EX and LUA is facilitated by a set of LUA modules.

4.1 The tex library

The `tex` table contains a large list of virtual internal T_EX parameters that are partially writable.

The designation virtual means that these items are not properly defined in LUA, but are only frontends that are handled by a metatable that operates on the actual T_EX values. As a result, most of the LUA table operators (like `pairs` and `#`) do not work on such items.

At the moment, it is possible to access almost every parameter that has these characteristics:

- You can use it after `\the`
- It is a single token.

This excludes parameters that need extra arguments, like `\the\scriptfont`.

The subset comprising simple integer and dimension registers are writable as well as readable (stuff like `\tracingcommands` and `\parindent`).

4.1.1 Integer parameters

The integer parameters accept and return LUA numbers.

Read-write:

<code>tex.adjdemerits</code>	<code>tex.globaldefs</code>
<code>tex.binoppenalty</code>	<code>tex.hangafter</code>
<code>tex.brokenpenalty</code>	<code>tex.hbadness</code>
<code>tex.catcodetable</code>	<code>tex.holdinginserts</code>
<code>tex.clubpenalty</code>	<code>tex.hyphenpenalty</code>
<code>tex.day</code>	<code>tex.interlinepenalty</code>
<code>tex.defaultthyphenchar</code>	<code>tex.language</code>
<code>tex.defaultskewchar</code>	<code>tex.lastlinefit</code>
<code>tex.delimiterfactor</code>	<code>tex.lefthyphenmin</code>
<code>tex.displaywidowpenalty</code>	<code>tex.linepenalty</code>
<code>tex.doublehyphndemerits</code>	<code>tex.localbrokenpenalty</code>
<code>tex.endlinechar</code>	<code>tex.localinterlinepenalty</code>
<code>tex.errorcontextlines</code>	<code>tex.looseness</code>
<code>tex.escapechar</code>	<code>tex.mag</code>
<code>tex.exhyphenpenalty</code>	<code>tex.maxdeadcycles</code>
<code>tex.fam</code>	<code>tex.month</code>
<code>tex.finalhyphndemerits</code>	<code>tex.newlinechar</code>
<code>tex.floatingpenalty</code>	<code>tex.outputpenalty</code>



<code>tex.pausing</code>	<code>tex.predisplaypenalty</code>
<code>tex.pdfadjustinterwordglue</code>	<code>tex.pretolerance</code>
<code>tex.pdfadjustspacing</code>	<code>tex.relpenalty</code>
<code>tex.pdfappendkern</code>	<code>tex.righthyphenmin</code>
<code>tex.pdfcompresslevel</code>	<code>tex.savinghyphcodes</code>
<code>tex.pdfdecimaldigits</code>	<code>tex.savingvdiscards</code>
<code>tex.pdfforcepagebox</code>	<code>tex.showboxbreadth</code>
<code>tex.pdfgamma</code>	<code>tex.showboxdepth</code>
<code>tex.pdfgentounicode</code>	<code>tex.time</code>
<code>tex.pdfimageapplygamma</code>	<code>tex.tolerance</code>
<code>tex.pdfimagegamma</code>	<code>tex.tracingassigns</code>
<code>tex.pdfimagehicolor</code>	<code>tex.tracingcommands</code>
<code>tex.pdfimageresolution</code>	<code>tex.tracinggroups</code>
<code>tex.pdfinclusionerrorlevel</code>	<code>tex.tracingifs</code>
<code>tex.pdfminorversion</code>	<code>tex.tracinglostchars</code>
<code>tex.pdfmovechars</code>	<code>tex.tracingmacros</code>
<code>tex.pdfobjcompresslevel</code>	<code>tex.tracingnesting</code>
<code>tex.pdfoptionalwaysusepdfpagebox</code>	<code>tex.tracingonline</code>
<code>tex.pdfoptionpdfinclusionerrorlevel</code>	<code>tex.tracingoutput</code>
<code>tex.pdfoptionpdfminorversion</code>	<code>tex.tracingpages</code>
<code>tex.pdfoutput</code>	<code>tex.tracingparagraphs</code>
<code>tex.pdfpagebox</code>	<code>tex.tracingrestores</code>
<code>tex.pdfpkresolution</code>	<code>tex.tracingscantokens</code>
<code>tex.pdfprependkern</code>	<code>tex.tracingstats</code>
<code>tex.pdfprotrudechars</code>	<code>tex.uchyph</code>
<code>tex.pdftracingfonts</code>	<code>tex.vbadness</code>
<code>tex.pdfuniqueresname</code>	<code>tex.widowpenalty</code>
<code>tex.postdisplaypenalty</code>	<code>tex.year</code>
<code>tex.predisplaydirection</code>	



Read-only:

<code>tex.deadcycles</code>	<code>tex.parshape</code>	<code>tex.spacefactor</code>
<code>tex.insertpenalties</code>	<code>tex.prevgraf</code>	

4.1.2 Dimension parameters

The dimension parameters accept LUA numbers (signifying scaled points) or strings (with included dimension). The result is always a string.

Read-write:

<code>tex.boxmaxdepth</code>	<code>tex.overfullrule</code>	<code>tex.pdfpageheight</code>
<code>tex.delimitershortfall</code>	<code>tex.pagebottomoffset</code>	<code>tex.pdfpagewidth</code>
<code>tex.displayindent</code>	<code>tex.pageheight</code>	<code>tex.pdfpxdimen</code>
<code>tex.displaywidth</code>	<code>tex.pagerightoffset</code>	<code>tex.pdfthreadmargin</code>
<code>tex.emergencystretch</code>	<code>tex.pagewidth</code>	<code>tex.pdfvorigin</code>
<code>tex.hangindent</code>	<code>tex.parindent</code>	<code>tex.predisplaysize</code>
<code>tex.hfuzz</code>	<code>tex.pdfdestmargin</code>	<code>tex.scriptspace</code>
<code>tex.hoffset</code>	<code>tex.pdfeachlinedepth</code>	<code>tex.splitmaxdepth</code>
<code>tex.hsize</code>	<code>tex.pdfeachlineheight</code>	<code>tex.vfuzz</code>
<code>tex.lineskiplimit</code>	<code>tex.pdffirstlineheight</code>	<code>tex.voffset</code>
<code>tex.mathsurround</code>	<code>tex.pdfhorigin</code>	<code>tex.vsize</code>
<code>tex.maxdepth</code>	<code>tex.pdflastlinedepth</code>	
<code>tex.nulldelimiterspace</code>	<code>tex.pdflinkmargin</code>	

Read-only:

<code>tex.pagedepth</code>	<code>tex.pagegoal</code>	<code>tex.prevdepth</code>
<code>tex.pagefilllstretch</code>	<code>tex.pageshrink</code>	
<code>tex.pagefillstretch</code>	<code>tex.pagestretch</code>	
<code>tex.pagefilstretch</code>	<code>tex.pagetotal</code>	

4.1.3 Direction parameters

All direction parameters are read-only and return a LUA string

<code>tex.bodydir</code>	<code>tex.pagedir</code>	<code>tex.textdir</code>
<code>tex.mathdir</code>	<code>tex.pardir</code>	

4.1.4 Glue parameters

All glue parameters are read-only and return a LUA string

<code>tex.abovedisplayshortskip</code>	<code>tex.belowdisplayskip</code>	<code>tex.parskip</code>
<code>tex.abovedisplayskip</code>	<code>tex.leftskip</code>	<code>tex.rightskip</code>
<code>tex.baselineskip</code>	<code>tex.lineskip</code>	<code>tex.spaceskip</code>
<code>tex.belowdisplayshortskip</code>	<code>tex.parfillskip</code>	<code>tex.splittopskip</code>



<code>tex.tabskip</code>	<code>tex.xspaceskip</code>
<code>tex.topskip</code>	

4.1.5 Muglue parameters

All muglue parameters are read-only and return a Lua string

<code>tex.medmuskip</code>	<code>tex.thinmuskip</code>
<code>tex.thickmuskip</code>	

4.1.6 Tokenlist parameters

All tokenlist parameters are read-only and return a Lua string

<code>tex.errhelp</code>	<code>tex.everyjob</code>	<code>tex.pdfpageattr</code>
<code>tex.everycr</code>	<code>tex.verymath</code>	<code>tex.pdfpageresources</code>
<code>tex.everydisplay</code>	<code>tex.verypar</code>	<code>tex.pdfpagesattr</code>
<code>tex.everyeof</code>	<code>tex.veryvbox</code>	<code>tex.pdfpkmode</code>
<code>tex.veryhbox</code>	<code>tex.output</code>	

4.1.7 Convert commands

The supported commands at this moment are:

<code>tex.AlephVersion</code>	<code>tex.eTeXVersion</code>	<code>tex.pdfnormaldeviate</code>
<code>tex.Alephrevision</code>	<code>tex.eTeXrevision</code>	<code>tex.pdftebanner</code>
<code>tex.OmegaVersion</code>	<code>tex.formatname</code>	<code>tex.pdfterevision</code>
<code>tex.Omegarevision</code>	<code>tex.jobname</code>	

All convert commands are read-only and return a Lua string

This list looks haphazard, but it really is not. These are all the cases of the convert internal command that do not require an argument.

4.1.8 attribute, count, dimension and token registers

T_EX's attributes (`\attribute`), counters (`\count`), dimensions (`\dimen`) and token (`\toks`) registers can be accessed and written to using four virtual sub-tables of the `tex` table:

<code>tex.attribute</code>	<code>tex.dimen</code>
<code>tex.count</code>	<code>tex.toks</code>

It is possible to use the names of relevant `\attributedef`, `\countdef`, `\dimendef`, or `\toksdef` control sequences as indices to these tables:

```
tex.count.scratchcounter = 0
enormous = tex.dimen["maxdimen"]
```



In this case, L^AT_EX looks up the value for you on the fly. You have to use a valid `\countdef` (or `\attributedef`, or `\dimendef`, or `\toksdef`), anything else will generate an error (the intent is to eventually also allow `<chardef tokens>` and even macros that expand into a number)

The attribute and count registers accept and return Lua numbers.

The dimension registers accept Lua numbers (in scaled points) or strings (with an included absolute dimension; `em` and `ex` and `px` are forbidden). The result is always a number in scaled points.

The token registers accept and return Lua strings. Lua strings are converted to token lists using `\the\toks` style expansion.

As an alternative to array addressing, there are also accessor functions defined:

```
tex.setdimen(number n, string s)
tex.setdimen(string s, string s)
tex.setdimen(number n, number n)
tex.setdimen(string s, number n)
number n = tex.getdimen(number n)
number n = tex.getdimen(string s)
```

```
tex.setcount(number n, number n)
tex.setcount(string s, number n)
number n = tex.getcount(number n)
number n = tex.getcount(string s)
```

```
tex.settoks (number n, string s)
tex.settoks (string s, string s)
string s = tex.gettoks (number n)
string s = tex.gettoks (string s)
```

4.1.9 Box registers

The current dimensions of `\box` registers can be read and altered using three other virtual sub-tables :

```
tex.wd
tex.ht
tex.dp
```

These are indexed strictly by number.

The box size registers accept Lua numbers (in scaled points) or strings (with included dimension). The result is always a number in scaled points.

As an alternative to array addressing, there are also accessor functions defined:



```
tex.setboxwd(number n, number n)
number n = tex.getboxwd(number n)
```

```
tex.setboxht(number n, number n)
number n = tex.getboxht(number n)
```

```
tex.setboxdp(number n, number n)
number n = tex.getboxdp(number n)
```

It is also possible to set and query actual boxes, using the node interface as defined in the [node](#) library:

```
tex.box
```

for array access, or

```
tex.setbox(number n, node s)
node n = tex.getbox(number n)
```

for function-based access

Be warned that an assignment like

```
tex.box[0] = tex.box[2]
```

does not copy the node list, it just duplicates a node pointer. If `\box2` will be cleared by T_EX commands later on, the contents of `\box0` becomes invalid as well. To prevent this from happening, always use `node.copy_list()` unless you are assigning to a temporary variable:

```
tex.box[0] = node.copy_list(tex.box[2])
```

4.1.10 Print functions

The `tex` table also contains the three print functions that are the major interface from LUA scripting to T_EX.

The arguments to these three functions are all stored in an in-memory virtual file that is fed to the T_EX scanner as the result of the expansion of `\directlua`.

The total amount of returnable text from a `\directlua` command is only limited by available system RAM. However, each separate printed string has to fit completely in T_EX's input buffer.

4.1.10.1 tex.print

```
tex.print(<string s>, ...)
tex.print(<number n>, <string s>, ...)
```

Each string argument is treated by T_EX as a separate input line.



The optional parameter can be used to print the strings using the catcode regime defined by `\catcodetable n`. If `n` is not a valid catcode table, then it is ignored, and the currently active catcode regime is used instead.

The very last string of the very last `tex.print()` command in a `\directlua` will not have the `\endlinechar` appended, all others do.

4.1.10.2 `tex.sprint`

```
tex.sprint(<string s>, ...)
tex.sprint(<number n>, <string s>, ...)
```

Each string argument is treated by T_EX as a special kind of input line that makes it suitable for use as a partial line input mechanism:

- T_EX does not switch to the new line state, so that leading spaces are not ignored.
- No `\endlinechar` is inserted.
- Trailing spaces are not removed.

4.1.10.3 `tex.write`

```
tex.write(<string s>, ...)
```

Each string argument is treated by T_EX as a special kind of input line that makes it suitable for use as a quick way to dump information:

- All catcodes on that line are either space (for " ") or character (for all others).
- There is no `\endlinechar` appended.

4.1.11 Helper functions

4.1.11.1 `tex.round`

```
number n = tex.round(number o)
```

Rounds lua number `o`, and returns a number that is in the range of a valid T_EX register value. If the number starts out of range, it generates a ‘Number too big’ error as well.

4.1.11.2 `tex.scale`

```
number n = tex.scale(number o, number delta)
table n = tex.scale(table o, number delta)
```



Multiplies the lua numbers `o` and `delta`, and returns a rounded number that is in the range of a valid T_EX register value. In the table version, it creates a copy of the table with all numeric top-level values scaled in that manner. If the multiplied number(s) are of range, it generates ‘Number to big’ error(s) as well.

4.2 The token library

The `token` table contains interface functions to T_EX’s handling of tokens. These functions are most useful when combined with the `token_filter` callback, but they could be used standalone as well.

A token is represented in LUA as a small table. For the moment, this table consists of three numeric entries:

nr	meaning	description
1	command code	this is a value between 0 and 130 (approximately)
2	command modifier	this is a value between 0 and 2^{21}
3	control sequence id	for commands that are not the result of control sequences, like letters and characters, it is zero, otherwise, it is number pointing into the equivalence table

4.2.1 `token.get_next`

```
token t = token.get_next()
```

This fetches the next input token from the current input source, without expansion.

4.2.2 `token.is_expandable`

```
boolean b = token.is_expandable(token t)
```

This tests if the token `t` could be expanded.

4.2.3 `token.expand`

```
token.expand()
```

If a token is expandable, this will expand one level of it, so that the first token of the expansion will now be the next token to be read by `tex.get_next()`.

4.2.4 `token.is_activechar`

```
boolean b = token.is_activechar(token t)
```



This is a special test that is sometimes handy. Discovering whether some token is the result of an active character turned out to be very hard otherwise.

4.2.5 `token.create`

```
token t = token.create(string csname)
token t = token.create(number charcode)
token t = token.create(number charcode, number catcode)
```

This is the token factory. If you feed it a string, then it is a control sequence `csname`, and it will be looked up in the equivalence table.

If you feed it number, then this is assumed to be an input character, and an optional second number gives its category code. This means it is possible to overrule a character's category code, with a few exceptions: the category codes 0 (escape), 9 (ignored), 13 (active), 14 (comment), and 15 (invalid) cannot occur inside a token. The values 0, 9, 14 and 15 are therefore illegal as input to `token.create()`, and active characters will be resolved immediately.

Note: unknown string sequences and never defined active characters will result in a token representing an undefined control sequence with a near-random name. It is *not* possible to define brand new control sequences using `token.create`!

4.2.6 `token.command_name`

```
string commandname = token.command_name(token t)
```

This returns the name associated with the command value of the token in L^AT_EX. There is no direct connection between these names and primitives. For instance, all `\ifxxx` tests are grouped under `if_fest`, and the command modifier defines which test is to be run.

4.2.7 `token.command_id`

```
number i = token.command_name(string commandname)
```

This returns a number that is the inverse operation of the previous command, to be used as the first item in a token table.

4.2.8 `token.csname_name`

```
string csname = token.csname_name(token t)
```



This returns the name associated with the equivalence table value of the token in `LUATEX`. It returns the string value of the command used to create the current token, or an empty string if there is no associated control sequence.

4.2.9 `token.csname_id`

```
number i = token.csname_id(string csname)
```

This returns a number that is the inverse operation of the previous command, to be used as the third item in a token table.

4.3 The node library

The `node` library contains functions that facilitate dealing with (lists of) nodes and their values.

`LUATEX` nodes are represented in `LUA` as userdata with the metadata type `luatex.node`. The various fields within a node can be accessed using fields that are indexed by strings.

Each node has at least the three fields `next`, `id`, and `subtype`:

- The `next` field returns the userdata object for the next node in a linked list of nodes, or `nil`, if there is no next node.
- The `id` indicates `TEX`'s node type. The field `id` has a numeric value for efficiency reasons, but some of the library functions also accept a string value instead of `id`.
- The `subtype` is another number. It often gives further information about a node of a particular `id`, but it is most important when dealing with whatsits, because they are differentiated solely based on their `subtype`.

The other available fields depend on the `id` (and for whatsits, the `subtype`) of the node. Further details on the various fields and their meanings are given in the chapter 6.

`TEX`'s math nodes are not yet supported: there is not yet an interface to the internals of the math list and it is not possible to create them from `LUA`. Support for `unset` (alignment) nodes is partial: they can be queried and modified from `LUA` code, but not created.

Nodes can be compared to each other, but: you are actually comparing indices into the node memory. This means that equality tests can only be trusted under very limited conditions. It will not work correctly in any situation where one of the two nodes has been freed and/or reallocated: in that case, there will be false positives.

4.3.1 Node handling functions

4.3.1.1 `node.types`

```
table t = node.types()
```



This function returns an array that maps node id numbers to node type strings, providing an overview of the possible top-level `id` types.

4.3.1.2 `node.whatsits`

```
table t = node.whatsits()
```

TeX's whatsits all have the same `id`. The various subtypes are defined by their `subtype`. The function is much like `node.id`, except that it provides an array of `subtype` mappings.

4.3.1.3 `node.id`

```
number id = node.id(string type)
```

This converts a single type name to its internal numeric representation.

4.3.1.4 `node.subtype`

```
number subtype = node.subtype(string type)
```

This converts a single whatsit name to its internal numeric representation (`subtype`).

4.3.1.5 `node.type`

```
string type = node.type(number id)
```

This converts a internal numeric representation to an external string representation.

4.3.1.6 `node.fields`

```
table t = node.fields(string type)
table t = node.fields(string type, string subtype)
```

This function returns an array of valid field names for a particular type of node. If you want to get the valid fields for a whatsit, then you have to supply the second argument also. Otherwise, the second argument will be silently ignored.

This function accepts numeric `id` and `subtype` values as well.

4.3.1.7 `node.has_field`

```
boolean t = node.has_field(node n, string field)
```



This function returns a boolean that is only true if `n` is actually a node, and it has the field.

4.3.1.8 `node.new`

```
node n = node.new(string type)
node n = node.new(string type, string subtype)
```

Creates a new node. All fields are initialized to either zero or null (depending on whether the field is a number or a pointer), except for `id` and `subtype` (if supplied). If you need a new `whatsit`, then the second argument is required.

This function accepts numeric `id` and `subtype` values as well.

4.3.1.9 `node.free`

```
node.free(node n)
```

Removes the node `m` from $\text{T}_{\text{E}}\text{X}$'s memory. Be careful: no checks are done on whether this node is still pointed to from a register or some `next` field: it is up to you to make sure that the internal data structures remain correct.

4.3.1.10 `node.flush_list`

```
node.flush_list(node n)
```

Removes the node list `m` and the complete node list following `m` from $\text{T}_{\text{E}}\text{X}$'s memory. Be careful: no checks are done on whether any of these nodes is still pointed to from a register or some `next` field: it is up to you to make sure that the internal data structures remain correct.

4.3.1.11 `node.copy`

```
node m = node.copy(node n)
```

Creates a deep copy of node `n`. Only the `next` field is not copied.

4.3.1.12 `node.copy_list`

```
node m = node.copy_list(node n)
```

Creates a deep copy of the node list that starts at `n`.



4.3.1.13 node.hpack

```
node h = node.hpack(node n)
node h = node.hpack(node n, number w, string info)
```

This function creates a new hlist by packagin the list that with node `n` into a horizontal box. With only a single argument, this box is created using the natural width of it's components. In the three argument form, `info` must be either `additional` or `exactly`, and `w` is the additional (`\hbox spread`) or exact (`\hbox to`) width to be used.

Caveat: at this moment, there can be unintended side-effects to this function, like updating some of the `\marks` and `\inserts`.

4.3.1.14 node.slide

```
node m = node.slide(node n)
```

Returns the last node of the node list that starts at `n`.

4.3.1.15 node.length

```
number i = node.length(node n)
number i = node.length(node n, node m)
```

Returns the number of nodes contained in the node list that starts at `n`. If `m` is also supplied it stops at `m` instead of at the end of the list. The node `m` is not counted.

4.3.1.16 node.count

```
number i = node.count(string type, node n)
number i = node.count(string type, node n, node m)
```

Returns the number of nodes contained in the node list that starts at `n` that have an `id` matching `type`. If `m` is also supplied, counting stops at `m` instead of at the end of the list. The node `m` is not counted.

4.3.1.17 node.traverse

```
node t = node.traverse(node n)
node t = node.traverse(node n, node m)
```

This is an iterator that loops over the node list that starts at `n`. If `m` is also supplied, the iterator stops at `m` instead of at the end of the list. The node `m` is not processed.



4.3.1.18 node.traverse_id

```
node t = node.traverse_id(string type, node n, node m)
node t = node.traverse_id(string type, node n)
```

This is an iterator that loops over all the nodes in the list that starts at `n` that have an `id` matching `type`. If `m` is also supplied, the iterator stops at `m` instead of at the end of the list. The node `m` is not processed.

4.3.1.19 node.remove

```
node head, current = node.remove(node head, node current)
```

This function removes the node `current` from the list following `head`. It is your responsibility to make sure it is really part of that list. The return values are the new `head` and `current` nodes. The returned `current` is the node in the calling argument, and is only passed back as a convenience (it's `next` field will be cleared). The returned `head` is more important, because: if the function is called with `current` equal to `head`, it will be changed.

4.3.1.20 node.insert_before

```
node head, new = node.insert_before(node head, node current, node new)
```

This function inserts the node `new` before `current` into the list following `head`. It is your responsibility to make sure that `current` is really part of that list. The return values are the (potentially mutated) `head` and the `new`, set up to be part of the list (with correct `next` field). If `head` is initially `nil`, it will become `new`.

4.3.1.21 node.insert_after

```
node head, new = node.insert_after(node head, node current, node new)
```

This function inserts the node `new` after `current` into the list following `head`. It is your responsibility to make sure that `current` is really part of that list. The return values are the `head` and the `new`, set up to be part of the list (with correct `next` field). If `head` is initially `nil`, it will become `new`.

4.3.2 Attribute handling

Attributes appear as linked list of userdata objects in the `attr` field of individual nodes. They can be handled individually, but it much safer and more efficient to use the dedicated functions associated with them.



4.3.2.1 node.has_attribute

```
number v = node.has_attribute(node n, number id)
number v = node.has_attribute(node n, number id, number val)
```

Tests if a node has the attribute with number `id` set. If `val` is also supplied, also tests if the value matches `val`. It returns the value, or, if no match is found, `nil`.

4.3.2.2 node.set_attribute

```
node.set_attribute(node n, number id, number val)
```

Sets the attribute with number `id` to the value `val`. Duplicate assignments are ignored.

4.3.2.3 node.unset_attribute

```
number v = node.unset_attribute(node n, number id, number val)
number v = node.unset_attribute(node n, number id)
```

Unsets the attribute with number `id`. If `val` is also supplied, it will only perform this operation if the value matches `val`. Missing attributes or attribute-value pairs are ignored.

If the attribute was actually deleted, returns its old value. Otherwise, returns `nil`.

4.4 The texio library

This library takes care of the low-level I/O interface.

4.4.1 Printing functions

4.4.1.1 texio.write

```
texio.write(string target, string s)
texio.write(string s)
```

Without the `target` argument, Writes the string to the same location(s) `TEX` writes messages to at this moment. If `\batchmode` is in effect, it writes only to the log, otherwise it writes to the log and the terminal.

The optional `target` can be one of three possibilities: `term`, `log` or `term and log`.



4.4.1.2 tex.write_nl

```
texio.write_nl(string target, string s)
texio.write_nl(string s)
```

Like `texio.write`, but make sure that the string `s` will appear at the beginning of a line. You can use an empty string if you only want to move to the next line.

4.5 The pdf library

This table contains the current `h` en `v` values that define the location on the output page. The values can be queried and set using scaled points as units.

```
pdf.v
pdf.h
```

The associated function calls are

```
pdf.setv(number n)
number n = pdf.getv()
pdf.seth(number n)
number n = pdf.geth()
```

It also holds a print function to write stuff to the PDF document, that can be used from within a `\lualatex` argument. This function is not to be used inside `\directlua` unless you know *exactly* what you are doing.

```
pdf.print
```

```
pdf.print(<string s>)
pdf.print(<string type>, <string s>)
```

The optional parameter can be used to mimic the behaviour of `\pdfliteral`: the `type` is `direct` or `page`.

4.6 The callback library

This library has functions that register, find and list callbacks.

The `callback` library is only available in LUA state zero (0).

```
callback.register(string <callback name>,function <callback_func>)
callback.register(string <callback name>,nil)
```

where the `<callback name>` is a predefined callback name, see below.



L^AT_EX internalizes the callback function in such a way that it does not matter if you redefine a function accidentally.

Callback assignments are always global. You can use the special value `nil` instead of a function for clearing the callback.

```
table <info> = callback.list()
```

The keys in the table are the known callback names, the value is a boolean where `true` means that the callback is currently set (active).

```
function <f> = callback.find(<callback name>)
```

If the callback is not set, `callback.find` returns `nil`.

4.6.1 File discovery callbacks

4.6.1.1 find_read_file and find_write_file

Your callback function should have the following conventions:

```
string <actual_name> = function (number <id_number>, string <asked_name>)
```

Arguments:

`id_number`

This number is zero for the log or `\input` files. For T_EX's `\read` or `\write` the number is incremented by one, so `\read0` becomes 1.

`asked_name`

This is the user-supplied filename, as found by `\input`, `\openin` or `\openout`.

Return value:

`actual_name`

This is the filename used. For the very first file that is read in by T_EX, you have to make sure you return an `actual_name` that has an extension and that is suitable for use as `jobname`. If you don't, you will have to manually fix the name for the log file and output file, and an eventual format filename will become mangled, since these depend on the `jobname`.

You have to return `nil` if the file cannot be found.

4.6.1.2 find_font_file

Your callback function should have the following conventions:

```
string <actual_name> = function (string <asked_name>)
```



The `asked_name` is an OTF or TFM font metrics file.

Return `nil` if the file cannot be found.

4.6.1.3 `find_output_file`

You callback function should have the following conventions:

```
string <actual_name> = function (string <asked_name>)
```

The `asked_name` is the PDF or DVI file for writing.

4.6.1.4 `find_format_file`

You callback function should have the following conventions:

```
string <actual_name> = function (string <asked_name>)
```

The `asked_name` is a format file for reading (the format file for writing is always opened in the current directory).

4.6.1.5 `find_vf_file`

Like `find_font_file`, but for virtual fonts. This applies to both ALEPH's ovf files and traditional Knuthian vf files.

4.6.1.6 `find_ocp_file`

Like `find_font_file`, but for ocp files.

4.6.1.7 `find_map_file`

Like `find_font_file`, but for map files.

4.6.1.8 `find_enc_file`

Like `find_font_file`, but for enc files.

4.6.1.9 `find_sfd_file`

Like `find_font_file`, but for subfont definition files.



4.6.1.10 find_pk_file

Like `find_font_file`, but for pk bitmap files. The argument `<name>` is a bit special in this case. It's form is

```
<base res>dpi/<fontname>.<actual res>pk
```

So you may be asked for `600dpi/manfnt.720pk`. It is up to you to find a reasonable bitmap file to go with that specification.

4.6.1.11 find_data_file

Like `find_font_file`, but for embedded files (`\pdfobj file "..."`).

4.6.1.12 find_opentype_file

Like `find_font_file`, but for `OPENTYPE` font files.

4.6.1.13 find_truetype_file and find_type1_file

Your callback function should have the following conventions:

```
string <actual_name> = function (string <asked_name>)
```

The `asked_name` is a font file. This callback is called while `LUATEX` is building its internal list of needed font files, so the actual timing may surprise you. Your return value is later fed back into the matching `read_file` callback.

Strangely enough, `find_type1_file` is also used for `OPENTYPE` (OTF) fonts.

4.6.1.14 find_image_file

Your callback function should have the following conventions:

```
string <actual_name> = function (string <asked_name>)
```

The `asked_name` is an image file. Your return value is used to open a file from the harddisk, so make sure you return something that is considered the name of a valid file by your operating system.

4.6.2 File reading callbacks

4.6.2.1 open_read_file

Your callback function should have the following conventions:



```
table <env> = function (string <file_name>)
```

Argument:

`file_name`

the filename returned by a previous `find_read_file` or the return value of `kpse_find_file()` if there was no such callback defined.

Return value:

`env`

this is a table containing at least one required and one optional callback functions for this file. The required field is `reader` and the associated function will be called once for each new line to be read, the optional one is `close` that will be called once when L^AT_EX is done with the file.

L^AT_EX never looks at the rest of the table, so you can use it to store your private per-file data. Both the callback functions will receive the table as their only argument.

4.6.2.1.1 reader

L^AT_EX will run this function whenever it needs a new input line from the file.

```
function(table <env>)  
    return string <line>  
end
```

Your function should return either a string or `nil`. The value `nil` signals that the end of file has occurred, and will make T_EX call the optional `close` function next.

4.6.2.1.2 close

L^AT_EX will optionally run this function when it needs to close the file.

```
function(table <env>)  
    return  
end
```

Your function should not return any value.

4.6.2.2 read_font_file

This function is called when T_EX needs to read a `ofm` or `tfm` file.

```
function(string <name>)  
    return boolean <success>, string <data>, number <data_size>  
end
```



success

return false when a fatal error occurred (e.g. when the file cannot be found, after all).

data

the bytes comprising the file.

data_size

the length of the `data`, in bytes.

return an empty string and zero if the file was found but there was a reading problem.

4.6.2.3 `read_vf_file`

Like `read_font_file`, but for virtual fonts.

4.6.2.4 `read_ocp_file`

Like `read_font_file`, but for ocp files.

4.6.2.5 `read_map_file`

Like `read_font_file`, but for map files.

4.6.2.6 `read_enc_file`

Like `read_font_file`, but for enc files.

4.6.2.7 `read_sfd_file`

Like `read_font_file`, but for subfont definition files.

4.6.2.8 `read_pk_file`

Like `read_font_file`, but for pk bitmap files.

4.6.2.9 `read_data_file`

Like `read_font_file`, but for embedded files (`\pdfobj file "..."`).

4.6.2.10 `read_truetype_file`

Like `read_font_file`, but for TRuetype font files. The `name` is a path name as returned by `find_truetype_file` or `kpse_find_file`.



4.6.2.11 read_type1_file

Like `read_font_file`, but for type1 font files. The `name` is a path name as returned by `find_type1_file` or `kpse_find_file`.

4.6.2.12 read_opentype_file

Like `read_font_file`, but for OPENTYPE font files. The `name` is a path name as returned by `find_type1_file` or `kpse_find_file`.

4.6.3 Data processing callbacks

4.6.3.1 process_input_buffer

This callback allows you to change the contents of the line input buffer just before L^AT_EX actually starts looking at it.

```
function(string <buffer>)  
    return string <adjusted_buffer>  
end
```

If you return `nil`, L^AT_EX will pretend like your callback never happened. You can gain a small amount of processing time from that.

4.6.3.2 token_filter

This callback allows you to change the fetch and preprocess any lexical token that enters L^AT_EX, before L^AT_EX executes or expands the associated command.

```
function()  
    return table <token>  
end
```

The calling convention for this callback is bit more complicated than for most other callbacks. The function should either return a Lua table representing a valid to-be-processed token or tokenlist, or something else like `nil` or an empty table.

If your Lua function does not return a table representing a valid token, it will be immediately called again, until it eventually does return a useful token or tokenlist (or until you reset the callback value to `nil`). See the description of `token` for some handy functions to be used in conjunction with this callback.

If your function returns a single usable token, then that token will be processed by L^AT_EX immediately. If the function returns a token list (a table consisting of a list of consecutive token tables), then that list



will be pushed to the input stack as completely new token list level, with it's token type set to inserted. In either case, the returned token(s) will not be fed back into the callback function.

4.6.4 Node list processing callbacks

The description of nodes and node lists is in the chapter 6.

4.6.4.1 `buildpage_filter`

This callback is called whenever L^AT_EX is ready to move stuff to the main vertical list. You can use this callback to do specialized manipulation of the page building stage like imposition or column balancing.

```
function(node <head>, string <extrainfo>)  
    return true | false | node <newhead>  
end
```

As for all the callbacks that deal with nodes, the return value can be one of three things:

- `boolean true` signals succesful processing
- `node` signals that the head node should be replaced by this node
- `boolean false` signals that the head node list should be ignored and flushed from memory

The string `extrainfo` gives some additional information about what T_EX's state is with respect to the 'current page'. The possible values are:

value	explanation
<code>alignment</code>	a (partial) alignment is being added
<code>box</code>	a typeset box is being added
<code>begin_of_par</code>	the beginning of a new paragraph
<code>vmode_par</code>	<code>\par</code> was found in vertical mode
<code>hmode_par</code>	<code>\par</code> was found in horizontal mode
<code>insert</code>	an insert is added
<code>penalty</code>	a penalty (in vertical mode)
<code>before_display</code>	immediately before a display starts
<code>after_display</code>	a display is finished

4.6.4.2 `pre_linebreak_filter`

This callback is called just before L^AT_EX starts converting a list of nodes into a stack of `\hboxes`. The removal of any final skip and the subsequent insertion of `\parfillskip` has not happened yet at that moment.

```
function(node <head>, string <groupcode>, int <glyph_count>)  
    return true | false | node <newhead>  
end
```



The string called `groupcode` identifies the nodelist's context within T_EX's processing. The range of possibilities is given in the table below, but not all of those can actually appear in `pre_linebreak_filter`, some are for the `Xpack_filter` callbacks that will be explained in the next two paragraphs.

value	explanation
<code>hbox</code>	<code>\hbox</code> in horizontal mode
<code>adjusted_hbox</code>	<code>\hbox</code> in vertical mode
<code>vbox</code>	<code>\vbox</code>
<code>vtop</code>	<code>\vtop</code>
<code>align</code>	<code>\halign</code> or <code>\valign</code>
<code>disc</code>	discretionaries
<code>insert</code>	packaging an insert
<code>vcenter</code>	<code>\vcenter</code>
<code>local_box</code>	<code>\localleftbox</code> or <code>\localrightbox</code>
<code>split_off</code>	top of a <code>\vsplit</code>
<code>split_keep</code>	remainder of a <code>\vsplit</code>
<code>preamble</code>	alignment preamble
<code>align_set</code>	alignment cell
<code>fin_row</code>	alignment row

4.6.4.3 `hpack_filter`

This callback is called when T_EX is ready to start boxing some horizontal mode material. Math items are ignored at the moment.

```
function(node <head>, number <size>, string <packtype>,
         string <groupcode>, number <glyphcount>)
    return true | false | node <newhead>
end
```

The `packtype` is either `additional` or `exactly`. If `additional`, then the `size` is a `\hbox spread ...` argument. If `exactly`, then the `size` is a `\hbox to ...`. In both cases, the number is in scaled points.

4.6.4.4 `vpack_filter`

This callback is called when T_EX is ready to start boxing some vertical mode material. Math displays are ignored at the moment.

This function is very similar to the `hpack_filter`. Besides the fact that it is called at different moments, there is an extra variable that matches T_EX's `\maxdepth` setting, but there is no glyph count.



```
function(node <head>, number <size>, string <packtype>,
         number <maxdepth>, string <groupcode>)
  return true | false | node <newhead>
end
```

4.6.4.5 pre_output_filter

This callback is called when T_EX is ready to start boxing the box 255 for `\output`.

```
function(node <head>, number <size>, string <packtype>,
         number <maxdepth>, string <groupcode>)
  return true | false | node <newhead>
end
```

4.6.5 Information reporting callbacks

4.6.5.1 start_run

```
function()
```

Replaces the code that prints L^AT_EX's banner

4.6.5.2 stop_run

```
function()
```

Replaces the code that prints L^AT_EX's statistics and output written to messages.

4.6.5.3 start_page_number

```
function()
```

Replaces the code that prints the `[` and the page number at the begin of `\shipout`. This callback will also override the printing of box information that normally takes place when `\tracingoutput` is positive.

4.6.5.4 stop_page_number

```
function()
```



Replaces the code that prints the] at the end of \shipout

4.6.5.5 show_error_hook

```
function()  
    return  
end
```

This callback is run from inside the T_EX error function, and the idea is to allow you to do some extra reporting on top of what T_EX already does (none of the normal actions are removed). You may find some of the values in the `status` table useful.

message

is the formal error message T_EX has given to the user (the line after the "!")

indicator

is either a filename (when it is a string) or a location indicator (a number) that can mean lots of different things like a token list id or a \read number.

lineno

is the current line number

This is an investigative item only, only for 'testing the water'.

The final goal is the total replacement of T_EX's error handling routines, but that needs lots of adjustments in the web source because T_EX deals with errors in a somewhat haphazard fashion.

4.6.6 Font-related callbacks

4.6.6.1 define_font

```
function(string <name>, number <size>, number <id>)  
    return table <font>  
end
```

The string `<name>` is the filename part of the font specification, as given by the user.

The number `<size>` is a bit special:

- if it is positive, it specifies an at size in scaled points.
- if it is negative, its absolute value represents a scaled setting relative to the designsize of the font.

The internal structure of the `` table that is to be returned is explained in chapter 5. That table is saved internally, so you can put extra fields in the table for your later Lua code to use.



4.7 The lua library

This library contains two read-only items:

4.7.1 Variables

```
number n = lua.id
```

This sets the id number of the instance.

```
string s = lua.version
```

This sets a L^AT_EX version identifier string (currently "0.1").

4.7.2 LUA bytecode registers

LUA registers can be used to communicate LUA functions across LUA states. The accepted values for assignments are functions and nil. Likewise, the retrieved value is either a function or nil.

```
lua.bytecode[n] = function () .. end  
lua.bytecode[n]()
```

The contents of the `lua.bytecode` array is stored inside the format file as actual LUA bytecode, so it can also be used to preload lua code.

The associated function calls are

```
function f = lua.getbytecode(number n)  
lua.setbytecode(number n, function f)
```

4.8 The kpse library

4.8.1 kpse.find_file

The most important function in the library is `find_file`:

```
string f = kpse.find_file(string filename)  
string f = kpse.find_file(string filename, string ftype)  
string f = kpse.find_file(string filename, boolean mustexist)  
string f = kpse.find_file(string filename, string ftype, boolean mustexist)  
string f = kpse.find_file(string filename, string ftype, number dpi)
```

Arguments:

filename

the name of the file you want to find, with or without extension.

type

maps to the `-format` argument of `KPSEWHICH`. The supported values are:



gf	TeX system documentation
pk	texpool
bitmap font	TeX system sources
tfm	PostScript header
afm	Troff fonts
base	type1 fonts
bib	vf
bst	dvips config
cnf	ist
ls-R	truetype fonts
fnt	type42 fonts
map	web2c files
mem	other text files
mf	other binary files
mfpool	misc fonts
mft	web
mp	cweb
mppool	enc files
MetaPost support	cmap files
ocp	subfont definition files
ofm	opentype fonts
opl	pdftex config
otp	lig files
ovf	texmfscripts
ovp	
graphic/figure	
tex	



The default type is `tex`.

`mustexist`

is similar to `KPSEWHICH`'s `-must-exist`, and the default is `false`. If you specify `true` (or a non-zero integer), then the `KPSE` library will search the disk as well as the `ls-R` databases.

`dpi`

This is used for the size argument of the formats `pk`, `gf`, and `bitmap font`.

4.8.2 `kpse.set_program_name`

Sets the `KPATHSEA` executable (and optionally program) name

```
kpse.set_program_name(string name)
```

```
kpse.set_program_name(string name, string progname)
```

The second argument controls the use of the dotted values in the `texmf.cnf` configuration file, and defaults to the first argument.

4.8.3 `kpse.init_prog`

Extra initialization for programs that need to generate bitmap fonts.

```
kpse.init_prog(string prefix, number base_dpi, string mfmode)
```

```
kpse.init_prog(string prefix, number base_dpi, string mfmode, string fall-back)
```

4.8.4 `kpse.readable_file`

Test if a (absolute) file name is a readable file

```
string f = kpse.readable_file(string name)
```

The return value is the actual absolute filename you should use, because the disk name is not always the same as the requested name, due to aliases and system-specific handling under e.g. `MSDOS`.

Returns `nil` if the file does not exist or is not readable.

4.8.5 `kpse.expand_path`

Like `kpsewhich`'s `-expand-path`:

```
string r = kpse.expand_path(string s)
```



4.8.6 `kpse.expand_var`

Like `kpsewhich`'s `-expand-var`:

```
string r = kpse.expand_var(string s)
```

4.8.7 `kpse.expand_braces`

Like `kpsewhich`'s `-expand-braces`:

```
string r = kpse.expand_braces(string s)
```

4.8.8 `kpse.var_value`

Like `kpsewhich`'s `-var-value`:

```
string r = kpse.var_value(string s)
```

4.9 The status library

This contains a number of run-time configuration items that you may find useful in message reporting, as well as an iterator function that gets all of the names and values as a table.

```
table <info> = status.list()
```

The keys in the table are the known items, the value is the current value.

Almost all of the values in `status` are fetched through a metatable at run-time whenever they are accessed, so you cannot use `pairs` on `status`, but you *can* use `pairs` on `<info>`, of course.

If you do not need the full list, you can also ask for a single item by using it's name as an index into `status`.

The current list is:

key	explanation
<code>pdf_gone</code>	written PDF bytes
<code>pdf_ptr</code>	not yet written pdf bytes
<code>dvi_gone</code>	written dvi bytes
<code>dvi_ptr</code>	not yet written dvi bytes
<code>total_pages</code>	number of written pages
<code>output_file_name</code>	name of the PDF or DVI file
<code>log_name</code>	name of the log file
<code>banner</code>	terminal display banner
<code>pdftex_banner</code>	



<code>var_used</code>	variable (one-word) memory in use
<code>dyn_used</code>	token (multi-word) memory in use
<code>str_ptr</code>	number of strings
<code>init_str_ptr</code>	number of INITEX strings
<code>max_strings</code>	maximum allowed strings
<code>pool_ptr</code>	string pool index
<code>init_pool_ptr</code>	INITEX string pool index
<code>pool_size</code>	maximum allowed string characters
<code>var_mem_max</code>	number of allocated words for nodes
<code>fix_mem_max</code>	number of allocated words for tokens
<code>fix_mem_end</code>	maximum number of used tokens
<code>cs_count</code>	number of control sequences
<code>hash_size</code>	size of hash
<code>hash_extra</code>	extra allowed hash
<code>font_ptr</code>	number of active fonts
<code>hyph_count</code>	hyphenation exceptions
<code>hyph_size</code>	max used hyphenation exceptions
<code>max_in_stack</code>	max used input stack entries
<code>max_nest_stack</code>	max used nesting stack entries
<code>max_param_stack</code>	max used parameter stack entries
<code>max_buf_stack</code>	max used buffer position
<code>max_save_stack</code>	max used save stack entries
<code>stack_size</code>	input stack size
<code>nest_size</code>	nesting stack size
<code>param_size</code>	parameter stack size
<code>buf_size</code>	line buffer size
<code>save_size</code>	save stack size
<code>obj_ptr</code>	max PDF object pointer
<code>obj_tab_size</code>	PDF object table size
<code>pdf_os_cntr</code>	max PDF object stream pointer
<code>pdf_os_objidx</code>	PDF object stream index
<code>pdf_dest_names_ptr</code>	max PDF destination pointer
<code>dest_names_size</code>	PDF destination table size
<code>pdf_mem_ptr</code>	max PDF memory used
<code>pdf_mem_size</code>	PDF memory size
<code>largest_used_mark</code>	max referenced marks class
<code>filename</code>	name of the current input file
<code>inputid</code>	numeric id of the current input
<code>linenumber</code>	location in the current input file
<code>lasterrorstring</code>	last error string
<code>luabytecodes</code>	number of active LUA bytecode registers
<code>luabytecode_bytes</code>	number of bytes in LUA bytecode registers
<code>luastates</code>	number of active LUA interpreters



luastate_bytes	number of bytes in use by LUA interpreters
output_active	<code>true</code> if the <code>\output</code> routine is active

4.10 The texconfig table

This is a table that is created empty. A startup LUA script could fill this table with a number of settings that are read out by the executable after loading and executing the startup file.

key	type	default	explanation
string_vacancies	number	75000	cf. web2c docs
pool_free	number	5000	cf. web2c docs
max_strings	number	15000	cf. web2c docs
strings_free	number	100	cf. web2c docs
trie_size	number	20000	cf. web2c docs
hyph_size	number	659	cf. web2c docs
nest_size	number	50	cf. web2c docs
max_in_open	number	15	cf. web2c docs
param_size	number	60	cf. web2c docs
save_size	number	4000	cf. web2c docs
stack_size	number	300	cf. web2c docs
dvi_buf_size	number	16384	cf. web2c docs
error_line	number	79	cf. web2c docs
half_error_line	number	50	cf. web2c docs
max_print_line	number	79	cf. web2c docs
ocp_list_size	number	1000	cf. web2c docs
ocp_buf_size	number	1000	cf. web2c docs
ocp_stack_size	number	1000	cf. web2c docs
hash_extra	number	0	cf. web2c docs
pk_dpi	number	72	cf. web2c docs
kpse_init	boolean	true	<code>false</code> totally disables KPATHSEA initialisation (only ever unset this if you implement <i>all</i> file find callbacks!)
trace_file_names	boolean	true	<code>false</code> disables T _E X's normal file open-close feedback (the assumption is that callbacks will take care of that)
src_special_auto	boolean	false	source specials sub-item
src_special_everypar	boolean	false	source specials sub-item
src_special_everyparend	boolean	false	source specials sub-item
src_special_everyocr	boolean	false	source specials sub-item
src_special_everymath	boolean	false	source specials sub-item
src_special_everyhbox	boolean	false	source specials sub-item
src_special_everyvbox	boolean	false	source specials sub-item
src_special_everydisplay	boolean	false	source specials sub-item



<code>file_line_error</code>	boolean	false	Do <code>file:line</code> style error messages
<code>halt_on_error</code>	boolean	false	Abort run on the first encountered error
<code>formatname</code>	string		If no format name was given on the commandline, this key will be tested first instead of simply quitting
<code>jobname</code>	string		If no input file name was given on the commandline, this key will be tested first instead of simply giving up

4.11 The font library

The font library will provide the interface into the internals of the font system, as well as contain some binary font loaders.

4.11.1 Loading a TFM file

```
table fnt = font.read_tfm(string name, number s)
```

The number is a bit special:

- if it is positive, it specifies an at size in scaled points.
- if it is negative, its absolute value represents a scaled setting relative to the designsize of the font.

The internal structure of the virtual font table that is returned is explained in chapter 5.

4.11.2 Loading a VF file

```
table vf_fnt = font.read_vf(string name, number s)
```

The number is a bit special:

- if it is positive, it specifies an at size in scaled points.
- if it is negative, its absolute value represents a scaled setting relative to the designsize of the font.

4.11.3 The fonts array

The whole table of T_EX fonts is accessible from lua using a virtual array.

```
font.fonts[n] = { ... }
table f = font.fonts[n]
```

See chapter 5 for the structure of the tables. Because this is a virtual array, you cannot call `pairs` on it, but see below for the `font.each` iterator.



The two metatable functions implementing the virtual array are:

```
table f = font.getfont(number n)
font.setfont(number n, table f)
```

Also note the following: Assignments can only be made to fonts that have already been defined in T_EX, but have not been accessed *at all* since that definition. This limits the usability of the write access to `font.fonts` quite a lot, a less stringent ruleset will likely be implemented later.

4.11.4 Checking a font's status

You can test for the status of a font by calling this function:

```
boolean f = font.frozen(number n)
```

The return value is one of true (unassignable), false (can be changed) or nil (not a valid font at all).

4.11.5 Defining a font directly

You can define your own font into `font.fonts`

```
number i = font.define(table f)
```

The return value is the internal id number of the defined font (the index into `font.fonts`). If the font creation fails, an error is raised. The table is a font structure, as explained in chapter 5.

4.11.6 Currently active font

```
number i = font.current();
```

This is the currently used font number.

4.11.7 Maximum font id

```
number i = font.max();
```

This is the largest used index in `font.fonts`.

4.11.8 Iterating over all fonts

```
for i,v in font.each() do
    ...
end
```



This is an iterator over each of the defined T_EX fonts. The first returned value is the index in `font.fonts`, the second the font itself, as a lua table. The indices are listed incrementally, but they do not always form an array of consecutive numbers: in some cases there can be holes in the sequence.

4.12 The fontforge library

4.12.1 Getting quick information on a font

```
local info = fontforge.info("filename")
```

This function returns either `nil`, or a `table` containing six fairly interesting information items from the font file:

key	type	explanation
fontname	string	
fullname	string	
familyname	string	
weight	string	
version	string	
italicangle	float	

Getting information through this function is (sometimes much) more efficient than loading the font properly, and is therefore handy when you want to create a dictionary of available fonts based on a directory contents.

4.12.2 Loading an OPENTYPE or TRUETYPE file

If you want to use an OPENTYPE font, you have to get the metric information from somewhere. Using the `fontforge` library, the basic way to get that information is thus:

```
local font = fontforge.open("filename")
local metrics = fontforge.to_table(font)
fontforge.close(font)
```

the `filename` font file is parsed and even partially interpreted by the OPENTYPE/TRUETYPE loading routines from FONTFORGE. There are a few advantages to this approach compared to reading the actual font file ourselves:

- The font is automatically re-encoded, so that the `ttf_metrics` table is using UNICODE for the character indices.
- Many features are pre-processed into a format that is easier to handle than just the bare tables would be.



- POSTSCRIPT-based OPENTYPE fonts do not store the character height and depth in the font file, so the character boundingbox has to be calculated in some way.
- In the future, it may be interesting to allow LUA scripts access to the font program itself, perhaps even creating or changing the font.

The top-level keys in the returned table are (this documentation is not yet finished):

key	type	explanation
table_version	number	indicates the metrics version
fontname	string	
fullname	string	
familynname	string	
weight	string	
copyright	string	
filename	string	
defbasefilename	string	
version	string	
italicangle	float	
upos	float	
uwidth	float	
units_per_em	number	
ascent	number	
descent	number	
vertical_origin	number	
uniqueid	number	
glyphcnt	number	
glyphmax	number	
glyphs	array	
changed	number	
hasvmetrics	number	
order2	number	
strokedfont	number	
weight_width_slope_only	number	
head_optimized_for_cleartype	number	
uni_interp	enum	unset , none , adobe , greek , japanese , trad_chinese , simp_chinese , korean , ams
map	table	
origname	string	
private	table	
xuid	string	
pfminfo	table	
names	table	
cidinfo	table	
subfonts	array	
cidmaster	array	



comments	string
anchor_classes	table
ttf_tables	table
kerns	table
vkerns	table
texdata	table
lookups	table
gpos	table
gsub	table
chosename	string
macstyle	number
fondname	string
design_size	number
fontstyle_id	number
fontstyle_name	table
design_range_bottom	number
design_range_top	number
strokewidth	float
mark_classes	array
mark_class_names	array
creationtime	number
modificationtime	number
os2_version	number

4.12.2.1 Glyph items

The `glyphs` is an array containing the per-character information (quite a few of these are only present if nonzero).

key	type	explanation
name	string	
unicodeenc	number	
boundingbox	array	array of four numbers
width	number	(only for horizontal fonts)
vwidth	number	(only for vertical fonts)
lsidebearing	number	(only if nonzero)
glyph_class	number	(only if nonzero)
kerns	array	(only for horizontal fonts, if set)
vkerns	array	(only for vertical fonts, if set)
dependents	array	linear array of glyph name strings (only if nonempty)
lookups	table	(only if nonempty)
ligofme	table	(only if nonempty)
comment	string	(only if set)
anchor	table	(only if set)



<code>tex_height</code>	number	(only if set)
<code>tex_depth</code>	number	(only if set)
<code>tex_sub_pos</code>	number	(only if set)
<code>tex_super_pos</code>	number	(only if set)

The `ker`ns and `vk`erns are linear arrays of small hashes:

key	type	explanation
<code>char</code>	string	
<code>off</code>	number	
<code>lookup</code>	string	

The `lookups` is a hash based on lookup subtable names, with the value of each key inside that a linear array of small hashes:

key	type	explanation
<code>type</code>	enum	<code>position</code> , <code>pair</code> , <code>substitution</code> , <code>alternate</code> , <code>multiple</code> , <code>ligature</code> , <code>lcaret</code> , <code>ker</code> ning, <code>vk</code> erning, <code>anchors</code> , <code>contextpos</code> , <code>contextsub</code> , <code>chainpos</code> , <code>chainsub</code> , <code>reversesub</code> , <code>max</code> , <code>kernback</code> , <code>vkernback</code>
<code>specification</code>	table	extra data

For the first seven values of `type`, there can be additional sub-information, stored in the sub-table `specification`:

value	type	explanation
<code>position</code>	table	<code>vr</code> table
<code>pair</code>	table	one string: <code>paired</code> , and a <code>vr</code> (sub)table
<code>substitution</code>	table	one string: <code>variant</code>
<code>alternate</code>	table	one string: <code>components</code>
<code>multiple</code>	table	one string: <code>components</code>
<code>ligature</code>	table	two strings: <code>components</code> , <code>char</code>
<code>lcaret</code>	array	linear array of numbers

The `vr` table contains for number-valued fields: `xoff`, `yoff`, `h_adv_off` and `v_adv_off`.

The other values of `type` could probably use some extra information as well, but I do not know which case of the union is supposed to be selected.

The `ligofme` is a linear array of small hashes:

key	type	explanation
<code>lig</code>	table	uses the same substructure as a single <code>possub</code> item
<code>char</code>	string	
<code>components</code>	array	linear array of named components
<code>ccnt</code>	number	

The `anchor` table is indexed by a string signifying the anchor type, which is one of



key	type	explanation
mark	string	placement mark
basechar	string	mark for attaching combining items to a base char
baselig	string	mark for attaching combining items to a ligature
basemark	string	generic mark for attaching combining items to connect to
centry	string	cursive entry point
cexit	string	cursive exit point

Each of those, if present is another hash, with the entries anchors. The anchor name (class) is the key, and the values are the actual data:

key	type	explanation
x	number	x location
y	number	y location
ttf_pt_index	number	truetype point index, only if given
lig_index	number	only if non-zero

For clarification, an anchor table could for example look like this :

```
["anchor"]={
  ["basemark"]={
    ["Anchor-7"]={ ["x"]=170, ["y"]=1080 }
  },
  ["mark"]={
    ["Anchor-1"] = { ["x"]=160, ["y"]=810 },
    ["Anchor-4"] = { ["x"]=160, ["y"]=800 }
  }
},
```

4.12.2.2 map table

The top-level map is a list of encoding mappings. Each of those is a table itself.

key	type	explanation
enccount	number	
encmax	number	
backmax	number	
remap	table	
map	array	non-linear array of mappings
backmap	array	non-linear array of backward mappings
enc	table	

The `remap` table is very small:

key	type	explanation
firstenc	number	



lastenc	number
infont	number

The `enc` table is a bit more verbose:

key	type	explanation
enc_name	string	
char_cnt	number	
char_max	number	
unicode	array	of UNICODE position numbers
psnames	array	of POSTSCRIPT glyph names
builtin	number	
hidden	number	
only_1byte	number	
has_1byte	number	
has_2byte	number	
is_unicodebmp	number	(only if nonzero)
is_unicodedefull	number	(only if nonzero)
is_custom	number	(only if nonzero)
is_original	number	(only if nonzero)
is_compact	number	(only if nonzero)
is_japanese	number	(only if nonzero)
is_korean	number	(only if nonzero)
is_tradchinese	number	(only if nonzero)
is_simplechinese	number	
low_page	number	
high_page	number	
iconv_name	string	
iso_2022_escape	string	

4.12.2.3 private table

This is the font's private POSTSCRIPT dictionary, if any. Keys and values are both strings.

4.12.2.4 cidinfo table

registry	string
ordering	string
supplement	number
version	number

4.12.2.5 pfminfo table

The `pfminfo` table contains most of the OS/2 information:



key	type	explanation
pfmset	number	
winascent_add	number	
windescent_add	number	
hheadascent_add	number	
hheaddescent_add	number	
typoascent_add	number	
typodescent_add	number	
subsuper_set	number	
panose_set	number	
hheadset	number	
vheadset	number	
pfmfamily	number	
weight	number	
width	number	
avgwidth	number	
firstchar	number	
lastchar	number	
fstype	number	
linegap	number	
vlinegap	number	
hhead_ascent	number	
hhead_descent	number	
hhead_descent	number	
os2_typoascent	number	
os2_typodescent	number	
os2_typolinegap	number	
os2_winascent	number	
os2_windescent	number	
os2_subxsize	number	
os2_subysize	number	
os2_subxoff	number	
os2_subyoff	number	
os2_supxsize	number	
os2_supysize	number	
os2_supxoff	number	
os2_supyoff	number	
os2_strikeysize	number	
os2_strikeypos	number	
os2_family_class	number	
os2_xheight	number	
os2_capheight	number	
os2_defaultchar	number	
os2_breakchar	number	



os2_vendor	string
panose	table

The [panose](#) subtable has exactly 10 string keys:

	key	type	
familytype	string		Values as in the OPENType font specification: Any , No Fit , Text and Display , Script
serifstyle	string		See the OPENType font specification for values
weight	string	id.	
proportion	string	id.	
contrast	string	id.	
strokevariation	string	id.	
armstyle	string	id.	
letterform	string	id.	
midline	string	id.	
xheight	string	id.	

4.12.2.6 names table

Each item has two top-level keys:

	key	type	explanation
lang	string		language for this entry
names table			

The [names](#) keys are the actual TRUEType name strings. The possible keys are:

key	type	explanation
copyright		
family		
subfamily		
uniqueid		
fullname		
version		
postscriptname		
trademark		
manufacturer		
designer		
descriptor		
venderurl		
designerurl		
license		
licenseurl		
idontknow		
preffamilyname		



```

prefmodifiers
compatfull
sampletext
cidfindfontname

```

4.12.2.7 anchor_classes table

The anchor_classes classes:

key	type	explanation
name	string	A descriptive id of this anchor class
lookup	string	
type	string	One of "mark", "mkmk", "curs", "mklg"

4.12.2.8 gpos table

Th gpos table has one array entry for each lookup.

key	type	explanation
type	string	One of "gpos_start", "gpos_single", "gpos_pair", "gpos_cursive", "gpos_mark2base", "gpos_mark2ligature", "gpos_mark2mark", "gpos_context", "gpos_contextchain"
flags	table	
name	string	
features	array	
subtables	array	

The flags table has a true value for each of the lookup flags that is actually set:

key	type	explanation
r2l	boolean	
ignorebaseglyphs	boolean	
ignoreligatures	boolean	
ignorecombiningmarks	boolean	

The features table has:

key	type	explanation
tag	string	
scripts	table	
ismax	number	(only if true)

The scripts table within features has:

key	type	explanation
script	string	
langs	array of strings	



The subtables table has:

key	type	explanation
name	string	
suffix	string	(only if used)
anchor_classes	number	(only if used)
vertical_kerning	number	(only if used)
kernclass	table	(only if used)

The kernclass with subtables table has:

key	type	explanation
first_cnt	number	
second_cnt	number	
firsts	array of strings	
seconds	array of strings	
lookup	string	associated lookup
offsets	array of numbers	
kcid	number	?

4.12.2.9 gsub table

This has identical layout to the [gpos](#) table, except for the type:

key	type	explanation
type	string	One of "gsub_start", "gsub_single", "gsub_multiple", "gsub_alternate", "gsub_ligature", "gsub_context", "gsub_contextchain", "gsub_reversecontextchain"

4.12.2.10 ttf_tables table

key	type	explanation
tag	string	
len	number	
maxlen	number	
data	number	

4.12.2.11 kerns table

Substructure is identical to the per-glyph subtable.

4.12.2.12 vkerns table

Substructure is identical to the per-glyph subtable.



4.12.2.13 texdata table

key	type	explanation
type	string	unset , text , math , mathext
params	array	22 font numeric parameters

4.12.2.14 lookups table

Top-level [lookups](#) is quite different from the ones at character level. The keys in this hash are strings, the values the actual lookups, represented as dictionary tables.

key	type	explanation
type	number	
format	enum	glyphs , class , coverage , reversecoverage
tag	string	
current_class	array	
before_class	array	
after_class	array	
rules	array	an array of rule items

Rule items have one common item and one specialized item:

key	type	explanation
lookups	array	A list of lookup items
glyph	array	Only if the parent's format is glyph
class	array	Only if the parent's format is glyph
coverage	array	Only if the parent's format is glyph
reversecoverage	array	Only if the parent's format is glyph

Each of the lookup item is:

key	type	explanation
seq	number	
lookup	string	

glyph:

key	type	explanation
names	string	
back	string	
fore	string	

class:

key	type	explanation
current	array	of numbers



before	array	of numbers
after	array	of numbers

coverage:

key	type	explanation
current	array	of strings
before	array	of strings
after	array	of strings

reversecoverage:

key	type	explanation
current	array	of strings
before	array	of strings
after	array	of strings
replacements	string	



5 Font structure

All T_EX fonts are represented to LUA code as tables, and internally as C structures. All keys in the table below are saved in the internal font structure if they are present in the table returned by the `define_font` callback, or if they result from the normal TFM/VF reading routines if there is no `define_font` callback defined.

The column from vf means that this key will be created by the `font.read_vf()` routine, from TFM means that the key will be created by the `font.read_tfm()` routine, and used means whether or not the L^AT_EX engine itself will do something with the key.

The top-level keys in the table are as follows:

key	from vf	from tfm	used	value type	description
name	yes	yes	yes	string	metric (file) name
area	no	yes	yes	string	(directory)location, typically empty
used	no	yes	yes	boolean	used already? (initial: false)
characters	yes	yes	yes	table	the defined glyphs of this font
checksum	yes	yes	no	number	default: 0
designsize	no	yes	yes	number	expected size (default: 655360 == 10pt)
direction	no	yes	yes	number	default: 0 (LTR)
encodingbytes	no	no	yes	number	default: depends on <code>format</code>
encodingname	no	no	yes	string	encoding name
fonts	yes	no	yes	table	locally used fonts
fullname	no	no	yes	string	actual (POSTSCRIPT) name
header	yes	no	no	string	header comments, if any
hyphenchar	no	no	yes	number	default: TeX's <code>\hyphenchar</code>
parameters	no	yes	yes	hash	default: 7 parameters, all zero
size	no	yes	yes	number	loaded (at) size. (default: same as designsize)
skewchar	no	no	yes	number	default: TeX's <code>\skewchar</code>
type	yes	no	yes	string	basic type of this font
format	no	no	yes	string	disk format type
embedding	no	no	yes	string	PDF inclusion
filename	no	no	yes	string	disk file name

The key `name` is always required.

The key `used` is set by the engine when a font is actively in use, this makes sure that the font's definition is written to the output file (DVI or PDF). The TFM reader sets it to false.

The `direction` is a number signalling the normal direction for this font. There are sixteen possibilities:

number	meaning	number	meaning
0	LT	8	TT
1	LL	9	TL
2	LB	10	TB



3	LR	11	TR
4	RT	12	BT
5	RL	13	BL
6	RB	14	BB
7	RR	15	BR

These are OMEGA-style direction abbreviations: the first character indicates the first edge of the character glyphs (the edge that is seen first in the writing direction), the second the top side.

The `parameters` is a hash with mixed key types. There are seven possible string keys, as well as a number of integer indices (these start from 8 up). The seven strings are actually used instead of the bottom seven indices, because that gives a nicer user interface.

The names and their internal remapping:

name	internal remapped number
slant	1
space	2
space_stretch	3
space_shrink	4
x_height	5
quad	6
extra_space	7

The keys `type`, `format`, `embedding`, `fullname` and `filename` are used to embed OPENTYPE fonts in the result PDF.

The `characters` table is a list of character hashes indexed by integer number. The number is the internal code. T_EX knows this character by.

Two very special string indexes can be used also: `left_boundary` is a virtual character whose ligatures and kerns are used to handle word boundary processing. `right_boundary` is similar but not actually used for anything (yet!).

Other index keys are ignored.

Each character hash itself is a hash. For example, here is the character f (decimal 102) in the font cmr10 at 10 points:

```
[102] = {
  ["kerns"] = {
    [63] = 50973,
    [93] = 50973,
    [39] = 50973,
    [33] = 50973,
    [41] = 50973
  },
  ["italic"] = 50973,
  ["height"] = 455111,
```



```

["depth"] = 0,
["ligatures"] = {
  [102] = {
    ["char"] = 11,
    ["type"] = 0
  },
  [108] = {
    ["char"] = 13,
    ["type"] = 0
  },
  [105] = {
    ["char"] = 12,
    ["type"] = 0
  }
},
["width"] = 200250
}

```

The following top-level keys can be present inside a character hash:

key	from vf	from tfm	used	value type	description
width	yes	yes	yes	number	character's width, in sp (default 0)
height	no	yes	yes	number	character's height, in sp (default 0)
depth	no	yes	yes	number	character's depth, in sp (default 0)
italic	no	yes	yes	number	character's italic correction, in sp (default zero)
next	no	yes	yes	number	the next larger character index
extensible	no	yes	yes	table	the constituent bits of an extensible recipe
kerns	no	yes	yes	table	kertering information
ligatures	no	yes	yes	table	ligaturing information
commands	yes	no	yes	array	virtual font commands
name	no	no	no	string	the character (POSTSCRIPT) name
index	no	no	yes	number	the (OPENTYPE or TRUETYPE) font glyph index
used	no	yes	yes	boolean	typeset already (default: false)?

The presence of **extensible** will overrule **next**, if that is also present.

The **extensible** table is very simple:

key	type	description
top	number	top character index
mid	number	middle character index
bot	number	bottom character index
rep	number	repeatable character index

The **kerns** table is a hash indexed by character index (and character index is defined as either a non-negative integer or the string value **right_boundary**), with the values the kerning to be applied, in scaled points.



The `ligatures` table is a hash indexed by character index (and character index is defined as either a non-negative integer or the string value `right_boundary`), with the values being yet another small hash, with two fields:

key	type	description
type	number	the type of this ligature command, default 0
char	number	the character index of the resultant ligature

The `char` field in a ligature is required.

The `type` field inside a ligature is the numerical or string value of one of the eight possible ligature types supported by \TeX . When \TeX inserts a new ligature, it puts the new glyph in the middle of the left and right glyphs. The original left and right glyphs can optionally be retained, and when at least one of them is kept, it is also possible to move the new insertion point forward one or two places. The glyph that ends up to the right of the insertion point will become the next left.

textual (Knuth)	number	string	result
$l + r =: n$	0	<code>=:</code>	$ n$
$l + r =: n$	1	<code>=: </code>	$ nr$
$l + r =: n$	2	<code> =:</code>	$ ln$
$l + r =: n$	3	<code> =: </code>	$ lnr$
$l + r =: > n$	5	<code>=: ></code>	$n r$
$l + r =:> n$	6	<code> =: ></code>	$l n$
$l + r =: > n$	7	<code> =: ></code>	$l nr$
$l + r =: >> n$	11	<code> =: >></code>	$ln r$

The default value is 0, and can be left out. That signifies a normal ligature where the ligature replaces both original glyphs. In this table the `|` indicates the final insertion point.

The `commands` array is explained below.

5.1 Real fonts

Whether or not a \TeX font is a real font that should be written to the \PDF document is decided by the `type` value in the top-level font structure. If the value is `real`, then this is a proper font, and the inclusion mechanism will attempt to add the needed font object definitions to the \PDF .

Values for `type`:

value	description
real	this is a base font
virtual	this is a virtual font

The actions to be taken depend on a number of different variables:

- Whether the used font fits in an 8-bit encoding scheme or not
- The type of the disk font file
- The level of embedding requested



A font that uses anything other than an 8-bit encoding vector has to be written to the PDF in a different way.

The rule is: if the font table has `encodingbytes` set to 2, then this is a wide font, in all other cases it isn't. The value 2 is the default for `OPENTYPE` and `TRUETYPE` fonts loaded via `LUA`.

If no special care is needed, `LUATEX` currently falls back to the mapfile-based solution used by `PDFTEX` and `DVIPS`. This behaviour will be removed in the future, when the existing code becomes integrated in the new subsystem.

But if this is a wide font, then the new subsystem kicks in, and some extra fields have to be present in the font structure. In this case, `LUATEX` does not use a map file at all.

The extra fields are: `format`, `embedding`, `fullname`, `cidinfo` (as explained above), `filename`, and the `index` key in the separate characters.

Values for `format`:

value	description
type1	this is a <code>POSTSCRIPT TYPE1</code> font
type3	this is a bitmapped (<code>PK</code>) font
truetype	this is a <code>TRUETYPE</code> or <code>TRUETYPE</code> -based <code>OPENTYPE</code> font
opentype	this is a <code>POSTSCRIPT</code> -based <code>OPENTYPE</code> font

Currently, only `TRUETYPE` and `OPENTYPE` fonts can be wide fonts (`TYPE1` `POSTSCRIPT` fonts are not supported).

Values for `embedding`:

value	description
no	don't embed the font at all
subset	include and attempt to subset the font
full	include this font in its entirety

At the moment, `subset` only works for `POSTSCRIPT`-based non-CID `OPENTYPE` fonts, every other font format essentially is treated as `full`.

It is not possible to artificially modify the transformation matrix for the font at the moment.

The other fields are used as follows: The `fullname` will be the `POSTSCRIPT/PDF` font name. The `cidinfo` will be used as the character set (the CID `/Ordering` and `/Registry` keys). The `filename` points to the actual font file. If you include the full path in the `filename` or if the file is in the local directory, `LUATEX` will run a little bit more efficient because it will not have to re-run the `find_xxx_file` callback in that case.

Be careful: when mixing old and new fonts in one document, it is possible to create name `POSTSCRIPT` name clashes that can result in printing errors. When this happens, you have to change the `fullname` of the font.

Typeset strings are written out in a wide format using 2 bytes per glyph, using the `index` key in the character information as value. The overall effect is like having an encoding based on numbers instead of traditional (`POSTSCRIPT`) name-based reencoding.



This type of reencoding means that there is no longer a clear connection between the text in your input file and the strings in the output PDF file. Dealing with this is on the agenda.

5.2 Virtual fonts

You have to take the following steps if you want L^AT_EX to treat the returned table from `define_font` as a virtual font:

- Set the top-level key `type` to `virtual`.
- Make sure there is at least one valid entry in `fonts` (see below)
- Give a `commands` array to every character (see below)

The presence of the toplevel `type` key with the specific value `virtual` will trigger handling of the rest of the special virtual font fields in the table, but the mere existence of 'type' is enough to prevent L^AT_EX from looking for a virtual font on its own.

Therefore, this also works in reverse: if you are absolutely certain that a font is not a virtual font, assigning the value `base` or `real` to `type` will inhibit L^AT_EX from looking for a virtual font file, thereby saving you a disk search.

The `fonts` is another LUA array. The values are one- or two-key hashes themselves, each entry indicating one of the base fonts in a virtual font. An example makes this easy to understand

```
fonts = { { name = "ptmr8a", size = 655360},
          { name = "psyr", size = 600000},
          { id = 38 } }
```

says that the first referenced font (index 1) in this virtual font is `ptrmr8a` loaded at 10pt, and the second is `psyr` loaded at a little over 9pt. The third one is previously defined font that is known to L^AT_EX as fontid 38.

The array index numbers are used by the character command definitions that are part of each character.

The `commands` array is a hash here each item is another small array, with first entry representing a command and the extra items the parameters to that command. The allowed commands and their arguments are:

command name	arguments	arg type	description
font	1	number	select a new font from the local <code>fonts</code> table
char	1	number	typeset this character number from the current
node	1	node	output this node (list), and move right
slot	2	number	a shortcut for a font, char set
push	0		save current position
nop	0		do nothing
pop	0		pop position
rule	2	2 numbers	output a rule $w * h$, and move right
down	1	number	move down on the page



right	1	number	move right on the page
special	1	string	output a <code>\special</code> command
comment	any	any	the rest of the command is ignored

Here is a rather elaborate example:

```
...
commands = {
  {"push"},           -- remember where we are
  {"right", 5000},    -- move right about 0.08pt
  {"font", 1},        -- select the fonts[1] entry
  {"char", 97},       -- place character 97 (a)
  {"pop"},            -- go all the way back
  {"down", -200000},  -- move *up* about 3pt
  {"special", "pdf: 1 0 0 rg"} -- switch to red color
  {"rule", 500000, 20000} -- draw a bar
  {'special',"pdf: 0 g"} -- back to black
}
...
```

The default value for `font` is always 1, for each character anew. If the virtual font is essentially only a re-encoding, then you do usually do not have create an explicit font entry.

Regardless of the amount of movement you create within the `commands`, the output pointer will always move by exactly the width as given in the `width` key of the character hash, after running the `commands`.

Even in a real font, there can be virtual characters. When L^AT_EX encounters a `commands` field inside a character when it becomes time to typeset the character, it will interpret the commands, just like for a true virtual character. In this case, if you have created no fonts array, then the default and only base font is taken to be the current font itself. In practise, this means that you can create virtual duplicates of existing characters.

Note: this feature does *not* work the other way around. There can not be real characters in a virtual font!

Finally, here is a plain T_EX input file with a demonstration:

```
% start of virtual-demo.tex

\pdfoutput=1
\directlua0 {
  callback.register("define_font",
    function (name,area,size)
      if name == 'cmr10-red' then
        f = font.read_tfm('cmr10',size)
        f.name = 'cmr10-red'
        f.type = 'virtual'
        f.fonts = {'cmr10', size}}
}
```



```

    for i,v in pairs(f.characters) do
        if (string.char(i)):find("[tacohanshartmut]") then
            v.commands = {
                {'special','pdf: 1 0 0 rg'},
                {'char',i},
                {'special','pdf: 0 g'},
            }
        else
            v.commands = {'char',i}
        end
    end
end
else
    f = font.read_tfm(name,size)
end
return f
end
)
}

\font\myfont = cmr10-red \myfont This is a line of text \par
\font\myfontx= cmr10 \myfontx Here is another line of text \par

\bye
% end of virtual-demo.tex

```



6 Nodes

6.1 LUA node representation

T_EX's nodes are represented in LUA as userdata object with a variable set of fields.

The current return value of `node.types()` is: `vlist` (1), `rule` (2), `ins` (3), `mark` (4), `adjust` (5), `disc` (7), `whatsit` (8), `math` (9), `glue` (10), `kern` (11), `penalty` (12), `unset` (13), `style` (14), `choice` (15), `ord` (16), `op` (17), `bin` (18), `rel` (19), `open` (20), `close` (21), `punct` (22), `inner` (23), `radical` (24), `fraction` (25), `under` (26), `over` (27), `accent` (28), `vcenter` (29), `left` (30), `right` (31), `action` (39), `margin_kern` (40), `glyph` (41), `attribute` (42), `glue_spec` (43), `attribute_list` (44), `hlist` (0), but as already mentioned, the math and alignment nodes in this list are not supported at the moment. The useful list is described in the next sections.

6.1.1 Auxiliary items

A few node-typed userdata objects do not occur in the normal list of nodes, but can be pointed to from within that list. They are not quite really nodes, but it is easier for the library routines to treat them as if they were.

6.1.1.1 glue_spec items

Skips are about the only type of data objects in traditional T_EX that are not a simple value. The structure that represents the glue components of a skip is called a `glue_spec`, and it has the following accessible fields:

key	type	explanation
<code>width</code>	number	
<code>stretch</code>	number	
<code>stretch_order</code>	number	
<code>shrink</code>	number	
<code>shrink_order</code>	number	

These objects are reference counted, so there is actually an extra field named `ref_count` as well, but that will become completely hidden in the near future.

6.1.1.2 attribute_list items

The newly introduced attribute registers are also non-trivial, because the value that is attached to a node is essentially a sparse array of key-value pairs.

It is generally easiest to deal with attributes by using the dedicated functions in the `node` library, but for completeness, here is the low-level interface:



field	type	explanation
next	node	pointer to the first attribute

There are no extra fields, this kind of item is only used as a head pointer for attribute items, making them easier to handle.

A normal node's `attribute` field will point to a field of this type, and the `next` field in that field will then point to the first defined attribute item, whose `next` will point to the second attribute item, etc.

6.1.1.3 attribute item

Valid fields:

field	type	explanation
next	node	pointer to the next attribute
number	number	the LHS
value	number	the RHS

6.1.2 Main text nodes

These are the nodes that comprise actual typesetting commands.

A few fields are present in all nodes regardless of their type, these are:

field	type	explanation
next	node	The next node in a list, or nil
id	number	The node's type (<code>id</code>) number
subtype	number	The node <code>subtype</code> identifier

The subtype is sometimes just a stub entry. Almost all nodes also have an `attr` field. In the following tables `next` and `id` are not explicitly mentioned.

6.1.2.1 hlist nodes

Valid fields: `attr`, `width`, `depth`, `height`, `dir`, `shift`, `glue_order`, `glue_sign`, `glue_set`, `list`

field	type	explanation
subtype	number	unused
attr	node	The head of the associated attribute list
width	number	
height	number	
depth	number	
shift	number	a displacement perpendicular to the character progression direction
glue_order	number	a number in the range 0–4, indicating the glue order
glue_set	number	the calculated glue ratio



<code>glue_sign</code>	number	
<code>list</code>	node	the body of this list
<code>dir</code>	number	the direction of this box

6.1.2.2 vlist nodes

Valid fields: As for hlist, except that shift is a displacement perpendicular to the line progression direction.

6.1.2.3 rule nodes

Valid fields: `attr`, `width`, `depth`, `height`, `dir`

field	type	explanation
<code>subtype</code>	number	unused
<code>attr</code>	node	
<code>width</code>	number	Rule size. The special value <code>-1073741824</code> is used for running glue dimensions
<code>height</code>	number	<code>' '</code>
<code>depth</code>	number	<code>' '</code>
<code>dir</code>	number	the direction of this rule

6.1.2.4 ins nodes

Valid fields: `attr`, `cost`, `depth`, `height`, `top_skip`, `list`

field	type	explanation
<code>subtype</code>	number	the insertion class
<code>attr</code>	node	
<code>cost</code>	number	The penalty associated with this insert
<code>height</code>	number	
<code>depth</code>	number	
<code>list</code>	node	the body of this insert
<code>top_skip</code>	node	a pointer to the <code>\splittopskip</code> glue spec

6.1.2.5 mark nodes

Valid fields: `attr`, `class`, `mark`

field	type	explanation
<code>subtype</code>	number	unused
<code>attr</code>	node	
<code>class</code>	number	the mark class
<code>mark</code>	table	a table representing a token list



6.1.2.6 adjust nodes

Valid fields: `attr`, `list`

field	type	explanation
subtype	number	1 = pre
attr	node	
list	node	adjusted material

6.1.2.7 disc nodes

Valid fields: `attr`, `pre`, `post`, `replace`

field	type	explanation
subtype	number	unused
attr	node	
pre	node	pointer to the pre-break text
post	node	pointer to the post-break text
replace	number	the number of nodes to skip if this discretionary is chosen as a breakpoint

6.1.2.8 math nodes

Valid fields: `attr`, `surround`

field	type	explanation
subtype	number	0 = on, 1 = off
attr	node	
surround	number	width of the <code>\mathsurround</code> kern

6.1.2.9 glue nodes

Valid fields: `attr`, `spec`, `leader`

field	type	explanation
subtype	number	0 = <code>\skip</code> , 1–18 = internal glue parameters, 100 = <code>\leaders</code> , 101 = <code>\cleaders</code> , 102 = <code>\xleaders</code>
attr	node	
spec	node	pointer to a <code>glue_spec</code> item
leader	node	pointer to a box or rule for leaders

6.1.2.10 kern nodes

Valid fields: `attr`, `kern`



field	type	explanation
subtype	number	0 = from font, 1 = from \kern or \/, 2 = from \accent
attr	node	
width	number	

6.1.2.11 penalty nodes

Valid fields: `attr`, `penalty`

field	type	explanation
subtype	number	not used
attr	node	
penalty	number	

6.1.2.12 glyph nodes

Valid fields: `attr`, `char`, `font`, `components`, `xoffset`, `yoffset`

field	type	explanation
subtype	number	0 = normal, 1 = ligature, 2 = leftboundary ligature, 3 = rightboundary ligature
attr	node	
char	number	
font	number	
components	node	pointer to ligature components
xoffset	number	
yoffset	number	

6.1.2.13 margin_kern nodes

Valid fields: `attr`, `width`, `glyph`

field	type	explanation
subtype	number	0 = left side, 1 = right side
attr	node	
width	number	
glyph	node	

6.1.3 whatsit nodes

Whatsit nodes come in many subtypes, that you can ask for my running `node.whatsits()`: `write` (1), `close` (2), `special` (3), `language` (4), `local_par` (6), `dir` (7), `pdf_literal` (8), `pdf_refobj` (10), `pdf_refxform` (12), `pdf_refximage` (14), `pdf_annot` (15), `pdf_start_link` (16),



`pdf_end_link` (17), `pdf_dest` (19), `pdf_thread` (20), `pdf_start_thread` (21), `pdf_end_thread` (22), `pdf_save_pos` (23), `late_lua` (35), `close_lua` (36), `user_defined` (43), `pdf_restore` (42), `pdf_colorstack` (39), `pdf_setmatrix` (40), `pdf_save` (41), `open` (0),

6.1.3.1 open nodes

Valid fields: `attr`, `stream`, `name`, `area`, `ext`

field	type	explanation
<code>attr</code>	node	
<code>stream</code>	number	T _E X's stream id number
<code>name</code>	string	file name
<code>ext</code>	string	file extension
<code>area</code>	string	file area

6.1.3.2 write nodes

Valid fields: `attr`, `stream`, `data`

field	type	explanation
<code>attr</code>	node	
<code>stream</code>	number	T _E X's stream id number
<code>data</code>	table	a table representing the token list to be written

6.1.3.3 close nodes

Valid fields: `attr`, `stream`

field	type	explanation
<code>attr</code>	node	
<code>stream</code>	number	T _E X's stream id number

6.1.3.4 special nodes

Valid fields: `attr`, `data`

field	type	explanation
<code>attr</code>	node	
<code>data</code>	string	The <code>\special</code> information

6.1.3.5 language nodes

Valid fields: `attr`, `lang`, `left`, `right`



field	type	explanation
attr	node	
lang	number	language id number
left	number	value of \lefthyphenmin
right	number	value of \righthyphenmin

6.1.3.6 local_par nodes

Valid fields: `attr`, `pen_inter`, `pen_broken`, `dir`, `box_left`, `box_left_width`, `box_right`, `box_right_width`

field	type	explanation
attr	node	
pen_inter	number	interline penalty
pen_broken	number	broken penalty
dir	number	the direction of this par
box_left	node	the \localleftbox
box_left_width	number	width of the \localleftbox
box_right	node	the \localrightbox
box_right_width	number	width of the \localrightbox

6.1.3.7 dir nodes

Valid fields: `attr`, `dir`, `level`, `dvi_ptr`, `dvi_h`

field	type	explanation
attr	node	
dir	number	the direction
level	number	nesting level of this direction whatsit
dvi_ptr	number	a saved dvi buffer byte offset
dir_h	number	a saved dvi position

6.1.3.8 pdf_literal nodes

Valid fields: `attr`, `mode`, `data`

field	type	explanation
attr	node	
mode	number	The mode setting of this literal
data	string	The \pdfliteral information



6.1.3.9 pdf_refobj nodes

Valid fields: [attr](#), [objnum](#)

field	type	explanation
attr	node	
objnum	number	the referenced PDF object number

6.1.3.10 pdf_refxform nodes

Valid fields: [attr](#), [width](#), [height](#), [depth](#), [objnum](#)

field	type	explanation
attr	node	
width	number	
height	number	
depth	number	
objnum	number	the referenced PDF object number

6.1.3.11 pdf_refximage nodes

Valid fields: [attr](#), [width](#), [height](#), [depth](#), [objnum](#)

field	type	explanation
attr	node	
width	number	
height	number	
depth	number	
objnum	number	the referenced PDF object number

6.1.3.12 pdf_annot nodes

Valid fields: [attr](#), [width](#), [height](#), [depth](#), [objnum](#), [data](#)

field	type	explanation
attr	node	
width	number	
height	number	
depth	number	
objnum	number	the referenced PDF object number
data	string	the annotation data



6.1.3.13 pdf_start_link nodes

Valid fields: `attr`, `width`, `height`, `depth`, `objnum`, `link_attr`, `action`

field	type	explanation
<code>attr</code>	node	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>objnum</code>	number	the referenced PDF object number
<code>link_attr</code>	table	the link attribute token list
<code>action</code>	node	the action to perform

6.1.3.14 action items

Valid fields: `action_type`, `named_id`, `action_id`, `file`, `new_window`, `data`, `ref_count`

These are a special kind of item that only appears inside pdf start link objects.

6.1.3.15 pdf_end_link nodes

Valid fields: `attr`

field	type	explanation
<code>attr</code>	node	

6.1.3.16 pdf_dest nodes

Valid fields: `attr`, `width`, `height`, `depth`, `named_id`, `dest_id`, `dest_type`, `xyz_zoom`, `objnum`

field	type	explanation
<code>attr</code>	node	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>named_id</code>	number	is the <code>dest_id</code> a string value?
<code>dest_id</code>	number or string	the destination id
<code>dest_type</code>	number	type of destination
<code>xyz_zoom</code>	number	
<code>objnum</code>	number	the PDF object number

6.1.3.17 pdf_thread nodes

Valid fields: `attr`, `width`, `height`, `depth`, `named_id`, `thread_id`, `thread_attr`



field	type	explanation
attr	node	
width	number	
height	number	
depth	number	
named_id	number	is the tread_id a string value?
tread_id	number or string	the thread id
thread_attr	number	extra thread information

6.1.3.18 pdf_start_thread nodes

Valid fields: `attr`, `width`, `height`, `depth`, `named_id`, `thread_id`, `thread_attr`

field	type	explanation
attr	node	
width	number	
height	number	
depth	number	
named_id	number	is the tread_id a string value?
tread_id	number or string	the thread id
thread_attr	number	extra thread information

6.1.3.19 pdf_end_thread nodes

Valid fields: `attr`

field	type	explanation
attr	node	

6.1.3.20 pdf_save_pos nodes

Valid fields: `attr`

field	type	explanation
attr	node	

6.1.3.21 late_lua nodes

Valid fields: `attr`, `reg`, `data`

field	type	explanation
attr	node	



reg	number	LUA state id number
data	string	data to execute

6.1.3.22 close_lua nodes

Valid fields: `attr`, `reg`

field	type	explanation
attr	node	
reg	number	LUA state id number

6.1.3.23 pdf_colorstack nodes

Valid fields: `attr`, `stack`, `cmd`, `data`

field	type	explanation
attr	node	
stack	number	colorstack id number
cmd	number	command to execute
data	string	data

6.1.3.24 pdf_setmatrix nodes

Valid fields: `attr`, `data`

field	type	explanation
attr	node	
data	string	data

6.1.3.25 pdf_save nodes

Valid fields: `attr`

field	type	explanation
attr	node	

6.1.3.26 pdf_restore nodes

Valid fields: `attr`

field	type	explanation
attr	node	



6.1.3.27 user_defined nodes

Valid fields: `attr`, `user_id`, `type`, `value`

field	type	explanation
<code>attr</code>	node	
<code>user_id</code>	number	id number
<code>type</code>	number	type of the value
<code>value</code>	number string node table	



7 Modifications

Besides the expected changes caused by new functionality, there are a number of not-so-expected changes. These are sometimes a side-effect of a new (conflicting) feature, or, more often than not, a change necessary to clean up the internal interfaces.

7.1 Changes from T_EX 3.141592

- There is no pool file, all strings are embedded during compilation.
- `plus 1 filllll` does not generate an error. The extra `l` is simply typeset.
- The `\endlinechar` can be either added (values 0 or more), or not (negative values). If it is added, the character is always decimal 13 a/k/a `^^M` a/k/a carriage return (This change may be temporary).

7.2 Changes from ϵ -T_EX 2.2

- The ϵ -T_EX functionality is always present and enabled (but see below about T_EXX_ET), so the prepended asterisk or `-etex` switch for INIT_EX is not needed.
- T_EXX_ET is not present, so the primitives

```
\TeXxETstate
\beginR
\beginL
\endR
\endL
```

are missing

7.3 Changes from PDFT_EX 1.40

- The (experimental) support for snap nodes has been removed, because it much more natural to build this functionality on top of node processing and attributes. The associated primitives that are now gone are: `\pdfsnaprefpoint`, `\pdfsnapy`, and `\pdfsnapycomp`.
- A number of utility functions is removed:

<code>\pdfelapsedtime</code>	<code>\pdffilesize</code>	<code>\pdfstrcmp</code>
<code>\pdfescapehex</code>	<code>\pdflastmatch</code>	<code>\pdfunescapehex</code>
<code>\pdfescapename</code>	<code>\pdfmatch</code>	
<code>\pdfescapestring</code>	<code>\pdfmdfivesum</code>	
<code>\pdffiledump</code>	<code>\pdfresettimer</code>	
<code>\pdffilemoddate</code>	<code>\pdfshellescape</code>	

- A few other experimental primitives are also provided without the extra `pdf` prefix, so they can also be called like this:



```
\primitive          \ifabsnum
\ifprimitive        \ifabsdim
```

- The definitions for new didot and new cicero are patched.
- The `\pdfprimitive` is bugfixed.
- The `\pdftexversion` is set to 200.

7.4 Changes from ALEPH RC4

- The input translations from ALEPH are not implemented, the related primitives are not available


```
\DefaultInputMode          \noDefaultInputTranslation
\noDefaultInputMode        \noInputTranslation
\noInputMode               \InputTranslation
\InputMode                 \DefaultOutputTranslation
\DefaultOutputMode         \noDefaultOutputTranslation
\noDefaultOutputMode       \noOutputTranslation
\noOutputMode              \OutputTranslation
\OutputMode
\DefaultInputTranslation
```
- A small series of bounds checking fixes to `\ocp` and `\ocplist` has been added to prevent the system from crashing due to array indexes running out of bounds.
- The `\hoffset` bug when `\pagedir` TRT is fixed, removing the need for an explicit fix to `\hoffset`
- A bug causing `\fam` to fail for family numbers above 15 is fixed.
- Some bits of ALEPH assumed 0 and `null` were identical. This resulted for instance in a bug that sometimes caused an eternal loop when trying to `\show` a box.
- A fair amount of minor bugs are fixed as well, most of these related to `\tracingcommands` output.
- The number of possible fonts, ocps and ocplists is smaller than their maximum ALEPH value (around 5000 fonts and 30000 ocps / ocplists).
- The internal function `scan_dir()` has been renamed to `scan_direction()` to prevent a naming clash.
- The `^^` notation can come in five and six item repetitions also, to insert characters that do not fit in the BMP.

7.5 Changes from standard WEB2C

- There is no `mltex`
- There is no `enctex`
- The following commandline switches are silently ignored, even in non-LUA mode:

```
-8bit
-translate-file=TCXNAME
-mltex
```



`-enc`
`-etex`

- `\openout` whatsits are not written to the log file.
- Some of the so-called web2c extensions are hard to set up in non-KPSE mode because `texmf.cnf` is not read: `shell-escape` is off (but that is not a problem because of Lua's `os.execute`), and the paranoia checks on `openin` and `openout` do not happen (however, it is easy for a Lua script to do this itself by overloading `io.open`).





8 Implementation notes

8.1 Primitives overlap

The primitives

<code>\pdfpagewidth</code>	<code>\pagewidth</code>
<code>\pdfpageheight</code>	<code>\pageheight</code>
<code>\fontcharwd</code>	<code>\charwd</code>
<code>\fontcharht</code>	<code>\charht</code>
<code>\fontchardp</code>	<code>\chardp</code>
<code>\fontcharic</code>	<code>\charic</code>

are all aliases of each other.

8.2 Memory allocation

The single internal memory heap that traditional T_EX used for tokens and nodes is split into two wholly separate arrays. Each of those can grow dynamically as needed.

The `texmf.cnf` settings related to main memory are no longer used (these are: `main_memory`, `mem_bot`, `extra_mem_top` and `extra_mem_bot`). Out of main memory errors can still occur, but the limiting factor is now solely the amount of RAM in your system.

Also, the memory (de)allocation routines for nodes are completely rewritten. The relevant code now lives in the C file `luanode.c`, and now normally uses a dozen or so avail lists instead of a doubly-linked model. At this moment, speed is still a little suboptimal because separate helper structures are maintained for debugging checks.

Because of the split in two arrays and the resulting differences in the data structures, some of the pascal web macros have been split. For instance, there are now `vlink` and `vinfo` as well as `link` and `info`. All access to the variable memory array is now hidden behind a macro called `vmem`.

The implementation of the growth of two arrays (via reallocation) introduces a potential pitfall: the memory arrays should never be used as the left hand side of a statement that can modify the array in question.

The input line buffer and pool size are now also reallocated when needed, and the `texmf.cnf` settings `buf_size` and `pool_size` are silently ignored.

8.3 Sparse arrays

The `\mathcode`, `\delcode`, `\catcode`, `\sfcode`, `\lccode` and `\uccode` tables are now sparse arrays that are implemented in C. They are no longer part of the T_EX equivalence table and because



each had 1.1 million entries with a few memory words each, this makes a major difference in memory usage.

These assignments do not yet show up when using the etex tracing routines `\tracingassigns` and `\tracingrestores` (code simply not written yet)

A side-effect of the current implementation is that `\global` is now more expensive in terms of processing than non-global assignments.

See [mathcodes.c](#) and [textcodes.c](#) if you are interested in the gory details.

Also, the glyph ids within a font are now managed by means of a sparse array and glyph ids can go up to index $2^{21} - 1$.

8.4 Simple single-character csnames

Single-character commands are no longer treated aspecialy in the internals, they are stored in the hash just like the multiletter csnames.

The code that displays control sequences explicitly checks if the length is one when it has to decide whether or not to add a trailing space.

8.5 Compressed format

The format is passed through zlib, allowing it to shrink to roughly half of the size it would have had in uncompressed form. This takes a bit more CPU cycles but much less disk I/O, so it should still be faster.

8.6 Binary file reading

All of the internal code is changed in such a way that if one of the `read_XXX_file` callbacks is not set, then the file is read by a C function using basically the same convention as the callback: a single read into a buffer big enough to hold the entire file contents. While this uses more memory than the previous code (that mostly used `getc` calls), it can be quite a bit faster (depending on your I/O subsystem).



9 Known bugs and limitations

The bugs below are going to be fixed eventually.

The top ones will be fixed soon, but in the later items either the actual problem is hard to find, or the code that causes the bug is going to be replaced by a new subsystem soon anyway, or it may not be worth the hassle and the limitations will eventually be documented.

- Searching of directly-generated PDF documents that use wide fonts is impossible, due to explicit referencing (by L^AT_EX) of font glyph indices without a ToUnicode map being included.
- Hyphenation can only deal with the Base Multilingual Plane (BMP)
- `tex.print()` and `tex.sprint()` do not work if `\directlua` is used in an OTF file (in the output of an `expression` rule).
- Handling of attributes in math mode is not complete. The data structures in math mode are quite different from those in text mode, so this will take some extra effort to implement correctly.
- When used inside `\directlua`, `pdf.print()` should create a literal node instead of flushing immediately.
- At the moment, only characters in plane 0 and plane 1 can be assigned catcode 13 (i.e. turned into active characters). This is a temporary measure to reduce the memory requirements of L^AT_EX. In general, L^AT_EX's memory footprint is a bit larger than we would like (with `plain.fmt` preloaded it needs about 55MB).
- Not all of ALEPH's direction commands are handled properly in PDF mode, and especially the vertical scripts support is missing almost completely (only TRT and TLT are actually tested).
- Letter spacing (`\letterspacefont`) is currently non-functional due to massive changes in the virtual font handling. This functionality may actually be removed completely in the future, because it is straightforward to set up letterspacing using the Lua `define_font` interface.





10 TODO

On top of the normal extensions that are planned, there are some more specific small feature requests. Whether these will all be included is not certain yet, (and new requests are welcome).

- Implement the T_EX primitive `\dimension`, cf. `\number`
- Change the Lua table `typetex.dimen` to accept and return float values instead of strings
- Do something about `\withoutpt` and/or a new register type `\real`?
- Create callback for the automatic creation of missing characters in fonts
- Implement the T_EX primitive `\htdp`?
- Do boxes with dual baselines.
- A way to (re?)calculate the width of a `\vbox`, taking only the natural width of the included items into account.
- Make the number of the output box configurable.
- Complete the attributes in math and switch all the nodes to a double-linked list.
- Finish the interface from Lua to T_EX's internals, specially the hash and equivalence table (a small subpart is implementing `\csname` lookups for `tex.box` access).
- Integrate the various PDF_TE_X extended font codes for hz en protruding into the font table
- Use of Type1C for embedded PostScript font subsets.
- Support font reencoding of 8-bit fonts via char index.
- Attempt to parse OFM level 0 fonts that are masquerading as level 1.
- Add line numbers to the lua errors



