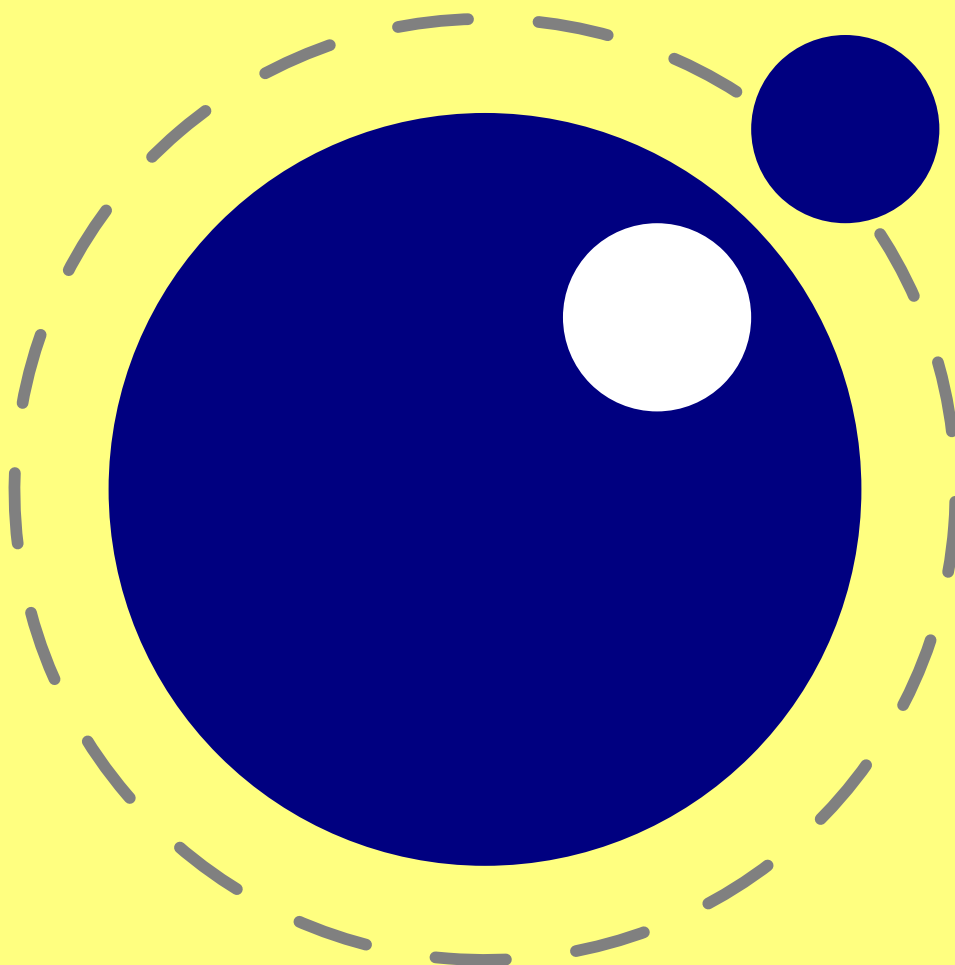


# LuaT<sub>E</sub>X

# Reference

beta 0.80.1





# **LuaT<sub>E</sub>X**

# **Reference**

# **Manual**

copyright: LuaT<sub>E</sub>X development team  
more info: [www.luatex.org](http://www.luatex.org)  
version: September 22, 2015



# Contents

1	Introduction	15
2	Basic T <sub>E</sub> X enhancements	17
2.1	Introduction	17
2.2	Version information	17
2.3	UNICODE text support	18
2.4	Extended tables	18
2.5	Attribute registers	19
2.5.1	Box attributes	19
2.6	LUA related primitives	20
2.6.1	<code>\directlua</code>	20
2.6.2	<code>\luafunction</code>	21
2.6.3	<code>\latelua</code>	22
2.6.4	<code>\luaescapestring</code>	22
2.7	New $\varepsilon$ -T <sub>E</sub> X primitives	23
2.7.1	<code>\clearmarks</code>	23
2.7.2	<code>\noligs</code> and <code>\nokerns</code>	23
2.7.3	<code>\formatname</code>	23
2.7.4	<code>\scantextokens</code>	23
2.7.5	Verbose versions of single-character alignments commands (0.45)	24
2.7.6	Catcode tables	24
2.7.6.1	<code>\catcodetable</code>	24
2.7.6.2	<code>\initcatcodetable</code>	24
2.7.6.3	<code>\savecatcodetable</code>	25
2.7.7	<code>\suppressfontnotfounderror</code> (0.11)	25
2.7.8	<code>\suppresslongerror</code> (0.36)	25
2.7.9	<code>\suppressifcsnameerror</code> (0.36)	25
2.7.10	<code>\suppressoutererror</code> (0.36)	25
2.7.11	<code>\suppressmathparerror</code> (0.80)	26
2.7.12	<code>\matheqnogapstep</code> (0.81)	26
2.7.13	<code>\outputbox</code> (0.37)	26



2.7.14	<code>\fontid</code>	26
2.7.15	Font syntax	27
2.7.16	File syntax (0.45)	27
2.7.17	Images and Forms	27
2.8	Debugging	27
2.9	Global leaders	27
2.10	Expandable character codes (0.75)	28
3	LUA general	29
3.1	Initialization	29
3.1.1	LUAT <sub>E</sub> X as a LUA interpreter	29
3.1.2	LUAT <sub>E</sub> X as a LUA byte compiler	29
3.1.3	Other commandline processing	29
3.2	LUA changes	32
3.3	LUA modules	35
4	LUAT <sub>E</sub> X LUA Libraries	37
4.1	The <code>callback</code> library	37
4.1.1	File discovery callbacks	38
4.1.1.1	<code>find_read_file</code> and <code>find_write_file</code>	38
4.1.1.2	<code>find_font_file</code>	38
4.1.1.3	<code>find_output_file</code>	39
4.1.1.4	<code>find_format_file</code>	39
4.1.1.5	<code>find_vf_file</code>	39
4.1.1.6	<code>find_map_file</code>	39
4.1.1.7	<code>find_enc_file</code>	39
4.1.1.8	<code>find_sfd_file</code>	39
4.1.1.9	<code>find_pk_file</code>	39
4.1.1.10	<code>find_data_file</code>	40
4.1.1.11	<code>find_opentype_file</code>	40
4.1.1.12	<code>find_truetype_file</code> and <code>find_type1_file</code>	40
4.1.1.13	<code>find_image_file</code>	40
4.1.2	File reading callbacks	40
4.1.2.1	<code>open_read_file</code>	40



4.1.2.1.1	<code>reader</code>	41
4.1.2.1.2	<code>close</code>	41
4.1.2.2	General file readers	41
4.1.3	Data processing callbacks	42
4.1.3.1	<code>process_input_buffer</code>	42
4.1.3.2	<code>process_output_buffer</code> (0.43)	42
4.1.3.3	<code>process_jobname</code> (0.71)	43
4.1.3.4	<code>token_filter</code>	43
4.1.4	Node list processing callbacks	43
4.1.4.1	<code>buildpage_filter</code>	43
4.1.4.2	<code>pre_linebreak_filter</code>	44
4.1.4.3	<code>linebreak_filter</code>	45
4.1.4.4	<code>post_linebreak_filter</code>	45
4.1.4.5	<code>hpack_filter</code>	45
4.1.4.6	<code>vpack_filter</code>	46
4.1.4.7	<code>pre_output_filter</code>	46
4.1.4.8	<code>hyphenate</code>	46
4.1.4.9	<code>ligaturing</code>	47
4.1.4.10	<code>kerning</code>	47
4.1.4.11	<code>mlist_to_hlist</code>	47
4.1.5	Information reporting callbacks	47
4.1.5.1	<code>pre_dump</code> (0.61)	47
4.1.5.2	<code>start_run</code>	48
4.1.5.3	<code>stop_run</code>	48
4.1.5.4	<code>start_page_number</code>	48
4.1.5.5	<code>stop_page_number</code>	48
4.1.5.6	<code>show_error_hook</code>	48
4.1.5.7	<code>show_error_message</code>	49
4.1.5.8	<code>show_lua_error_hook</code>	49
4.1.5.9	<code>start_file</code>	49
4.1.5.10	<code>stop_file</code>	49
4.1.6	PDF-related callbacks	49
4.1.6.1	<code>finish_pdffile</code>	49



4.1.6.2	<code>finish_pdfpage</code>	50
4.1.7	Font-related callbacks	50
4.1.7.1	<code>define_font</code>	50
4.2	The <code>epdf</code> library	50
4.3	The <code>font</code> library	59
4.3.1	Loading a TFM file	59
4.3.2	Loading a VF file	59
4.3.3	The fonts array	60
4.3.4	Checking a font's status	60
4.3.5	Defining a font directly	60
4.3.6	Projected next font id	61
4.3.7	Font id (0.47)	61
4.3.8	Currently active font	61
4.3.9	Maximum font id	61
4.3.10	Iterating over all fonts	61
4.4	The <code>fontloader</code> library (0.36)	61
4.4.1	Getting quick information on a font	61
4.4.2	Loading an OPENTYPE or TRUETYPE file	62
4.4.3	Applying a 'feature file'	63
4.4.4	Applying an 'AFM file'	64
4.4.5	Fontloader font tables	64
4.4.5.1	Table types	65
4.4.5.1.1	Top-level	65
4.4.5.1.2	Glyph items	67
4.4.5.1.3	map table	70
4.4.5.1.4	private table	71
4.4.5.1.5	cidinfo table	71
4.4.5.1.6	pfminfo table	71
4.4.5.1.7	names table	72
4.4.5.1.8	anchor_classes table	73
4.4.5.1.9	gpos table	73
4.4.5.1.10	gsub table	74
4.4.5.1.11	ttf_tables and ttf_tab_saved tables	75





4.4.5.1.12	mm table	75
4.4.5.1.13	mark_classes table (0.44)	75
4.4.5.1.14	math table	75
4.4.5.1.15	validation_state table	77
4.4.5.1.16	horiz_base and vert_base table	77
4.4.5.1.17	altuni table	78
4.4.5.1.18	vert_variants and horiz_variants table	78
4.4.5.1.19	mathkern table	78
4.4.5.1.20	kerns table	78
4.4.5.1.21	vkerns table	79
4.4.5.1.22	texdata table	79
4.4.5.1.23	lookups table	79
4.5	The <code>img</code> library	80
4.5.1	<code>img.new</code>	80
4.5.2	<code>img.keys</code>	81
4.5.3	<code>img.scan</code>	82
4.5.4	<code>img.copy</code>	83
4.5.5	<code>img.write</code>	83
4.5.6	<code>img.immediatewrite</code>	83
4.5.7	<code>img.node</code>	83
4.5.8	<code>img.types</code>	84
4.5.9	<code>img.bboxes</code>	84
4.6	The <code>kpse</code> library	84
4.6.1	<code>kpse.set_program_name</code> and <code>kpse.new</code>	84
4.6.2	<code>find_file</code>	85
4.6.3	<code>lookup</code>	86
4.6.4	<code>init_prog</code>	87
4.6.5	<code>readable_file</code>	87
4.6.6	<code>expand_path</code>	87
4.6.7	<code>expand_var</code>	87
4.6.8	<code>expand_braces</code>	88
4.6.9	<code>show_path</code>	88
4.6.10	<code>var_value</code>	88



4.6.11	<code>version</code>	88
4.7	The <code>lang</code> library	88
4.8	The <code>lua</code> library	90
4.8.1	LUA bytecode registers	90
4.8.2	LUA chunk name registers	91
4.9	The <code>mplib</code> library	91
4.9.1	<code>mplib.new</code>	91
4.9.2	<code>mp:statistics</code>	92
4.9.3	<code>mp:execute</code>	92
4.9.4	<code>mp:finish</code>	92
4.9.5	Result table	93
4.9.5.1	fill	94
4.9.5.2	outline	94
4.9.5.3	text	95
4.9.5.4	special	95
4.9.5.5	start_bounds, start_clip	95
4.9.5.6	stop_bounds, stop_clip	95
4.9.6	Subsidiary table formats	95
4.9.6.1	Paths and pens	95
4.9.6.2	Colors	96
4.9.6.3	Transforms	96
4.9.6.4	Dashes	96
4.9.7	Character size information	96
4.9.7.1	<code>mp:char_width</code>	96
4.9.7.2	<code>mp:char_height</code>	97
4.9.7.3	<code>mp:char_depth</code>	97
4.10	The <code>node</code> library	97
4.10.1	Node handling functions	98
4.10.1.1	<code>node.is_node</code>	98
4.10.1.2	<code>node.types</code>	98
4.10.1.3	<code>node.whatsits</code>	98
4.10.1.4	<code>node.id</code>	98
4.10.1.5	<code>node.subtype</code>	98



4.10.1.6	<code>node.type</code>	98
4.10.1.7	<code>node.fields</code>	99
4.10.1.8	<code>node.has_field</code>	99
4.10.1.9	<code>node.new</code>	99
4.10.1.10	<code>node.free</code>	99
4.10.1.11	<code>node.flush_list</code>	99
4.10.1.12	<code>node.copy</code>	100
4.10.1.13	<code>node.copy_list</code>	100
4.10.1.14	<code>node.next</code> (0.65)	100
4.10.1.15	<code>node.prev</code> (0.65)	100
4.10.1.16	<code>node.current_attr</code> (0.66)	100
4.10.1.17	<code>node.hpack</code>	101
4.10.1.18	<code>node.vpack</code> (since 0.36)	101
4.10.1.19	<code>node.dimensions</code> (0.43)	102
4.10.1.20	<code>node.mlist_to_hlist</code>	103
4.10.1.21	<code>node.slide</code>	103
4.10.1.22	<code>node.tail</code>	103
4.10.1.23	<code>node.length</code>	103
4.10.1.24	<code>node.count</code>	103
4.10.1.25	<code>node.traverse</code>	103
4.10.1.26	<code>node.traverse_id</code>	104
4.10.1.27	<code>node.end_of_math</code> (0.76)	105
4.10.1.28	<code>node.remove</code>	105
4.10.1.29	<code>node.insert_before</code>	105
4.10.1.30	<code>node.insert_after</code>	105
4.10.1.31	<code>node.first_glyph</code> (0.65)	106
4.10.1.32	<code>node.ligaturing</code>	106
4.10.1.33	<code>node.kerning</code>	106
4.10.1.34	<code>node.unprotect_glyphs</code>	106
4.10.1.35	<code>node.protect_glyphs</code>	106
4.10.1.36	<code>node.last_node</code>	107
4.10.1.37	<code>node.write</code>	107
4.10.1.38	<code>node.protrusion_skipable</code> (0.60.1)	107



4.10.2	Attribute handling	107
4.10.2.1	<code>node.has_attribute</code>	107
4.10.2.2	<code>node.set_attribute</code>	107
4.10.2.3	<code>node.unset_attribute</code>	108
4.11	The <code>pdf</code> library	108
4.11.1	<code>pdf.mapfile</code> , <code>pdf.mapline</code> (new in 0.53.0)	108
4.11.2	<code>pdf.catalog</code> , <code>pdf.info</code> , <code>pdf.names</code> , <code>pdf.trailer</code> (new in 0.53.0)	108
4.11.3	<code>pdf.&lt;set/get&gt;pageattributes</code> , <code>pdf.&lt;set/get&gt;pageresources</code> , <code>pdf.&lt;set/get&gt;pagesattributes</code>	108
4.11.4	<code>pdf.h</code> , <code>pdf.v</code>	109
4.11.5	<code>pdf.getpos</code> , <code>pdf.gethpos</code> , <code>pdf.getvpos</code>	109
4.11.6	<code>pdf.hasmatrix</code> , <code>pdf.getmatrix</code>	109
4.11.7	<code>pdf.print</code>	109
4.11.8	<code>pdf.immediateobj</code>	110
4.11.9	<code>pdf.obj</code>	110
4.11.10	<code>pdf.refobj</code>	111
4.11.11	<code>pdf.reserveobj</code>	111
4.11.12	<code>pdf.registerannot</code> (new in 0.47.0)	112
4.12	The <code>pdfscanner</code> library (new in 0.72.0)	112
4.13	The <code>status</code> library	114
4.14	The <code>tex</code> library	116
4.14.1	Internal parameter values	116
4.14.1.1	Integer parameters	117
4.14.1.2	Dimension parameters	118
4.14.1.3	Direction parameters	119
4.14.1.4	Glue parameters	120
4.14.1.5	Muglue parameters	120
4.14.1.6	Tokenlist parameters	120
4.14.2	Convert commands	120
4.14.3	Last item commands	121
4.14.4	Attribute, count, dimension, skip and token registers	121
4.14.5	Character code registers (0.63)	122
4.14.6	Box registers	123



4.14.7	Math parameters	124
4.14.8	Special list heads	125
4.14.9	Semantic nest levels (0.51)	125
4.14.10	Print functions	126
4.14.10.1	<code>tex.print</code>	127
4.14.10.2	<code>tex.sprint</code>	127
4.14.10.3	<code>tex.tprint</code>	127
4.14.10.4	<code>tex.write</code>	128
4.14.11	Helper functions	128
4.14.11.1	<code>tex.round</code>	128
4.14.11.2	<code>tex.scale</code>	128
4.14.11.3	<code>tex.sp</code> (0.51)	128
4.14.11.4	<code>tex.definefont</code>	129
4.14.11.5	<code>tex.error</code> (0.61)	129
4.14.11.6	<code>tex.hashtokens</code> (0.25)	129
4.14.12	Functions for dealing with primitives	129
4.14.12.1	<code>tex.enableprimitives</code>	129
4.14.12.2	<code>tex.extraprimitives</code>	130
4.14.12.3	<code>tex.primitives</code>	134
4.14.13	Core functionality interfaces	134
4.14.13.1	<code>tex.badness</code> (0.53)	134
4.14.13.2	<code>tex.linebreak</code> (0.53)	134
4.14.13.3	<code>tex.shipout</code> (0.51)	136
4.15	The <code>texconfig</code> table	136
4.16	The <code>texio</code> library	137
4.16.1	Printing functions	137
4.16.1.1	<code>texio.write</code>	137
4.16.1.2	<code>texio.write_nl</code>	137
4.17	The <code>token</code> library	138
4.17.1	<code>token.get_next</code>	138
4.17.2	<code>token.is_expandable</code>	138
4.17.3	<code>token.expand</code>	138
4.17.4	<code>token.is_activechar</code>	138



4.17.5	<code>token.create</code>	139
4.17.6	<code>token.command_name</code>	139
4.17.7	<code>token.command_id</code>	139
4.17.8	<code>token.csname_name</code>	139
4.17.9	<code>token.csname_id</code>	140
4.17.10	The <code>newtoken</code> library	140
5	Math	145
5.1	The current math style	145
5.1.1	<code>\mathstyle</code>	145
5.1.2	<code>\Ustack</code>	145
5.2	Unicode math characters	146
5.3	Cramped math styles	147
5.4	Math parameter settings	147
5.5	Font-based Math Parameters	149
5.6	Math spacing setting	151
5.7	Math accent handling	152
5.8	Math root extension	153
5.9	Math kerning in super- and subscripts	153
5.10	Scripts on horizontally extensible items like arrows	154
5.11	Extensible delimiters	154
5.12	Other Math changes	154
5.12.1	Verbose versions of single-character math commands	154
5.12.2	Allowed math commands in non-math modes	155
5.13	Math todo	155
6	Languages and characters, fonts and glyphs	157
6.1	Characters and glyphs	157
6.2	The main control loop	158
6.3	Loading patterns and exceptions	159
6.4	Applying hyphenation	160
6.5	Applying ligatures and kerning	162
6.6	Breaking paragraphs into lines	164



7	Font structure	165
7.1	Real fonts	170
7.2	Virtual fonts	172
7.2.1	Artificial fonts	174
7.2.2	Example virtual font	174
8	Nodes	177
8.1	LUA node representation	177
8.1.1	Auxiliary items	177
8.1.1.1	glue_spec items	177
8.1.1.2	attribute_list and attribute items	178
8.1.1.3	action item	178
8.1.2	Main text nodes	179
8.1.2.1	hlist nodes	179
8.1.2.2	vlist nodes	180
8.1.2.3	rule nodes	180
8.1.2.4	ins nodes	180
8.1.2.5	mark nodes	180
8.1.2.6	adjust nodes	181
8.1.2.7	disc nodes	181
8.1.2.8	math nodes	181
8.1.2.9	glue nodes	182
8.1.2.10	kern nodes	182
8.1.2.11	penalty nodes	183
8.1.2.12	glyph nodes	183
8.1.2.13	margin_kern nodes	184
8.1.3	Math nodes	184
8.1.3.1	Math kernel subnodes	184
8.1.3.1.1	math_char and math_text_char subnodes	184
8.1.3.1.2	sub_box and sub_mlist subnodes	185
8.1.3.2	Math delimiter subnode	185
8.1.3.2.1	delim subnodes	185
8.1.3.3	Math core nodes	186
8.1.3.3.1	simple nodes	186



8.1.3.3.2	accent nodes	186
8.1.3.3.3	style nodes	187
8.1.3.3.4	choice nodes	187
8.1.3.3.5	radical nodes	187
8.1.3.3.6	fraction nodes	188
8.1.3.3.7	fence nodes	188
8.1.4	whatsit nodes	188
8.1.4.1	open nodes	188
8.1.4.2	write nodes	189
8.1.4.3	close nodes	189
8.1.4.4	special nodes	189
8.1.4.5	language nodes	189
8.1.4.6	local_par nodes	189
8.1.4.7	dir nodes	190
8.1.4.8	pdf_literal nodes	190
8.1.4.9	pdf_refobj nodes	191
8.1.4.10	pdf_refxform nodes	191
8.1.4.11	pdf_refximage nodes	191
8.1.4.12	pdf_annot nodes	192
8.1.4.13	pdf_start_link nodes	192
8.1.4.14	pdf_end_link nodes	192
8.1.4.15	pdf_dest nodes	192
8.1.4.16	pdf_thread nodes	193
8.1.4.17	pdf_start_thread nodes	193
8.1.4.18	pdf_end_thread nodes	193
8.1.4.19	pdf_save_pos nodes	194
8.1.4.20	late_lua nodes	194
8.1.4.21	pdf_colorstack nodes	194
8.1.4.22	pdf_setmatrix nodes	194
8.1.4.23	pdf_save nodes	195
8.1.4.24	pdf_restore nodes	195
8.1.4.25	user_defined nodes	195
8.2	Two access models	196





9	Modifications	201
9.1	Changes from T <sub>E</sub> X 3.1415926	201
9.2	Changes from $\varepsilon$ -T <sub>E</sub> X 2.2	201
9.3	Changes from PDFT <sub>E</sub> X 1.40	201
9.4	Changes from ALEPH RC4	202
9.5	Changes from standard WEB2C	203
10	Implementation notes	205
10.1	Primitives overlap	205
10.2	Memory allocation	205
10.3	Sparse arrays	205
10.4	Simple single-character csnames	206
10.5	Compressed format	206
10.6	Binary file reading	206
11	Known bugs and limitations, TODO	207





# 1 Introduction

This book will eventually become the reference manual of L<sup>A</sup>T<sub>E</sub>X. At the moment, it simply reports the behavior of the executable matching the snapshot or beta release date in the title page.

Features may come and go. The current version of L<sup>A</sup>T<sub>E</sub>X is not meant for production and users cannot depend on stability, nor on functionality staying the same.

Nothing is considered stable just yet. This manual therefore simply reflects the current state of the executable. ***Absolutely nothing*** on the following pages is set in stone. When the need arises, anything can (and will) be changed.

L<sup>A</sup>T<sub>E</sub>X consists of a number of interrelated but (still) distinguishable parts:

- PDF<sub>T</sub>E<sub>X</sub> version 1.40.9, converted to C (with patches from later releases).
- The direction model and some other bits from ALEPH RC4 converted to C.
- L<sup>A</sup>UA 5.2.1
- dedicated L<sup>A</sup>UA libraries
- various T<sub>E</sub>X extensions
- parts of FONTFORGE 2008.11.17
- the METAPOST library
- newly written compiled source code to glue it all together

Neither ALEPH's I/O translation processes, nor tcx files, nor ENC<sub>T</sub>E<sub>X</sub> can be used, these encoding-related functions are superseded by a L<sup>A</sup>UA-based solution (reader callbacks). Also, some experimental PDF<sub>T</sub>E<sub>X</sub> features are removed. These can be implemented in L<sup>A</sup>UA instead.





## 2 Basic T<sub>E</sub>X enhancements

### 2.1 Introduction

From day one, L<sup>A</sup>T<sub>E</sub>X has offered extra functionality when compared to the superset of P<sub>D</sub>F<sub>T</sub>E<sub>X</sub> and A<sub>L</sub>E<sub>P</sub>H. That has not been limited to the possibility to execute L<sup>A</sup>U<sub>A</sub> code via `\directlua`, but L<sup>A</sup>T<sub>E</sub>X also adds functionality via new T<sub>E</sub>X-side primitives.

However, starting with beta 0.39.0, most of that functionality is hidden by default. When L<sup>A</sup>T<sub>E</sub>X 0.40.0 starts up in ‘iniluatex’ mode (`luatex -ini`), it defines only the primitive commands known by T<sub>E</sub>X82 and the one extra command `\directlua`.

As is fitting, a L<sup>A</sup>U<sub>A</sub> function has to be called to add the extra primitives to the user environment. The simplest method to get access to all of the new primitive commands is by adding this line to the format generation file:

```
\directlua { tex.enableprimitives('',tex.extraprimitives()) }
```

But be aware that the curly braces may not have the proper `\catcode` assigned to them at this early time (giving a ‘Missing number’ error), so it may be needed to put these assignments

```
\catcode `\{=1
\catcode `\}=2
```

before the above line. More fine-grained primitives control is possible, you can look up the details in [section 4.14.12](#). For simplicity’s sake, this manual assumes that you have executed the `\directlua` command as given above.

The startup behavior documented above is considered stable in the sense that there will not be backward-incompatible changes any more.

### 2.2 Version information

There are three new primitives to test the version of L<sup>A</sup>T<sub>E</sub>X:

primitive	explanation	value
<code>\luatexbanner</code>	the banner as reported on the command line	This is LuaTeX, Version beta-0.80.1 (TeX Live 2016/dev)
<code>\luatexversion</code>	a combination of major and minor number	80
<code>\luatexrevision</code>	the revision number, the current value is	1

The official L<sup>A</sup>T<sub>E</sub>X version is defined as follows:



- The major version is the integer result of `\luatexversion` divided by 100. The primitive is an ‘internal variable’, so you may need to prefix its use with `\the` depending on the context.
- The minor version is the two-digit result of `\luatexversion` modulo 100.
- The revision is the given by `\luatexrevision`. This primitive expands to a positive integer.
- The full version number consists of the major version, minor version and revision, separated by dots.

## 2.3 UNICODE text support

Text input and output is now considered to be UNICODE text, so input characters can use the full range of UNICODE ( $2^{20} + 2^{16} - 1 = 0x10FFFF$ ).

Later chapters will talk of characters and glyphs. Although these are not interchangeable, they are closely related. During typesetting, a character is always converted to a suitable graphic representation of that character in a specific font. However, while processing a list of to-be-typeset nodes, its contents may still be seen as a character. Inside L<sup>A</sup>T<sub>E</sub>X there is not yet a clear separation between the two concepts. Until this is implemented, please do not be too harsh on us if we make errors in the usage of the terms.

A few primitives are affected by this, all in a similar fashion: each of them has to accommodate for a larger range of acceptable numbers. For instance, `\char` now accepts values between 0 and 1,114,111. This should not be a problem for well-behaved input files, but it could create incompatibilities for input that would have generated an error when processed by older T<sub>E</sub>X-based engines. The affected commands with an altered initial (left of the equals sign) or secondary (right of the equals sign) value are: `\char`, `\lccode`, `\uccode`, `\catcode`, `\sfcode`, `\efcode`, `\lpcode`, `\rprcode`, `\chardef`.

As far as the core engine is concerned, all input and output to text files is UTF-8 encoded. Input files can be pre-processed using the `reader` callback. This will be explained in a later chapter.

Output in byte-sized chunks can be achieved by using characters just outside of the valid UNICODE range, starting at the value 1,114,112 (0x110000). When the time comes to print a character  $c \geq 1,114,112$ , L<sup>A</sup>T<sub>E</sub>X will actually print the single byte corresponding to  $c$  minus 1,114,112.

Output to the terminal uses `^^` notation for the lower control range ( $c < 32$ ), with the exception of `^^I`, `^^J` and `^^M`. These are considered ‘safe’ and therefore printed as-is.

Normalization of the UNICODE input can be handled by a macro package during callback processing (this will be explained in [section 4.1.2](#)).

## 2.4 Extended tables

All traditional T<sub>E</sub>X and  $\epsilon$ -T<sub>E</sub>X registers can be 16-bit numbers as in ALEPH. The affected commands are:

<code>\count</code>	<code>\toks</code>	<code>\toksdef</code>	<code>\copy</code>
<code>\dimen</code>	<code>\countdef</code>	<code>\insert</code>	<code>\unhcopy</code>
<code>\skip</code>	<code>\dimendef</code>	<code>\box</code>	<code>\unvcopy</code>
<code>\muskip</code>	<code>\skipdef</code>	<code>\unhbox</code>	<code>\wd</code>
<code>\marks</code>	<code>\muskipdef</code>	<code>\unvbox</code>	<code>\ht</code>



`\dp`

`\setbox`

`\vsplit`

The glyph properties (like `\efcode`) introduced in `PDFTEX` that deal with font expansion (hz) and character protruding are also 16-bit. Because font memory management has been rewritten, these character properties are no longer shared among fonts instances that originate from the same metric file.

The behavior documented in the above section is considered stable in the sense that there will not be backward-incompatible changes any more.

## 2.5 Attribute registers

Attributes are a completely new concept in `LUATEX`. Syntactically, they behave a lot like counters: attributes obey `TEX`'s nesting stack and can be used after `\the` etc. just like the normal `\count` registers.

`\attribute` `<16-bit number>` `<optional equals>` `<32-bit number>`

`\attributedef` `<csname>` `<optional equals>` `<16-bit number>`

Conceptually, an attribute is either ‘set’ or ‘unset’. Unset attributes have a special negative value to indicate that they are unset, that value is the lowest legal value: `-0x7FFFFFFF` in hexadecimal, a.k.a. `-2147483647` in decimal. It follows that the value `-0x7FFFFFFF` cannot be used as a legal attribute value, but you *can* assign `-0x7FFFFFFF` to ‘unset’ an attribute. All attributes start out in this ‘unset’ state in `INITTEX` (prior to 0.37, there could not be valid negative attribute values, and the ‘unset’ value was `-1`).

Attributes can be used as extra counter values, but their usefulness comes mostly from the fact that the numbers and values of all ‘set’ attributes are attached to all nodes created in their scope. These can then be queried from any `LUA` code that deals with node processing. Further information about how to use attributes for node list processing from `LUA` is given in **chapter 8**.

The behavior documented in the above subsection is considered stable in the sense that there will not be backward-incompatible changes any more.

### 2.5.1 Box attributes

Nodes typically receive the list of attributes that is in effect when they are created. This moment can be quite asynchronous. For example: in paragraph building, the individual line boxes are created after the `\par` command has been processed, so they will receive the list of attributes that is in effect then, not the attributes that were in effect in, say, the first or third line of the paragraph.

Similar situations happen in `LUATEX` regularly. A few of the more obvious problematic cases are dealt with: the attributes for nodes that are created during hyphenation, kerning and ligaturing borrow their attributes from their surrounding glyphs, and it is possible to influence box attributes directly.

When you assemble a box in a register, the attributes of the nodes contained in the box are unchanged when such a box is placed, unboxed, or copied. In this respect attributes act the same as characters that have been converted to references to glyphs in fonts. For instance, when you use attributes to implement color support, each node carries information about its eventual color. In that case, unless you implement mechanisms that deal with it, applying a color to already boxed material will have no effect. Keep in



mind that this incompatibility is mostly due to the fact that separate specials and literals are a more unnatural approach to colors than attributes.

It is possible to fine-tune the list of attributes that are applied to a `hbox`, `vbox` or `vtop` by the use of the keyword `attr`. An example:

```
\attribute2=5
\setbox0=\hbox {Hello}
\setbox2=\hbox attr1=12 attr2=-"7FFFFFFF{Hello}
```

This will set the attribute list of box 2 to 1 = 12, and the attributes of box 0 will be 2 = 5. As you can see, assigning the maximum negative value causes an attribute to be ignored.

The `attr` keyword(s) should come before a `to` or `spread`, if that is also specified.

## 2.6 LUA related primitives

In order to merge LUA code with T<sub>E</sub>X input, a few new primitives are needed.

### 2.6.1 `\directlua`

The primitive `\directlua` is used to execute LUA code immediately. The syntax is

```
\directlua <general text>
\directlua name <general text> <general text>
\directlua <16-bit number> <general text>
```

The last `<general text>` is expanded fully, and then fed into the LUA interpreter. After reading and expansion has been applied to the `<general text>`, the resulting token list is converted to a string as if it was displayed using `\the\toks`. On the LUA side, each `\directlua` block is treated as a separate chunk. In such a chunk you can use the `local` directive to keep your variables from interfering with those used by the macro package.

The conversion to and from a token list means that you normally can not use LUA line comments (starting with `--`) within the argument. As there typically will be only one ‘line’ the first line comment will run on until the end of the input. You will either need to use T<sub>E</sub>X-style line comments (starting with `%`), or change the T<sub>E</sub>X category codes locally. Another possibility is to say:

```
\begingroup
\endlinechar=10
\directlua ...
\endgroup
```

Then LUA line comments can be used, since T<sub>E</sub>X does not replace line endings with spaces.

The `name <general text>` specifies the name of the LUA chunk, mainly shown in the stack backtrace of error messages created by LUA code. The `<general text>` is expanded fully, thus macros can be used to generate the chunk name, i.e.





`\directlua name{\jobname:\the\inputlineno} ...`

to include the name of the input file as well as the input line into the chunk name.

Likewise, the `<16-bit number>` designates a name of a Lua chunk, but in this case the name will be taken from the `lua.name` array (see the documentation of the `lua` table further in this manual). This syntax is new in version 0.36.0.

The chunk name should not start with a `@`, or it will be displayed as a file name (this is a quirk in the current Lua implementation).

The `\directlua` command is expandable. Since it passes Lua code to the Lua interpreter its expansion from the T<sub>E</sub>X viewpoint is usually empty. However, there are some Lua functions that produce material to be read by T<sub>E</sub>X, the so called print functions. The most simple use of these is `tex.print(<string>s)`. The characters of the string `s` will be placed on the T<sub>E</sub>X input buffer, that is, ‘before T<sub>E</sub>X’s eyes’ to be read by T<sub>E</sub>X immediately. For example:

```
\count10=20
a\directlua{tex.print(tex.count[10]+5)}b
```

expands to

a25b

Here is another example:

```
$\pi = \directlua{tex.print(math.pi)}$
```

will result in

$\pi = 3.1415926535898$

Note that the expansion of `\directlua` is a sequence of characters, not of tokens, contrary to all T<sub>E</sub>X commands. So formally speaking its expansion is null, but it places material on a pseudo-file to be immediately read by T<sub>E</sub>X, as  $\epsilon$ -T<sub>E</sub>X’s `\scantokens`.

For a description of print functions look at **section 4.14.10**.

Because the `<general text>` is a chunk, the normal Lua error handling is triggered if there is a problem in the included code. The Lua error messages should be clear enough, but the contextual information is still pretty bad. Often, you will only see the line number of the right brace at the end of the code.

While on the subject of errors: some of the things you can do inside Lua code can break up L<sup>A</sup>T<sub>E</sub>X pretty bad. If you are not careful while working with the node list interface, you may even end up with assertion errors from within the T<sub>E</sub>X portion of the executable.

The behavior documented in the above subsection is considered stable in the sense that there will not be backward-incompatible changes any more.

## 2.6.2 `\luafunction`

The `\directlua` commands involves tokenization of its argument (after picking up an optional name or number specification). The tokenlist is then converted into a string and given to Lua to turn into a function



that is called. The overhead is rather small but when you use this primitive hundreds or thousands of times, it can become noticeable. For this reason there is a variant call available: `\luafunction`. This command is used as follows:

```
\directlua {
  local t = lua.get_functions_table()
  t[1] = function() tex.print("!") end
  t[2] = function() tex.print("?") end
}
```

```
\luafunction1
\uafunction2
```

Of course the functions can also be defined in a separate file. There is no limit on the number of functions apart from normal LUA limitations. Of course there is the limitation of no arguments but that would involve parsing and thereby give no gain. The function, when called in fact gets one argument, being the index, so in:

```
\directlua {
  local t = lua.get_functions_table()
  t[8] = function(slot) tex.print(slot) end
}
```

the number 8 gets typeset.

### 2.6.3 `\latelua`

`\latelua` stores LUA code in a whatsit that will be processed at the time of shipping out. Its intended use is a cross between `\pdfliteral` and `\write`. Within the LUA code you can print PDF statements directly to the PDF file via `pdf.print`, or you can write to other output streams via `texio.write` or simply using lua's I/O routines.

```
\latelua <general text>
\latelua name <general text> <general text>
\latelua <16-bit number> <general text>
```

Expansion of macros etcetera in the final `<general text>` is delayed until just before the whatsit is executed (like in `\write`). With regard to PDF output stream `\latelua` behaves as `\pdfliteral page`.

The `name <general text>` and `<16-bit number>` behave in the same way as they do for `\directlua`

### 2.6.4 `\luaescapestring`

This primitive converts a T<sub>E</sub>X token sequence so that it can be safely used as the contents of a LUA string: embedded backslashes, double and single quotes, and newlines and carriage returns are escaped. This



is done by prepending an extra token consisting of a backslash with category code 12, and for the line endings, converting them to `n` and `r` respectively. The token sequence is fully expanded.

```
\luaescapestring <general text>
```

Most often, this command is not actually the best way to deal with the differences between the  $\TeX$  and  $\text{Lua}$ . In very short bits of  $\text{Lua}$  code it is often not needed, and for longer stretches of  $\text{Lua}$  code it is easier to keep the code in a separate file and load it using  $\text{Lua}$ 's `dofile`:

```
\directlua { dofile('mysetups.lua')}
```

## 2.7 New $\varepsilon$ - $\text{TeX}$ primitives

### 2.7.1 `\clearmarks`

This primitive clears a mark class completely, resetting all three connected mark texts to empty.

```
\clearmarks <16-bit number>
```

### 2.7.2 `\noligs` and `\nokerns`

These primitives prohibit ligature and kerning insertion at the time when the initial node list is built by  $\text{LuaTeX}$ 's main control loop. They are part of a temporary trick and will be removed in the near future. For now, you need to enable these primitives when you want to do node list processing of 'characters', where  $\text{TeX}$ 's normal processing would get in the way.

```
\noligs <integer>  
\nokerns <integer>
```

These primitives can now be implemented by overloading the ligature building and kerning functions, i.e. by assigning dummy functions to their associated callbacks.

### 2.7.3 `\formatname`

`\formatname`'s syntax is identical to `\jobname`.

In  $\text{iniTeX}$ , the expansion is empty. Otherwise, the expansion is the value that `\jobname` had during the  $\text{iniTeX}$  run that dumped the currently loaded format.

### 2.7.4 `\scantextokens`

The syntax of `\scantextokens` is identical to `\scantokens`. This primitive is a slightly adapted version of  $\varepsilon$ - $\text{TeX}$ 's `\scantokens`. The differences are:



- The last (and usually only) line does not have a `\endlinechar` appended
- `\scantextokens` never raises an EOF error, and it does not execute `\everyeof` tokens.
- The ‘... while end of file ...’ error tests are not executed, allowing the expansion to end on a different grouping level or while a conditional is still incomplete.

## 2.7.5 Verbose versions of single-character alignments commands (0.45)

L<sup>A</sup>T<sub>E</sub>X defines two new primitives that have the same function as `#` and `&` in alignments:

primitive	explanation
<code>\alignmark</code>	Duplicates the functionality of <code>#</code> inside alignment preambles
<code>\aligntab</code>	Duplicates the functionality of <code>&amp;</code> inside alignments (and preambles)

## 2.7.6 Catcode tables

Catcode tables are a new feature that allows you to switch to a predefined catcode regime in a single statement. You can have a practically unlimited number of different tables.

The subsystem is backward compatible: if you never use the following commands, your document will not notice any difference in behavior compared to traditional T<sub>E</sub>X.

The contents of each catcode table is independent from any other catcode tables, and their contents is stored and retrieved from the format file.

### 2.7.6.1 `\catcodetable`

`\catcodetable` <15-bit number>

The primitive `\catcodetable` switches to a different catcode table. Such a table has to be previously created using one of the two primitives below, or it has to be zero. Table zero is initialized by `INITEX`.

### 2.7.6.2 `\initcatcodetable`

`\initcatcodetable` <15-bit number>

The primitive `\initcatcodetable` creates a new table with catcodes identical to those defined by `INITEX`:

0	<code>\</code>		escape	
5	<code>^^M</code>	return	<code>car_ret</code>	(this name may change)
9	<code>^^@</code>	null	ignore	
10	<code>&lt;space&gt;</code>	space	spacer	
11	<code>a – z</code>		letter	
11	<code>A – Z</code>		letter	
12	everything else		other	



```

14 %                comment
15 ^^?            delete  invalid_char

```

The new catcode table is allocated globally: it will not go away after the current group has ended. If the supplied number is identical to the currently active table, an error is raised.

### 2.7.6.3 `\savecatcodetable`

`\savecatcodetable` <15-bit number>

`\savecatcodetable` copies the current set of catcodes to a new table with the requested number. The definitions in this new table are all treated as if they were made in the outermost level.

The new table is allocated globally: it will not go away after the current group has ended. If the supplied number is the currently active table, an error is raised.

### 2.7.7 `\suppressfontnotfounderror` (0.11)

`\suppressfontnotfounderror = 1`

If this new integer parameter is non-zero, then L<sup>A</sup>T<sub>E</sub>X will not complain about font metrics that are not found. Instead it will silently skip the font assignment, making the requested csname for the font `\ifx` equal to `\nullfont`, so that it can be tested against that without bothering the user.

### 2.7.8 `\suppresslongerror` (0.36)

`\suppresslongerror = 1`

If this new integer parameter is non-zero, then L<sup>A</sup>T<sub>E</sub>X will not complain about `\par` commands encountered in contexts where that is normally prohibited (most prominently in the arguments of non-long macros).

### 2.7.9 `\suppressifcsnameerror` (0.36)

`\suppressifcsnameerror = 1`

If this new integer parameter is non-zero, then L<sup>A</sup>T<sub>E</sub>X will not complain about non-expandable commands appearing in the middle of a `\ifcsname` expansion. Instead, it will keep getting expanded tokens from the input until it encounters an `\endcsname` command. Use with care! This command is experimental: if the input expansion is unbalanced wrt. `\csname ... \endcsname` pairs, the L<sup>A</sup>T<sub>E</sub>X process may hang indefinitely.

### 2.7.10 `\suppressoutererror` (0.36)

`\suppressoutererror = 1`



If this new integer parameter is non-zero, then `LUATEX` will not complain about `\outer` commands encountered in contexts where that is normally prohibited.

The addition of this command coincides with a change in the `LUATEX` engine: ever since the snapshot of 20060915, `\outer` was simply ignored. That behavior has now reverted back to be `TEX82`-compatible by default.

### 2.7.11 `\suppressmathparerror` (0.80)

The following setting will permit tokens in a math formula:

```
\suppressmathparerror = 1
```

So, the next code is valid then:

```
$ x + 1 =
```

```
a $
```

### 2.7.12 `\matheqnogapstep` (0.81)

By default `TEX` will add one quad between the equation and the number. This is hardcoded. A new primitive can control this:

```
\matheqnogapstep = 1000
```

Because a math quad from the math text font is used instead of a dimension, we use a step to control the size. A value of zero will suppress the gap. The step is divided by 1000 which is the usual way to mimick floating point factors in `TEX`.

### 2.7.13 `\outputbox` (0.37)

```
\outputbox = 65535
```

This new integer parameter allows you to alter the number of the box that will be used to store the page sent to the output routine. Its default value is 255, and the acceptable range is from 0 to 65535.

### 2.7.14 `\fontid`

```
\fontid\font
```

This primitive expands into a number. It is not a register so there is no need to prefix with `\number` (and using `\the` gives an error). The currently used font id is 1. Here are some more:

```
\bf 93  
\it 95  
\bi 144
```



These numbers depend on the macro package used because each one has its own way of dealing with fonts. They can also differ per run, as they can depend on the order of loading fonts. For instance, when in `CONTEXT` virtual math `UNICODE` fonts are used, we can easily get over a hundred ids in use. Not all ids have to be bound to a real font, after all it's just a number.

### 2.7.15 Font syntax

L<sup>A</sup>T<sub>E</sub>X will accept a braced argument as a font name:

```
\font\myfont = {cmr10}
```

This allows for embedded spaces, without the need for double quotes. Macro expansion takes place inside the argument.

### 2.7.16 File syntax (0.45)

L<sup>A</sup>T<sub>E</sub>X will accept a braced argument as a file name:

```
\input {plain}  
\openin 0 {plain}
```

This allows for embedded spaces, without the need for double quotes. Macro expansion takes place inside the argument.

### 2.7.17 Images and Forms

L<sup>A</sup>T<sub>E</sub>X accepts optional dimension parameters for `\pdfrefximage` and `\pdfrefxform` in the same format as for `\pdfximage`. With images, these dimensions are then used instead of the ones given to `\pdfximage`; but the original dimensions are not overwritten, so that a `\pdfrefximage` without dimensions still provides the image with dimensions defined by `\pdfximage`. These optional parameters are not implemented for `\pdfxform`.

```
\pdfrefximage width 20mm height 10mm depth 5mm \pdflastximage  
\pdfrefxform width 20mm height 10mm depth 5mm \pdflastxform
```

## 2.8 Debugging

If `\tracingonline` is larger than 2, the node list display will also print the node number of the nodes.

## 2.9 Global leaders

There is a new experimental primitive: `\gleaders` (a L<sup>A</sup>T<sub>E</sub>X extension, added in 0.43). This type of leaders is anchored to the origin of the box to be shipped out. So they are like normal `\leaders` in



that they align nicely, except that the alignment is based on the *largest* enclosing box instead of the *smallest*.

## 2.10 Expandable character codes (0.75)

The new expandable command `\Uchar` reads a number between 0 and 1,114,111 and expands to the associated Unicode character.





## 3 LUA general

### 3.1 Initialization

#### 3.1.1 L<sup>A</sup>T<sub>E</sub>X as a LUA interpreter

There are some situations that make L<sup>A</sup>T<sub>E</sub>X behave like a standalone LUA interpreter:

- if a `--luaonly` option is given on the commandline, or
- if the executable is named `texlua` (or `luatexlua`), or
- if the only non-option argument (file) on the commandline has the extension `lua` or `luc`.

In this mode, it will set LUA's `arg[0]` to the found script name, pushing preceding options in negative values and the rest of the commandline in the positive values, just like the LUA interpreter.

L<sup>A</sup>T<sub>E</sub>X will exit immediately after executing the specified LUA script and is, in effect, a somewhat bulky standalone LUA interpreter with a bunch of extra preloaded libraries.

#### 3.1.2 L<sup>A</sup>T<sub>E</sub>X as a LUA byte compiler

There are two situations that make L<sup>A</sup>T<sub>E</sub>X behave like the LUA byte compiler:

- if a `--luaonly` option is given on the commandline, or
- if the executable is named `texluac`

In this mode, L<sup>A</sup>T<sub>E</sub>X is exactly like `luac` from the standalone LUA distribution, except that it does not have the `-l` switch, and that it accepts (but ignores) the `--luaonly` switch.

#### 3.1.3 Other commandline processing

When the L<sup>A</sup>T<sub>E</sub>X executable starts, it looks for the `--lua` commandline option. If there is no `--lua` option, the commandline is interpreted in a similar fashion as in traditional PDF<sub>T</sub>E<sub>X</sub> and ALEPH.

The following command-line switches are understood.

<code>--fmt=FORMAT</code>	load the format file <code>FORMAT</code>
<code>--lua=FILE</code>	load and execute a LUA initialization script
<code>--safer</code>	disable easily exploitable LUA commands
<code>--nosocket</code>	disable the LUA socket library
<code>--help</code>	display help and exit
<code>--ini</code>	be iniluatex, for dumping formats
<code>--interaction=STRING</code>	set interaction mode (STRING=batchmode/nonstopmode/scrollmode/errorstopmode)



<code>--halt-on-error</code>	stop processing at the first error
<code>--kpathsea-debug=NUMBER</code>	set path searching debugging flags according to the bits of NUMBER
<code>--progname=STRING</code>	set the program name to STRING
<code>--version</code>	display version and exit
<code>--credits</code>	display credits and exit
<code>--recorder</code>	enable filename recorder
<code>--etex</code>	ignored
<code>--output-comment=STRING</code>	use STRING for DVI file comment instead of date (no effect for PDF)
<code>--output-directory=DIR</code>	use DIR as the directory to write files to
<code>--draftmode</code>	switch on draft mode (generates no output PDF)
<code>--output-format=FORMAT</code>	use FORMAT for job output; FORMAT is 'dvi' or 'pdf'
<code>--[no-]shell-escape</code>	disable/enable <code>\write 18{SHELL COMMAND}</code>
<code>--enable-write18</code>	enable <code>\write 18{SHELL COMMAND}</code>
<code>--disable-write18</code>	disable <code>\write 18{SHELL COMMAND}</code>
<code>--shell-restricted</code>	restrict <code>\write 18</code> to a list of commands given in texmf.cnf
<code>--debug-format</code>	enable format debugging
<code>--[no-]file-line-error</code>	disable/enable file:line:error style messages
<code>--[no-]file-line-error-style</code>	aliases of <code>--[no-]file-line-error</code>
<code>--jobname=STRING</code>	set the job name to STRING
<code>--[no-]parse-first-line</code>	ignored
<code>--translate-file=</code>	ignored
<code>--default-translate-file=</code>	ignored
<code>--8bit</code>	ignored
<code>--[no-]mktex=FMT</code>	disable/enable mktexFMT generation (FMT=tex/tfm)
<code>--synctex=NUMBER</code>	enable synctex

A note on the creation of the various temporary files and the `\jobname`. The value to use for `\jobname` is decided as follows:

- If `--jobname` is given on the command line, its argument will be the value for `\jobname`, without any changes. The argument will not be used for actual input so it need not exist. The `--jobname` switch only controls the `\jobname` setting.
- Otherwise, `\jobname` will be the name of the first file that is read from the file system, with any path components and the last extension (the part following the last `.`) stripped off.
- An exception to the previous point: if the command line goes into interactive mode (by starting with a command) and there are no files input via `\everyjob` either, then the `\jobname` is set to `texput` as a last resort.

The file names for output files that are generated automatically are created by attaching the proper extension (`.log`, `.pdf`, etc.) to the found `\jobname`. These files are created in the directory pointed to by `--output-directory`, or in the current directory, if that switch is not present.

Without the `--lua` option, command line processing works like it does in any other web2c-based type-setting engine, except that L<sup>A</sup>T<sub>E</sub>X has a few extra switches.



If the `--lua` option is present, L<sup>A</sup>T<sub>E</sub>X will enter an alternative mode of commandline processing in comparison to the standard web2c programs.

In this mode, a small series of actions is taken in order. First, it will parse the commandline as usual, but it will only interpret a small subset of the options immediately: `--safer`, `--nosocket`, `--[no-]shell-escape`, `--enable-write18`, `--disable-write18`, `--shell-restricted`, `--help`, `--version`, and `--credits`.

Now it searches for the requested LUA initialization script. If it cannot be found using the actual name given on the commandline, a second attempt is made by prepending the value of the environment variable `LUATEXDIR`, if that variable is defined in the environment.

Then it checks the various safety switches. You can use those to disable some LUA commands that can easily be abused by a malicious document. At the moment, `--safer` nils the following functions:

#### library functions

os	execute exec setenv rename remove tmpdir
io	popen output tmpfile
lfs	rmdir mkdir chdir lock touch

Furthermore, it disables loading of compiled LUA libraries (support for these was added in 0.46.0), and it makes `io.open()` fail on files that are opened for anything besides reading.

`--nosocket` makes the socket library unavailable, so that LUA cannot use networking.

The switches `--[no-]shell-escape`, `--[enable|disable]-write18`, and `--shell-restricted` have the same effects as in P<sup>D</sup>F<sup>T</sup>E<sub>X</sub>, and additionally make `io.popen()`, `os.execute`, `os.exec` and `os.spawn` adhere to the requested option.

Next the initialization script is loaded and executed. From within the script, the entire commandline is available in the LUA table `arg`, beginning with `arg[0]`, containing the name of the executable. As consequence, the warning about unrecognized option is suppressed.

Commandline processing happens very early on. So early, in fact, that none of T<sub>E</sub>X's initializations have taken place yet. For that reason, the tables that deal with typesetting, like `tex`, `token`, `node` and `pdf`, are off-limits during the execution of the startup file (they are nilled). Special care is taken that `texio.write` and `texio.write_nl` function properly, so that you can at least report your actions to the log file when (and if) it eventually becomes opened (note that T<sub>E</sub>X does not even know its `\jobname` yet at this point). See chapter 4 for more information about the L<sup>A</sup>T<sub>E</sub>X-specific LUA extension tables.

Everything you do in the LUA initialization script will remain visible during the rest of the run, with the exception of the aforementioned `tex`, `token`, `node` and `pdf` tables: those will be initialized to their documented state after the execution of the script. You should not store anything in variables or within tables with these four global names, as they will be overwritten completely.

We recommend you use the startup file only for your own T<sub>E</sub>X-independent initializations (if you need any), to parse the commandline, set values in the `texconfig` table, and register the callbacks you need.

L<sup>A</sup>T<sub>E</sub>X allows some of the commandline options to be overridden by reading values from the `texconfig` table at the end of script execution (see the description of the `texconfig` table later on in this document for more details on which ones exactly).



Unless the `texconfig` table tells L<sup>A</sup>T<sub>E</sub>X not to initialize KPATHSEA at all (set `texconfig.kpse_init` to `false` for that), L<sup>A</sup>T<sub>E</sub>X acts on some more commandline options after the initialization script is finished: in order to initialize the built-in KPATHSEA library properly, L<sup>A</sup>T<sub>E</sub>X needs to know the correct program name to use, and for that it needs to check `--progname`, or `--ini` and `--fmt`, if `--progname` is missing.

## 3.2 LUA changes

**NOTE:** L<sup>A</sup>T<sub>E</sub>X 0.74.0 is the first version with Lua 5.2, and this is used without any patches to the core, which has some side effects. In particular, Lua's `tonumber()` may return values in scientific notation, thereby confusing the T<sub>E</sub>X end of things when it is used as the right-hand side of an assignment to a `\dimen` or `\count`.

**NOTE:** Also in L<sup>A</sup>T<sub>E</sub>X 0.74.0 (this is a change in Lua 5.2), loading dynamic Lua libraries will fail if there are two Lua libraries loaded at the same time (which will typically happen on Win32, because there is one Lua 5.2 inside luatex, and another will likely be linked to the `dll` file of the module itself). We plan to fix that later by switching L<sup>A</sup>T<sub>E</sub>X itself to using the DLL version of Lua 5.2 inside L<sup>A</sup>T<sub>E</sub>X instead of including a static version in the binary.

Starting from version 0.45, L<sup>A</sup>T<sub>E</sub>X is able to use the kpathsea library to find `require()`d modules. For this purpose, `package.searchers[2]` is replaced by a different loader function, that decides at runtime whether to use kpathsea or the built-in core lua function. It uses KPATHSEA when that is already initialized at that point in time, otherwise it reverts to using the normal `package.path` loader.

Initialization of KPATHSEA can happen either implicitly (when L<sup>A</sup>T<sub>E</sub>X starts up and the startup script has not set `texconfig.kpse_init` to false), or explicitly by calling the LUA function `kpse.set_program_name()`.

Starting from version 0.46.0 L<sup>A</sup>T<sub>E</sub>X is also able to use dynamically loadable LUA libraries, unless `--safer` was given as an option on the command line.

For this purpose, `package.searchers[3]` is replaced by a different loader function, that decides at runtime whether to use kpathsea or the built-in core lua function. As in the previous paragraph, it uses KPATHSEA when that is already initialized at that point in time, otherwise it reverts to using the normal `package.cpath` loader.

This functionality required an extension to kpathsea:

There is a new kpathsea file format: `kpse_clua_format` that searches for files with extension `.dll` and `.so`. The `texmf.cnf` setting for this variable is `CLUAINPUTS`, and by default it has this value:

```
CLUAINPUTS=.:$SELFAUTOLOC/lib/{$progname,$engine,}/lua//
```

This path is imperfect (it requires a TDS subtree below the binaries directory), but the architecture has to be in the path somewhere, and the currently simplest way to do that is to search below the binaries directory only.

One level up (a `lib` directory parallel to `bin`) would have been nicer, but that is not doable because T<sub>E</sub>XLIVE uses a `bin/<arch>` structure.



In keeping with the other T<sub>E</sub>X-like programs in T<sub>E</sub>XLive, the two Lua functions `os.execute` and `io.popen` (as well as the two new functions `os.exec` and `os.spawn` that are explained below) take the value of `shell_escape` and/or `shell_escape_commands` in account. Whenever L<sup>A</sup>T<sub>E</sub>X is run with the assumed intention to typeset a document (and by that I mean that it is called as `luatex`, as opposed to `texlua`, and that the commandline option `--luaonly` was not given), it will only run the four functions above if the matching `texmf.cnf` variable(s) or their `texconfig` (see [section 4.15](#)) counterparts allow execution of the requested system command. In ‘script interpreter’ runs of L<sup>A</sup>T<sub>E</sub>X, these settings have no effect, and all four functions function as normal. This change is new in 0.37.0.

The `f:read(*line)` and `f:lines()` functions from the `io` library have been adjusted so that they are line-ending neutral: any of `LF`, `CR` or `CR+LF` are acceptable line endings.

`luafilesystem` has been extended: there are two extra boolean functions (`lfs.isdir(filename)` and `lfs.isfile(filename)`) and one extra string field in its attributes table (`permissions`). There is an additional function (added in 0.51) `lfs.shortname()` which takes a file name and returns its short name on WIN32 platforms. On other platforms, it just returns the given argument. The file name is not tested for existence. Finally, for non-WIN32 platforms only, there is the new function `lfs.readlink()` (added in 0.51) that takes an existing symbolic link as argument and returns its content. It returns an error on WIN32.

The `string` library has an extra function: `string.explode(s[,m])`. This function returns an array containing the string argument `s` split into sub-strings based on the value of the string argument `m`. The second argument is a string that is either empty (this splits the string into characters), a single character (this splits on each occurrence of that character, possibly introducing empty strings), or a single character followed by the plus sign `+` (this special version does not create empty sub-strings). The default value for `m` is `' + '` (multiple spaces).

Note: `m` is not hidden by surrounding braces (as it would be if this function was written in T<sub>E</sub>X macros).

The `string` library also has six extra iterators that return strings piecemeal:

- `string.utfvalues(s)` (returns an integer value in the UNICODE range)
- `string.utfcharacters(s)` (returns a string with a single UTF-8 token in it)
- `string.characters(s)` (a string containing one byte)
- `string.characterpairs(s)` (two strings each containing one byte) will produce an empty second string if the string length was odd.
- `string.bytes(s)` (a single byte value)
- `string.bytepairs(s)` (two byte values) Will produce nil instead of a number as its second return value if the string length was odd.

The `string.characterpairs()` and `string.bytepairs()` are useful especially in the conversion of UTF-16 encoded data into UTF-8.

Starting with L<sup>A</sup>T<sub>E</sub>X 0.74, there is also a two-argument form of `string.dump()`. The second argument is a boolean which, if true, strips the symbols from the dumped data. This matches an extension made in `luajit`.

Note: The `string` library functions `len`, `lower`, `sub` etc. are not UNICODE-aware. For strings in the UTF-8 encoding, i.e., strings containing characters above code point 127, the corresponding functions from the `slnunicode` library can be used, e.g., `unicode.utf8.len`, `unicode.utf8.lower`



etc. The exceptions are `unicode.utf8.find`, that always returns byte positions in a string, and `unicode.utf8.match` and `unicode.utf8.gmatch`. While the latter two functions in general *are* UNICODE-aware, they fall-back to non-UNICODE-aware behavior when using the empty capture `()` (other captures work as expected). For the interpretation of character classes in `unicode.utf8` functions refer to the library sources at <http://luaforge.net/projects/sln>. The `slnunicode` library will be replaced by an internal UNICODE library in a future L<sup>A</sup>T<sub>E</sub>X version.

The `os` library has a few extra functions and variables:

- `os.selfdir` is a variable that holds the directory path of the actual executable. For example: `/opt/luatex/standalone-mkiv-new/tex/texmf-linux-64/bin` (present since 0.27.0).
- `os.exec(commandline)` is a variation on `os.execute`.

The `commandline` can be either a single string or a single table.

If the argument is a table: L<sup>A</sup>T<sub>E</sub>X first checks if there is a value at integer index zero. If there is, this is the command to be executed. Otherwise, it will use the value at integer index one. (if neither are present, nothing at all happens).

The set of consecutive values starting at integer 1 in the table are the arguments that are passed on to the command (the value at index 1 becomes `arg[0]`). The command is searched for in the execution path, so there is normally no need to pass on a fully qualified pathname.

If the argument is a string, then it is automatically converted into a table by splitting on whitespace. In this case, it is impossible for the command and first argument to differ from each other.

In the string argument format, whitespace can be protected by putting (part of) an argument inside single or double quotes. One layer of quotes is interpreted by L<sup>A</sup>T<sub>E</sub>X, and all occurrences of `\`, `'` or `\\` within the quoted text are un-escaped. In the table format, there is no string handling taking place.

This function normally does not return control back to the Lua script: the command will replace the current process. However, it will return the two values `nil` and `'error'` if there was a problem while attempting to execute the command.

On Windows, the current process is actually kept in memory until after the execution of the command has finished. This prevents crashes in situations where T<sub>E</sub>X<sub>L</sub>UA scripts are run inside integrated T<sub>E</sub>X environments.

The original reason for this command is that it cleans out the current process before starting the new one, making it especially useful for use in T<sub>E</sub>X<sub>L</sub>UA.

- `os.spawn(commandline)` is a returning version of `os.exec`, with otherwise identical calling conventions.

If the command ran ok, then the return value is the exit status of the command. Otherwise, it will return the two values `nil` and `'error'`.

- `os.setenv('key', 'value')` This sets a variable in the environment. Passing `nil` instead of a value string will remove the variable.
- `os.env` This is a hash table containing a dump of the variables and values in the process environment at the start of the run. It is writeable, but the actual environment is *not* updated automatically.
- `os.gettimeofday()` Returns the current 'UNIX time', but as a float. This function is not available on the SUNOS platforms, so do not use this function for portable documents.



- `os.times()` Returns the current process times according to the UNIX C library function 'times'. This function is not available on the MS Windows and SUNOS platforms, so do not use this function for portable documents.
- `os.tmpdir()` This will create a directory in the 'current directory' with the name `luatex.XXXXXX` where the X-es are replaced by a unique string. The function also returns this string, so you can `lfs.chdir()` into it, or `nil` if it failed to create the directory. The user is responsible for cleaning up at the end of the run, it does not happen automatically.
- `os.type` This is a string that gives a global indication of the class of operating system. The possible values are currently `windows`, `unix`, and `msdos` (you are unlikely to find this value 'in the wild').
- `os.name` This is a string that gives a more precise indication of the operating system. These possible values are not yet fixed, and for `os.type` values `windows` and `msdos`, the `os.name` values are simply `windows` and `msdos`.  
The list for the type `unix` is more precise: `linux`, `freebsd`, `kfreebsd` (since 0.51), `cygwin` (since 0.53), `openbsd`, `solaris`, `sunos` (pre-solaris), `hpux`, `irix`, `macosx`, `gnu` (hurd), `bsd` (unknown, but BSD-like), `sysv` (unknown, but sysv-like), `generic` (unknown).  
(`os.version` is planned as a future extension)
- `os.uname()` This function returns a table with specific operating system information acquired at runtime. The keys in the returned table are all string valued, and their names are: `sysname`, `machine`, `release`, `version`, and `nodename`.

In stock LUA, many things depend on the current locale. In L<sup>A</sup>T<sub>E</sub>X, we can't do that, because it makes documents unportable. While L<sup>A</sup>T<sub>E</sub>X is running it forces the following locale settings:

```
LC_CTYPE=C
LC_COLLATE=C
LC_NUMERIC=C
```

## 3.3 LUA modules

NOTE: Starting with L<sup>A</sup>T<sub>E</sub>X 0.74, the implied use of the built-in Lua modules in this section is deprecated. If you want to use one of these libraries, please start your source file with a proper `require` line. In the near future, L<sup>A</sup>T<sub>E</sub>X will switch to loading these modules on demand.

Some modules that are normally external to LUA are statically linked in with L<sup>A</sup>T<sub>E</sub>X, because they offer useful functionality:

- `slnunicode`, from the `Selene` libraries, <http://luaforge.net/projects/sln>. (version 1.1)  
This library has been slightly extended so that the `unicode.utf8.*` functions also accept the first 256 values of plane 18. This is the range L<sup>A</sup>T<sub>E</sub>X uses for raw binary output, as explained above.
- `luazip`, from the kepler project, <http://www.keplerproject.org/luazip/>. (version 1.2.1, but patched for compilation with LUA 5.2)
- `luafilesystem`, also from the kepler project, <http://www.keplerproject.org/luafilesystem/>. (version 1.5.0)
- `lpeg`, by Roberto Ierusalimschy, <http://www.inf.puc-rio.br/~roberto/lpeg/lpeg.html>. (version 0.10.2)



Note: `lpeg` is not UNICODE-aware, but interprets strings on a byte-per-byte basis. This mainly means that `lpeg.S` cannot be used with characters above code point 127, since those characters are encoded using two bytes, and thus `lpeg.S` will look for one of those two bytes when matching, not the combination of the two.

The same is true for `lpeg.R`, although the latter will display an error message if used with characters above code point 127: I.e. `lpeg.R('ää')` results in the message `bad argument #1 to 'R' (range must have two characters)`, since to `lpeg`, `ä` is two 'characters' (bytes), so `ää` totals three.

- `lzlib`, by Tiago Dionizio, <http://luaforge.net/projects/lzlib/>. (version 0.2)
- `md5`, by Roberto Ierusalimsky <http://www.inf.puc-rio.br/~roberto/md5/md5-5/md5.html>.
- `luasocket`, by Diego Nehab <http://w3.impa.br/~diego/software/luasocket/> (version 2.0.2).

Note: the `.lua` support modules from `luasocket` are also preloaded inside the executable, there are no external file dependencies.





## 4 L<sup>A</sup>T<sub>E</sub>X LUA Libraries

**NOTE:** Starting with L<sup>A</sup>T<sub>E</sub>X 0.74, the implied use of the built-in Lua modules `epdf`, `fontloader`, `mpbib`, and `pdfscanner` is deprecated. If you want to use these, please start your source file with a proper `require` line. In the near future, L<sup>A</sup>T<sub>E</sub>X will switch to loading these modules on demand.

The interfacing between T<sub>E</sub>X and Lua is facilitated by a set of library modules. The Lua libraries in this chapter are all defined and initialized by the L<sup>A</sup>T<sub>E</sub>X executable. Together, they allow Lua scripts to query and change a number of T<sub>E</sub>X's internal variables, run various internal T<sub>E</sub>X functions, and set up L<sup>A</sup>T<sub>E</sub>X's hooks to execute Lua code.

The following sections are in alphabetical order.

### 4.1 The callback library

This library has functions that register, find and list callbacks.

A quick note on what callbacks are (thanks, Paul!):

Callbacks are entry points to L<sup>A</sup>T<sub>E</sub>X's internal operations, which can be interspersed with additional Lua code, and even replaced altogether. In the first case, T<sub>E</sub>X is simply augmented with new operations (for instance, a manipulation of the nodes resulting from the paragraph builder); in the second case, its hard-coded behavior (for instance, the paragraph builder itself) is ignored and processing relies on user code only.

More precisely, the code to be inserted at a given callback is a function (an anonymous function or the name of a function variable); it will receive the arguments associated with the callback, if any, and must frequently return some other arguments for T<sub>E</sub>X to resume its operations.

The first task is registering a callback:

```
id, error = callback.register (<string> callback_name, <function> func)
id, error = callback.register (<string> callback_name, nil)
id, error = callback.register (<string> callback_name, false)
```

where the `callback_name` is a predefined callback name, see below. The function returns the internal `id` of the callback or `nil`, if the callback could not be registered. In the latter case, `error` contains an error message, otherwise it is `nil`.

L<sup>A</sup>T<sub>E</sub>X internalizes the callback function in such a way that it does not matter if you redefine a function accidentally.

Callback assignments are always global. You can use the special value `nil` instead of a function for clearing the callback.

For some minor speed gain, you can assign the boolean `false` to the non-file related callbacks, doing so will prevent L<sup>A</sup>T<sub>E</sub>X from executing whatever it would execute by default (when no callback function is registered at all). Be warned: this may cause all sorts of grief unless you know *exactly* what you are doing! This functionality is present since version 0.38.



Currently, callbacks are not dumped into the format file.

```
<table> info = callback.list()
```

The keys in the table are the known callback names, the value is a boolean where `true` means that the callback is currently set (active).

```
<function> f = callback.find (callback_name)
```

If the callback is not set, `callback.find` returns `nil`.

## 4.1.1 File discovery callbacks

The behavior documented in this subsection is considered stable in the sense that there will not be backward-incompatible changes any more.

### 4.1.1.1 find\_read\_file and find\_write\_file

Your callback function should have the following conventions:

```
<string> actual_name = function (<number> id_number, <string> asked_name)
```

Arguments:

`id_number`

This number is zero for the log or `\input` files. For T<sub>E</sub>X's `\read` or `\write` the number is incremented by one, so `\read0` becomes 1.

`asked_name`

This is the user-supplied filename, as found by `\input`, `\openin` or `\openout`.

Return value:

`actual_name`

This is the filename used. For the very first file that is read in by T<sub>E</sub>X, you have to make sure you return an `actual_name` that has an extension and that is suitable for use as `jobname`. If you don't, you will have to manually fix the name of the log file and output file after L<sup>A</sup>T<sub>E</sub>X is finished, and an eventual format filename will become mangled. That is because these file names depend on the `jobname`.

You have to return `nil` if the file cannot be found.

### 4.1.1.2 find\_font\_file

Your callback function should have the following conventions:

```
<string> actual_name = function (<string> asked_name)
```

The `asked_name` is an OTF or TFM font metrics file.



Return `nil` if the file cannot be found.

#### 4.1.1.3 `find_output_file`

Your callback function should have the following conventions:

```
<string> actual_name = function (<string> asked_name)
```

The `asked_name` is the PDF or DVI file for writing.

#### 4.1.1.4 `find_format_file`

Your callback function should have the following conventions:

```
<string> actual_name = function (<string> asked_name)
```

The `asked_name` is a format file for reading (the format file for writing is always opened in the current directory).

#### 4.1.1.5 `find_vf_file`

Like `find_font_file`, but for virtual fonts. This applies to both ALEPH's ovf files and traditional Knuthian vf files.

#### 4.1.1.6 `find_map_file`

Like `find_font_file`, but for map files.

#### 4.1.1.7 `find_enc_file`

Like `find_font_file`, but for enc files.

#### 4.1.1.8 `find_sfd_file`

Like `find_font_file`, but for subfont definition files.

#### 4.1.1.9 `find_pk_file`

Like `find_font_file`, but for pk bitmap files. The argument `asked_name` is a bit special in this case. Its form is

```
<base res>dpi/<fontname>.<actual res>pk
```

So you may be asked for `600dpi/manfnt.720pk`. It is up to you to find a 'reasonable' bitmap file to go with that specification.



#### 4.1.1.10 find\_data\_file

Like `find_font_file`, but for embedded files (`\pdfobj file '...'`).

#### 4.1.1.11 find\_opentype\_file

Like `find_font_file`, but for `OPENTYPE` font files.

#### 4.1.1.12 find\_truetype\_file and find\_type1\_file

Your callback function should have the following conventions:

```
<string> actual_name = function (<string> asked_name)
```

The `asked_name` is a font file. This callback is called while `LUATEX` is building its internal list of needed font files, so the actual timing may surprise you. Your return value is later fed back into the matching `read_file` callback.

Strangely enough, `find_type1_file` is also used for `OPENTYPE` (OTF) fonts.

#### 4.1.1.13 find\_image\_file

Your callback function should have the following conventions:

```
<string> actual_name = function (<string> asked_name)
```

The `asked_name` is an image file. Your return value is used to open a file from the harddisk, so make sure you return something that is considered the name of a valid file by your operating system.

### 4.1.2 File reading callbacks

The behavior documented in this subsection is considered stable in the sense that there will not be backward-incompatible changes any more.

#### 4.1.2.1 open\_read\_file

Your callback function should have the following conventions:

```
<table> env = function (<string> file_name)
```

Argument:

`file_name`

The filename returned by a previous `find_read_file` or the return value of `kpse.find_file()` if there was no such callback defined.



Return value:

`env`

This is a table containing at least one required and one optional callback function for this file. The required field is `reader` and the associated function will be called once for each new line to be read, the optional one is `close` that will be called once when L<sup>A</sup>T<sub>E</sub>X is done with the file.

L<sup>A</sup>T<sub>E</sub>X never looks at the rest of the table, so you can use it to store your private per-file data. Both the callback functions will receive the table as their only argument.

#### 4.1.2.1.1 `reader`

L<sup>A</sup>T<sub>E</sub>X will run this function whenever it needs a new input line from the file.

```
function(<table> env)
    return <string> line
end
```

Your function should return either a string or `nil`. The value `nil` signals that the end of file has occurred, and will make T<sub>E</sub>X call the optional `close` function next.

#### 4.1.2.1.2 `close`

L<sup>A</sup>T<sub>E</sub>X will run this optional function when it decides to close the file.

```
function(<table> env)
end
```

Your function should not return any value.

#### 4.1.2.2 General file readers

There is a set of callbacks for the loading of binary data files. These all use the same interface:

```
function(<string> name)
    return <boolean> success, <string> data, <number> data_size
end
```

The `name` will normally be a full path name as it is returned by either one of the file discovery callbacks or the internal version of `kpse.find_file()`.

`success`

Return `false` when a fatal error occurred (e. g. when the file cannot be found, after all).

`data`

The bytes comprising the file.

`data_size`

The length of the `data`, in bytes.



Return an empty string and zero if the file was found but there was a reading problem.

The list of functions is as follows:

<code>read_font_file</code>	ofm or tfm files
<code>read_vf_file</code>	virtual fonts
<code>read_map_file</code>	map files
<code>read_enc_file</code>	encoding files
<code>read_sfd_file</code>	subfont definition files
<code>read_pk_file</code>	pk bitmap files
<code>read_data_file</code>	embedded files ( <code>\pdfobj file ...</code> )
<code>read_truetype_file</code>	TRUEType font files
<code>read_type1_file</code>	TYPE1 font files
<code>read_opentype_file</code>	OPENTYPE font files

## 4.1.3 Data processing callbacks

### 4.1.3.1 `process_input_buffer`

This callback allows you to change the contents of the line input buffer just before L<sup>A</sup>T<sub>E</sub>X actually starts looking at it.

```
function(<string> buffer)
    return <string> adjusted_buffer
end
```

If you return `nil`, L<sup>A</sup>T<sub>E</sub>X will pretend like your callback never happened. You can gain a small amount of processing time from that.

This callback does not replace any internal code.

### 4.1.3.2 `process_output_buffer` (0.43)

This callback allows you to change the contents of the line output buffer just before L<sup>A</sup>T<sub>E</sub>X actually starts writing it to a file as the result of a `\write` command. It is only called for output to an actual file (that is, excluding the log, the terminal, and `\write18` calls).

```
function(<string> buffer)
    return <string> adjusted_buffer
end
```

If you return `nil`, L<sup>A</sup>T<sub>E</sub>X will pretend like your callback never happened. You can gain a small amount of processing time from that.

This callback does not replace any internal code.



### 4.1.3.3 process\_jobname (0.71)

This callback allows you to change the jobname given by `\jobname` in T<sub>E</sub>X and `tex.jobname` in Lua. It does not affect the internal job name or the name of the output or log files.

```
function(<string> jobname)
    return <string> adjusted_jobname
end
```

The only argument is the actual job name; you should not use `tex.jobname` inside this function or infinite recursion may occur. If you return `nil`, L<sup>A</sup>T<sub>E</sub>X will pretend your callback never happened.

This callback does not replace any internal code.

### 4.1.3.4 token\_filter

This callback allows you to replace the way L<sup>A</sup>T<sub>E</sub>X fetches lexical tokens.

```
function()
    return <table> token
end
```

The calling convention for this callback is a bit more complicated than for most other callbacks. The function should either return a LUA table representing a valid to-be-processed token or tokenlist, or something else like `nil` or an empty table.

If your LUA function does not return a table representing a valid token, it will be immediately called again, until it eventually does return a useful token or tokenlist (or until you reset the callback value to `nil`). See the description of `token` for some handy functions to be used in conjunction with this callback.

If your function returns a single usable token, then that token will be processed by L<sup>A</sup>T<sub>E</sub>X immediately. If the function returns a token list (a table consisting of a list of consecutive token tables), then that list will be pushed to the input stack at a completely new token list level, with its token type set to ‘inserted’. In either case, the returned token(s) will not be fed back into the callback function.

Setting this callback to `false` has no effect (because otherwise nothing would happen, forever).

## 4.1.4 Node list processing callbacks

The description of nodes and node lists is in [chapter 8](#).

### 4.1.4.1 buildpage\_filter

This callback is called whenever L<sup>A</sup>T<sub>E</sub>X is ready to move stuff to the main vertical list. You can use this callback to do specialized manipulation of the page building stage like imposition or column balancing.

```
function(<string> extrainfo)
```



`end`

The string `extrainfo` gives some additional information about what T<sub>E</sub>X's state is with respect to the 'current page'. The possible values are:

value	explanation
<code>alignment</code>	a (partial) alignment is being added
<code>after_output</code>	an output routine has just finished
<code>box</code>	a typeset box is being added
<code>new_graf</code>	the beginning of a new paragraph
<code>vmode_par</code>	<code>\par</code> was found in vertical mode
<code>hmode_par</code>	<code>\par</code> was found in horizontal mode
<code>insert</code>	an insert is added
<code>penalty</code>	a penalty (in vertical mode)
<code>before_display</code>	immediately before a display starts
<code>after_display</code>	a display is finished
<code>end</code>	LUA <sub>T</sub> E <sub>X</sub> is terminating (it's all over)

This callback does not replace any internal code.

#### 4.1.4.2 `pre_linebreak_filter`

This callback is called just before LUA<sub>T</sub>E<sub>X</sub> starts converting a list of nodes into a stack of `\hboxes`, after the addition of `\parfillskip`.

```
function(<node> head, <string> groupcode)
    return true | false | <node> newhead
end
```

The string called `groupcode` identifies the nodelist's context within T<sub>E</sub>X's processing. The range of possibilities is given in the table below, but not all of those can actually appear in `pre_linebreak_filter`, some are for the `hpack_filter` and `vpack_filter` callbacks that will be explained in the next two paragraphs.

value	explanation
<code>&lt;empty&gt;</code>	main vertical list
<code>hbox</code>	<code>\hbox</code> in horizontal mode
<code>adjusted_hbox</code>	<code>\hbox</code> in vertical mode
<code>vbox</code>	<code>\vbox</code>
<code>vtop</code>	<code>\vtop</code>
<code>align</code>	<code>\halign</code> or <code>\valign</code>
<code>disc</code>	discretionaries
<code>insert</code>	packaging an insert
<code>vcenter</code>	<code>\vcenter</code>
<code>local_box</code>	<code>\localleftbox</code> or <code>\localrightbox</code>
<code>split_off</code>	top of a <code>\vsplit</code>





<code>split_keep</code>	remainder of a <code>\vsplit</code>
<code>align_set</code>	alignment cell
<code>fin_row</code>	alignment row

As for all the callbacks that deal with nodes, the return value can be one of three things:

- boolean `true` signals succesful processing
- `<node>` signals that the ‘head’ node should be replaced by the returned node
- boolean `false` signals that the ‘head’ node list should be ignored and flushed from memory

This callback does not replace any internal code.

#### 4.1.4.3 `linebreak_filter`

This callback replaces L<sup>A</sup>T<sub>E</sub>X’s line breaking algorithm.

```
function(<node> head, <boolean> is_display)
    return <node> newhead
end
```

The returned node is the head of the list that will be added to the main vertical list, the boolean argument is true if this paragraph is interrupted by a following math display.

If you return something that is not a `<node>`, L<sup>A</sup>T<sub>E</sub>X will apply the internal linebreak algorithm on the list that starts at `<head>`. Otherwise, the `<node>` you return is supposed to be the head of a list of nodes that are all allowed in vertical mode, and at least one of those has to represent a hbox. Failure to do so will result in a fatal error.

Setting this callback to `false` is possible, but dangerous, because it is possible you will end up in an unfixable ‘deadcycles loop’.

#### 4.1.4.4 `post_linebreak_filter`

This callback is called just after L<sup>A</sup>T<sub>E</sub>X has converted a list of nodes into a stack of `\hboxes`.

```
function(<node> head, <string> groupcode)
    return true | false | <node> newhead
end
```

This callback does not replace any internal code.

#### 4.1.4.5 `hpack_filter`

This callback is called when T<sub>E</sub>X is ready to start boxing some horizontal mode material. Math items and line boxes are ignored at the moment.

```
function(<node> head, <string> groupcode, <number> size,
```



```

    <string> packtype [, <string> direction])
  return true | false | <node> newhead
end

```

The `packtype` is either `additional` or `exactly`. If `additional`, then the `size` is a `\hbox spread ...` argument. If `exactly`, then the `size` is a `\hbox to ....` In both cases, the number is in scaled points.

The `direction` is either one of the three-letter direction specifier strings, or `nil` (added in 0.45).

This callback does not replace any internal code.

#### 4.1.4.6 `vpack_filter`

This callback is called when T<sub>E</sub>X is ready to start boxing some vertical mode material. Math displays are ignored at the moment.

This function is very similar to the `hpack_filter`. Besides the fact that it is called at different moments, there is an extra variable that matches T<sub>E</sub>X's `\maxdepth` setting.

```

function(<node> head, <string> groupcode, <number> size, <string>
    packtype, <number> maxdepth [, <string> direction])
  return true | false | <node> newhead
end

```

This callback does not replace any internal code.

#### 4.1.4.7 `pre_output_filter`

This callback is called when T<sub>E</sub>X is ready to start boxing the box 255 for `\output`.

```

function(<node> head, <string> groupcode, <number> size, <string> packtype,
    <number> maxdepth [, <string> direction])
  return true | false | <node> newhead
end

```

This callback does not replace any internal code.

#### 4.1.4.8 `hyphenate`

```

function(<node> head, <node> tail)
end

```

No return values. This callback has to insert discretionary nodes in the node list it receives.

Setting this callback to `false` will prevent the internal discretionary insertion pass.



#### 4.1.4.9 ligaturing

```
function(<node> head, <node> tail)
end
```

No return values. This callback has to apply ligaturing to the node list it receives.

You don't have to worry about return values because the `head` node that is passed on to the callback is guaranteed not to be a `glyph_node` (if need be, a temporary node will be prepended), and therefore it cannot be affected by the mutations that take place. After the callback, the internal value of the 'tail of the list' will be recalculated.

The `next` of `head` is guaranteed to be non-nil.

The `next` of `tail` is guaranteed to be nil, and therefore the second callback argument can often be ignored. It is provided for orthogonality, and because it can sometimes be handy when special processing has to take place.

Setting this callback to `false` will prevent the internal ligature creation pass.

#### 4.1.4.10 kerning

```
function(<node> head, <node> tail)
end
```

No return values. This callback has to apply kerning between the nodes in the node list it receives. See [ligaturing](#) for calling conventions.

Setting this callback to `false` will prevent the internal kern insertion pass.

#### 4.1.4.11 mlist\_to\_hlist

This callback replaces L<sup>A</sup>T<sub>E</sub>X's math list to node list conversion algorithm.

```
function(<node> head, <string> display_type, <boolean> need_penalties)
    return <node> newhead
end
```

The returned node is the head of the list that will be added to the vertical or horizontal list, the string argument is either 'text' or 'display' depending on the current math mode, the boolean argument is `true` if penalties have to be inserted in this list, `false` otherwise.

Setting this callback to `false` is bad, it will almost certainly result in an endless loop.

### 4.1.5 Information reporting callbacks

#### 4.1.5.1 pre\_dump (0.61)

```
function()
```



```
end
```

This function is called just before dumping to a format file starts. It does not replace any code and there are neither arguments nor return values.

#### 4.1.5.2 `start_run`

```
function()  
end
```

This callback replaces the code that prints L<sup>A</sup>T<sub>E</sub>X's banner. Note that for successful use, this callback has to be set in the lua initialization script, otherwise it will be seen only after the run has already started.

#### 4.1.5.3 `stop_run`

```
function()  
end
```

This callback replaces the code that prints L<sup>A</sup>T<sub>E</sub>X's statistics and 'output written to' messages.

#### 4.1.5.4 `start_page_number`

```
function()  
end
```

Replaces the code that prints the [ and the page number at the begin of `\shipout`. This callback will also override the printing of box information that normally takes place when `\tracingoutput` is positive.

#### 4.1.5.5 `stop_page_number`

```
function()  
end
```

Replaces the code that prints the ] at the end of `\shipout`.

#### 4.1.5.6 `show_error_hook`

```
function()  
end
```

This callback is run from inside the T<sub>E</sub>X error function, and the idea is to allow you to do some extra reporting on top of what T<sub>E</sub>X already does (none of the normal actions are removed). You may find some of the values in the `status` table useful.



This callback does not replace any internal code.

#### 4.1.5.7 `show_error_message`

```
function()  
end
```

This callback replaces the code that prints the error message. The usual interaction after the message is not affected.

#### 4.1.5.8 `show_lua_error_hook`

```
function()  
end
```

This callback replaces the code that prints the extra lua error message.

#### 4.1.5.9 `start_file`

```
function(category,filename)  
end
```

This callback replaces the code that prints L<sup>A</sup>T<sub>E</sub>X's when a file is opened like (`filename` for regular files. The category is a number:

- 1 a normal data file, like a T<sub>E</sub>X source
- 2 a font map coupling font names to resources
- 3 an image file (`png`, `pdf`, etc)
- 4 an embedded font subset
- 5 a fully embedded font

#### 4.1.5.10 `stop_file`

```
function(category)  
end
```

This callback replaces the code that prints L<sup>A</sup>T<sub>E</sub>X's when a file is closed like the `)` for regular files.

### 4.1.6 PDF-related callbacks

#### 4.1.6.1 `finish_pdffile`

```
function()
```



`end`

This callback is called when all document pages are already written to the PDF file and L<sup>A</sup>T<sub>E</sub>X is about to finalize the output document structure. Its intended use is final update of PDF dictionaries such as `/Catalog` or `/Info`. The callback does not replace any code. There are neither arguments nor return values.

#### 4.1.6.2 `finish_pdfpage`

```
function(shippingout)
end
```

This callback is called after the pdf page stream has been assembled and before the page object gets finalized. This callback is available in L<sup>A</sup>T<sub>E</sub>X 0.78.4 and later.

### 4.1.7 Font-related callbacks

#### 4.1.7.1 `define_font`

```
function(<string> name, <number> size, <number> id)
    return <table> font | <number> id
end
```

The string `name` is the filename part of the font specification, as given by the user.

The number `size` is a bit special:

- if it is positive, it specifies an ‘at size’ in scaled points.
- if it is negative, its absolute value represents a ‘scaled’ setting relative to the designsize of the font.

The `id` is the internal number assigned to the font.

The internal structure of the `font` table that is to be returned is explained in **chapter 7**. That table is saved internally, so you can put extra fields in the table for your later Lua code to use. In alternative, `retval` can be a previously defined fontid. This is useful if a previous definition can be reused instead of creating a whole new font structure.

Setting this callback to `false` is pointless as it will prevent font loading completely but will nevertheless generate errors.

## 4.2 The `epdf` library

The `epdf` library provides Lua bindings to many PDF access functions that are defined by the poppler pdf viewer library (written in C++ by Kristian Høgsberg, based on xpdf by Derek Noonburg). Within L<sup>A</sup>T<sub>E</sub>X (and P<sup>D</sup>F<sub>T</sub>E<sub>X</sub>), xpdf functionality is being used since long time to embed PDF files. The `epdf`



library shall allow to scrutinize an external PDF file. It gives access to its document structure, e.g., catalog, cross-reference table, individual pages, objects, annotations, info, and metadata. The L<sup>A</sup>T<sub>E</sub>X team is evaluating the possibility of reducing the binding to a basic low level PDF primitives and delegate the complete set of functions to an external shared object module.

The `epdf` library is still in alpha state: PDF access is currently read-only (it's not yet possible to alter a PDF file or to assemble it from scratch), and many function bindings are still missing.

For a start, a PDF file is opened by `epdf.open()` with file name, e.g.:

```
doc = epdf.open("foo.pdf")
```

This normally returns a `PDFDoc` userdata variable; but if the file could not be opened successfully, instead of a fatal error just the value `nil` is returned.

All Lua functions in the `epdf` library are named after the poppler functions listed in the poppler header files for the various classes, e.g., files `PDFDoc.h`, `Dict.h`, and `Array.h`. These files can be found in the poppler subdirectory within the L<sup>A</sup>T<sub>E</sub>X sources. Which functions are already implemented in the `epdf` library can be found in the L<sup>A</sup>T<sub>E</sub>X source file `lepdflib.cc`. For using the `epdf` library, knowledge of the PDF file architecture is indispensable.

There are many different userdata types defined by the `epdf` library, currently these are `AnnotBorderStyle`, `AnnotBorder`, `Annots`, `Annot`, `Array`, `Attribute`, `Catalog`, `Dict`, `EmbFile`, `GString`, `LinkDest`, `Links`, `Link`, `ObjectStream`, `Object`, `PDFDoc`, `PDFRectangle`, `Page`, `Ref`, `Stream`, `StructElement`, `StructTreeRoot` `TextSpan`, `XRefEntry` and `XRef`

All these userdata names and the Lua access functions closely resemble the classes naming from the poppler header files, including the choice of mixed upper and lower case letters. The Lua function calls use object-oriented syntax, e.g., the following calls return the `Page` object for page 1:

```
pageref = doc:getCatalog():getPageRef(1)
pageobj = doc:getXRef():fetch(pageref.num, pageref.gen)
```

But writing such chained calls is risky, as an intermediate function may return `nil` on error; therefore between function calls there should be Lua type checks (e.g., against `nil`) done. If a non-object item is requested (e.g., a `Dict` item by calling `page:getPieceInfo()`, cf. `Page.h`) but not available, the Lua functions return `nil` (without error). If a function should return an `Object`, but it's not existing, a `Null` object is returned instead (also without error; this is in-line with poppler behavior).

All library objects have a `__gc` metamethod for garbage collection. The `__tostring` metamethod gives the type name for each object.

All object constructors:

```
<PDFDoc>      = epdf.open(<string> PDF filename)
<Annot>       = epdf.Annot(<XRef>, <Dict>, <Catalog>, <Ref>)
<Annots>      = epdf.Annots(<XRef>, <Catalog>, <Object>)
<Array>       = epdf.Array(<XRef>)
<Attribute>   = epdf.Attribute(<Type>,<Object>)| epdf.Attribute(<string>,
<int>, <Object>)
```



```

<Dict>          = epdf.Dict(<XRef>)
<Object>        = epdf.Object()
<PDFRectangle> = epdf.PDFRectangle()

```

The functions `StructElement_Type`, `Attribute_Type` and `AttributeOwner_Type` return a hash table `{<string>, <integer>}`.

`Annot` methods:

```

<boolean>      = <Annot>:isOk()
<Object>       = <Annot>:getAppearance()
<AnnotBorder>  = <Annot>:getBorder()
<boolean>      = <Annot>:match(<Ref>)

```

`AnnotBorderStyle` methods:

```

<number> = <AnnotBorderStyle>:getWidth()

```

`Annots` methods:

```

<integer> = <Annots>:getNumAnnots()
<Annot>   = <Annots>:getAnnot(<integer>)

```

`Array` methods:

```

          <Array>:incRef()
          <Array>:decRef()
<integer> = <Array>:getLength()
          <Array>:add(<Object>)
<Object>  = <Array>:get(<integer>)
<Object>  = <Array>:getNF(<integer>)
<string>  = <Array>:getString(<integer>)

```

`Attribute` methods:

```

<boolean> = <Attribute>:isOk()
<integer> = <Attribute>:getType()
<integer> = <Attribute>:getOwner()
<string>  = <Attribute>:getTypeName()
<string>  = <Attribute>:getOwnerName()
<Object>  = <Attribute>:getValue()
<Object>  = <Attribute>:getDefaultValue()
<string>  = <Attribute>:getName()
<integer> = <Attribute>:getRevision()
          <Attribute>:setRevision(<unsigned integer>)
<boolean> = <Attribute>:isHidden()
          <Attribute>:setHidden(<boolean>)
<string>  = <Attribute>:getFormattedValue()

```





```
<string> = <Attribute>:setFormattedValue(<string>)
```

Catalog methods:

```
<boolean> = <Catalog>:isOK()
<integer> = <Catalog>:getNumPages()
<Page>     = <Catalog>:getPage(<integer>)
<Ref>      = <Catalog>:getPageRef(<integer>)
<string>   = <Catalog>:getBaseURI()
<string>   = <Catalog>:readMetadata()
<Object>   = <Catalog>:getStructTreeRoot()
<integer> = <Catalog>:findPage(<integer> object number, <integer> object
generation)
<LinkDest> = <Catalog>:findDest(<string> name)
<Object>   = <Catalog>:getDests()
<integer> = <Catalog>:numEmbeddedFiles()
<EmbFile> = <Catalog>:embeddedFile(<integer>)
<integer> = <Catalog>:numJS()
<string> = <Catalog>:getJS(<integer>)
<Object> = <Catalog>:getOutline()
<Object> = <Catalog>:getAcroForm()
```

EmbFile methods:

```
<string> = <EmbFile>:name()
<string> = <EmbFile>:description()
<integer> = <EmbFile>:size()
<string> = <EmbFile>:modDate()
<string> = <EmbFile>:createDate()
<string> = <EmbFile>:checksum()
<string> = <EmbFile>:mimeType()
<Object> = <EmbFile>:streamObject()
<boolean> = <EmbFile>:isOk()
```

Dict methods:

```
<Dict>:incRef()
<Dict>:decRef()
<integer> = <Dict>:getLength()
<Dict>:add(<string>, <Object>)
<Dict>:set(<string>, <Object>)
<Dict>:remove(<string>)
<boolean> = <Dict>:is(<string>)
<Object> = <Dict>:lookup(<string>)
<Object> = <Dict>:lookupNF(<string>)
<integer> = <Dict>:lookupInt(<string>, <string>)
```



```

<string> = <Dict>:getKey(<integer>)
<Object> = <Dict>:getVal(<integer>)
<Object> = <Dict>:getValNF(<integer>)
<boolean> = <Dict>:hasKey(<string>)

```

Link methods:

```

<boolean> = <Link>:isOK()
<boolean> = <Link>:inRect(<number>, <number>)

```

LinkDest methods:

```

<boolean> = <LinkDest>:isOK()
<integer> = <LinkDest>:getKind()
<string> = <LinkDest>:getKindName()
<boolean> = <LinkDest>:isPageRef()
<integer> = <LinkDest>:getPageNum()
<Ref> = <LinkDest>:getPageRef()
<number> = <LinkDest>:getLeft()
<number> = <LinkDest>:getBottom()
<number> = <LinkDest>:getRight()
<number> = <LinkDest>:getTop()
<number> = <LinkDest>:getZoom()
<boolean> = <LinkDest>:getChangeLeft()
<boolean> = <LinkDest>:getChangeTop()
<boolean> = <LinkDest>:getChangeZoom()

```

Links methods:

```

<integer> = <Links>:getNumLinks()
<Link> = <Links>:getLink(<integer>)

```

Object methods:

```

    <Object>:initBool(<boolean>)
    <Object>:initInt(<integer>)
    <Object>:initReal(<number>)
    <Object>:initString(<string>)
    <Object>:initName(<string>)
    <Object>:initNull()
    <Object>:initArray(<XRef>)
    <Object>:initDict(<XRef>)
    <Object>:initStream(<Stream>)
    <Object>:initRef(<integer> object number, <integer> object gen-
eration)
    <Object>:initCmd(<string>)
    <Object>:initError()

```



```

        <Object>:initEOF()
<Object>  = <Object>:fetch(<XRef>)
<integer> = <Object>:getType()
<string>  = <Object>:getTypeName()
<boolean> = <Object>:isBool()
<boolean> = <Object>:isInt()
<boolean> = <Object>:isReal()
<boolean> = <Object>:isNum()
<boolean> = <Object>:isString()
<boolean> = <Object>:isName()
<boolean> = <Object>:isNull()
<boolean> = <Object>:isArray()
<boolean> = <Object>:isDict()
<boolean> = <Object>:isStream()
<boolean> = <Object>:isRef()
<boolean> = <Object>:isCmd()
<boolean> = <Object>:isError()
<boolean> = <Object>:isEOF()
<boolean> = <Object>:isNone()
<boolean> = <Object>:getBool()
<integer> = <Object>:getInt()
<number>  = <Object>:getReal()
<number>  = <Object>:getNum()
<string>  = <Object>:getString()
<string>  = <Object>:getName()
<Array>   = <Object>:getArray()
<Dict>    = <Object>:getDict()
<Stream>  = <Object>:getStream()
<Ref>     = <Object>:getRef()
<integer> = <Object>:getRefNum()
<integer> = <Object>:getRefGen()
<string>  = <Object>:getCmd()
<integer> = <Object>:arrayGetLength()
           = <Object>:arrayAdd(<Object>)
<Object>  = <Object>:arrayGet(<integer>)
<Object>  = <Object>:arrayGetNF(<integer>)
<integer> = <Object>:dictGetLength(<integer>)
           = <Object>:dictAdd(<string>, <Object>)
           = <Object>:dictSet(<string>, <Object>)
<Object>  = <Object>:dictLookup(<string>)
<Object>  = <Object>:dictLookupNF(<string>)
<string>  = <Object>:dictgetKey(<integer>)
<Object>  = <Object>:dictgetVal(<integer>)
<Object>  = <Object>:dictgetValNF(<integer>)

```



```

<boolean> = <Object>:streamIs(<string>)
           = <Object>:streamReset()
<integer> = <Object>:streamGetChar()
<integer> = <Object>:streamLookChar()
<integer> = <Object>:streamGetPos()
           = <Object>:streamSetPos(<integer>)
<Dict>    = <Object>:streamGetDict()

```

Page methods:

```

<boolean>    = <Page>:isOk()
<integer>    = <Page>:getNum()
<PDFRectangle> = <Page>:getMediaBox()
<PDFRectangle> = <Page>:getCropBox()
<boolean>    = <Page>:isCropped()
<number>     = <Page>:getMediaWidth()
<number>     = <Page>:getMediaHeight()
<number>     = <Page>:getCropWidth()
<number>     = <Page>:getCropHeight()
<PDFRectangle> = <Page>:getBleedBox()
<PDFRectangle> = <Page>:getTrimBox()
<PDFRectangle> = <Page>:getArtBox()
<integer>     = <Page>:getRotate()
<string>     = <Page>:getLastModified()
<Dict>       = <Page>:getBoxColorInfo()
<Dict>       = <Page>:getGroup()
<Stream>     = <Page>:getMetadata()
<Dict>       = <Page>:getPieceInfo()
<Dict>       = <Page>:getSeparationInfo()
<Dict>       = <Page>:getResourceDict()
<Object>     = <Page>:getAnnots()
<Links>      = <Page>:getLinks(<Catalog>)
<Object>     = <Page>:getContents()

```

PDFDoc methods:

```

<boolean>    = <PDFDoc>:isOk()
<integer>    = <PDFDoc>:getErrorCode()
<string>     = <PDFDoc>:getErrorCodeName()
<string>     = <PDFDoc>:getFileName()
<XRef>      = <PDFDoc>:getXRef()
<Catalog>   = <PDFDoc>:getCatalog()
<number>     = <PDFDoc>:getPageMediaWidth()
<number>     = <PDFDoc>:getPageMediaHeight()
<number>     = <PDFDoc>:getPageCropWidth()

```



```

<number>    = <PDFDoc>:getPageCropHeight()
<integer>    = <PDFDoc>:getNumPages()
<string>     = <PDFDoc>:readMetadata()
<Object>     = <PDFDoc>:getStructTreeRoot()
<integer>    = <PDFDoc>:findPage(<integer> object number, <integer> object
generation)
<Links>      = <PDFDoc>:getLinks(<integer>)
<LinkDest>   = <PDFDoc>:findDest(<string>)
<boolean>    = <PDFDoc>:isEncrypted()
<boolean>    = <PDFDoc>:okToPrint()
<boolean>    = <PDFDoc>:okToChange()
<boolean>    = <PDFDoc>:okToCopy()
<boolean>    = <PDFDoc>:okToAddNotes()
<boolean>    = <PDFDoc>:isLinearized()
<Object>     = <PDFDoc>:getDocInfo()
<Object>     = <PDFDoc>:getDocInfoNF()
<integer>    = <PDFDoc>:getPDFMajorVersion()
<integer>    = <PDFDoc>:getPDFMinorVersion()

```

PDFRectangle methods:

```

<boolean>    = <PDFRectangle>:isValid()

```

Stream methods:

```

<integer>    = <Stream>:getKind()
<string>     = <Stream>:getKindName()
              = <Stream>:reset()
              = <Stream>:close()
<integer>    = <Stream>:getChar()
<integer>    = <Stream>:lookChar()
<integer>    = <Stream>:getRawChar()
<integer>    = <Stream>:getUnfilteredChar()
              = <Stream>:unfilteredReset()
<integer>    = <Stream>:getPos()
<boolean>    = <Stream>:isBinary()
<Stream>     = <Stream>:getUndecodedStream()
<Dict>      = <Stream>:getDict()

```

StructElement methods:

```

<string>     = <StructElement>:getTypeName()
<integer>    = <StructElement>:getType()
<boolean>    = <StructElement>:isOk()
<boolean>    = <StructElement>:isBlock()
<boolean>    = <StructElement>:isInline()

```



```

<boolean>      = <StructElement>:isGrouping()
<boolean>      = <StructElement>:isContent()
<boolean>      = <StructElement>:isObjectRef()
<integer>      = <StructElement>:getMCID()
<Ref>          = <StructElement>:getObjectRef()
<Ref>          = <StructElement>:getParentRef()
<boolean>      = <StructElement>:hasPageRef()
<Ref>          = <StructElement>:getPageRef()
<StructTreeRoot> = <StructElement>:getStructTreeRoot()
<string>       = <StructElement>:getID()
<string>       = <StructElement>:getLanguage()
<integer>      = <StructElement>:getRevision()
               = <StructElement>:setRevision(<unsigned integer>)
<string>       = <StructElement>:getTitle()
<string>       = <StructElement>:getExpandedAbbr()
<integer>      = <StructElement>:getNumChildren()
<StructElement> = <StructElement>:getChild()
               = <StructElement>:appendChild<StructElement>()
<integer>      = <StructElement>:getNumAttributes()
<Attribute>    = <StructElement>:getAttribute(<integer>)
<string>       = <StructElement>:appendAttribute(<Attribute>)
<Attribute>    = <StructElement>:findAttribute(<Attribute::Type>,boolean,At-
tribute::Owner)
<string>       = <StructElement>:getAltText()
<string>       = <StructElement>:getActualText()
<string>       = <StructElement>:getText(<boolean>)
<table>        = <StructElement>:getTextSpans()

```

StructTreeRoot methods:

```

<StructElement> = <StructTreeRoot>:findParentElement
<PDFDoc>        = <StructTreeRoot>:getDoc
<Dict>          = <StructTreeRoot>:getRoleMap
<Dict>          = <StructTreeRoot>:getClassMap
<integer>       = <StructTreeRoot>:getNumChildren
<StructElement> = <StructTreeRoot>:getChild
               = <StructTreeRoot>:appendChild
<StructElement> = <StructTreeRoot>:findParentElement

```

TextSpan han only one method:

```

<string> = <TextSpan>:getText()

```

XRef methods:

```

<boolean> = <XRef>:isOk()

```



```

<integer> = <XRef>:getErrorCode()
<boolean> = <XRef>:isEncrypted()
<boolean> = <XRef>:okToPrint()
<boolean> = <XRef>:okToPrintHighRes()
<boolean> = <XRef>:okToChange()
<boolean> = <XRef>:okToCopy()
<boolean> = <XRef>:okToAddNotes()
<boolean> = <XRef>:okToFillForm()
<boolean> = <XRef>:okToAccessibility()
<boolean> = <XRef>:okToAssemble()
<Object> = <XRef>:getCatalog()
<Object> = <XRef>:fetch(<integer> object number, <integer> object genera-
tion)
<Object> = <XRef>:getDocInfo()
<Object> = <XRef>:getDocInfoNF()
<integer> = <XRef>:getNumObjects()
<integer> = <XRef>:getRootNum()
<integer> = <XRef>:getRootGen()
<integer> = <XRef>:getSize()
<Object> = <XRef>:getTrailerDict()

```

## 4.3 The font library

The font library provides the interface into the internals of the font system, and also it contains helper functions to load traditional T<sub>E</sub>X font metrics formats. Other font loading functionality is provided by the `fontloader` library that will be discussed in the next section.

### 4.3.1 Loading a TFM file

The behavior documented in this subsection is considered stable in the sense that there will not be backward-incompatible changes any more.

```
<table> fnt = font.read_tfm(<string> name, <number> s)
```

The number is a bit special:

- if it is positive, it specifies an ‘at size’ in scaled points.
- if it is negative, its absolute value represents a ‘scaled’ setting relative to the designsize of the font.

The internal structure of the metrics font table that is returned is explained in **chapter 7**.

### 4.3.2 Loading a VF file

The behavior documented in this subsection is considered stable in the sense that there will not be backward-incompatible changes any more.



```
<table> vf_fnt = font.read_vf(<string> name, <number> s)
```

The meaning of the number `s` and the format of the returned table are similar to the ones in the `read_tfm()` function.

### 4.3.3 The fonts array

The whole table of T<sub>E</sub>X fonts is accessible from LUA using a virtual array.

```
font.fonts[n] = { ... }  
<table> f = font.fonts[n]
```

See **chapter 7** for the structure of the tables. Because this is a virtual array, you cannot call `pairs` on it, but see below for the `font.each` iterator.

The two metatable functions implementing the virtual array are:

```
<table> f = font.getfont(<number> n)  
font.setfont(<number> n, <table> f)
```

Note that at the moment, each access to the `font.fonts` or call to `font.getfont` creates a lua table for the whole font. This process can be quite slow. In a later version of L<sup>A</sup>T<sub>E</sub>X, this interface will change (it will start using userdata objects instead of actual tables).

Also note the following: assignments can only be made to fonts that have already been defined in T<sub>E</sub>X, but have not been accessed *at all* since that definition. This limits the usability of the write access to `font.fonts` quite a lot, a less stringent ruleset will likely be implemented later.

### 4.3.4 Checking a font's status

You can test for the status of a font by calling this function:

```
<boolean> f = font.frozen(<number> n)
```

The return value is one of `true` (unassignable), `false` (can be changed) or `nil` (not a valid font at all).

### 4.3.5 Defining a font directly

You can define your own font into `font.fonts` by calling this function:

```
<number> i = font.define(<table> f)
```

The return value is the internal id number of the defined font (the index into `font.fonts`). If the font creation fails, an error is raised. The table is a font structure, as explained in **chapter 7**.





### 4.3.6 Projected next font id

```
<number> i = font.nextid()
```

This returns the font id number that would be returned by a `font.define` call if it was executed at this spot in the code flow. This is useful for virtual fonts that need to reference themselves.

### 4.3.7 Font id (0.47)

```
<number> i = font.id(<string> csname)
```

This returns the font id associated with `csname` string, or `-1` if `csname` is not defined; new in 0.47.

### 4.3.8 Currently active font

```
<number> i = font.current()  
font.current(<number> i)
```

This gets or sets the currently used font number.

### 4.3.9 Maximum font id

```
<number> i = font.max()
```

This is the largest used index in `font.fonts`.

### 4.3.10 Iterating over all fonts

```
for i,v in font.each() do  
    ...  
end
```

This is an iterator over each of the defined T<sub>E</sub>X fonts. The first returned value is the index in `font.fonts`, the second the font itself, as a Lua table. The indices are listed incrementally, but they do not always form an array of consecutive numbers: in some cases there can be holes in the sequence.

## 4.4 The fontloader library (0.36)

### 4.4.1 Getting quick information on a font

```
<table> info = fontloader.info(<string> filename)
```



This function returns either `nil`, or a `table`, or an array of small tables (in the case of a TrueType collection). The returned table(s) will contain some fairly interesting information items from the font(s) defined by the file:

key	type	explanation
fontname	string	the POSTSCRIPT name of the font
fullname	string	the formal name of the font
familyname	string	the family name this font belongs to
weight	string	a string indicating the color value of the font
version	string	the internal font version
italicangle	float	the slant angle
units_per_em	number	(since 0.78.2) 1000 for POSTSCRIPT-based fonts, usually 2048 for TRUETYPE
pfminfo	table	(since 0.78.2) (see <a href="#">section 4.4.5.1.6</a> )

Getting information through this function is (sometimes much) more efficient than loading the font properly, and is therefore handy when you want to create a dictionary of available fonts based on a directory contents.

## 4.4.2 Loading an OPENTYPE or TRUETYPE file

If you want to use an OPENTYPE font, you have to get the metric information from somewhere. Using the `fontloader` library, the simplest way to get that information is thus:

```
function load_font (filename)
    local metrics = nil
    local font = fontloader.open(filename)
    if font then
        metrics = fontloader.to_table(font)
        fontloader.close(font)
    end
    return metrics
end

myfont = load_font('/opt/tex/texmf/fonts/data/arial.ttf')
```

The main function call is

```
<userdata> f, <table> w = fontloader.open(<string> filename)
<userdata> f, <table> w = fontloader.open(<string> filename, <string> font-
name)
```

The first return value is a userdata representation of the font. The second return value is a table containing any warnings and errors reported by fontloader while opening the font. In normal typesetting, you would probably ignore the second argument, but it can be useful for debugging purposes.



For TRUEType collections (when filename ends in 'ttc') and DFont collections, you have to use a second string argument to specify which font you want from the collection. Use the `fontname` strings that are returned by `fontloader.info` for that.

To turn the font into a table, `fontloader.to_table` is used on the font returned by `fontloader.open`.

```
<table> f = fontloader.to_table(<userdata> font)
```

This table cannot be used directly by L<sup>A</sup>T<sub>E</sub>X and should be turned into another one as described in [chapter 7](#). Do not forget to store the `fontname` value in the `psname` field of the metrics table to be returned to L<sup>A</sup>T<sub>E</sub>X, otherwise the font inclusion backend will not be able to find the correct font in the collection.

See [section 4.4.5](#) for details on the userdata object returned by `fontloader.open()` and the layout of the `metrics` table returned by `fontloader.to_table()`.

The font file is parsed and partially interpreted by the font loading routines from FONTFORGE. The file format can be OPENType, TRUEType, TRUEType Collection, CFF, or TYPE1.

There are a few advantages to this approach compared to reading the actual font file ourselves:

- The font is automatically re-encoded, so that the `metrics` table for TRUEType and OPENType fonts is using UNICODE for the character indices.
- Many features are pre-processed into a format that is easier to handle than just the bare tables would be.
- POSTSCRIPT-based OPENType fonts do not store the character height and depth in the font file, so the character boundingbox has to be calculated in some way.
- In the future, it may be interesting to allow LUA scripts access to the font program itself, perhaps even creating or changing the font.

A loaded font is discarded with:

```
fontloader.close(<userdata> font)
```

### 4.4.3 Applying a 'feature file'

You can apply a 'feature file' to a loaded font:

```
<table> errors = fontloader.apply_featurefile(<userdata> font, <string>  
filename)
```

A 'feature file' is a textual representation of the features in an OPENType font. See

[http://www.adobe.com/devnet/opentype/afdko/topic\\_feature\\_file\\_syntax.html](http://www.adobe.com/devnet/opentype/afdko/topic_feature_file_syntax.html)

and

<http://fontforge.sourceforge.net/featurefile.html>

for a more detailed description of feature files.



If the function fails, the return value is a table containing any errors reported by fontloader while applying the feature file. On success, `nil` is returned. (the return value is new in 0.65)

#### 4.4.4 Applying an ‘AFM file’

You can apply an ‘AFM file’ to a loaded font:

```
<table> errors = fontloader.apply_afmfile(<userdata> font, <string> file-  
name)
```

An AFM file is a textual representation of (some of) the meta information in a TYPE1 font. See [ftp://ftp.math.utah.edu/u/ma/hohn/linux/postscript/5004.AFM\\_Spec.pdf](ftp://ftp.math.utah.edu/u/ma/hohn/linux/postscript/5004.AFM_Spec.pdf) for more information about afm files.

Note: If you `fontloader.open()` a TYPE1 file named `font.pfb`, the library will automatically search for and apply `font.afm` if it exists in the same directory as the file `font.pfb`. In that case, there is no need for an explicit call to `apply_afmfile()`.

If the function fails, the return value is a table containing any errors reported by fontloader while applying the AFM file. On success, `nil` is returned. (the return value is new in 0.65)

#### 4.4.5 Fontloader font tables

As mentioned earlier, the return value of `fontloader.open()` is a userdata object. In L<sup>A</sup>T<sub>E</sub>X versions before 0.63, the only way to have access to the actual metrics was to call `fontloader.to_table()` on this object, returning the table structure that is explained in the following subsections.

However, it turns out that the result from `fontloader.to_table()` sometimes needs very large amounts of memory (depending on the font’s complexity and size) so starting with L<sup>A</sup>T<sub>E</sub>X 0.63, it is possible to access the userdata object directly.

In the L<sup>A</sup>T<sub>E</sub>X 0.63.0, the following is implemented:

- all top-level keys that would be returned by `to_table()` can also be accessed directly.
- the top-level key ‘glyphs’ returns a *virtual* array that allows indices from `f.glyphmin` to `(f.glyph-max)`.
- the items in that virtual array (the actual glyphs) are themselves also userdata objects, and each has accessors for all of the keys explained in the section ‘Glyph items’ below.
- the top-level key ‘subfonts’ returns an *actual* array of userdata objects, one for each of the subfonts (or `nil`, if there are no subfonts).

A short example may be helpful. This code generates a printout of all the glyph names in the font `PunkNova.kern.otf`:

```
local f = fontloader.open('PunkNova.kern.otf')  
print (f.fontname)  
local i = 0
```



```

if f.glyphcnt > 0 then
  for i=f.glyphmin,f.glyphmax do
    local g = f.glyphs[i]
    if g then
      print(g.name)
    end
    i = i + 1
  end
end
fontloader.close(f)

```

In this case, the L<sup>A</sup>T<sub>E</sub>X memory requirement stays below 100MB on the test computer, while the internal stucture generated by `to_table()` needs more than 2GB of memory (the font itself is 6.9MB in disk size).

In L<sup>A</sup>T<sub>E</sub>X 0.63 only the top-level font, the subfont table entries, and the glyphs are virtual objects, everything else still produces normal lua values and tables. In future versions, more return values may be replaced by userdata objects (as much as needed to keep the memory requirements in check).

If you want to know the valid fields in a font or glyph structure, call the `fields` function on an object of a particular type (either glyph or font for now, more will be implemented later):

```

<table> fields = fontloader.fields(<userdata> font)
<table> fields = fontloader.fields(<userdata> font_glyph)

```

For instance:

```

local fields = fontloader.fields(f)
local fields = fontloader.fields(f.glyphs[0])

```

### 4.4.5.1 Table types

#### 4.4.5.1.1 Top-level

The top-level keys in the returned table are (the explanations in this part of the documentation are not yet finished):

key	type	explanation
table_version	number	indicates the metrics version (currently 0.3)
fontname	string	POSTSCRIPT font name
fullname	string	official (human-oriented) font name
famillyname	string	family name
weight	string	weight indicator
copyright	string	copyright information
filename	string	the file name
version	string	font version



<code>italicangle</code>	float	slant angle
<code>units_per_em</code>	number	1000 for PostSCRIPT-based fonts, usually 2048 for TRUEType
<code>ascent</code>	number	height of ascender in <code>units_per_em</code>
<code>descent</code>	number	depth of descender in <code>units_per_em</code>
<code>upos</code>	float	
<code>uwidth</code>	float	
<code>uniqueid</code>	number	
<code>glyphs</code>	array	
<code>glyphcnt</code>	number	number of included glyphs
<code>glyphmax</code>	number	maximum used index the glyphs array
<code>glyphmin</code>	number	minimum used index the glyphs array
<code>hasvmetrics</code>	number	
<code>onlybitmaps</code>	number	
<code>serifcheck</code>	number	
<code>isserif</code>	number	
<code>issans</code>	number	
<code>encodingchanged</code>	number	
<code>strokedfont</code>	number	
<code>use_typo_metrics</code>	number	
<code>weight_width_slope_only</code>	number	
<code>head_optimized_for_cleartype</code>	number	
<code>uni_interp</code>	enum	<code>unset</code> , <code>none</code> , <code>adobe</code> , <code>greek</code> , <code>japanese</code> , <code>trad_chinese</code> , <code>simp_chinese</code> , <code>korean</code> , <code>ams</code>
<code>origname</code>	string	the file name, as supplied by the user
<code>map</code>	table	
<code>private</code>	table	
<code>xuid</code>	string	
<code>pfminfo</code>	table	
<code>names</code>	table	
<code>cidinfo</code>	table	
<code>subfonts</code>	array	
<code>comments</code>	string	
<code>fontlog</code>	string	
<code>cvt_names</code>	string	
<code>anchor_classes</code>	table	
<code>ttf_tables</code>	table	
<code>ttf_tab_saved</code>	table	
<code>kerns</code>	table	
<code>vkerns</code>	table	
<code>texdata</code>	table	
<code>lookups</code>	table	
<code>gpos</code>	table	
<code>gsub</code>	table	



mm	table
chosename	string
macstyle	number
fondname	string
fontstyle_id	number
fontstyle_name	table
strokewidth	float
mark_classes	table
creationtime	number
modificationtime	number
os2_version	number
sfd_version	number
math	table
validation_state	table
horiz_base	table
vert_base	table
extrema_bound	number

#### 4.4.5.1.2 Glyph items

The `glyphs` is an array containing the per-character information (quite a few of these are only present if nonzero).

key	type	explanation
name	string	the glyph name
unicode	number	unicode code point, or -1
boundingbox	array	array of four numbers, see note below
width	number	only for horizontal fonts
vwidth	number	only for vertical fonts
tsidebearing	number	only for vertical ttf/otf fonts, and only if nonzero (0.79.0)
lsidebearing	number	only if nonzero and not equal to boundingbox[1]
class	string	one of "none", "base", "ligature", "mark", "component" (if not present, the glyph class is 'automatic')
kerns	array	only for horizontal fonts, if set
vkerns	array	only for vertical fonts, if set
dependents	array	linear array of glyph name strings, only if nonempty
lookups	table	only if nonempty
ligatures	table	only if nonempty
anchors	table	only if set
comment	string	only if set
tex_height	number	only if set
tex_depth	number	only if set
italic_correction	number	only if set
top_accent	number	only if set



<code>is_extended_shape</code>	number	only if this character is part of a math extension list
<code>altuni</code>	table	alternate UNICODE items
<code>vert_variants</code>	table	
<code>horiz_variants</code>	table	
<code>mathkern</code>	table	

On **boundingbox**: The boundingbox information for TRUEType fonts and TRUEType-based OTF fonts is read directly from the font file. POSTSCRIPT-based fonts do not have this information, so the boundingbox of traditional PostSCRIPT fonts is generated by interpreting the actual bezier curves to find the exact boundingbox. This can be a slow process, so starting from L<sup>A</sup>T<sub>E</sub>X 0.45, the boundingboxes of PostSCRIPT-based OTF fonts (and raw CFF fonts) are calculated using an approximation of the glyph shape based on the actual glyph points only, instead of taking the whole curve into account. This means that glyphs that have missing points at extrema will have a too-tight boundingbox, but the processing is so much faster that in our opinion the tradeoff is worth it.

The **kerns** and **vkerns** are linear arrays of small hashes:

key	type	explanation
<code>char</code>	string	
<code>off</code>	number	
<code>lookup</code>	string	

The **lookups** is a hash, based on lookup subtable names, with the value of each key inside that a linear array of small hashes:

key	type	explanation
<code>type</code>	enum	<a href="#">position</a> , <a href="#">pair</a> , <a href="#">substitution</a> , <a href="#">alternate</a> , <a href="#">multiple</a> , <a href="#">ligature</a> , <a href="#">lcaret</a> , <a href="#">kerning</a> , <a href="#">vkerning</a> , <a href="#">anchors</a> , <a href="#">contextpos</a> , <a href="#">contextsub</a> , <a href="#">chainpos</a> , <a href="#">chainsub</a> , <a href="#">reversesub</a> , <a href="#">max</a> , <a href="#">kernback</a> , <a href="#">vkernback</a>
<code>specification</code>	table	extra data

For the first seven values of **type**, there can be additional sub-information, stored in the sub-table **specification**:

value	type	explanation
<code>position</code>	table	a table of the <a href="#">offset_specs</a> type
<code>pair</code>	table	one string: <a href="#">paired</a> , and an array of one or two <a href="#">offset_specs</a> tables: <a href="#">offsets</a>
<code>substitution</code>	table	one string: <a href="#">variant</a>
<code>alternate</code>	table	one string: <a href="#">components</a>
<code>multiple</code>	table	one string: <a href="#">components</a>
<code>ligature</code>	table	two strings: <a href="#">components</a> , <a href="#">char</a>
<code>lcaret</code>	array	linear array of numbers

Tables for **offset\_specs** contain up to four number-valued fields: **x** (a horizontal offset), **y** (a vertical offset), **h** (an advance width correction) and **v** (an advance height correction).

The **ligatures** is a linear array of small hashes:





key	type	explanation
lig	table	uses the same substructure as a single item in the <a href="#">lookups</a> table explained above
char	string	
components	array	linear array of named components
ccnt	number	

The [anchor](#) table is indexed by a string signifying the anchor type, which is one of

key	type	explanation
mark	table	placement mark
basechar	table	mark for attaching combining items to a base char
baselig	table	mark for attaching combining items to a ligature
basemark	table	generic mark for attaching combining items to connect to
centry	table	cursive entry point
cexit	table	cursive exit point

The content of these is a short array of defined anchors, with the entry keys being the anchor names. For all except [baselig](#), the value is a single table with this definition:

key	type	explanation
x	number	x location
y	number	y location
ttf_pt_index	number	truetype point index, only if given

For [baselig](#), the value is a small array of such anchor sets, one for each constituent item of the ligature.

For clarification, an anchor table could for example look like this :

```
['anchor'] = {
  ['basemark'] = {
    ['Anchor-7'] = { ['x']=170, ['y']=1080 }
  },
  ['mark'] = {
    ['Anchor-1'] = { ['x']=160, ['y']=810 },
    ['Anchor-4'] = { ['x']=160, ['y']=800 }
  },
  ['baselig'] = {
    [1] = { ['Anchor-2'] = { ['x']=160, ['y']=650 } },
    [2] = { ['Anchor-2'] = { ['x']=460, ['y']=640 } }
  }
}
```

Note: The [baselig](#) table can be sparse!



### 4.4.5.1.3 map table

The top-level map is a list of encoding mappings. Each of those is a table itself.

key	type	explanation
enccount	number	
encmax	number	
backmax	number	
remap	table	
map	array	non-linear array of mappings
backmap	array	non-linear array of backward mappings
enc	table	

The [remap](#) table is very small:

key	type	explanation
firstenc	number	
lastenc	number	
infont	number	

The [enc](#) table is a bit more verbose:

key	type	explanation
enc_name	string	
char_cnt	number	
char_max	number	
unicode	array	of UNICODE position numbers
psnames	array	of POSTSCRIPT glyph names
builtin	number	
hidden	number	
only_1byte	number	
has_1byte	number	
has_2byte	number	
is_unicodebmp	number	only if nonzero
is_unicodedefull	number	only if nonzero
is_custom	number	only if nonzero
is_original	number	only if nonzero
is_compact	number	only if nonzero
is_japanese	number	only if nonzero
is_korean	number	only if nonzero
is_tradchinese	number	only if nonzero [name?]
is_simplechinese	number	only if nonzero
low_page	number	
high_page	number	
iconv_name	string	
iso_2022_escape	string	



#### 4.4.5.1.4 private table

This is the font's private POSTSCRIPT dictionary, if any. Keys and values are both strings.

#### 4.4.5.1.5 cidinfo table

key	type	explanation
registry	string	
ordering	string	
supplement	number	
version	number	

#### 4.4.5.1.6 pfminfo table

The `pfminfo` table contains most of the OS/2 information:

key	type	explanation
pfmset	number	
winascent_add	number	
windescent_add	number	
hheadascent_add	number	
hheaddescent_add	number	
typoascent_add	number	
typodescent_add	number	
subsuper_set	number	
panose_set	number	
hheadset	number	
vheadset	number	
pfmfamily	number	
weight	number	
width	number	
avgwidth	number	
firstchar	number	
lastchar	number	
fstype	number	
linegap	number	
vlinegap	number	
hhead_ascent	number	
hhead_descent	number	
os2_typoascent	number	
os2_typodescent	number	
os2_typolinegap	number	
os2_winascent	number	
os2_windescent	number	



os2_subxsize	number	
os2_subysize	number	
os2_subxoff	number	
os2_subyoff	number	
os2_supxsize	number	
os2_supysize	number	
os2_supxoff	number	
os2_supyoff	number	
os2_strikeysize	number	
os2_strikeypos	number	
os2_family_class	number	
os2_xheight	number	
os2_capheight	number	
os2_defaultchar	number	
os2_breakchar	number	
os2_vendor	string	
codepages	table	A two-number array of encoded code pages
unicoderanges	table	A four-number array of encoded unicode ranges
panose	table	

The [panose](#) subtable has exactly 10 string keys:

key	type	explanation
familytype	string	Values as in the OPENType font specification: <a href="#">Any</a> , <a href="#">No Fit</a> , <a href="#">Text</a> and <a href="#">Display</a> , <a href="#">Script</a> , <a href="#">Decorative</a> , <a href="#">Pictorial</a>
serifstyle	string	See the OPENType font specification for values
weight	string	id.
proportion	string	id.
contrast	string	id.
strokevariation	string	id.
armstyle	string	id.
letterform	string	id.
midline	string	id.
xheight	string	id.

#### 4.4.5.1.7 names table

Each item has two top-level keys:

key	type	explanation
lang	string	language for this entry
names	table	

The [names](#) keys are the actual TRUETYPE name strings. The possible keys are:



key	explanation
copyright	
family	
subfamily	
uniqueid	
fullname	
version	
postscriptname	
trademark	
manufacturer	
designer	
descriptor	
venderurl	
designerurl	
license	
licenseurl	
idontknow	
preffamilyname	
prefmodifiers	
compatfull	
sampletext	
cidfindfontname	
wwsfamily	
wwssubfamily	

#### 4.4.5.1.8 anchor\_classes table

The anchor\_classes classes:

key	type	explanation
name	string	a descriptive id of this anchor class
lookup	string	
type	string	one of <a href="#">mark</a> , <a href="#">mkmk</a> , <a href="#">curs</a> , <a href="#">mklg</a>

#### 4.4.5.1.9 gpos table

Th gpos table has one array entry for each lookup. (The [gpos\\_](#) prefix is somewhat redundant.)

key	type	explanation
type	string	one of <a href="#">gpos_single</a> , <a href="#">gpos_pair</a> , <a href="#">gpos_cursive</a> , <a href="#">gpos_mark2base</a> , <a href="#">gpos_mark2ligature</a> , <a href="#">gpos_mark2mark</a> , <a href="#">gpos_context</a> , <a href="#">gpos_contextchain</a>
flags	table	
name	string	



features array  
subtables array

The flags table has a true value for each of the lookup flags that is actually set:

key	type	explanation
r2l	boolean	
ignorebaseglyphs	boolean	
ignoreligatures	boolean	
ignorecombiningmarks	boolean	
mark_class	string	(new in 0.44)

The features subtable items of gpos have:

key	type	explanation
tag	string	
scripts	table	

The scripts table within features has:

key	type	explanation
script	string	
langs	array of strings	

The subtables table has:

key	type	explanation
name	string	
suffix	string	(only if used)
anchor_classes	number	(only if used)
vertical_kerning	number	(only if used)
kernclass	table	(only if used)

The kernclass with subtables table has:

key	type	explanation
firsts	array of strings	
seconds	array of strings	
lookup	string or array	associated lookup(s)
offsets	array of numbers	

#### 4.4.5.1.10 gsub table

This has identical layout to the [gpos](#) table, except for the type:

key	type	explanation
type	string	one of <a href="#">gsub_single</a> , <a href="#">gsub_multiple</a> , <a href="#">gsub_alternate</a> , <a href="#">gsub_ligature</a> , <a href="#">gsub_context</a> , <a href="#">gsub_contextchain</a> , <a href="#">gsub_reversecontextchain</a>



#### 4.4.5.1.11 ttf\_tables and ttf\_tab\_saved tables

key	type	explanation
tag	string	
len	number	
maxlen	number	
data	number	

#### 4.4.5.1.12 mm table

key	type	explanation
axes	table	array of axis names
instance_count	number	
positions	table	array of instance positions (#axes * instances )
defweights	table	array of default weights for instances
cdv	string	
ndv	string	
axismaps	table	

The **axismaps**:

key	type	explanation
blends	table	an array of blend points
designs	table	an array of design values
min	number	
def	number	
max	number	

#### 4.4.5.1.13 mark\_classes table (0.44)

The keys in this table are mark class names, and the values are a space-separated string of glyph names in this class.

Note: This table is indeed new in 0.44. The manual said it existed before then, but in practise it was missing due to a bug.

#### 4.4.5.1.14 math table

ScriptPercentScaleDown  
ScriptScriptPercentScaleDown  
DelimitedSubFormulaMinHeight  
DisplayOperatorMinHeight  
MathLeading  
AxisHeight



AccentBaseHeight  
 FlattenedAccentBaseHeight  
 SubscriptShiftDown  
 SubscriptTopMax  
 SubscriptBaselineDropMin  
 SuperscriptShiftUp  
 SuperscriptShiftUpCramped  
 SuperscriptBottomMin  
 SuperscriptBaselineDropMax  
 SubSuperscriptGapMin  
 SuperscriptBottomMaxWithSubscript  
 SpaceAfterScript  
 UpperLimitGapMin  
 UpperLimitBaselineRiseMin  
 LowerLimitGapMin  
 LowerLimitBaselineDropMin  
 StackTopShiftUp  
 StackTopDisplayStyleShiftUp  
 StackBottomShiftDown  
 StackBottomDisplayStyleShiftDown  
 StackGapMin  
 StackDisplayStyleGapMin  
 StretchStackTopShiftUp  
 StretchStackBottomShiftDown  
 StretchStackGapAboveMin  
 StretchStackGapBelowMin  
 FractionNumeratorShiftUp  
 FractionNumeratorDisplayStyleShiftUp  
 FractionDenominatorShiftDown  
 FractionDenominatorDisplayStyleShiftDown  
 FractionNumeratorGapMin  
 FractionNumeratorDisplayStyleGapMin  
 FractionRuleThickness  
 FractionDenominatorGapMin  
 FractionDenominatorDisplayStyleGapMin  
 SkewedFractionHorizontalGap  
 SkewedFractionVerticalGap  
 OverbarVerticalGap  
 OverbarRuleThickness  
 OverbarExtraAscender  
 UnderbarVerticalGap  
 UnderbarRuleThickness  
 UnderbarExtraDescender  
 RadicalVerticalGap





RadicalDisplayStyleVerticalGap	
RadicalRuleThickness	
RadicalExtraAscender	
RadicalKernBeforeDegree	
RadicalKernAfterDegree	
RadicalDegreeBottomRaisePercent	
MinConnectorOverlap	
FractionDelimiterSize	(new in 0.47.0)
FractionDelimiterDisplayStyleSize	(new in 0.47.0)

#### 4.4.5.1.15 validation\_state table

key	explanation
bad_ps_fontname	
bad_glyph_table	
bad_cff_table	
bad_metrics_table	
bad_cmap_table	
bad_bitmaps_table	
bad_gx_table	
bad_ot_table	
bad_os2_version	
bad_sfnt_header	

#### 4.4.5.1.16 horiz\_base and vert\_base table

key	type	explanation
tags	table	an array of script list tags
scripts	table	

The `scripts` subtable:

key	type	explanation
baseline	table	
default_baseline	number	
lang	table	

The `lang` subtable:

key	type	explanation
tag	string	a script tag
ascent	number	
descent	number	
features	table	



The `features` points to an array of tables with the same layout except that in those nested tables, the tag represents a language.

#### 4.4.5.1.17 `altuni` table

An array of alternate UNICODE values. Inside that array are hashes with:

key	type	explanation
<code>unicode</code>	number	this glyph is also used for this unicode
<code>variant</code>	number	the alternative is driven by this unicode selector

#### 4.4.5.1.18 `vert_variants` and `horiz_variants` table

key	type	explanation
<code>variants</code>	string	
<code>italic_correction</code>	number	
<code>parts</code>	table	

The `parts` table is an array of smaller tables:

key	type	explanation
<code>component</code>	string	
<code>extender</code>	number	
<code>start</code>	number	
<code>end</code>	number	
<code>advance</code>	number	

#### 4.4.5.1.19 `mathkern` table

key	type	explanation
<code>top_right</code>	table	
<code>bottom_right</code>	table	
<code>top_left</code>	table	
<code>bottom_left</code>	table	

Each of the subtables is an array of small hashes with two keys:

key	type	explanation
<code>height</code>	number	
<code>kern</code>	number	

#### 4.4.5.1.20 `kerns` table

Substructure is identical to the per-glyph subtable.



#### 4.4.5.1.21 vkerns table

Substructure is identical to the per-glyph subtable.

#### 4.4.5.1.22 texdata table

key	type	explanation
type	string	<a href="#">unset</a> , <a href="#">text</a> , <a href="#">math</a> , <a href="#">mathext</a>
params	array	22 font numeric parameters

#### 4.4.5.1.23 lookups table

Top-level [lookups](#) is quite different from the ones at character level. The keys in this hash are strings, the values the actual lookups, represented as dictionary tables.

key	type	explanation
type	string	
format	enum	one of <a href="#">glyphs</a> , <a href="#">class</a> , <a href="#">coverage</a> , <a href="#">reversecoverage</a>
tag	string	
current_class	array	
before_class	array	
after_class	array	
rules	array	an array of rule items

Rule items have one common item and one specialized item:

key	type	explanation
lookups	array	a linear array of lookup names
glyphs	array	only if the parent's format is <a href="#">glyphs</a>
class	array	only if the parent's format is <a href="#">class</a>
coverage	array	only if the parent's format is <a href="#">coverage</a>
reversecoverage	array	only if the parent's format is <a href="#">reversecoverage</a>

A glyph table is:

key	type	explanation
names	string	
back	string	
fore	string	

A class table is:

key	type	explanation
current	array	of numbers
before	array	of numbers
after	array	of numbers



coverage:

key	type	explanation
current	array	of strings
before	array	of strings
after	array	of strings

reversecoverage:

key	type	explanation
current	array	of strings
before	array	of strings
after	array	of strings
replacements	string	

## 4.5 The `img` library

The `img` library can be used as an alternative to `\pdfximage` and `\pdfrefximage`, and the associated ‘satellite’ commands like `\pdfximagebbox`. Image objects can also be used within virtual fonts via the `image` command listed in [section 7.2](#).

### 4.5.1 `img.new`

```
<image> var = img.new()  
<image> var = img.new(<table> image_spec)
```

This function creates a userdata object of type ‘image’. The `image_spec` argument is optional. If it is given, it must be a table, and that table must contain a `filename` key. A number of other keys can also be useful, these are explained below.

You can either say

```
a = img.new()
```

followed by

```
a.filename = "foo.png"
```

or you can put the file name (and some or all of the other keys) into a table directly, like so:

```
a = img.new({filename='foo.pdf', page=1})
```

The generated `<image>` userdata object allows access to a set of user-specified values as well as a set of values that are normally filled in and updated automatically by L<sup>A</sup>T<sub>E</sub>X itself. Some of those are derived from the actual image file, others are updated to reflect the PDF output status of the object.



There is one required user-specified field: the file name (`filename`). It can optionally be augmented by the requested image dimensions (`width`, `depth`, `height`), user-specified image attributes (`attr`), the requested PDF page identifier (`page`), the requested boundingbox (`pagebox`) for PDF inclusion, the requested color space object (`colorspace`).

The function `img.new` does not access the actual image file, it just creates the `<image>` userdata object and initializes some memory structures. The `<image>` object and its internal structures are automatically garbage collected.

Once the image is scanned, all the values in the `<image>` except `width`, `height` and `depth`, become frozen, and you cannot change them any more.

## 4.5.2 `img.keys`

```
<table> keys = img.keys()
```

This function returns a list of all the possible `image_spec` keys, both user-supplied and automatic ones.

field name	type	description
<code>attr</code>	string	the image attributes for L <sup>A</sup> T <sub>E</sub> X
<code>bbox</code>	table	table with 4 boundingbox dimensions <code>llx</code> , <code>lly</code> , <code>urx</code> , and <code>ury</code> overruling the <code>pagebox</code> entry
<code>colordepth</code>	number	the number of bits used by the color space
<code>colorspace</code>	number	the color space object number
<code>depth</code>	number	the image depth for L <sup>A</sup> T <sub>E</sub> X (in scaled points)
<code>filename</code>	string	the image file name
<code>filepath</code>	string	the full (expanded) file name of the image
<code>height</code>	number	the image height for L <sup>A</sup> T <sub>E</sub> X (in scaled points)
<code>imagetype</code>	string	one of <code>pdf</code> , <code>png</code> , <code>jpg</code> , <code>jp2</code> , <code>jbig2</code> , or <code>nil</code>
<code>index</code>	number	the PDF image name suffix
<code>objnum</code>	number	the PDF image object number
<code>page</code>	??	the identifier for the requested image page (type is number or string, default is the number 1)
<code>pagebox</code>	string	the requested bounding box, one of <code>none</code> , <code>media</code> , <code>crop</code> , <code>bleed</code> , <code>trim</code> , <code>art</code>
<code>pages</code>	number	the total number of available pages
<code>rotation</code>	number	the image rotation from included PDF file, in multiples of 90 deg.
<code>stream</code>	string	the raw stream data for an <code>/XObject</code> <code>/Form</code> object
<code>transform</code>	number	the image transform, integer number 0..7
<code>width</code>	number	the image width for L <sup>A</sup> T <sub>E</sub> X (in scaled points)
<code>xres</code>	number	the horizontal natural image resolution (in DPI)
<code>xsize</code>	number	the natural image width
<code>yres</code>	number	the vertical natural image resolution (in DPI)
<code>ysize</code>	number	the natural image height

A running (undefined) dimension in `width`, `height`, or `depth` is represented as `nil` in Lua, so if you want to load an image at its ‘natural’ size, you do not have to specify any of those three fields.



The `stream` parameter allows to fabricate an `/XObject /Form` object from a string giving the stream contents, e. g., for a filled rectangle:

```
a.stream = "0 0 20 10 re f"
```

When writing the image, an `/XObject /Form` object is created, like with embedded PDF file writing. The object is written out only once. The `stream` key requires that also the `bbox` table is given. The `stream` key conflicts with the `filename` key. The `transform` key works as usual also with `stream`.

The `bbox` key needs a table with four boundingbox values, e. g.:

```
a.bbox = {"30bp", 0, "225bp", "200bp"}
```

This replaces and overrules any given `pagebox` value; with given `bbox` the box dimensions coming with an embedded PDF file are ignored. The `xsize` and `ysize` dimensions are set accordingly, when the image is scaled. The `bbox` parameter is ignored for non-PDF images.

The `transform` allows to mirror and rotate the image in steps of 90 deg. The default value 0 gives an unmirrored, unrotated image. Values 1–3 give counterclockwise rotation by 90, 180, or 270 degrees, whereas with values 4–7 the image is first mirrored and then rotated counterclockwise by 90, 180, or 270 degrees. The `transform` operation gives the same visual result as if you would externally preprocess the image by a graphics tool and then use it by L<sup>A</sup>T<sub>E</sub>X. If a PDF file to be embedded already contains a `/Rotate` specification, the rotation result is the combination of the `/Rotate` rotation followed by the `transform` operation.

### 4.5.3 `img.scan`

```
<image> var = img.scan(<image> var)
<image> var = img.scan(<table> image_spec)
```

When you say `img.scan(a)` for a new image, the file is scanned, and variables such as `xsize`, `ysize`, image `type`, number of `pages`, and the resolution are extracted. Each of the `width`, `height`, `depth` fields are set up according to the image dimensions, if they were not given an explicit value already. An image file will never be scanned more than once for a given image variable. With all subsequent `img.scan(a)` calls only the dimensions are again set up (if they have been changed by the user in the meantime).

For ease of use, you can do right-away a

```
<image> a = img.scan ({ filename = "foo.png" })
```

without a prior `img.new`.

Nothing is written yet at this point, so you can do `a=img.scan`, retrieve the available info like image width and height, and then throw away `a` again by saying `a=nil`. In that case no image object will be reserved in the PDF, and the used memory will be cleaned up automatically.



## 4.5.4 `img.copy`

```
<image> var = img.copy(<image> var)
<image> var = img.copy(<table> image_spec)
```

If you say `a = b`, then both variables point to the same `<image>` object. if you want to write out an image with different sizes, you can do a `b=img.copy(a)`.

Afterwards, `a` and `b` still reference the same actual image dictionary, but the dimensions for `b` can now be changed from their initial values that were just copies from `a`.

## 4.5.5 `img.write`

```
<image> var = img.write(<image> var)
<image> var = img.write(<table> image_spec)
```

By `img.write(a)` a PDF object number is allocated, and a whatsit node of subtype `pdf_refximage` is generated and put into the output list. By this the image `a` is placed into the page stream, and the image file is written out into an image stream object after the shipping of the current page is finished.

Again you can do a terse call like

```
img.write ({ filename = "foo.png" })
```

The `<image>` variable is returned in case you want it for later processing.

## 4.5.6 `img.immediatewrite`

```
<image> var = img.immediatewrite(<image> var)
<image> var = img.immediatewrite(<table> image_spec)
```

By `img.immediatewrite(a)` a PDF object number is allocated, and the image file for image `a` is written out immediately into the PDF file as an image stream object (like with `\immediate\pdfximage`). The object number of the image stream dictionary is then available by the `objnum` key. No `pdf_refximage` whatsit node is generated. You will need an `img.write(a)` or `img.node(a)` call to let the image appear on the page, or reference it by another trick; else you will have a dangling image object in the PDF file.

Also here you can do a terse call like

```
a = img.immediatewrite ({ filename = "foo.png" })
```

The `<image>` variable is returned and you will most likely need it.

## 4.5.7 `img.node`

```
<node> n = img.node(<image> var)
```



```
<node> n = img.node(<table> image_spec)
```

This function allocates a PDF object number and returns a `whatsit` node of subtype `pdf_refximage`, filled with the image parameters `width`, `height`, `depth`, and `objnum`. Also here you can do a terse call like:

```
n = img.node ({ filename = "foo.png" })
```

This example outputs an image:

```
node.write(img.node{filename="foo.png"})
```

### 4.5.8 `img.types`

```
<table> types = img.types()
```

This function returns a list with the supported image file type names, currently these are `pdf`, `png`, `jpg`, `jp2` (JPEG 2000), and `jbig2`.

### 4.5.9 `img.bboxes`

```
<table> bboxes = img.bboxes()
```

This function returns a list with the supported PDF page box names, currently these are `media`, `crop`, `bleed`, `trim`, and `art` (all in lowercase letters).

## 4.6 The `kpse` library

This library provides two separate, but nearly identical interfaces to the `KPATHSEA` file search functionality: there is a ‘normal’ procedural interface that shares its `kpathsea` instance with `LUATEX` itself, and an object oriented interface that is completely on its own. The object oriented interface and `kpse.new` have been added in `LUATEX` 0.37.

### 4.6.1 `kpse.set_program_name` and `kpse.new`

Before the search library can be used at all, its database has to be initialized. There are three possibilities, two of which belong to the procedural interface.

First, when `LUATEX` is used to typeset documents, this initialization happens automatically and the `KPATHSEA` executable and program names are set to `luatex` (that is, unless explicitly prohibited by the user’s startup script. See [section 3.1](#) for more details).

Second, in `TEXLUA` mode, the initialization has to be done explicitly via the `kpse.set_program_name` function, which sets the `KPATHSEA` executable (and optionally program) name.





```
kpse.set_program_name(<string> name)
kpse.set_program_name(<string> name, <string> progame)
```

The second argument controls the use of the ‘dotted’ values in the `texmf.cnf` configuration file, and defaults to the first argument.

Third, if you prefer the object oriented interface, you have to call a different function. It has the same arguments, but it returns a userdata variable.

```
local kpathsea = kpse.new(<string> name)
local kpathsea = kpse.new(<string> name, <string> progame)
```

Apart from these two functions, the calling conventions of the interfaces are identical. Depending on the chosen interface, you either call `kpse.find_file()` or `kpathsea.find_file()`, with identical arguments and return values.

## 4.6.2 find\_file

The most often used function in the library is `find_file`:

```
<string> f = kpse.find_file(<string> filename)
<string> f = kpse.find_file(<string> filename, <string> ftype)
<string> f = kpse.find_file(<string> filename, <boolean> mustexist)
<string> f = kpse.find_file(<string> filename, <string> ftype, <boolean>
mustexist)
<string> f = kpse.find_file(<string> filename, <string> ftype, <number> dpi)
```

Arguments:

`filename`

the name of the file you want to find, with or without extension.

`ftype`

maps to the `-format` argument of `KPSEWHICH`. The supported `ftype` values are the same as the ones supported by the standalone `kpsewhich` program:



'gf'	'TeX system sources'
'pk'	'PostScript header'
'bitmap font'	'Troff fonts'
'tfm'	'type1 fonts'
'afm'	'vf'
'base'	'dvips config'
'bib'	'ist'
'bst'	'truetype fonts'
'cnf'	'type42 fonts'
'ls-R'	'web2c files'
'fmt'	'other text files'
'map'	'other binary files'
'mem'	'misc fonts'
'mf'	'web'
'mfpool'	'cweb'
'mft'	'enc files'
'mp'	'cmap files'
'mppool'	'subfont definition files'
'MetaPost support'	'opentype fonts'
'ocp'	'pdftex config'
'ofm'	'lig files'
'opl'	'texmfscripts'
'otp'	'lua',
'ovf'	'font feature files',
'ovp'	'cid maps',
'graphic/figure'	'mlbib',
'tex'	'mlbst',
'TeX system documentation'	'clua',
'texpool'	

The default type is `tex`. Note: this is different from `KPSEWHICH`, which tries to deduce the file type itself from looking at the supplied extension. The last four types: 'font feature files', 'cid maps', 'mlbib', 'mlbst' were new additions in L<sup>A</sup>T<sub>E</sub>X 0.40.2.

`mustexist`

is similar to `KPSEWHICH`'s `-must-exist`, and the default is `false`. If you specify `true` (or a non-zero integer), then the `KPSE` library will search the disk as well as the `ls-R` databases.

`dpi`

This is used for the size argument of the formats `pk`, `gf`, and `bitmap font`.

## 4.6.3 lookup

A more powerful (but slower) generic method for finding files is also available (since 0.51). It returns a string for each found file.

```
<string> f, ... = kpse.lookup(<string> filename, <table> options)
```



The options match commandline arguments from [kpsewhich](#):

key	type	description
debug	number	set debugging flags for this lookup
format	string	use specific file type (see list above)
dpi	number	use this resolution for this lookup; default 600
path	string	search in the given path
all	boolean	output all matches, not just the first
mustexist	boolean	(0.65 and higher) search the disk as well as ls-R if necessary
must-exist	boolean	(0.64 and lower) search the disk as well as ls-R if necessary
mktxpk	boolean	disable/enable mktexpk generation for this lookup
mktextex	boolean	disable/enable mktextex generation for this lookup
mktxmf	boolean	disable/enable mktexmf generation for this lookup
mktextfm	boolean	disable/enable mktextfm generation for this lookup
subdir	string or table	only output matches whose directory part ends with the given string(s)

## 4.6.4 `init_prog`

Extra initialization for programs that need to generate bitmap fonts.

```
kpse.init_prog(<string> prefix, <number> base_dpi, <string> mfmode)
kpse.init_prog(<string> prefix, <number> base_dpi, <string> mfmode, <string>
fallback)
```

## 4.6.5 `readable_file`

Test if an (absolute) file name is a readable file.

```
<string> f = kpse.readable_file(<string> name)
```

The return value is the actual absolute filename you should use, because the disk name is not always the same as the requested name, due to aliases and system-specific handling under e.g. MSDOS.

Returns `nil` if the file does not exist or is not readable.

## 4.6.6 `expand_path`

Like `kpsewhich`'s `-expand-path`:

```
<string> r = kpse.expand_path(<string> s)
```

## 4.6.7 `expand_var`

Like `kpsewhich`'s `-expand-var`:



```
<string> r = kpse.expand_var(<string> s)
```

### 4.6.8 expand\_braces

Like kpsewhich's `-expand-braces`:

```
<string> r = kpse.expand_braces(<string> s)
```

### 4.6.9 show\_path

Like kpsewhich's `-show-path`:

```
<string> r = kpse.show_path(<string> ftype)
```

### 4.6.10 var\_value

Like kpsewhich's `-var-value`:

```
<string> r = kpse.var_value(<string> s)
```

### 4.6.11 version

Returns the kpathsea version string (new in 0.51)

```
<string> r = kpse.version()
```

## 4.7 The lang library

This library provides the interface to L<sup>A</sup>T<sub>E</sub>X's structure representing a language, and the associated functions.

```
<language> l = lang.new()  
<language> l = lang.new(<number> id)
```

This function creates a new userdata object. An object of type `<language>` is the first argument to most of the other functions in the `lang` library. These functions can also be used as if they were object methods, using the colon syntax.

Without an argument, the next available internal id number will be assigned to this object. With argument, an object will be created that links to the internal language with that id number.

```
<number> n = lang.id(<language> l)
```

returns the internal `\language` id number this object refers to.



```
<string> n = lang.hyphenation(<language> l)
lang.hyphenation(<language> l, <string> n)
```

Either returns the current hyphenation exceptions for this language, or adds new ones. The syntax of the string is explained in [section 6.3](#).

```
lang.clear_hyphenation(<language> l)
```

Clears the exception dictionary for this language.

```
<string> n = lang.clean(<string> o)
```

Creates a hyphenation key from the supplied hyphenation value. The syntax of the argument string is explained in [section 6.3](#). This function is useful if you want to do something else based on the words in a dictionary file, like spell-checking.

```
<string> n = lang.patterns(<language> l)
lang.patterns(<language> l, <string> n)
```

Adds additional patterns for this language object, or returns the current set. The syntax of this string is explained in [section 6.3](#).

```
lang.clear_patterns(<language> l)
```

Clears the pattern dictionary for this language.

```
<number> n = lang.prehyphenchar(<language> l)
lang.prehyphenchar(<language> l, <number> n)
```

Gets or sets the ‘pre-break’ hyphen character for implicit hyphenation in this language (initially the hyphen, decimal 45).

```
<number> n = lang.posthyphenchar(<language> l)
lang.posthyphenchar(<language> l, <number> n)
```

Gets or sets the ‘post-break’ hyphen character for implicit hyphenation in this language (initially null, decimal 0, indicating emptiness).

```
<number> n = lang.preexhyphenchar(<language> l)
lang.preexhyphenchar(<language> l, <number> n)
```

Gets or sets the ‘pre-break’ hyphen character for explicit hyphenation in this language (initially null, decimal 0, indicating emptiness).

```
<number> n = lang.postexhyphenchar(<language> l)
lang.postexhyphenchar(<language> l, <number> n)
```

Gets or sets the ‘post-break’ hyphen character for explicit hyphenation in this language (initially null, decimal 0, indicating emptiness).



```
<boolean> success = lang.hyphenate(<node> head)
<boolean> success = lang.hyphenate(<node> head, <node> tail)
```

Inserts hyphenation points (discretionary nodes) in a node list. If `tail` is given as argument, processing stops on that node. Currently, `success` is always true if `head` (and `tail`, if specified) are proper nodes, regardless of possible other errors.

Hyphenation works only on ‘characters’, a special subtype of all the glyph nodes with the node subtype having the value 1. Glyph nodes with different subtypes are not processed. See [section 6.1](#) for more details.

## 4.8 The lua library

This library contains one read-only item:

```
<string> s = lua.version
```

This returns the LUA version identifier string. The value is currently Lua 5.2.

### 4.8.1 LUA bytecode registers

LUA registers can be used to communicate LUA functions across LUA chunks. The accepted values for assignments are functions and `nil`. Likewise, the retrieved value is either a function or `nil`.

```
lua.bytecode[<number> n] = <function> f
lua.bytecode[<number> n]()
```

The contents of the `lua.bytecode` array is stored inside the format file as actual LUA bytecode, so it can also be used to preload LUA code.

Note: The function must not contain any upvalues. Currently, functions containing upvalues can be stored (and their upvalues are set to `nil`), but this is an artifact of the current LUA implementation and thus subject to change.

The associated function calls are

```
<function> f = lua.getbytecode(<number> n)
lua.setbytecode(<number> n, <function> f)
```

Note: Since a LUA file loaded using `loadfile(filename)` is essentially an anonymous function, a complete file can be stored in a bytecode register like this:

```
lua.bytecode[n] = loadfile(filename)
```

Now all definitions (functions, variables) contained in the file can be created by executing this bytecode register:

```
lua.bytecode[n]()
```



Note that the path of the file is stored in the Lua bytecode to be used in stack backtraces and therefore dumped into the format file if the above code is used in `INITEX`. If it contains private information, i.e. the user name, this information is then contained in the format file as well. This should be kept in mind when preloading files into a bytecode register in `INITEX`.

## 4.8.2 LUA chunk name registers

There is an array of 65536 (0–65535) potential chunk names for use with the `\directlua` and `\latelua` primitives.

```
lua.name[<number> n] = <string> s
<string> s = lua.name[<number> n]
```

If you want to unset a lua name, you can assign `nil` to it.

## 4.9 The mplib library

The METAPost library interface registers itself in the table `mplib`. It is based on MPLIB version 1.999.

### 4.9.1 mplib.new

To create a new METAPost instance, call

```
<mpinstance> mp = mplib.new({...})
```

This creates the `mp` instance object. The argument hash can have a number of different fields, as follows:

name	type	description	default
<code>error_line</code>	number	error line width	79
<code>print_line</code>	number	line length in ps output	100
<code>random_seed</code>	number	the initial random seed	variable
<code>interaction</code>	string	the interaction mode, one of <code>batch</code> , <code>nonstop</code> , <code>scroll</code> , <code>errorstop</code>	<code>errorstop</code>
<code>job_name</code>	string	<code>--jobname</code>	<code>mpout</code>
<code>find_file</code>	function	a function to find files	only local files

The `find_file` function should be of this form:

```
<string> found = finder (<string> name, <string> mode, <string> type)
```

with:

**name** the requested file

**mode** the file mode: `r` or `w`



`type` the kind of file, one of: `mp`, `tfm`, `map`, `pfb`, `enc`

Return either the full pathname of the found file, or `nil` if the file cannot be found.

Note that the new version of MPLIB no longer uses binary mem files, so the way to preload a set of macros is simply to start off with an `input` command in the first `mp:execute()` call.

## 4.9.2 `mp:statistics`

You can request statistics with:

```
<table> stats = mp:statistics()
```

This function returns the vital statistics for an MPLIB instance. There are four fields, giving the maximum number of used items in each of four allocated object classes:

<code>main_memory</code>	number	memory size
<code>hash_size</code>	number	hash size
<code>param_size</code>	number	simultaneous macro parameters
<code>max_in_open</code>	number	input file nesting levels

Note that in the new version of MPLIB, this is informational only. The objects are all allocated dynamically, so there is no chance of running out of space unless the available system memory is exhausted.

## 4.9.3 `mp:execute`

You can ask the METAPost interpreter to run a chunk of code by calling

```
<table> rettable = mp:execute('metapost language chunk')
```

for various bits of METAPost language input. Be sure to check the `rettable.status` (see below) because when a fatal METAPost error occurs the MPLIB instance will become unusable thereafter.

Generally speaking, it is best to keep your chunks small, but beware that all chunks have to obey proper syntax, like each of them is a small file. For instance, you cannot split a single statement over multiple chunks.

In contrast with the normal standalone `mpost` command, there is *no* implied ‘input’ at the start of the first chunk.

## 4.9.4 `mp:finish`

```
<table> rettable = mp:finish()
```





If for some reason you want to stop using an MPLIB instance while processing is not yet actually done, you can call `mp:finish`. Eventually, used memory will be freed and open files will be closed by the LUA garbage collector, but an explicit `mp:finish` is the only way to capture the final part of the output streams.

## 4.9.5 Result table

The return value of `mp:execute` and `mp:finish` is a table with a few possible keys (only `status` is always guaranteed to be present).

log	string	output to the 'log' stream
term	string	output to the 'term' stream
error	string	output to the 'error' stream (only used for 'out of memory')
status	number	the return value: 0=good, 1=warning, 2=errors, 3=fatal error
fig	table	an array of generated figures (if any)

When `status` equals 3, you should stop using this MPLIB instance immediately, it is no longer capable of processing input.

If it is present, each of the entries in the `fig` array is a userdata representing a figure object, and each of those has a number of object methods you can call:

boundingbox	function	returns the bounding box, as an array of 4 values
postscript	function	returns a string that is the ps output of the <code>fig</code> . this function accepts two optional integer arguments for specifying the values of <code>prologues</code> (first argument) and <code>procset</code> (second argument)
svg	function	returns a string that is the svg output of the <code>fig</code> . This function accepts an optional integer argument for specifying the value of <code>prologues</code>
objects	function	returns the actual array of graphic objects in this <code>fig</code>
copy_objects	function	returns a deep copy of the array of graphic objects in this <code>fig</code>
filename	function	the filename this <code>fig</code> 's PostSCRIPT output would have written to in standalone mode
width	function	the <code>charwd</code> value
height	function	the <code>charht</code> value
depth	function	the <code>chardp</code> value
italcorr	function	the <code>charit</code> value
charcode	function	the (rounded) <code>charcode</code> value

**NOTE:** you can call `fig:objects()` only once for any one `fig` object!

When the boundingbox represents a 'negated rectangle', i.e. when the first set of coordinates is larger than the second set, the picture is empty.

Graphical objects come in various types that each has a different list of accessible values. The types are: `fill`, `outline`, `text`, `start_clip`, `stop_clip`, `start_bounds`, `stop_bounds`, `special`.

There is helper function (`mplib.fields(obj)`) to get the list of accessible values for a particular object, but you can just as easily use the tables given below.



All graphical objects have a field `type` that gives the object type as a string value; it is not explicit mentioned in the following tables. In the following, `numbers` are PostScript points represented as a floating point number, unless stated otherwise. Field values that are of type `table` are explained in the next section.

#### 4.9.5.1 fill

<code>path</code>	table	the list of knots
<code>htap</code>	table	the list of knots for the reversed trajectory
<code>pen</code>	table	knots of the pen
<code>color</code>	table	the object's color
<code>linejoin</code>	number	line join style (bare number)
<code>miterlimit</code>	number	miterlimit
<code>prescript</code>	string	the prescript text
<code>postscript</code>	string	the postscript text

The entries `htap` and `pen` are optional.

There is helper function (`mplib.pen_info(obj)`) that returns a table containing a bunch of vital characteristics of the used pen (all values are floats):

<code>width</code>	number	width of the pen
<code>sx</code>	number	$x$ scale
<code>rx</code>	number	$xy$ multiplier
<code>ry</code>	number	$yx$ multiplier
<code>sy</code>	number	$y$ scale
<code>tx</code>	number	$x$ offset
<code>ty</code>	number	$y$ offset

#### 4.9.5.2 outline

<code>path</code>	table	the list of knots
<code>pen</code>	table	knots of the pen
<code>color</code>	table	the object's color
<code>linejoin</code>	number	line join style (bare number)
<code>miterlimit</code>	number	miterlimit
<code>linecap</code>	number	line cap style (bare number)
<code>dash</code>	table	representation of a dash list
<code>prescript</code>	string	the prescript text
<code>postscript</code>	string	the postscript text

The entry `dash` is optional.



### 4.9.5.3 text

text	string	the text
font	string	font tfm name
dsize	number	font size
color	table	the object's color
width	number	
height	number	
depth	number	
transform	table	a text transformation
prescript	string	the prescript text
postscript	string	the postscript text

### 4.9.5.4 special

prescript string special text

### 4.9.5.5 start\_bounds, start\_clip

path table the list of knots

### 4.9.5.6 stop\_bounds, stop\_clip

Here are no fields available.

## 4.9.6 Subsidiary table formats

### 4.9.6.1 Paths and pens

Paths and pens (that are really just a special type of paths as far as MPLIB is concerned) are represented by an array where each entry is a table that represents a knot.

left_type	string	when present: 'endpoint', but usually absent
right_type	string	like <a href="#">left_type</a>
x_coord	number	X coordinate of this knot
y_coord	number	Y coordinate of this knot
left_x	number	X coordinate of the precontrol point of this knot
left_y	number	Y coordinate of the precontrol point of this knot
right_x	number	X coordinate of the postcontrol point of this knot
right_y	number	Y coordinate of the postcontrol point of this knot

There is one special case: pens that are (possibly transformed) ellipses have an extra string-valued key [type](#) with value [elliptical](#) besides the array part containing the knot list.



### 4.9.6.2 Colors

A color is an integer array with 0, 1, 3 or 4 values:

0	marking only	no values
1	greyscale	one value in the range (0, 1), 'black' is 0
3	RGB	three values in the range (0, 1), 'black' is 0, 0, 0
4	CMYK	four values in the range (0, 1), 'black' is 0, 0, 0, 1

If the color model of the internal object was `uninitialized`, then it was initialized to the values representing 'black' in the colorspace `defaultcolormodel` that was in effect at the time of the `shipout`.

### 4.9.6.3 Transforms

Each transform is a six-item array.

1	number	represents x
2	number	represents y
3	number	represents xx
4	number	represents yx
5	number	represents xy
6	number	represents yy

Note that the translation (index 1 and 2) comes first. This differs from the ordering in `POSTSCRIPT`, where the translation comes last.

### 4.9.6.4 Dashes

Each `dash` is two-item hash, using the same model as `POSTSCRIPT` for the representation of the dashlist. `dashes` is an array of 'on' and 'off', values, and `offset` is the phase of the pattern.

dashes	hash	an array of on-off numbers
offset	number	the starting offset value

## 4.9.7 Character size information

These functions find the size of a glyph in a defined font. The `fontname` is the same name as the argument to `infont`; the `char` is a glyph id in the range 0 to 255; the returned `w` is in AFM units.

### 4.9.7.1 mp:char\_width

```
<number> w = mp:char_width(<string> fontname, <number> char)
```



### 4.9.7.2 mp:char\_height

```
<number> w = mp:char_height(<string> fontname, <number> char)
```

### 4.9.7.3 mp:char\_depth

```
<number> w = mp:char_depth(<string> fontname, <number> char)
```

## 4.10 The node library

The `node` library contains functions that facilitate dealing with (lists of) nodes and their values. They allow you to create, alter, copy, delete, and insert L<sup>A</sup>T<sub>E</sub>X node objects, the core objects within the typesetter.

L<sup>A</sup>T<sub>E</sub>X nodes are represented in Lua as userdata with the metadata type `luatex.node`. The various parts within a node can be accessed using named fields.

Each node has at least the three fields `next`, `id`, and `subtype`:

- The `next` field returns the userdata object for the next node in a linked list of nodes, or `nil`, if there is no next node.
- The `id` indicates T<sub>E</sub>X's 'node type'. The field `id` has a numeric value for efficiency reasons, but some of the library functions also accept a string value instead of `id`.
- The `subtype` is another number. It often gives further information about a node of a particular `id`, but it is most important when dealing with 'whatsits', because they are differentiated solely based on their `subtype`.

The other available fields depend on the `id` (and for 'whatsits', the `subtype`) of the node. Further details on the various fields and their meanings are given in **chapter 8**.

Support for `unset` (alignment) nodes is partial: they can be queried and modified from Lua code, but not created.

Nodes can be compared to each other, but: you are actually comparing indices into the node memory. This means that equality tests can only be trusted under very limited conditions. It will not work correctly in any situation where one of the two nodes has been freed and/or reallocated: in that case, there will be false positives.

At the moment, memory management of nodes should still be done explicitly by the user. Nodes are not 'seen' by the Lua garbage collector, so you have to call the node freeing functions yourself when you are no longer in need of a node (list). Nodes form linked lists without reference counting, so you have to be careful that when control returns back to L<sup>A</sup>T<sub>E</sub>X itself, you have not deleted nodes that are still referenced from a `next` pointer elsewhere, and that you did not create nodes that are referenced more than once.

There are statistics available with regards to the allocated node memory, which can be handy for tracing.



## 4.10.1 Node handling functions

### 4.10.1.1 `node.is_node`

```
<boolean> t = node.is_node(<any> item)
```

This function returns true if the argument is a userdata object of type `<node>`.

### 4.10.1.2 `node.types`

```
<table> t = node.types()
```

This function returns an array that maps node id numbers to node type strings, providing an overview of the possible top-level `id` types.

### 4.10.1.3 `node.whatsits`

```
<table> t = node.whatsits()
```

TeX's 'whatsits' all have the same `id`. The various subtypes are defined by their `subtype` fields. The function is much like `node.types`, except that it provides an array of `subtype` mappings.

### 4.10.1.4 `node.id`

```
<number> id = node.id(<string> type)
```

This converts a single type name to its internal numeric representation.

### 4.10.1.5 `node.subtype`

```
<number> subtype = node.subtype(<string> type)
```

This converts a single whatsit name to its internal numeric representation (`subtype`).

### 4.10.1.6 `node.type`

```
<string> type = node.type(<any> n)
```

If the argument is a number, then this function converts an internal numeric representation to an external string representation. Otherwise, it will return the string `node` if the object represents a node (this is new in 0.65), and `nil` otherwise.



#### 4.10.1.7 node.fields

```
<table> t = node.fields(<number> id)
<table> t = node.fields(<number> id, <number> subtype)
```

This function returns an array of valid field names for a particular type of node. If you want to get the valid fields for a ‘whatsit’, you have to supply the second argument also. In other cases, any given second argument will be silently ignored.

This function accepts string `id` and `subtype` values as well.

#### 4.10.1.8 node.has\_field

```
<boolean> t = node.has_field(<node> n, <string> field)
```

This function returns a boolean that is only true if `n` is actually a node, and it has the field.

#### 4.10.1.9 node.new

```
<node> n = node.new(<number> id)
<node> n = node.new(<number> id, <number> subtype)
```

Creates a new node. All of the new node’s fields are initialized to either zero or `nil` except for `id` and `subtype` (if supplied). If you want to create a new `whatsit`, then the second argument is required, otherwise it need not be present. As with all node functions, this function creates a node on the  $\text{T}_{\text{E}}\text{X}$  level.

This function accepts string `id` and `subtype` values as well.

#### 4.10.1.10 node.free

```
node.free(<node> n)
```

Removes the node `n` from  $\text{T}_{\text{E}}\text{X}$ ’s memory. Be careful: no checks are done on whether this node is still pointed to from a register or some `next` field: it is up to you to make sure that the internal data structures remain correct.

#### 4.10.1.11 node.flush\_list

```
node.flush_list(<node> n)
```

Removes the node list `n` and the complete node list following `n` from  $\text{T}_{\text{E}}\text{X}$ ’s memory. Be careful: no checks are done on whether any of these nodes is still pointed to from a register or some `next` field: it is up to you to make sure that the internal data structures remain correct.



#### 4.10.1.12 `node.copy`

```
<node> m = node.copy(<node> n)
```

Creates a deep copy of node `n`, including all nested lists as in the case of a `hlist` or `vlist` node. Only the `next` field is not copied.

#### 4.10.1.13 `node.copy_list`

```
<node> m = node.copy_list(<node> n)
<node> m = node.copy_list(<node> n, <node> m)
```

Creates a deep copy of the node list that starts at `n`. If `m` is also given, the copy stops just before node `m`.

Note that you cannot copy attribute lists this way, specialized functions for dealing with attribute lists will be provided later but are not there yet. However, there is normally no need to copy attribute lists as when you do assignments to the `attr` field or make changes to specific attributes, the needed copying and freeing takes place automatically.

#### 4.10.1.14 `node.next` (0.65)

```
<node> m = node.next(<node> n)
```

Returns the node following this node, or `nil` if there is no such node.

#### 4.10.1.15 `node.prev` (0.65)

```
<node> m = node.prev(<node> n)
```

Returns the node preceding this node, or `nil` if there is no such node.

#### 4.10.1.16 `node.current_attr` (0.66)

```
<node> m = node.current_attr()
```

Returns the currently active list of attributes, if there is one.

The intended usage of `current_attr` is as follows:

```
local x1 = node.new("glyph")
x1.attr = node.current_attr()
local x2 = node.new("glyph")
x2.attr = node.current_attr()
```

or:





```

local x1 = node.new("glyph")
local x2 = node.new("glyph")
local ca = node.current_attr()
x1.attr = ca
x2.attr = ca

```

The attribute lists are ref counted and the assignment takes care of incrementing the refcount. You cannot expect the value `ca` to be valid any more when you assign attributes (using `tex.setattribute`) or when control has been passed back to `TEX`.

Note: this function is somewhat experimental, and it returns the *actual* attribute list, not a copy thereof. Therefore, changing any of the attributes in the list will change these values for all nodes that have the current attribute list assigned to them.

#### 4.10.1.17 `node.hpack`

```

<node> h, <number> b = node.hpack(<node> n)
<node> h, <number> b = node.hpack(<node> n, <number> w, <string> info)
<node> h, <number> b = node.hpack(<node> n, <number> w, <string> info, <string>
dir)

```

This function creates a new hlist by packaging the list that begins at node `n` into a horizontal box. With only a single argument, this box is created using the natural width of its components. In the three argument form, `info` must be either `additional` or `exactly`, and `w` is the additional (`\hbox spread`) or exact (`\hbox to`) width to be used.

Direction support added in L<sup>A</sup>T<sub>E</sub>X 0.45.

The second return value is the badness of the generated box, this extension was added in 0.51.

Caveat: at this moment, there can be unexpected side-effects to this function, like updating some of the `\marks` and `\inserts`. Also note that the content of `h` is the original node list `n`: if you call `node.free(h)` you will also free the node list itself, unless you explicitly set the `list` field to `nil` beforehand. And in a similar way, calling `node.free(n)` will invalidate `h` as well!

#### 4.10.1.18 `node.vpack` (since 0.36)

```

<node> h, <number> b = node.vpack(<node> n)
<node> h, <number> b = node.vpack(<node> n, <number> w, <string> info)
<node> h, <number> b = node.vpack(<node> n, <number> w, <string> info, <string>
dir)

```

This function creates a new vlist by packaging the list that begins at node `n` into a vertical box. With only a single argument, this box is created using the natural height of its components. In the three argument form, `info` must be either `additional` or `exactly`, and `w` is the additional (`\vbox spread`) or exact (`\vbox to`) height to be used.

Direction support added in L<sup>A</sup>T<sub>E</sub>X 0.45.



The second return value is the badness of the generated box, this extension was added in 0.51. See the description of `node.hpack()` for a few memory allocation caveats.

#### 4.10.1.19 `node.dimensions` (0.43)

```
<number> w, <number> h, <number> d = node.dimensions(<node> n)
<number> w, <number> h, <number> d = node.dimensions(<node> n, <string>
dir)
<number> w, <number> h, <number> d = node.dimensions(<node> n, <node> t)
<number> w, <number> h, <number> d = node.dimensions(<node> n, <node> t,
<string> dir)
```

This function calculates the natural in-line dimensions of the node list starting at node `n` and terminating just before node `t` (or the end of the list, if there is no second argument). The return values are scaled points. An alternative format that starts with glue parameters as the first three arguments is also possible:

```
<number> w, <number> h, <number> d =
  node.dimensions(<number> glue_set, <number> glue_sign,
                  <number> glue_order, <node> n)
<number> w, <number> h, <number> d =
  node.dimensions(<number> glue_set, <number> glue_sign,
                  <number> glue_order, <node> n, <string> dir)
<number> w, <number> h, <number> d =
  node.dimensions(<number> glue_set, <number> glue_sign,
                  <number> glue_order, <node> n, <node> t)
<number> w, <number> h, <number> d =
  node.dimensions(<number> glue_set, <number> glue_sign,
                  <number> glue_order, <node> n, <node> t, <string> dir)
```

This calling method takes glue settings into account and is especially useful for finding the actual width of a sublist of nodes that are already boxed, for example in code like this, which prints the width of the space inbetween the `a` and `b` as it would be if `\box0` was used as-is:

```
\setbox0 = \hbox to 20pt {a b}

\directlua{print (node.dimensions(tex.box[0].glue_set,
                                tex.box[0].glue_sign,
                                tex.box[0].glue_order,
                                tex.box[0].head.next,
                                node.tail(tex.box[0].head))) }
```

Direction support added in L<sup>A</sup>T<sub>E</sub>X 0.45.



#### 4.10.1.20 `node.mlist_to_hlist`

```
<node> h = node.mlist_to_hlist(<node> n,  
                               <string> display_type, <boolean> penalties)
```

This runs the internal mlist to hlist conversion, converting the math list in `n` into the horizontal list `h`. The interface is exactly the same as for the callback `mlist_to_hlist`.

#### 4.10.1.21 `node.slide`

```
<node> m = node.slide(<node> n)
```

Returns the last node of the node list that starts at `n`. As a side-effect, it also creates a reverse chain of `prev` pointers between nodes.

#### 4.10.1.22 `node.tail`

```
<node> m = node.tail(<node> n)
```

Returns the last node of the node list that starts at `n`.

#### 4.10.1.23 `node.length`

```
<number> i = node.length(<node> n)  
<number> i = node.length(<node> n, <node> m)
```

Returns the number of nodes contained in the node list that starts at `n`. If `m` is also supplied it stops at `m` instead of at the end of the list. The node `m` is not counted.

#### 4.10.1.24 `node.count`

```
<number> i = node.count(<number> id, <node> n)  
<number> i = node.count(<number> id, <node> n, <node> m)
```

Returns the number of nodes contained in the node list that starts at `n` that have a matching `id` field. If `m` is also supplied, counting stops at `m` instead of at the end of the list. The node `m` is not counted.

This function also accept string `id`'s.

#### 4.10.1.25 `node.traverse`

```
<node> t = node.traverse(<node> n)
```

This is a lua iterator that loops over the node list that starts at `n`. Typical input code like this



```
for n in node.traverse(head) do
    ...
end
```

is functionally equivalent to:

```
do
    local n
    local function f (head,var)
        local t
        if var == nil then
            t = head
        else
            t = var.next
        end
        return t
    end
    while true do
        n = f (head, n)
        if n == nil then break end
        ...
    end
end
```

It should be clear from the definition of the function `f` that even though it is possible to add or remove nodes from the node list while traversing, you have to take great care to make sure all the `next` (and `prev`) pointers remain valid.

If the above is unclear to you, see the section ‘For Statement’ in the Lua Reference Manual.

#### 4.10.1.26 `node.traverse_id`

```
<node> t = node.traverse_id(<number> id, <node> n)
```

This is an iterator that loops over all the nodes in the list that starts at `n` that have a matching `id` field. See the previous section for details. The change is in the local function `f`, which now does an extra while loop checking against the upvalue `id`:

```
local function f (head,var)
    local t
    if var == nil then
        t = head
    else
        t = var.next
    end
```



```

while not t.id == id do
    t = t.next
end
return t
end

```

#### 4.10.1.27 node.end\_of\_math (0.76)

```
<node> t = node.end_of_math(<node> start)
```

Looks for and returns the next `math_node` following the `start`. If the given node is a math endnode this helper return that node, else it follows the list and return the next math endnote. If no such node is found nil is returned.

#### 4.10.1.28 node.remove

```
<node> head, current = node.remove(<node> head, <node> current)
```

This function removes the node `current` from the list following `head`. It is your responsibility to make sure it is really part of that list. The return values are the new `head` and `current` nodes. The returned `current` is the node following the `current` in the calling argument, and is only passed back as a convenience (or `nil`, if there is no such node). The returned `head` is more important, because if the function is called with `current` equal to `head`, it will be changed.

#### 4.10.1.29 node.insert\_before

```
<node> head, new = node.insert_before(<node> head, <node> current, <node>
new)
```

This function inserts the node `new` before `current` into the list following `head`. It is your responsibility to make sure that `current` is really part of that list. The return values are the (potentially mutated) `head` and the node `new`, set up to be part of the list (with correct `next` field). If `head` is initially `nil`, it will become `new`.

#### 4.10.1.30 node.insert\_after

```
<node> head, new = node.insert_after(<node> head, <node> current, <node>
new)
```

This function inserts the node `new` after `current` into the list following `head`. It is your responsibility to make sure that `current` is really part of that list. The return values are the `head` and the node `new`, set up to be part of the list (with correct `next` field). If `head` is initially `nil`, it will become `new`.



#### 4.10.1.31 `node.first_glyph` (0.65)

```
<node> n = node.first_glyph(<node> n)
<node> n = node.first_glyph(<node> n, <node> m)
```

Returns the first node in the list starting at `n` that is a glyph node with a subtype indicating it is a glyph, or `nil`. If `m` is given, processing stops at (but including) that node, otherwise processing stops at the end of the list.

Note: this function used to be called `first_character`. It has been renamed in L<sup>A</sup>T<sub>E</sub>X 0.65, and the old name is deprecated now.

#### 4.10.1.32 `node.ligaturing`

```
<node> h, <node> t, <boolean> success = node.ligaturing(<node> n)
<node> h, <node> t, <boolean> success = node.ligaturing(<node> n, <node> m)
```

Apply T<sub>E</sub>X-style ligaturing to the specified nodelist. The tail node `m` is optional. The two returned nodes `h` and `t` are the new head and tail (both `n` and `m` can change into a new ligature).

#### 4.10.1.33 `node.kerning`

```
<node> h, <node> t, <boolean> success = node.kerning(<node> n)
<node> h, <node> t, <boolean> success = node.kerning(<node> n, <node> m)
```

Apply T<sub>E</sub>X-style kerning to the specified nodelist. The tail node `m` is optional. The two returned nodes `h` and `t` are the head and tail (either one of these can be an inserted kern node, because special kernings with word boundaries are possible).

#### 4.10.1.34 `node.unprotect_glyphs`

```
node.unprotect_glyphs(<node> n)
```

Subtracts 256 from all glyph node subtypes. This and the next function are helpers to convert from `characters` to `glyphs` during node processing.

#### 4.10.1.35 `node.protect_glyphs`

```
node.protect_glyphs(<node> n)
```

Adds 256 to all glyph node subtypes in the node list starting at `n`, except that if the value is 1, it adds only 255. The special handling of 1 means that `characters` will become `glyphs` after subtraction of 256.



#### 4.10.1.36 `node.last_node`

```
<node> n = node.last_node()
```

This function pops the last node from T<sub>E</sub>X's 'current list'. It returns that node, or `nil` if the current list is empty.

#### 4.10.1.37 `node.write`

```
node.write(<node> n)
```

This is an experimental function that will append a node list to T<sub>E</sub>X's 'current list' (the node list is not deep-copied any more since version 0.38). There is no error checking yet!

#### 4.10.1.38 `node.protrusion_skippable` (0.60.1)

```
<boolean> skippable = node.protrusion_skippable(<node> n)
```

Returns `true` if, for the purpose of line boundary discovery when character protrusion is active, this node can be skipped.

### 4.10.2 Attribute handling

Attributes appear as linked list of userdata objects in the `attr` field of individual nodes. They can be handled individually, but it is much safer and more efficient to use the dedicated functions associated with them.

#### 4.10.2.1 `node.has_attribute`

```
<number> v = node.has_attribute(<node> n, <number> id)
<number> v = node.has_attribute(<node> n, <number> id, <number> val)
```

Tests if a node has the attribute with number `id` set. If `val` is also supplied, also tests if the value matches `val`. It returns the value, or, if no match is found, `nil`.

#### 4.10.2.2 `node.set_attribute`

```
node.set_attribute(<node> n, <number> id, <number> val)
```

Sets the attribute with number `id` to the value `val`. Duplicate assignments are ignored. *[needs explanation]*



### 4.10.2.3 `node.unset_attribute`

```
<number> v = node.unset_attribute(<node> n, <number> id)
<number> v = node.unset_attribute(<node> n, <number> id, <number> val)
```

Unsets the attribute with number `id`. If `val` is also supplied, it will only perform this operation if the value matches `val`. Missing attributes or attribute-value pairs are ignored.

If the attribute was actually deleted, returns its old value. Otherwise, returns `nil`.

## 4.11 The pdf library

This contains variables and functions that are related to the PDF backend.

### 4.11.1 `pdf.mapfile`, `pdf.mapline` (new in 0.53.0)

```
pdf.mapfile(<string> map file)
pdf.mapline(<string> map line)
```

These two functions can be used to replace primitives `\pdfmapfile` and `\pdfmapline` from PDFTEX. They expect a string as only parameter and have no return value.

The also functions replace the former variables `pdf.pdfmapfile` and `pdf.pdfmapline`.

### 4.11.2 `pdf.catalog`, `pdf.info`, `pdf.names`, `pdf.trailer` (new in 0.53.0)

These variables offer a read-write interface to the corresponding PDFTEX token lists. The value types are strings and they are written out to the PDF file directly after the PDFTEX token registers.

The preferred interface is now `pdf.setcatalog`, `pdf.setinfo`, `pdf.setnames` and `pdf.settrailer` for setting these properties and `pdf.getcatalog`, `pdf.getinfo`, `pdf.getnames` and `pdf.gettrailer` for querying them,

The corresponding ‘pdf’ parameter names `pdf.pdfcatalog`, `pdf.pdfinfo`, `pdf.pdfnames`, and `pdf.pdftrailer` are removed in 0.79.0.

### 4.11.3 `pdf.<set/get>pageattributes`, `pdf.<set/get>pagere-sources`, `pdf.<set/get>pagesattributes`

These variables offer a read-write interface to related token lists. The value types are strings. The variables have no interaction with the corresponding PDFTEX token registers `\pdfpageattr`, `\pdf-pageresources`, and `\pdfpagesattr`. They are written out to the PDF file directly after the PDFTEX token registers.





The preferred interface is now `pdf.setpageattributes`, `pdf.setpagesattributes` and `pdf.setpageresources` for setting these properties and `pdf.getpageattributes`, `pdf.getpageattributes` and `pdf.getpageresources` for querying them.

#### 4.11.4 `pdf.h`, `pdf.v`

These are the `h` and `v` values that define the current location on the output page, measured from its lower left corner. The values can be queried using scaled points as units.

```
local h = pdf.h
local v = pdf.v
```

#### 4.11.5 `pdf.getpos`, `pdf.gethpos`, `pdf.getvpos`

These are the function variants of `pdf.h` and `pdf.v`. Sometimes using a function is preferred over a key so this saves wrapping. Also, these functions are faster than the key based access, as `h` and `v` keys are not real variables but looked up using a metatable call. The `getpos` function returns two values, the other return one.

```
local h, v = pdf.getpos()
```

#### 4.11.6 `pdf.hasmatrix`, `pdf.getmatrix`

The current matrix transformation is available via the `getmatrix` command, which returns 6 values: `sx`, `rx`, `ry`, `sy`, `tx`, and `ty`. The `hasmatrix` function returns `true` when a matrix is applied.

```
if pdf.hasmatrix() then
    local sx, rx, ry, sy, tx, ty = pdf.getmatrix()
    -- do something useful or not
end
```

#### 4.11.7 `pdf.print`

A print function to write stuff to the PDF document that can be used from within a `\latelua` argument. This function is not to be used inside `\directlua` unless you know *exactly* what you are doing.

```
pdf.print(<string> s)
pdf.print(<string> type, <string> s)
```

The optional parameter can be used to mimic the behavior of `\pdfliteral`: the `type` is `direct` or `page`.



### 4.11.8 pdf.immediateobj

This function creates a PDF object and immediately writes it to the PDF file. It is modelled after PDF<sub>T</sub>E<sub>X</sub>'s `\immediate\pdfobj` primitives. All function variants return the object number of the newly generated object.

```
<number> n = pdf.immediateobj(<string> objtext)
<number> n = pdf.immediateobj("file", <string> filename)
<number> n = pdf.immediateobj("stream", <string> streamtext, <string> attr-
text)
<number> n = pdf.immediateobj("streamfile", <string> filename, <string> at-
trtext)
```

The first version puts the `objtext` raw into an object. Only the object wrapper is automatically generated, but any internal structure (like `<< >>` dictionary markers) needs to be provided by the user. The second version with keyword `"file"` as 1st argument puts the contents of the file with name `filename` raw into the object. The third version with keyword `"stream"` creates a stream object and puts the `streamtext` raw into the stream. The stream length is automatically calculated. The optional `attrtext` goes into the dictionary of that object. The fourth version with keyword `"streamfile"` does the same as the 3rd one, it just reads the stream data raw from a file.

An optional first argument can be given to make the function use a previously reserved PDF object.

```
<number> n = pdf.immediateobj(<integer> n, <string> objtext)
<number> n = pdf.immediateobj(<integer> n, "file", <string> filename)
<number> n = pdf.immediateobj(<integer> n, "stream", <string> streamtext,
<string> attrtext)
<number> n = pdf.immediateobj(<integer> n, "streamfile", <string> filename,
<string> attrtext)
```

### 4.11.9 pdf.obj

This function creates a PDF object, which is written to the PDF file only when referenced, e.g., by `pdf.refobj()`.

All function variants return the object number of the newly generated object, and there are two separate calling modes.

The first mode is modelled after PDF<sub>T</sub>E<sub>X</sub>'s `\pdfobj` primitive.

```
<number> n = pdf.obj(<string> objtext)
<number> n = pdf.obj("file", <string> filename)
<number> n = pdf.obj("stream", <string> streamtext, <string> attrtext)
<number> n = pdf.obj("streamfile", <string> filename, <string> attrtext)
```

An optional first argument can be given to make the function use a previously reserved PDF object.

```
<number> n = pdf.obj(<integer> n, <string> objtext)
```



```

<number> n = pdf.obj(<integer> n, "file", <string> filename)
<number> n = pdf.obj(<integer> n, "stream", <string> streamtext, <string>
attrtext)
<number> n = pdf.obj(<integer> n, "streamfile", <string> filename, <string>
attrtext)

```

The second mode accepts a single argument table with key–value pairs.

```

<number> n = pdf.obj{ type = <string>,
                      immediate = <boolean>,
                      objnum = <number>,
                      attr = <string>,
                      compresslevel = <number>,
                      objcompression = <boolean>,
                      file = <string>,
                      string = <string>}

```

The `type` field can have the values `raw` and `stream`, this field is required, the others are optional (within constraints).

Note: this mode makes `pdf.obj` look more flexible than it actually is: the constraints from the separate parameter version still apply, so for example you can't have both `string` and `file` at the same time.

#### 4.11.10 pdf.refobj

This function, the Lua version of the `\pdfrefobj` primitive, references an object by its object number, so that the object will be written out.

```
pdf.refobj(<integer> n)
```

This function works in both the `\directlua` and `\latelua` environment. Inside `\directlua` a new whatsit node 'pdf\_refobj' is created, which will be marked for flushing during page output and the object is then written directly after the page, when also the resources objects are written out. Inside `\latelua` the object will be marked for flushing.

This function has no return values.

#### 4.11.11 pdf.reserveobj

This function creates an empty PDF object and returns its number.

```

<number> n = pdf.reserveobj()
<number> n = pdf.reserveobj("annot")

```



### 4.11.12 pdf.registerannot (new in 0.47.0)

This function adds an object number to the `/Annots` array for the current page without doing anything else. This function can only be used from within `\lualatex`.

```
pdf.registerannot (<number> objnum)
```

## 4.12 The pdfscanner library (new in 0.72.0)

The `pdfscanner` library allows interpretation of PDF content streams and `/ToUnicode` (cmap) streams. You can get those streams from the `epdf` library, as explained in an earlier section. There is only a single top-level function in this library:

```
pdfscanner.scan (<Object> stream, <table> operatortable, <table> info)
```

The first argument, `stream`, should be either a PDF stream object, or a PDF array of PDF stream objects (those options comprise the possible return values of `<Page>:getContents()` and `<Object>:getStream()` in the `epdf` library).

The second argument, `operatortable`, should be a Lua table where the keys are PDF operator name strings and the values are Lua functions (defined by you) that are used to process those operators. The functions are called whenever the scanner finds one of these PDF operators in the content stream(s). The functions are called with two arguments: the `scanner` object itself, and the `info` table that was passed as the third argument to `pdfscanner.scan`.

Internally, `pdfscanner.scan` loops over the PDF operators in the stream(s), collecting operands on an internal stack until it finds a PDF operator. If that PDF operator's name exists in `operatortable`, then the associated function is executed. After the function has run (or when there is no function to execute) the internal operand stack is cleared in preparation for the next operator, and processing continues.

The `scanner` argument to the processing functions is needed because it offers various methods to get the actual operands from the internal operand stack. The most important of those functions is

A simple example of processing a PDF's document stream could look like this:

```
function Do (scanner, info)
  local val = scanner:pop()
  local name = val[2] -- val[1] == 'name'
  print (info.space .. 'Use XObject ' .. name)
  local resources = info.resources
  local xobject = resources:lookup("XObject"):getDict():lookup(name)
  if (xobject and xobject:isStream()) then
    local dict = xobject:getStream():getDict()
    if dict then
      local name = dict:lookup('Subtype')
      if name:getName() == 'Form' then
        local newinfo = { space = info.space .. " " ,
```



```

        resources = dict:lookup('Resources'):getDict() }
    pdfscanner.scan(xobject, operatortable, newinfo)
end
end
end
end
operatortable = {Do = Do}

doc = epdf.open(arg[1])
pagenum = 1
while pagenum <= doc:getNumPages() do
    local page = doc:getCatalog():getPage(pagenum)
    local info = { space = "  ", resources = page:getResourceDict() }
    print ('Page ' .. pagenum)
    pdfscanner.scan(page:getContents(), operatortable, info)
    pagenum = pagenum + 1
end

```

This example iterates over all the actual content in the PDF, and prints out the found XObject names. While the code demonstrates quite some of the `epdf` functions, let's focus on the type `pdfscanner` specific code instead.

From the bottom up, the line

```
pdfscanner.scan(page:getContents(), operatortable, info)
```

runs the scanner with the PDF page's top-level content.

The third argument, `info`, contains two entries: `space` is used to indent the printed output, and `resources` is needed so that embedded `XForms` can find their own content.

The second argument, `operatortable` defines a processing function for a single PDF operator, `Do`.

The function `Do` prints the name of the current XObject, and then starts a new scanner for that object's content stream, under the condition that the XObject is in fact a `/Form`. That nested scanner is called with new `info` argument with an updated `space` value so that the indentation of the output nicely nests, and with an new `resources` field to help the next iteration down to properly process any other, embedded XObjects.

Of course, this is not a very useful example in practise, but for the purpose of demonstrating `pdfscanner`, it is just long enough. It makes use of only one `scanner` method: `scanner:pop()`. That function pops the top operand of the internal stack, and returns a lua table where the object at index one is a string representing the type of the operand, and object two is its value.

The list of possible operand types and associated lua value types is:

```

integer    <number>
real       <number>
boolean    <boolean>

```



name	<string>
operator	<string>
string	<string>
array	<table>
dict	<table>

In case of **integer** or **real**, the value is always a Lua (floating point) number.

In case of **name**, the leading slash is always stripped.

In case of **string**, please bear in mind that PDF actually supports different types of strings (with different encodings) in different parts of the PDF document, so may need to reencode some of the results; **pdfscanner** always outputs the byte stream without reencoding anything. **pdfscanner** does not differentiate between literal strings and hexadecimal strings (the hexadecimal values are decoded), and it treats the stream data for inline images as a string that is the single operand for **EI**.

In case of **array**, the table content is a list of **pop** return values.

In case of **dict**, the table keys are PDF name strings and the values are **pop** return values.

There are few more methods defined that you can ask **scanner**:

pop	as explained above
popNumber	return only the value of a <b>real</b> or <b>integer</b>
popName	return only the value of a <b>name</b>
popString	return only the value of a <b>string</b>
popArray	return only the value of a <b>array</b>
popDict	return only the value of a <b>dict</b>
popBool	return only the value of a <b>boolean</b>
done	abort further processing of this <b>scan()</b> call

The **popXXX** are convenience functions, and come in handy when you know the type of the operands beforehand (which you usually do, in PDF). For example, the **Do** function could have used **local name = scanner:popName()** instead, because the single operand to the **Do** operator is always a PDF name object.

The **done** function allows you to abort processing of a stream once you have learned everything you want to learn. This comes in handy while parsing **/ToUnicode**, because there usually is trailing garbage that you are not interested in. Without **done**, processing only end at the end of the stream, possibly wasting CPU cycles.

## 4.13 The status library

This contains a number of run-time configuration items that you may find useful in message reporting, as well as an iterator function that gets all of the names and values as a table.

```
<table> info = status.list()
```



The keys in the table are the known items, the value is the current value. Almost all of the values in `status` are fetched through a metatable at run-time whenever they are accessed, so you cannot use `pairs` on `status`, but you *can* use `pairs` on `info`, of course. If you do not need the full list, you can also ask for a single item by using its name as an index into `status`.

The current list is:

key	explanation
<code>pdf_gone</code>	written PDF bytes
<code>pdf_ptr</code>	not yet written PDF bytes
<code>dvi_gone</code>	written DVI bytes
<code>dvi_ptr</code>	not yet written DVI bytes
<code>total_pages</code>	number of written pages
<code>output_file_name</code>	name of the PDF or DVI file
<code>log_name</code>	name of the log file
<code>banner</code>	terminal display banner
<code>var_used</code>	variable (one-word) memory in use
<code>dyn_used</code>	token (multi-word) memory in use
<code>str_ptr</code>	number of strings
<code>init_str_ptr</code>	number of INITEX strings
<code>max_strings</code>	maximum allowed strings
<code>pool_ptr</code>	string pool index
<code>init_pool_ptr</code>	INITEX string pool index
<code>pool_size</code>	current size allocated for string characters
<code>node_mem_usage</code>	a string giving insight into currently used nodes
<code>var_mem_max</code>	number of allocated words for nodes
<code>fix_mem_max</code>	number of allocated words for tokens
<code>fix_mem_end</code>	maximum number of used tokens
<code>cs_count</code>	number of control sequences
<code>hash_size</code>	size of hash
<code>hash_extra</code>	extra allowed hash
<code>font_ptr</code>	number of active fonts
<code>max_in_stack</code>	max used input stack entries
<code>max_nest_stack</code>	max used nesting stack entries
<code>max_param_stack</code>	max used parameter stack entries
<code>max_buf_stack</code>	max used buffer position
<code>max_save_stack</code>	max used save stack entries
<code>stack_size</code>	input stack size
<code>nest_size</code>	nesting stack size
<code>param_size</code>	parameter stack size
<code>buf_size</code>	current allocated size of the line buffer
<code>save_size</code>	save stack size
<code>obj_ptr</code>	max PDF object pointer
<code>obj_tab_size</code>	PDF object table size
<code>pdf_os_cntr</code>	max PDF object stream pointer



<code>pdf_os_objidx</code>	PDF object stream index
<code>pdf_dest_names_ptr</code>	max PDF destination pointer
<code>dest_names_size</code>	PDF destination table size
<code>pdf_mem_ptr</code>	max PDF memory used
<code>pdf_mem_size</code>	PDF memory size
<code>largest_used_mark</code>	max referenced marks class
<code>filename</code>	name of the current input file
<code>inputid</code>	numeric id of the current input
<code>linenumber</code>	location in the current input file
<code>lasterrorstring</code>	last error string
<code>luabytecodes</code>	number of active LUA bytecode registers
<code>luabytecode_bytes</code>	number of bytes in LUA bytecode registers
<code>luastate_bytes</code>	number of bytes in use by LUA interpreters
<code>output_active</code>	<code>true</code> if the <code>\output</code> routine is active
<code>callbacks</code>	total number of executed callbacks so far
<code>indirect_callbacks</code>	number of those that were themselves a result of other callbacks (e.g. file readers)
<code>luatex_svn</code>	the luatex repository id (added in 0.51)
<code>luatex_version</code>	the luatex version number (added in 0.38)
<code>luatex_revision</code>	the luatex revision string (added in 0.38)
<code>ini_version</code>	<code>true</code> if this is an <code>INITEX</code> run (added in 0.38)

## 4.14 The tex library

The `tex` table contains a large list of virtual internal T<sub>E</sub>X parameters that are partially writable.

The designation ‘virtual’ means that these items are not properly defined in LUA, but are only frontends that are handled by a metatable that operates on the actual T<sub>E</sub>X values. As a result, most of the LUA table operators (like `pairs` and `#`) do not work on such items.

At the moment, it is possible to access almost every parameter that has these characteristics:

- You can use it after `\the`
- It is a single token.
- Some special others, see the list below

This excludes parameters that need extra arguments, like `\the\scriptfont`.

The subset comprising simple integer and dimension registers are writable as well as readable (stuff like `\tracingcommands` and `\parindent`).

### 4.14.1 Internal parameter values

For all the parameters in this section, it is possible to access them directly using their names as index in the `tex` table, or by using one of the functions `tex.get()` and `tex.set()`.





The exact parameters and return values differ depending on the actual parameter, and so does whether `tex.set` has any effect. For the parameters that *can* be set, it is possible to use `'global'` as the first argument to `tex.set`; this makes the assignment global instead of local.

```
tex.set (<string> n, ...)
tex.set ('global', <string> n, ...)
... = tex.get (<string> n)
```

#### 4.14.1.1 Integer parameters

The integer parameters accept and return LUA numbers.

Read-write:

<code>tex.adjdemerits</code>	<code>tex.mag</code>
<code>tex.binoppenalty</code>	<code>tex.maxdeadcycles</code>
<code>tex.brokenpenalty</code>	<code>tex.month</code>
<code>tex.catcodetable</code>	<code>tex.newlinechar</code>
<code>tex.clubpenalty</code>	<code>tex.outputpenalty</code>
<code>tex.day</code>	<code>tex.pausing</code>
<code>tex.defaultthyphenchar</code>	<code>tex.pdfadjustspacing</code>
<code>tex.defaultskewchar</code>	<code>tex.pdfcompresslevel</code>
<code>tex.delimiterfactor</code>	<code>tex.pdfdecimaldigits</code>
<code>tex.displaywidowpenalty</code>	<code>tex.pdfgamma</code>
<code>tex.doublehyphndemerits</code>	<code>tex.pdfgentounicode</code>
<code>tex.endlinechar</code>	<code>tex.pdfimageapplygamma</code>
<code>tex.errorcontextlines</code>	<code>tex.pdfimagegamma</code>
<code>tex.escapechar</code>	<code>tex.pdfimagehicolor</code>
<code>tex.exhyphenpenalty</code>	<code>tex.pdfimageresolution</code>
<code>tex.fam</code>	<code>tex.pdfinclusionerrorlevel</code>
<code>tex.finalhyphndemerits</code>	<code>tex.pdfminorversion</code>
<code>tex.floatingpenalty</code>	<code>tex.pdfobjcompresslevel</code>
<code>tex.globaldefs</code>	<code>tex.pdfoutput</code>
<code>tex.hangafter</code>	<code>tex.pdfpagebox</code>
<code>tex.hbadness</code>	<code>tex.pdfpkresolution</code>
<code>tex.holdinginserts</code>	<code>tex.pdfprotrudechars</code>
<code>tex.hyphenpenalty</code>	<code>tex.pdftracingfonts</code>
<code>tex.interlinepenalty</code>	<code>tex.pdfuniqueresname</code>
<code>tex.language</code>	<code>tex.postdisplaypenalty</code>
<code>tex.lastlinefit</code>	<code>tex.predisplaydirection</code>
<code>tex.lefthyphenmin</code>	<code>tex.predisplaypenalty</code>
<code>tex.linepenalty</code>	<code>tex.pretolerance</code>
<code>tex.localbrokenpenalty</code>	<code>tex.relpenalty</code>
<code>tex.localinterlinepenalty</code>	<code>tex.righthyphenmin</code>
<code>tex.looseness</code>	<code>tex.savinghyphcodes</code>



<code>tex.savingvdiscards</code>	<code>tex.tracingonline</code>
<code>tex.showboxbreadth</code>	<code>tex.tracingoutput</code>
<code>tex.showboxdepth</code>	<code>tex.tracingpages</code>
<code>tex.time</code>	<code>tex.tracingparagraphs</code>
<code>tex.tolerance</code>	<code>tex.tracingrestores</code>
<code>tex.tracingassigns</code>	<code>tex.tracingscantokens</code>
<code>tex.tracingcommands</code>	<code>tex.tracingstats</code>
<code>tex.tracinggroups</code>	<code>tex.uchyph</code>
<code>tex.tracingifs</code>	<code>tex.vbadness</code>
<code>tex.tracinglostchars</code>	<code>tex.widowpenalty</code>
<code>tex.tracingmacros</code>	<code>tex.year</code>
<code>tex.tracingnesting</code>	

Read-only:

<code>tex.deadcycles</code>	<code>tex.parshape</code>	<code>tex.spacefactor</code>
<code>tex.insertpenalties</code>	<code>tex.prevgraf</code>	

#### 4.14.1.2 Dimension parameters

The dimension parameters accept LUA numbers (signifying scaled points) or strings (with included dimension). The result is always a number in scaled points.

Read-write:

<code>tex.boxmaxdepth</code>	<code>tex.nulldelimiter-</code>	<code>tex.pdfpageheight</code>
<code>tex.delimitershort-</code>	<code>space</code>	<code>tex.pdfpagewidth</code>
<code>fall</code>	<code>tex.overfullrule</code>	<code>tex.pdfpxdimen</code>
<code>tex.displayindent</code>	<code>tex.pagebottomoffset</code>	<code>tex.pdfthreadmargin</code>
<code>tex.displaywidth</code>	<code>tex.pageheight</code>	<code>tex.pdfvorigin</code>
<code>tex.emergencystretch</code>	<code>tex.pageleftoffset</code>	<code>tex.predisplaysize</code>
<code>tex.hangindent</code>	<code>tex.pagerightoffset</code>	<code>tex.scriptspace</code>
<code>tex.hfuzz</code>	<code>tex.pagetopoffset</code>	<code>tex.splitmaxdepth</code>
<code>tex.hoffset</code>	<code>tex.pagewidth</code>	<code>tex.vfuzz</code>
<code>tex.hsize</code>	<code>tex.parindent</code>	<code>tex.voffset</code>
<code>tex.lineskiplimit</code>	<code>tex.pdfdestmargin</code>	<code>tex.vsize</code>
<code>tex.mathsurround</code>	<code>tex.pdfhorigin</code>	
<code>tex.maxdepth</code>	<code>tex.pdflinkmargin</code>	

Read-only:

<code>tex.pagedepth</code>	<code>tex.pagefilstretch</code>	<code>tex.pagestretch</code>
<code>tex.pagefillllstretch</code>	<code>tex.pagegoal</code>	<code>tex.pagetotal</code>
<code>tex.pagefillstretch</code>	<code>tex.pageshrink</code>	<code>tex.prevdepth</code>



#### 4.14.1.3 Direction parameters

The direction parameters are read-only and return a Lua string.



<code>tex.bodydir</code>	<code>tex.pagedir</code>	<code>tex.textdir</code>
<code>tex.mathdir</code>	<code>tex.pardir</code>	

#### 4.14.1.4 Glue parameters

The glue parameters accept and return a userdata object that represents a `glue_spec` node.

<code>tex.abovedisplayshort-</code> <code>skip</code>	<code>tex.belowdisplayskip</code>	<code>tex.spaceskip</code>
<code>tex.abovedisplayskip</code>	<code>tex.leftskip</code>	<code>tex.splittopskip</code>
<code>tex.baselineskip</code>	<code>tex.lineskip</code>	<code>tex.tabskip</code>
<code>tex.belowdisplayshort-</code> <code>skip</code>	<code>tex.parfillskip</code>	<code>tex.topskip</code>
	<code>tex.parskip</code>	<code>tex.xspaceskip</code>
	<code>tex.rightskip</code>	

#### 4.14.1.5 Muglue parameters

All muglue parameters are to be used read-only and return a Lua string.

<code>tex.medmuskip</code>	<code>tex.thickmuskip</code>	<code>tex.thinmuskip</code>
----------------------------	------------------------------	-----------------------------

#### 4.14.1.6 Tokenlist parameters

The tokenlist parameters accept and return Lua strings. Lua strings are converted to and from token lists using `\the\toks` style expansion: all category codes are either space (10) or other (12). It follows that assigning to some of these, like ‘`tex.output`’, is actually useless, but it feels bad to make exceptions in view of a coming extension that will accept full-blown token strings.

<code>tex.errhelp</code>	<code>tex.everyjob</code>	<code>tex.pdfpageattr</code>
<code>tex.everycr</code>	<code>tex.everymath</code>	<code>tex.pdfpageresources</code>
<code>tex.everydisplay</code>	<code>tex.everypar</code>	<code>tex.pdfpagesattr</code>
<code>tex.everyeof</code>	<code>tex.everyvbox</code>	<code>tex.pdfpkmode</code>
<code>tex.everyhbox</code>	<code>tex.output</code>	

### 4.14.2 Convert commands

All ‘convert’ commands are read-only and return a Lua string. The supported commands at this moment are:

<code>tex.eTeXVersion</code>	<code>tex.formatname</code>
<code>tex.eTeXrevision</code>	<code>tex.jobname</code>



<code>tex.luatexbanner</code>	<code>tex.uniformdeviate(number)</code>
<code>tex.luatexrevision</code>	<code>tex.number(number)</code>
<code>tex.pdfnormaldeviate</code>	<code>tex.romannumeral(number)</code>
<code>tex.fontname(number)</code>	<code>tex.pdfpageref(number)</code>
<code>tex.pdffontname(number)</code>	<code>tex.pdfxformname(number)</code>
<code>tex.pdffontobjnum(number)</code>	<code>tex.fontidentifier(number)</code>
<code>tex.pdffontsize(number)</code>	

If you are wondering why this list looks haphazard; these are all the cases of the ‘convert’ internal command that do not require an argument, as well as the ones that require only a simple numeric value.

The special (lua-only) case of `tex.fontidentifier` returns the `csname` string that matches a font id number (if there is one).

The PDF<sub>T</sub>E<sub>X</sub> related primitives will be removed from versions after 0.80. All the version related details are available in the `status` table at the LUA end, for instance one can say:

```
\let\pdfTeXbanner \luatexbanner
\let\pdfTeXversion \luatexversion
\let\pdfTeXrevision\luatexrevision
```

if these are really needed in a macro package.

### 4.14.3 Last item commands

All ‘last item’ commands are read-only and return a number.

The supported commands at this moment are:

<code>tex.lastpenalty</code>	<code>tex.pdflastximagepages</code>	<code>tex.eTeXversion</code>
<code>tex.lastkern</code>	<code>tex.pdflastannot</code>	<code>tex.currentgrouplevel</code>
<code>tex.lastskip</code>	<code>tex.pdflastxpos</code>	<code>tex.currentgrouptype</code>
<code>tex.lastnodetype</code>	<code>tex.pdflastypos</code>	<code>tex.currentiflevel</code>
<code>tex.inputlineno</code>	<code>tex.pdfrandomseed</code>	<code>tex.currentiftype</code>
<code>tex.pdflastobj</code>	<code>tex.pdflastlink</code>	<code>tex.currentifbranch</code>
<code>tex.pdflastxform</code>	<code>tex.luatexversion</code>	<code>tex.pdflastximagecol-</code>
<code>tex.pdflastximage</code>	<code>tex.eTeXminorversion</code>	<code>ordepth</code>

### 4.14.4 Attribute, count, dimension, skip and token registers

T<sub>E</sub>X’s attributes (`\attribute`), counters (`\count`), dimensions (`\dimen`), skips (`\skip`) and token (`\toks`) registers can be accessed and written to using two times five virtual sub-tables of the `tex` table:

<code>tex.attribute</code>	<code>tex.count</code>	<code>tex.dimen</code>
----------------------------	------------------------	------------------------



```
tex.skip                tex.toks
```

It is possible to use the names of relevant `\attributedef`, `\countdef`, `\dimendef`, `\skipdef`, or `\toksdef` control sequences as indices to these tables:

```
tex.count.scratchcounter = 0
enormous = tex.dimen['maxdimen']
```

In this case, L<sup>A</sup>T<sub>E</sub>X looks up the value for you on the fly. You have to use a valid `\countdef` (or `\attributedef`, or `\dimendef`, or `\skipdef`, or `\toksdef`), anything else will generate an error (the intent is to eventually also allow `<chardef tokens>` and even macros that expand into a number).

The attribute and count registers accept and return Lua numbers.

The dimension registers accept Lua numbers (in scaled points) or strings (with an included absolute dimension; `em` and `ex` and `px` are forbidden). The result is always a number in scaled points.

The token registers accept and return Lua strings. Lua strings are converted to and from token lists using `\the\toks` style expansion: all category codes are either space (10) or other (12).

The skip registers accept and return `glue_spec` userdata node objects (see the description of the node interface elsewhere in this manual).

As an alternative to array addressing, there are also accessor functions defined for all cases, for example, here is the set of possibilities for `\skip` registers:

```
tex.setskip (<number> n, <node> s)
tex.setskip (<string> s, <node> s)
tex.setskip ('global', <number> n, <node> s)
tex.setskip ('global', <string> s, <node> s)
<node> s = tex.getskip (<number> n)
<node> s = tex.getskip (<string> s)
```

In the function-based interface, it is possible to define values globally by using the string `'global'` as the first function argument.

### 4.14.5 Character code registers (0.63)

T<sub>E</sub>X's character code tables (`\lccode`, `\uccode`, `\sfcode`, `\catcode`, `\mathcode`, `\delcode`) can be accessed and written to using six virtual subtables of the `tex` table

```
tex.lccode                tex.sfcode                tex.mathcode
tex.uccode                tex.catcode               tex.delcode
```

The function call interfaces are roughly as above, but there are a few twists. `sfcodes` are the simple ones:

```
tex.setsfcode (<number> n, <number> s)
```



```
tex.setsfcodes ('global', <number> n, <number> s)
<number> s = tex.getsfcodes (<number> n)
```

The function call interface for `lccode` and `uccode` additionally allows you to set the associated sibling at the same time:

```
tex.setlccode (['global'], <number> n, <number> lc)
tex.setlccode (['global'], <number> n, <number> lc, <number> uc)
<number> lc = tex.getlccode (<number> n)
tex.setuccode (['global'], <number> n, <number> uc)
tex.setuccode (['global'], <number> n, <number> uc, <number> lc)
<number> uc = tex.getuccode (<number> n)
```

The function call interface for `catcode` also allows you to specify a category table to use on assignment or on query (default in both cases is the current one):

```
tex.setcatcode (['global'], <number> n, <number> c)
tex.setcatcode (['global'], <number> cattable, <number> n, <number> c)
<number> lc = tex.getcatcode (<number> n)
<number> lc = tex.getcatcode (<number> cattable, <number> n)
```

The interfaces for `delcode` and `mathcode` use small array tables to set and retrieve values:

```
tex.setmathcode (['global'], <number> n, <table> mval )
<table> mval = tex.getmathcode (<number> n)
tex.setdelcode (['global'], <number> n, <table> dval )
<table> dval = tex.getdelcode (<number> n)
```

Where the table for `mathcode` is an array of 3 numbers, like this:

```
{<number> mathclass, <number> family, <number> character}
```

And the table for `delcode` is an array with 4 numbers, like this:

```
{<number> small_fam, <number> small_char, <number> large_fam, <number> large_char}
```

Normally, the third and fourth values in a delimiter code assignment will be zero according to `\Udelcode` usage, but the returned table can have values there (if the delimiter code was set using `\delcode`, for example). Unset `delcode`'s can be recognized because `dval[1]` is `-1`.

## 4.14.6 Box registers

It is possible to set and query actual boxes, using the node interface as defined in the `node` library:

```
tex.box
```

for array access, or



```

tex.setbox(<number> n, <node> s)
tex.setbox(<string> cs, <node> s)
tex.setbox('global', <number> n, <node> s)
tex.setbox('global', <string> cs, <node> s)
<node> n = tex.getbox(<number> n)
<node> n = tex.getbox(<string> cs)

```

for function-based access. In the function-based interface, it is possible to define values globally by using the string 'global' as the first function argument.

Be warned that an assignment like

```
tex.box[0] = tex.box[2]
```

does not copy the node list, it just duplicates a node pointer. If `\box2` will be cleared by  $\TeX$  commands later on, the contents of `\box0` becomes invalid as well. To prevent this from happening, always use `node.copy_list()` unless you are assigning to a temporary variable:

```
tex.box[0] = node.copy_list(tex.box[2])
```

## 4.14.7 Math parameters

It is possible to set and query the internal math parameters using:

```

tex.setmath(<string> n, <string> t, <number> n)
tex.setmath('global', <string> n, <string> t, <number> n)
<number> n = tex.getmath(<string> n, <string> t)

```

As before an optional first parameter 'global' indicates a global assignment.

The first string is the parameter name minus the leading 'Umath', and the second string is the style name minus the trailing 'style'.

Just to be complete, the values for the math parameter name are:

quad	axis	operatorsize
overbarkern	overbarrule	overbarvgap
underbarkern	underbarrule	underbarvgap
radicalkern	radicalrule	radicalvgap
radicaldegreebefore	radicaldegreeafter	radicaldegreeraise
stackvgap	stacknumup	stackdenomdown
fractionrule	fractionnumvgap	fractionnumup
fractiondenomvgap	fractiondenomdown	fractiondelsize
limitabovevgap	limitabovebgap	limitabovekern
limitbelowvgap	limitbelowbgap	limitbelowkern
underdelimitervgap	underdelimiterbgap	
overdelimitervgap	overdelimiterbgap	
subshiftdrop	supshiftdrop	subshiftdown





subsupshiftdown	subtopmax	supshiftup	
supbottommin	supsubbottommax	subsupvgap	
spaceafterscript	connectoroverlapmin		
ordordspacing	ordopspacing	ordbinspacing	ordrelspacing
ordopenspacing	ordclosespacing	ordpunctspacing	ordinnerspacing
opordspacing	opopspacing	opbinspacing	oprelspacing
opopenspacing	opclosespacing	oppunctspacing	opinnerspacing
binordspacing	binopspacing	binbinspacing	binrelspacing
binopenspacing	binclosespacing	binpunctspacing	bininnerspacing
relordspacing	relopspacing	relbinspacing	relrelspacing
reloppspacing	relclosespacing	relpunctspacing	relinnerspacing
openordspacing	openopspacing	openbinspacing	openrelspacing
openopenspacing	openclosespacing	openpunctspacing	openinnerspacing
closeordspacing	closeopspacing	closebinspacing	closerelspacing
closeopenspacing	closeclosespacing	closepunctspacing	closeinnerspacing
punctordspacing	punctopspacing	punctbinspacing	punctrelspacing
punctopenspacing	punctclosespacing	punctpunctspacing	punctinnerspacing
innerordspacing	inneropspacing	innerbinspacing	innerrelspacing
inneropenspacing	innerclosespacing	innerpunctspacing	innerinnerspacing

The values for the style parameter name are:

display	crampeddisplay
text	crampedtext
script	crampedscript
scriptscript	crampedscriptscript

## 4.14.8 Special list heads

The virtual table `tex.lists` contains the set of internal registers that keep track of building page lists.

field	description
<code>page_ins_head</code>	circular list of pending insertions
<code>contrib_head</code>	the recent contributions
<code>page_head</code>	the current page content
<code>hold_head</code>	used for held-over items for next page
<code>adjust_head</code>	head of the current <code>\vadjust</code> list
<code>pre_adjust_head</code>	head of the current <code>\vadjust pre</code> list

## 4.14.9 Semantic nest levels (0.51)

The virtual table `tex.nest` contains the currently active semantic nesting state. It has two main parts: a zero-based array of userdata for the semantic nest itself, and the numerical value `tex.nest.ptr`, which gives the highest available index. Neither the array items in `tex.nest[]` nor `tex.nest.ptr`



can be assigned to (as this would confuse the typesetting engine beyond repair), but you can assign to the individual values inside the array items, e.g. `tex.nest[tex.nest.ptr].prevdepth`.

`tex.nest[tex.nest.ptr]` is the current nest state, `tex.nest[0]` the outermost (main vertical list) level.

The known fields are:

key	type	modes	explanation
mode	number	all	The current mode. This is a number representing the main mode at this level: 0 == no mode (this happens during <code>\write</code> ) 1 == vertical, 127 = horizontal, 253 = display math. -1 == internal vertical, -127 = restricted horizontal, -253 = inline math.
modeline	number	all	source input line where this mode was entered in, negative inside the output routine.
head	node	all	the head of the current list
tail	node	all	the tail of the current list
prevgraf	number	vmode	number of lines in the previous paragraph
prevdepth	number	vmode	depth of the previous paragraph (equal to <code>\pdfignoreddimen</code> when it is to be ignored)
spacefactor	number	hmode	the current space factor
dirs	node	hmode	used for temporary storage by the line break algorithm
noad	node	mmode	used for temporary storage of a pending fraction numerator, for <code>\over</code> etc.
delimptr	node	mmode	used for temporary storage of the previous math delimiter, for <code>\middle</code> .
mathdir	boolean	mmode	true when during math processing the <code>\mathdir</code> is not the same as the surrounding <code>\textdir</code>
mathstyle	number	mmode	the current <code>\mathstyle</code>

## 4.14.10 Print functions

The `tex` table also contains the three print functions that are the major interface from Lua scripting to T<sub>E</sub>X.

The arguments to these three functions are all stored in an in-memory virtual file that is fed to the T<sub>E</sub>X scanner as the result of the expansion of `\directlua`.

The total amount of returnable text from a `\directlua` command is only limited by available system RAM. However, each separate printed string has to fit completely in T<sub>E</sub>X's input buffer.

The result of using these functions from inside callbacks is undefined at the moment.



#### 4.14.10.1 `tex.print`

```
tex.print(<string> s, ...)
tex.print(<number> n, <string> s, ...)
tex.print(<table> t)
tex.print(<number> n, <table> t)
```

Each string argument is treated by T<sub>E</sub>X as a separate input line. If there is a table argument instead of a list of strings, this has to be a consecutive array of strings to print (the first non-string value will stop the printing process). This syntax was added in 0.36.

The optional parameter can be used to print the strings using the catcode regime defined by `\catcodetable n`. If `n` is `-1`, the currently active catcode regime is used. If `n` is `-2`, the resulting catcodes are the result of `\the\toks`: all category codes are 12 (other) except for the space character, that has category code 10 (space). Otherwise, if `n` is not a valid catcode table, then it is ignored, and the currently active catcode regime is used instead.

The very last string of the very last `tex.print()` command in a `\directlua` will not have the `\endlinechar` appended, all others do.

#### 4.14.10.2 `tex.sprint`

```
tex.sprint(<string> s, ...)
tex.sprint(<number> n, <string> s, ...)
tex.sprint(<table> t)
tex.sprint(<number> n, <table> t)
```

Each string argument is treated by T<sub>E</sub>X as a special kind of input line that makes it suitable for use as a partial line input mechanism:

- T<sub>E</sub>X does not switch to the ‘new line’ state, so that leading spaces are not ignored.
  - No `\endlinechar` is inserted.
  - Trailing spaces are not removed.
- Note that this does not prevent T<sub>E</sub>X itself from eating spaces as result of interpreting the line. For example, in

```
before\directlua{tex.sprint("\relax")tex.sprint(" inbetween")}after
```

the space before `inbetween` will be gobbled as a result of the ‘normal’ scanning of `\relax`.

If there is a table argument instead of a list of strings, this has to be a consecutive array of strings to print (the first non-string value will stop the printing process). This syntax was added in 0.36.

The optional argument sets the catcode regime, as with `tex.print()`.

#### 4.14.10.3 `tex.tprint`

```
tex.tprint({<number> n, <string> s, ...}, {...})
```



This function is basically a shortcut for repeated calls to `tex.sprint(<number> n, <string> s, ...)`, once for each of the supplied argument tables.

#### 4.14.10.4 `tex.write`

```
tex.write(<string> s, ...)
tex.write(<table> t)
```

Each string argument is treated by T<sub>E</sub>X as a special kind of input line that makes it suitable for use as a quick way to dump information:

- All catcodes on that line are either ‘space’ (for ‘ ’) or ‘character’ (for all others).
- There is no `\endlinechar` appended.

If there is a table argument instead of a list of strings, this has to be a consecutive array of strings to print (the first non-string value will stop the printing process). This syntax was added in 0.36.

### 4.14.11 Helper functions

#### 4.14.11.1 `tex.round`

```
<number> n = tex.round(<number> o)
```

Rounds Lua number `o`, and returns a number that is in the range of a valid T<sub>E</sub>X register value. If the number starts out of range, it generates a ‘number to big’ error as well.

#### 4.14.11.2 `tex.scale`

```
<number> n = tex.scale(<number> o, <number> delta)
<table> n = tex.scale(table o, <number> delta)
```

Multiplies the Lua numbers `o` and `delta`, and returns a rounded number that is in the range of a valid T<sub>E</sub>X register value. In the table version, it creates a copy of the table with all numeric top-level values scaled in that manner. If the multiplied number(s) are of range, it generates ‘number to big’ error(s) as well.

Note: the precision of the output of this function will depend on your computer’s architecture and operating system, so use with care! An interface to L<sup>A</sup>T<sub>E</sub>X’s internal, 100% portable scale function will be added at a later date.

#### 4.14.11.3 `tex.sp` (0.51)

```
<number> n = tex.sp(<number> o)
<number> n = tex.sp(<string> s)
```



Converts the number `o` or a string `s` that represents an explicit dimension into an integer number of scaled points.

For parsing the string, the same scanning and conversion rules are used that `LuATEX` would use if it was scanning a dimension specifier in its `TEX`-like input language (this includes generating errors for bad values), expect for the following:

1. only explicit values are allowed, control sequences are not handled
2. infinite dimension units (`fil...`) are forbidden
3. `mu` units do not generate an error (but may not be useful either)

#### 4.14.11.4 `tex.definefont`

```
tex.definefont(<string> csname, <number> fontid)
tex.definefont(<boolean> global, <string> csname, <number> fontid)
```

Associates `csname` with the internal font number `fontid`. The definition is global if (and only if) `global` is specified and true (the setting of `globaldefs` is not taken into account).

#### 4.14.11.5 `tex.error` (0.61)

```
tex.error(<string> s)
tex.error(<string> s, <table> help)
```

This creates an error somewhat like the combination of `\errhelp` and `\errmessage` would. During this error, deletions are disabled.

The array part of the `help` table has to contain strings, one for each line of error help.

#### 4.14.11.6 `tex.hashtokens` (0.25)

```
for i,v in pairs (tex.hashtokens()) do ... end
```

Returns a name and token table pair (see [section 4.17](#) about token tables) iterator for every non-zero entry in the hash table. This can be useful for debugging, but note that this also reports control sequences that may be unreachable at this moment due to local redefinitions: it is strictly a dump of the hash table.

### 4.14.12 Functions for dealing with primitives

#### 4.14.12.1 `tex.enableprimitives`

```
tex.enableprimitives(<string> prefix, <table> primitive names)
```

This function accepts a prefix string and an array of primitive names.



For each combination of ‘prefix’ and ‘name’, the `tex.enableprimitives` first verifies that ‘name’ is an actual primitive (it must be returned by one of the `tex.extraprimtives()` calls explained below, or part of T<sub>E</sub>X82, or `\directlua`). If it is not, `tex.enableprimitives` does nothing and skips to the next pair.

But if it is, then it will construct a csname variable by concatenating the ‘prefix’ and ‘name’, unless the ‘prefix’ is already the actual prefix of ‘name’. In the latter case, it will discard the ‘prefix’, and just use ‘name’.

Then it will check for the existence of the constructed csname. If the csname is currently undefined (note: that is not the same as `\relax`), it will globally define the csname to have the meaning: run code belonging to the primitive ‘name’. If for some reason the csname is already defined, it does nothing and tries the next pair.

An example:

```
tex.enableprimitives('LuaTeX', {'formatname'})
```

will define `\LuaTeXformatname` with the same intrinsic meaning as the documented primitive `\formatname`, provided that the control sequences `\LuaTeXformatname` is currently undefined.

Second example:

```
tex.enableprimitives('Omega',tex.extraprimtives ('omega'))
```

will define a whole series of csnames like `\Omegatextdir`, `\Omegapardir`, etc., but it will stick with `\OmegaVersion` instead of creating the doubly-prefixed `\OmegaOmegaVersion`.

Starting with version 0.39.0 (and this is why the above two functions are needed), L<sup>A</sup>T<sub>E</sub>X in `--ini` mode contains only the T<sub>E</sub>X82 primitives and `\directlua`, no extra primitives **at all**.

So, if you want to have all the new functionality available using their default names, as it is now, you will have to add

```
\ifx\directlua\undefined \else
  \directlua {tex.enableprimitives('',tex.extraprimtives ())}
\fi
```

near the beginning of your format generation file. Or you can choose different prefixes for different subsets, as you see fit.

Calling some form of `tex.enableprimitives()` is highly important though, because if you do not, you will end up with a T<sub>E</sub>X82-lookalike that can run lua code but not do much else. The defined csnames are (of course) saved in the format and will be available at runtime.

#### 4.14.12.2 `tex.extraprimtives`

```
<table> t = tex.extraprimtives(<string> s, ...)
```

This function returns a list of the primitives that originate from the engine(s) given by the requested string value(s). The possible values and their (current) return values are:



## name values

tex

□

- / above abovedisplayshortskip abovedisplayskip abovewithdelims accent adjdemerits advance afterassignment aftergroup atop atopwithdelims badness baselineskip batchmode begingroup belowdisplayshortskip belowdisplayskip binoppenalty botmark box boxmaxdepth brokenpenalty catcode char chardef cleaders closein closeout clubpenalty copy count countdef cr crcr csname day deadcycles def defaultthyphenchar defaultskewchar delcode delimiter delimiterfactor delimitershortfall dimen dimendef discretionary displayindent displaylimits displaystyle displaywidowpenalty displaywidth divide doublehyphendemerits dp dump edef else emergencystretch end endcsname endgroup endinput endlineskip eqno errhelp errmessage errorcontextlines errorstopmode escapechar everycr everydisplay everyhbox everyjob everymath everypar everyvbox exhyphenchar exhyphenpenalty expandafter fam fi finalhyphendemerits firstmark floatingpenalty font fontdimen fontname futurelet gdef global globaldefs halign hangafter hangindent hbadness hbox hfil hfill hfilneg hfuzz hoffset holdinginserts hrule hsize hskip hss ht hyphenation hyphenchar hyphenpenalty if ifcase ifcat ifdim ifeof iffalse ifhbox ifhmode ifinner ifmmode ifnum ifodd iftrue ifvbox ifvmode ifvoid ifx ignorespaces immediate indent input inputlineno insert insertpenalties interlinepenalty jobname kern language lastbox lastkern lastpenalty lastskip lcode leaders left lefthyphenmin leftskip legno let limits linepenalty lineskip lineskiplimit long looseness lower lowercase mag mark mathaccent mathbin mathchar mathchardef mathchoice mathclose mathcode mathinner mathop mathopen mathord mathpunct mathrel mathsurround maxdeadcycles maxdepth meaning medmuskip message middle mkern month moveleft moveright mskip multiply muskip muskipdef newlinechar noalign noboundary noexpand noindent nolimits nonscript nonstopmode nulldelimiterspace nullfont number omit openin openout or outer output outputpenalty over overfullrule overline overwithdelims pagedepth pagefillstretch pagefillstretch pagefilstretch pagegoal pageshrink pagestretch pagetotal par parfillskip parindent parshape parskip patterns pausing penalty postdisplaypenalty predisdisplaypenalty predisplaysize pretolerance prevdepth prevgraf radical raise read relax relpenalty right righthyphenmin rightskip romannumeral scriptfont scriptscriptfont scriptscriptstyle scriptspace scriptstyle scrollmode setbox setlanguage scode shipout show showbox showboxbreadth showboxdepth showlists showthe skewchar skip skipdef spacefactor spaceskip span special splitbotmark splitfirstmark splitmaxdepth splittopskip string tabskip textfont textstyle the thickmuskip thinmuskip time toks toksdef tolerance topmark topskip tracingcommands tracinglostchars tracingmacros tracingonline tracingoutput tracingpages tracingparagraphs tracingrestores tracingstats uccode uchyph underline unhbox unhcopy unkern unpenalty unskip unvbox unvcopy uppercase vadjust valign vbadness vbox vcenter vfil vfill vfilneg vfuzz voffset vrule vsize vskip vsplit vss vtop wd widowpenalty write xdef xleaders xspaceskip year

core

directlua

etex

botmarks clubpenalties currentgrouplevel currentgrouptype currentifbranch currentiflevel currentifttype detokenize dimexpr displaywidowpenalties eTeXVersion eTeXminorversion eTeXrevision eTeXversion everyeof firstmarks fontcharhp fontcharht fontcharic fontcharwd glueexpr glueshrink glueshrinkorder gluestretch gluestretchorder glueto muexpr mutoglu numexpr pagediscards parshapedimen parshapeindent parshapelength predisdisplaydirection



	protected readline savinghyphcodes savingvdiscards scantokens showgroups showifs showtokens splitbotmarks splitdiscards splitfirstmarks topmarks tracingassigns tracinggroups tracingifs tracingnesting tracingscantokens unexpanded unless widowpenalties
pdftex	efcode expanded ifincsname ifpdfabsdim ifpdfabsnum ifpdfprimitive leftmarginkern letterspacefont lpcode luatexbanner pdfadjustspacing pdfannot pdfcatalog pdfcolorstack pdfcolorstackinit pdfcompresslevel pdfcopyfont pdfcreationdate pdfdecimaldigits pdfdest pdfdestmargin pdfdraftmode pdfeachlinedepth pdfeachlineheight pdfendlink pdfendthread pdffirstlineheight pdffontattr pdffontexpand pdffontname pdffontobjnum pdffontsize pdfgamma pdfgentounicode pdfglyphtounicode pdfhorigin pdfignoreddimen pdfimageapplygamma pdfimagegamma pdfimagehicolor pdfimageresolution pdfincludechars pdfinclusioncopyfonts pdfinclusionerrorlevel pdfinfo pdfinsertht pdflastannot pdflastlinedepth pdflastlink pdflastobj pdflastxform pdflastximage pdflastximagecolordepth pdflastximagepages pdflastxpos pdflastypos pdflinkmargin pdfliteral pdfmapfile pdfmapline pdfminorversion pdfnames pdfnoligatures pdfnormaldeviate pdfobj pdfobjcompresslevel pdfoptionpdfminorversion pdfoutline pdfoutput pdfpageattr pdfpagebox pdfpageheight pdfpageref pdfpageresources pdfpagesattr pdfpagewidth pdfpkmode pdfpkresolution pdfprimitive pdfprotrudechars pdfpxdimen pdfrandomseed pdfrefobj pdfrefxform pdfrefximage pdfreplacefont pdfrestore pdfretval pdfsave pdfsavepos pdfsetmatrix pdfsetrandomseed pdfstartlink pdfstartthread pdftexbanner pdftexrevision pdftexversion pdfthread pdfthreadmargin pdftracingfonts pdftrailer pdfuniformdeviate pdfuniqueresname pdfvorigin pdfxform pdfxformattr pdfxformname pdfxformresources pdfximage pdfximagebbox quitvmode rightmarginkern rpcode tagcode
omega	bodydir chardp charht charit charwd leftghost localbrokenpenalty localinterlinepenalty localleftbox localrightbox mathdir pagedir pageheight pagewidth pardir rightghost textdir
aleph	boxdir pagebottomoffset pagerightoffset
luatex	Uchar Udelcode Udelcodenum Udelimiter Udelimiterover Udelimiterunder Umathaccent Umathaxis Umathbinbinspacing Umathbinclosespacing Umathbininnerspacing Umathbinopenspacing Umathbinopspacing Umathbinordspacing Umathbinpunctspacing Umathbinrelspacing Umathchar Umathchardef Umathcharnum Umathcharnumdef Umathclosebinspacing Umathcloseclosespacing Umathcloseinnerspacing Umathcloseopenspacing Umathcloseopspacing Umathcloseordspacing Umathclosepunctspacing Umathcloserelspacing Umathcode Umathcodenum Umathconnectoroverlapmin Umathfractiondelsize Umathfractiondenomdown Umathfractiondenomvgap Umathfractionnumup Umathfractionnumvgap Umathfractionrule Umathinnerbinspacing Umathinnerclosespacing Umathinnerinnerspacing Umathinneropenspacing Umathinneropspacing Umathinnerordspacing Umathinnerpunctspacing Umathinnerrelspacing Umathlimitabovebgap Umathlimitabovekern Umathlimitabovevgap Umathlimitbelowbgap Umathlimitbelowkern Umathlimitbelowvgap Umathopbinspacing Umathopclosespacing Umathopenbinspacing Umathopenclosespacing Umathopeninnerspacing Umathopenopenspacing Umathopenopspacing Umathopenordspacing Umathopenpunctspacing Umathopenrelspacing Umathoperatorsize Umathopinnerspacing Umathopopenspacing Umathopopspacing Umathopordspacing Umathoppunctspacing Umathoprelspacing Umathordbinspacing Umathordclosespacing Umathordinnerspacing Umathordopenspacing Umathordopspacing Umathordordspacing Umathordpunctspacing Umathordrelspacing Umathoverbarkern Umathoverbarrule Umathoverbarvgap Umathoverdelimiterbgap Umathoverdelimitervgap Umathpunctbinspacing





Umathpunctclosespacing      Umathpunctinnerspacing      Umathpunctopenspacing  
 Umathpunctopspacing   Umathpunctordspacing   Umathpunctpunctspacing   Umathpunctrelspacing  
 Umathquad   Umathradicaldegreeafter   Umathradicaldegreebefore   Umathradicaldegreeraise  
 Umathradicalkern      Umathradicalrule      Umathradicalvgap      Umathrelbinspacing  
 Umathrelclosespacing   Umathrelinnerspacing   Umathrelopenspacing   Umathrelopspacing  
 Umathrelordspacing   Umathrelpunctspacing   Umathrelrelspacing   Umathspaceafterscript  
 Umathstackdenomdown      Umathstacknumup      Umathstackvgap      Umathsubshiftdown  
 Umathsubshiftdrop      Umathsubsupshiftdown      Umathsubsupvgap      Umathsubtopmax  
 Umathsupbottommin      Umathsupshiftdrop      Umathsupshiftup      Umathsupsubbottommax  
 Umathunderbarkern   Umathunderbarrule   Umathunderbarvgap   Umathunderdelimiterbgap  
 Umathunderdelimitervgap   Uoverdelimiter   Uradical   Uroot   Ustack   Ustartdisplaymath  
 Ustartmath   Ustopdisplaymath   Ustopmath   Usubscript   Usuperscript   Uunderdelimiter  
 alignmark   aligntab   attribute   attributedef   catcodetable   clearmarks   crampeddisplaystyle  
 crampedscriptscriptstyle   crampedscriptstyle   crampedtextstyle   discpenalty   fontid  
 formatname   gleaders   hyphenationmin   ifabsdim   ifabsnum   ifprimitive   initcatcodetable  
 latelua   luaescapestring   luafunction   luastartup   luatexdatestamp   luatexrevision  
 luatexversion   matheqnogapstep   mathstyle   nokerns   noligs   outputbox   pageleftoffset  
 pagetopoffset   postexhyphenchar   posthyphenchar   preexhyphenchar   prehyphenchar  
 primitive   savecatcodetable   scantextokens   suppressfontnotfounderror   suppressifcsnameerror  
 suppresslongerror   suppressmathparerror   suppressoutererror   syntex  
 umath   Udelcode   Udelcodenum   Udelimiter   Udelimiterover   Udelimiterunder  
 Umathaccent      Umathaxis      Umathbinbinspacing      Umathbinclosespacing  
 Umathbininnerspacing   Umathbinopenspacing   Umathbinopspacing   Umathbinordspacing  
 Umathbinpunctspacing   Umathbinrelspacing   Umathchar   Umathchardef   Umathcharnum  
 Umathcharnumdef   Umathclosebinspacing   Umathcloseclosespacing   Umathcloseinnerspacing  
 Umathcloseopenspacing   Umathcloseopspacing   Umathcloseordspacing   Umathclosepunctspacing  
 Umathcloserelspacing      Umathcode      Umathcodenum      Umathconnectoroverlapmin  
 Umathfractiondelsize   Umathfractiondenomdown   Umathfractiondenomvgap   Umathfractionnumup  
 Umathfractionnumvgap   Umathfractionrule   Umathinnerbinspacing   Umathinnerclosespacing  
 Umathinnerinnerspacing   Umathinneropenspacing   Umathinneropspacing   Umathinnerordspacing  
 Umathinnerpunctspacing   Umathinnerrelspacing   Umathlimitabovebgap   Umathlimitabovekern  
 Umathlimitabovevgap   Umathlimitbelowbgap   Umathlimitbelowkern   Umathlimitbelowvgap  
 Umathopbinspacing   Umathopclosespacing   Umathopenbinspacing   Umathopenclosespacing  
 Umathopeninnerspacing   Umathopenopenspacing   Umathopenopspacing   Umathopenordspacing  
 Umathopenpunctspacing   Umathopenrelspacing   Umathoperatorsize   Umathopinnerspacing  
 Umathopopenspacing      Umathopopspacing      Umathopordspacing      Umathoppunctspacing  
 Umathoprelspacing   Umathordbinspacing   Umathordclosespacing   Umathordinnerspacing  
 Umathordopenspacing   Umathordopspacing   Umathordordspacing   Umathordpunctspacing  
 Umathordrelspacing      Umathoverbarkern      Umathoverbarrule      Umathoverbarvgap  
 Umathoverdelimiterbgap      Umathoverdelimitervgap      Umathpunctbinspacing  
 Umathpunctclosespacing      Umathpunctinnerspacing      Umathpunctopenspacing  
 Umathpunctopspacing   Umathpunctordspacing   Umathpunctpunctspacing   Umathpunctrelspacing  
 Umathquad   Umathradicaldegreeafter   Umathradicaldegreebefore   Umathradicaldegreeraise  
 Umathradicalkern      Umathradicalrule      Umathradicalvgap      Umathrelbinspacing



Umathrelclosespacing	Umathrelinnerspacing	Umathreloopenspacing	Umathrelopspacing
Umathrelordspacing	Umathrelpunctspacing	Umathrelrelspacing	Umathspaceafterscript
Umathstackdenomdown	Umathstacknumup	Umathstackvgap	Umathsubshiftdown
Umathsubshiftdrop	Umathsubsupshiftdown	Umathsubsupvgap	Umathsubtopmax
Umathsupbottommin	Umathsupshiftdrop	Umathsupshiftup	Umathsupsubbottommax
Umathunderbarkern	Umathunderbarrule	Umathunderbarvgap	Umathunderdelimeterbgap
Umathunderdelimitervgap	Uoverdelimeter	Uradical	Uroot
Ustartmath	Ustopdisplaymath	Ustopmath	Usubscript
		Usuperscript	Uunderdelimeter

Note that `'luatex'` does not contain `directlua`, as that is considered to be a core primitive, along with all the T<sub>E</sub>X82 primitives, so it is part of the list that is returned from `'core'`.

`'umath'` is a subset of `'luatex'` that covers the Unicode math primitives and have been added in L<sup>A</sup>T<sub>E</sub>X 0.75.0 as it might be desired to handle the prefixing of that subset differently.

Running `tex.extraprimitives()` will give you the complete list of primitives that are not defined at L<sup>A</sup>T<sub>E</sub>X 0.39.0 `-ini` startup. It is exactly equivalent to `tex.extraprimitives('etex', 'pdf-tex', 'omega', 'aleph', 'luatex')`

#### 4.14.12.3 `tex.primitives`

```
<table> t = tex.primitives()
```

This function returns a hash table listing all primitives that L<sup>A</sup>T<sub>E</sub>X knows about. The keys in the hash are primitives names, the values are tables representing tokens (see [section 4.17](#)). The third value is always zero.

### 4.14.13 Core functionality interfaces

#### 4.14.13.1 `tex.badness (0.53)`

```
<number> b = tex.badness(<number> t, <number> s)
```

This helper function is useful during linebreak calculations. `t` and `s` are scaled values; the function returns the badness for when total `t` is supposed to be made from amounts that sum to `s`. The returned number is a reasonable approximation of  $100(t/s)^3$ ;

#### 4.14.13.2 `tex.linebreak (0.53)`

```
local <node> nodelist, <table> info =
    tex.linebreak(<node> listhead, <table> parameters)
```

The understood parameters are as follows:



name	type	description
pardir	string	
pretolerance	number	
tracingparagraphs	number	
tolerance	number	
looseness	number	
hyphenpenalty	number	
exhyphenpenalty	number	
pdfadjustspacing	number	
adjdemerits	number	
pdfprotrudechars	number	
linepenalty	number	
lastlinefit	number	
doublehyphendemerits	number	
finalhyphendemerits	number	
hangafter	number	
interlinepenalty	number or table	if a table, then it is an array like <code>\interlinepenalties</code>
clubpenalty	number or table	if a table, then it is an array like <code>\clubpenalties</code>
widowpenalty	number or table	if a table, then it is an array like <code>\widowpenalties</code>
brokenpenalty	number	
emergencystretch	number	in scaled points
hangindent	number	in scaled points
hsize	number	in scaled points
leftskip	glue_spec node	
rightskip	glue_spec node	
pdfignoreddimen	number	in scaled points
parshape	table	

Note that there is no interface for `\displaywidowpenalties`, you have to pass the right choice for `widowpenalties` yourself.

The meaning of the various keys should be fairly obvious from the table (the names match the T<sub>E</sub>X and P<sub>D</sub>F<sub>T</sub>E<sub>X</sub> primitives) except for the last 5 entries. The four `pdf...line...` keys are ignored if their value equals `pdfignoreddimen`.

It is your own job to make sure that `listhead` is a proper paragraph list: this function does not add any nodes to it. To be exact, if you want to replace the core line breaking, you may have to do the following (when you are not actually working in the `pre_linebreak_filter` or `linebreak_filter` callbacks, or when the original list starting at `listhead` was generated in horizontal mode):

- add an ‘indent box’ and perhaps a `local_par` node at the start (only if you need them)
- replace any found final glue by an infinite penalty (or add such a penalty, if the last node is not a glue)
- add a glue node for the `\parfillskip` after that penalty node
- make sure all the `prev` pointers are OK

The result is a node list, it still needs to be vpacked if you want to assign it to a `\vbox`.



The returned `info` table contains four values that are all numbers:

<code>prevdepth</code>	depth of the last line in the broken paragraph
<code>prevgraf</code>	number of lines in the broken paragraph
<code>looseness</code>	the actual looseness value in the broken paragraph
<code>demerits</code>	the total demerits of the chosen solution

Note there are a few things you cannot interface using this function: You cannot influence font expansion other than via `pdfadjustspacing`, because the settings for that take place elsewhere. The same is true for `hbadness` and `hfuzz` etc. All these are in the `hpack()` routine, and that fetches its own variables via `globals`.

### 4.14.13.3 `tex.shipout` (0.51)

`tex.shipout(<number> n)`

Ships out box number `n` to the output file, and clears the box register.

## 4.15 The `texconfig` table

This is a table that is created empty. A startup Lua script could fill this table with a number of settings that are read out by the executable after loading and executing the startup file.

key	type	default	explanation
<code>kpse_init</code>	boolean	true	<code>false</code> totally disables KPATHSEA initialisation, and enables interpretation of the following numeric key–value pairs. (only ever unset this if you implement <i>all</i> file find callbacks!)
<code>shell_escape</code>	string	'f'	Use 'y' or 't' or '1' to enable <code>\write 18</code> unconditionally, 'p' to enable the commands that are listed in <code>shell_escape_commands</code> (new in 0.37)
<code>shell_escape_commands</code>	string		Comma-separated list of command names that may be executed by <code>\write 18</code> even if <code>shell_escape</code> is set to 'p'. Do <i>not</i> use spaces around commas, separate any required command arguments by using a space, and use the ASCII double quote (") for any needed argument or path quoting (new in 0.37)
<code>string_vacancies</code>	number	75000	cf. web2c docs
<code>pool_free</code>	number	5000	cf. web2c docs
<code>max_strings</code>	number	15000	cf. web2c docs
<code>strings_free</code>	number	100	cf. web2c docs
<code>nest_size</code>	number	50	cf. web2c docs
<code>max_in_open</code>	number	15	cf. web2c docs
<code>param_size</code>	number	60	cf. web2c docs



<code>save_size</code>	number	4000	cf. web2c docs
<code>stack_size</code>	number	300	cf. web2c docs
<code>dvi_buf_size</code>	number	16384	cf. web2c docs
<code>error_line</code>	number	79	cf. web2c docs
<code>half_error_line</code>	number	50	cf. web2c docs
<code>max_print_line</code>	number	79	cf. web2c docs
<code>hash_extra</code>	number	0	cf. web2c docs
<code>pk_dpi</code>	number	72	cf. web2c docs
<code>trace_file_names</code>	boolean	true	<code>false</code> disables T <sub>E</sub> X's normal file open-close feedback (the assumption is that callbacks will take care of that)
<code>file_line_error</code>	boolean	false	do <code>file:line</code> style error messages
<code>halt_on_error</code>	boolean	false	abort run on the first encountered error
<code>formatname</code>	string		if no format name was given on the commandline, this key will be tested first instead of simply quitting
<code>jobname</code>	string		if no input file name was given on the commandline, this key will be tested first instead of simply giving up

**Note:** the numeric values that match web2c parameters are only used if `kpse_init` is explicitly set to `false`. In all other cases, the normal values from `texmf.cnf` are used.

## 4.16 The texio library

This library takes care of the low-level I/O interface.

### 4.16.1 Printing functions

#### 4.16.1.1 `texio.write`

```
texio.write(<string> target, <string> s, ...)
texio.write(<string> s, ...)
```

Without the `target` argument, writes all given strings to the same location(s) T<sub>E</sub>X writes messages to at this moment. If `\batchmode` is in effect, it writes only to the log, otherwise it writes to the log and the terminal. The optional `target` can be one of three possibilities: `term`, `log` or `term and log`.

Note: If several strings are given, and if the first of these strings is or might be one of the targets above, the `target` must be specified explicitly to prevent LUA from interpreting the first string as the target.

#### 4.16.1.2 `texio.write_nl`

```
texio.write_nl(<string> target, <string> s, ...)
texio.write_nl(<string> s, ...)
```



This function behaves like `texio.write`, but make sure that the given strings will appear at the beginning of a new line. You can pass a single empty string if you only want to move to the next line.

## 4.17 The token library

The `token` table contains interface functions to T<sub>E</sub>X's handling of tokens. These functions are most useful when combined with the `token_filter` callback, but they could be used standalone as well.

A token is represented in LUA as a small table. For the moment, this table consists of three numeric entries:

index	meaning	description
1	command code	this is a value between 0 and 130 (approximately)
2	command modifier	this is a value between 0 and 2 <sup>21</sup>
3	control sequence id	for commands that are not the result of control sequences, like letters and characters, it is zero, otherwise, it is a number pointing into the 'equivalence table'

### 4.17.1 `token.get_next`

```
token t = token.get_next()
```

This fetches the next input token from the current input source, without expansion.

### 4.17.2 `token.is_expandable`

```
<boolean> b = token.is_expandable(<token> t)
```

This tests if the token `t` could be expanded.

### 4.17.3 `token.expand`

```
token.expand(<token> t)
```

If a token is expandable, this will expand one level of it, so that the first token of the expansion will now be the next token to be read by `token.get_next()`.

### 4.17.4 `token.is_activechar`

```
<boolean> b = token.is_activechar(<token> t)
```

This is a special test that is sometimes handy. Discovering whether some control sequence is the result of an active character turned out to be very hard otherwise.



### 4.17.5 token.create

```
token t = token.create(<string> csname)
token t = token.create(<number> charcode)
token t = token.create(<number> charcode, <number> catcode)
```

This is the token factory. If you feed it a string, then it is the name of a control sequence (without leading backslash), and it will be looked up in the equivalence table.

If you feed it number, then this is assumed to be an input character, and an optional second number gives its category code. This means it is possible to overrule a character's category code, with a few exceptions: the category codes 0 (escape), 9 (ignored), 13 (active), 14 (comment), and 15 (invalid) cannot occur inside a token. The values 0, 9, 14 and 15 are therefore illegal as input to `token.create()`, and active characters will be resolved immediately.

Note: unknown string sequences and never defined active characters will result in a token representing an 'undefined control sequence' with a near-random name. It is *not* possible to define brand new control sequences using `token.create`!

### 4.17.6 token.command\_name

```
<string> commandname = token.command_name(<token> t)
```

This returns the name associated with the 'command' value of the token in L<sup>A</sup>T<sub>E</sub>X. There is not always a direct connection between these names and primitives. For instance, all `\ifxxx` tests are grouped under `if_test`, and the 'command modifier' defines which test is to be run.

### 4.17.7 token.command\_id

```
<number> i = token.command_id(<string> commandname)
```

This returns a number that is the inverse operation of the previous command, to be used as the first item in a token table.

### 4.17.8 token.csname\_name

```
<string> csname = token.csname_name(<token> t)
```

This returns the name associated with the 'equivalence table' value of the token in L<sup>A</sup>T<sub>E</sub>X. It returns the string value of the command used to create the current token, or an empty string if there is no associated control sequence.

Keep in mind that there are potentially two control sequences that return the same csname string: single character control sequences and active characters have the same 'name'.



## 4.17.9 token.csname\_id

```
<number> i = token.csname_id(<string> csname)
```

This returns a number that is the inverse operation of the previous command, to be used as the third item in a token table.

### 4.17.10 The newtoken library

The current `token` library will be replaced by a new one that is more flexible and powerful. The transition takes place in steps. In version 0.80 we have `newtoken` and in version 0.85 the old lib will be replaced completely. So if you use this new mechanism in production code you need to be aware of incompatible updates between 0.80 and 0.90. Because the related in- and output code will also be cleaned up and rewritten you should be aware of incompatible logging and error reporting too.

The old library presents tokens as triplets or numbers, the new library presents a userdata object. The old library used a callback to intercept tokens in the input but the new library provides a basic scanner infrastructure that can be used to write macros that accept a wide range of arguments. This interface is on purpose kept general and as performance is quite ok one can build additional parsers without too much overhead. It's up to macro package writers to see how they can benefit from this as the main principle behind L<sup>A</sup>T<sub>E</sub>X is to provide a minimal set of tools and no solutions.

The current functions in the `newtoken` namespace are given in the next table:

function	argument	result
is_token	token	checks if the given argument is a token userdata
get_next		returns the next token in the input
scan_keyword	string	returns true if the given keyword is gobbled
scan_int		returns a number
scan_dimen	infinity, mu-units	returns a number representing a dimension and or two numbers being the filler and order
scan_glue	mu-units	returns a glue spec node
scan_toks	definer, expand	returns a table of tokens token list (this can become a linked list in later releases)
scan_code	bitset	returns a character if its category is in the given bitset (representing catcodes)
scan_string		returns a string given between {}, as <code>\macro</code> or as sequence of characters with catcode 11 or 12
scan_word		returns a sequence of characters with catcode 11 or 12 as string
create		returns a userdata token object of the given control sequence name (or character); this interface can change

The scanners can be considered stable apart from the one scanning for a token. This is because futures releases can return a linked list instead of a table (as with nodes). The `scan_code` function takes an optional number, the `keyword` function a normal LUA string. The `infinity` boolean signals that we





also permit `fill` as dimension and the `mu-units` flags the scanner that we expect math units. When scanning tokens we can indicate that we are defining a macro, in which case the result will also provide information about what arguments are expected and in the result this is separated from the meaning by a separator token. The `expand` flag determines if the list will be expanded.

The string scanner scans for something between curly braces and expands on the way, or when it sees a control sequence it will return its meaning. Otherwise it will scan characters with catcode `letter` or `other`. So, given the following definition:

```
\def\bar{bar}
\def\foo{foo-\bar}
```

we get:

```
\directlua{newtoken.scan_string()}{foo}  foo      full expansion
\directlua{newtoken.scan_string()}foo     foo      letters and others
\directlua{newtoken.scan_string()}\foo    foo-bar   meaning
```

The `\foo` case only gives the meaning, but one can pass an already expanded definition (`\edef`'d). In the case of the braced variant one can of course use the `\detokenize` and `\unexpanded` primitives as there we do expand.

The `scan_word` scanner can be used to implement for instance a number scanner:

```
function newtokens.scan_number(base)
    return tonumber(newtoken.scan_word(),base)
end
```

This scanner accepts any valid LUA number so it is a way to pick up floats in the input.

The creator function can be used as follows:

```
local t = newtoken("relax")
```

This gives back a token object that has the properties of the `\relax` primitive. The possible properties of tokens are:

<code>command</code>	a number representing the internal command number
<code>cmdname</code>	the type of the command (for instance the catcode in case of a character or the classifier that determines the internal treatment)
<code>csname</code>	the associated control sequence (if applicable)
<code>id</code>	the unique id of the token
<code>active</code>	a boolean indicating the active state of the token
<code>expandable</code>	a boolean indicating if the token (macro) is expandable
<code>protected</code>	a boolean indicating if the token (macro) is protected

The numbers that represent a catcode are the same as in T<sub>E</sub>X itself, so using this information assumes that you know a bit about T<sub>E</sub>X's internals. The other numbers and names are used consistently but are



not frozen. So, when you use them for comparing you can best query a known primitive or character first to see the values.

More interesting are the scanners. You can use the LUA interface as follows:

```
\directlua {  
    function mymacro(n)  
        ...  
    end  
}  
  
\def\mymacro#1{%  
    \directlua {  
        mymacro(\number\dimexpr#1)  
    }%  
}  
  
\mymacro{12pt}  
\mymacro{\dimen0}
```

But starting with version 0.80 you can also do this:

```
\directlua {  
    function mymacro()  
        local d = newtoken.scan_dimen()  
        ...  
    end  
}  
  
\def\mymacro{%  
    \directlua {  
        mymacro()  
    }%  
}  
  
\mymacro 12pt  
\mymacro \dimen0
```

It is quite clear from looking at the code what the first method needs as argument(s). For the second method you need to look at the LUA code to see what gets picked up. Instead of passing from T<sub>E</sub>X to LUA we let LUA fetch from the input stream.

In the first case the input is tokenized and then turned into a string when it's passed to LUA where it gets interpreted. In the second case only a function call gets interpreted but then the input is picked up by explicitly calling the scanner functions. These return proper LUA variables so no further conversion has to be done. This is more efficient but in practice (given what T<sub>E</sub>X has to do) this effect should not be overestimated. For numbers and dimensions it saves a bit but for passing strings conversion to and



from tokens has to be done anyway (although we can probably speed up the process in later versions if needed).

When the interface is stable and has replaced the old one completely we will add some more information here. By that time the internals have been cleaned up a bit more so we know then what will stay and go. A positive side effect of this transition is that we can simplify the input part because we no longer need to intercept using callbacks.





## 5 Math

The handling of mathematics in L<sup>A</sup>T<sub>E</sub>X differs quite a bit from how T<sub>E</sub>X82 (and therefore P<sub>D</sub>F<sub>T</sub>E<sub>X</sub>) handles math. First, L<sup>A</sup>T<sub>E</sub>X adds primitives and extends some others so that U<sup>N</sup>I<sup>C</sup>O<sup>D</sup>E input can be used easily. Second, all of T<sub>E</sub>X82's internal special values (for example for operator spacing) have been made accessible and changeable via control sequences. Third, there are extensions that make it easier to use O<sup>P</sup>E<sup>N</sup>T<sup>Y</sup>P<sup>E</sup> math fonts. And finally, there are some extensions that have been proposed in the past that are now added to the engine.

### 5.1 The current math style

Starting with L<sup>A</sup>T<sub>E</sub>X 0.39.0, it is possible to discover the math style that will be used for a formula in an expandable fashion (while the math list is still being read). To make this possible, L<sup>A</sup>T<sub>E</sub>X adds the new primitive: `\mathstyle`. This is a 'convert command' like e.g. `\romannumeral`: its value can only be read, not set.

#### 5.1.1 `\mathstyle`

The returned value is between 0 and 7 (in math mode), or  $-1$  (all other modes). For easy testing, the eight math style commands have been altered so that they can be used as numeric values, so you can write code like this:

```
\ifnum\mathstyle=\textstyle
  \message{normal text style}
\else \ifnum\mathstyle=\crampedtextstyle
  \message{cramped text style}
\fi \fi
```

#### 5.1.2 `\Ustack`

There are a few math commands in T<sub>E</sub>X where the style that will be used is not known straight from the start. These commands (`\over`, `\atop`, `\overwithdelims`, `\atopwithdelims`) would therefore normally return wrong values for `\mathstyle`. To fix this, L<sup>A</sup>T<sub>E</sub>X introduces a special prefix command: `\Ustack`:

```
 $\Ustack {a \over b}$ 
```

The `\Ustack` command will scan the next brace and start a new math group with the correct (numerator) math style.



## 5.2 Unicode math characters

Character handling is now extended up to the full UNICODE range (the `\U` prefix), which is compatible with  $\XTeX$ .

The math primitives from  $\TeX$  are kept as they are, except for the ones that convert from input to math commands: `\mathcode`, and `\delcode`. These two now allow for a 21-bit character argument on the left hand side of the equals sign.

Some of the new  $\text{LUA}\TeX$  primitives read more than one separate value. This is shown in the tables below by a plus sign in the second column.

The input for such primitives would look like this:

```
\def\overbrace {\Umathaccent 0 1 "23DE }
```

Altered  $\text{TEX}82$  primitives:

primitive	value range (in hex)
<code>\mathcode</code>	0–10FFFF = 0–8000
<code>\delcode</code>	0–10FFFF = 0–FFFFFF

Unaltered:

primitive	value range (in hex)
<code>\mathchardef</code>	0–8000
<code>\mathchar</code>	0–7FFF
<code>\mathaccent</code>	0–7FFF
<code>\delimiter</code>	0–7FFFFFFF
<code>\radical</code>	0–7FFFFFFF

New primitives that are compatible with  $\XTeX$ :

primitive	value range (in hex)
<code>\Umathchardef</code>	0+0+0–7+FF+10FFFF <sup>1</sup>
<code>\Umathcharnumdef</code> <sup>5</sup>	–80000000–7FFFFFFF <sup>3</sup>
<code>\Umathcode</code>	0–10FFFF = 0+0+0–7+FF+10FFFF <sup>1</sup>
<code>\Udelcode</code>	0–10FFFF = 0+0–FF+10FFFF <sup>2</sup>
<code>\Umathchar</code>	0+0+0–7+FF+10FFFF
<code>\Umathaccent</code>	0+0+0–7+FF+10FFFF <sup>2,4</sup>
<code>\Udelimiter</code>	0+0+0–7+FF+10FFFF <sup>2</sup>
<code>\Uradical</code>	0+0–FF+10FFFF <sup>2</sup>
<code>\Umathcharnum</code>	–80000000–7FFFFFFF <sup>3</sup>
<code>\Umathcodenum</code>	0–10FFFF = –80000000–7FFFFFFF <sup>3</sup>
<code>\Udelcodenum</code>	0–10FFFF = –80000000–7FFFFFFF <sup>3</sup>

Note 1: `\Umathchardef<csname>="8"0"0` and `\Umathchardef<number>="8"0"0` are also accepted.



Note 2: The new primitives that deal with delimiter-style objects do not set up a ‘large family’. Selecting a suitable size for display purposes is expected to be dealt with by the font via the `\Umathoperator-size` parameter (more information a following section).

Note 3: For these three primitives, all information is packed into a single signed integer. For the first two (`\Umathcharnum` and `\Umathcodenum`), the lowest 21 bits are the character code, the 3 bits above that represent the math class, and the family data is kept in the topmost bits (This means that the values for math families 128–255 are actually negative). For `\Udelcodenum` there is no math class; the math family information is stored in the bits directly on top of the character code. Using these three commands is not as natural as using the two- and three-value commands, so unless you know exactly what you are doing and absolutely require the speedup resulting from the faster input scanning, it is better to use the verbose commands instead.

Note 4: As of L<sup>A</sup>T<sub>E</sub>X 0.65, `\Umathaccent` accepts optional keywords to control various details regarding math accents. See [section 5.7](#) below for details.

Note 5: `\Umathcharnumdef` was added in release 0.72.

New primitives that exist in L<sup>A</sup>T<sub>E</sub>X only (all of these will be explained in following sections):

primitive	value range (in hex)
<code>\Uroot</code>	0+0–FF+10FFFF <sup>2</sup>
<code>\Uoverdelimiter</code>	0+0–FF+10FFFF <sup>2</sup>
<code>\Uunderdelimiter</code>	0+0–FF+10FFFF <sup>2</sup>
<code>\Udelimiterover</code>	0+0–FF+10FFFF <sup>2</sup>
<code>\Udelimiterunder</code>	0+0–FF+10FFFF <sup>2</sup>

## 5.3 Cramped math styles

L<sup>A</sup>T<sub>E</sub>X has four new primitives to set the cramped math styles directly:

```
\crampeddisplaystyle
\crampedtextstyle
\crampedscriptstyle
\crampedscriptscriptstyle
```

These additional commands are not all that valuable on their own, but they come in handy as arguments to the math parameter settings that will be added shortly.

## 5.4 Math parameter settings

In L<sup>A</sup>T<sub>E</sub>X, the font dimension parameters that T<sub>E</sub>X used in math typesetting are now accessible via primitive commands. In fact, refactoring of the math engine has resulted in many more parameters than were accessible before.

primitive name	description
<code>\Umathquad</code>	the width of 18mu’s



<code>\Umathaxis</code>	height of the vertical center axis of the math formula above the baseline
<code>\Umathoperatorsize</code>	minimum size of large operators in display mode
<code>\Umathoverbarkern</code>	vertical clearance above the rule
<code>\Umathoverbarrule</code>	the width of the rule
<code>\Umathoverbarvgap</code>	vertical clearance below the rule
<code>\Umathunderbarkern</code>	vertical clearance below the rule
<code>\Umathunderbarrule</code>	the width of the rule
<code>\Umathunderbarvgap</code>	vertical clearance above the rule
<code>\Umathradicalkern</code>	vertical clearance above the rule
<code>\Umathradicalrule</code>	the width of the rule
<code>\Umathradicalvgap</code>	vertical clearance below the rule
<code>\Umathradicaldegreebefore</code>	the forward kern that takes place before placement of the radical degree
<code>\Umathradicaldegreeafter</code>	the backward kern that takes place after placement of the radical degree
<code>\Umathradicaldegreeraise</code>	this is the percentage of the total height and depth of the radical sign that the degree is raised by. It is expressed in <b>percents</b> , so 60% is expressed as the integer 60.
<code>\Umathstackvgap</code>	vertical clearance between the two elements in a <code>\atop</code> stack
<code>\Umathstacknumup</code>	numerator shift upward in <code>\atop</code> stack
<code>\Umathstackdenomdown</code>	denominator shift downward in <code>\atop</code> stack
<code>\Umathfractionrule</code>	the width of the rule in a <code>\over</code>
<code>\Umathfractionnumvgap</code>	vertical clearance between the numerator and the rule
<code>\Umathfractionnumup</code>	numerator shift upward in <code>\over</code>
<code>\Umathfractiondenomvgap</code>	vertical clearance between the denominator and the rule
<code>\Umathfractiondenomdown</code>	denominator shift downward in <code>\over</code>
<code>\Umathfractiondelsize</code>	minimum delimiter size for <code>\dotswithdelims</code>
<code>\Umathlimitabovevgap</code>	vertical clearance for limits above operators
<code>\Umathlimitabovebgap</code>	vertical baseline clearance for limits above operators
<code>\Umathlimitabovekern</code>	space reserved at the top of the limit
<code>\Umathlimitbelowvgap</code>	vertical clearance for limits below operators
<code>\Umathlimitbelowbgap</code>	vertical baseline clearance for limits below operators
<code>\Umathlimitbelowkern</code>	space reserved at the bottom of the limit
<code>\Umathoverdelimitervgap</code>	vertical clearance for limits above delimiters
<code>\Umathoverdelimiterbgap</code>	vertical baseline clearance for limits above delimiters
<code>\Umathunderdelimitervgap</code>	vertical clearance for limits below delimiters
<code>\Umathunderdelimiterbgap</code>	vertical baseline clearance for limits below delimiters
<code>\Umathsubshiftdrop</code>	subscript drop for boxes and subformulas
<code>\Umathsubshiftdown</code>	subscript drop for characters
<code>\Umathsupshiftdrop</code>	superscript drop (raise, actually) for boxes and subformulas
<code>\Umathsupshiftdown</code>	superscript raise for characters
<code>\Umathsubsupshiftdown</code>	subscript drop in the presence of a superscript





<code>\Umathsubtopmax</code>	the top of standalone subscripts cannot be higher than this above the baseline
<code>\Umathsupbottommin</code>	the bottom of standalone superscripts cannot be less than this above the baseline
<code>\Umathsupsubbottommax</code>	the bottom of the superscript of a combined super- and subscript be at least as high as this above the baseline
<code>\Umathsubsupvgap</code>	vertical clearance between super- and subscript
<code>\Umathspaceafterscript</code>	additional space added after a super- or subscript
<code>\Umathconnectoroverlapmin</code>	minimum overlap between parts in an extensible recipe

Each of the parameters in this section can be set by a command like this:

```
\Umathquad\displaystyle=1em
```

they obey grouping, and you can use `\the\Umathquad\displaystyle` if needed.

## 5.5 Font-based Math Parameters

While it is nice to have these math parameters available for tweaking, it would be tedious to have to set each of them by hand. For this reason, L<sup>A</sup>T<sub>E</sub>X initializes a bunch of these parameters whenever you assign a font identifier to a math family based on either the traditional math font dimensions in the font (for assignments to math family 2 and 3 using TFM-based fonts like `cmsy` and `cmex`), or based on the named values in a potential `MathConstants` table when the font is loaded via Lua. If there is a `MathConstants` table, this takes precedence over font dimensions, and in that case no attention is paid to which family is being assigned to: the `MathConstants` tables in the last assigned family sets all parameters.

In the table below, the one-letter style abbreviations and symbolic tfm font dimension names match those using in the T<sub>E</sub>Xbook. Assignments to `\textfont` set the values for the cramped and uncramped display and text styles. Use `\scriptfont` for the script styles, and `\scriptscriptfont` for the scriptscript styles (totalling eight parameters for three font sizes). In the TFM case, assignments only happen in family 2 and family 3 (and of course only for the parameters for which there are font dimensions).

Besides the parameters below, L<sup>A</sup>T<sub>E</sub>X also looks at the ‘space’ font dimension parameter. For math fonts, this should be set to zero.

variable	style	default value opentype	default value tfm
<code>\Umathaxis</code>	–	AxisHeight	axis_height
<code>\Umathoperatorsiz</code>	D, D’	DisplayOperatorMinHeight	<sup>6</sup>
<code>\Umathfractiondelsize</code>	D, D’	FractionDelimiterDisplayStyleSize <sup>9</sup>	delim1
"	T, T’, S, S’, SS, SS’	FractionDelimiterSize <sup>9</sup>	delim2
<code>\Umathfractiondenomdown</code>	D, D’	FractionDenominatorDisplayStyleShiftDown	denom1
"	T, T’, S, S’, SS, SS’	FractionDenominatorShiftDown	denom2
<code>\Umathfractiondenomvgap</code>	D, D’	FractionDenominatorDisplayStyleGapMin	3*default_rule_thickness
"	T, T’, S, S’, SS, SS’	FractionDenominatorGapMin	default_rule_thickness
<code>\Umathfractionnumup</code>	D, D’	FractionNumeratorDisplayStyleShiftUp	num1
"	T, T’, S, S’, SS, SS’	FractionNumeratorShiftUp	num2
<code>\Umathfractionnumvgap</code>	D, D’	FractionNumeratorDisplayStyleGapMin	3*default_rule_thickness
"	T, T’, S, S’, SS, SS’	FractionNumeratorGapMin	default_rule_thickness
<code>\Umathfractionrule</code>	–	FractionRuleThickness	default_rule_thickness



<code>\Umathlimitabovebgap</code>	–	UpperLimitBaselineRiseMin	<code>big_op_spacing3</code>
<code>\Umathlimitabovekern</code>	–	$0^1$	<code>big_op_spacing5</code>
<code>\Umathlimitabovevgap</code>	–	UpperLimitGapMin	<code>big_op_spacing1</code>
<code>\Umathlimitbelowbgap</code>	–	LowerLimitBaselineDropMin	<code>big_op_spacing4</code>
<code>\Umathlimitbelowkern</code>	–	$0^1$	<code>big_op_spacing5</code>
<code>\Umathlimitbelowvgap</code>	–	LowerLimitGapMin	<code>big_op_spacing2</code>
<code>\Umathoverdelimitervgap</code>	–	StretchStackGapBelowMin	<code>big_op_spacing1</code>
<code>\Umathoverdelimiterbgap</code>	–	StretchStackTopShiftUp	<code>big_op_spacing3</code>
<code>\Umathunderdelimitervgap</code>	–	StretchStackGapAboveMin	<code>big_op_spacing2</code>
<code>\Umathunderdelimeterbgap</code>	–	StretchStackBottomShiftDown	<code>big_op_spacing4</code>
<code>\Umathoverbarkern</code>	–	OverbarExtraAscender	<code>default_rule_thickness</code>
<code>\Umathoverbarrule</code>	–	OverbarRuleThickness	<code>default_rule_thickness</code>
<code>\Umathoverbarvgap</code>	–	OverbarVerticalGap	$3 * \text{default\_rule\_thickness}$
<code>\Umathquad</code>	–	$\langle \text{font\_size}(f) \rangle^1$	<code>math_quad</code>
<code>\Umathradicalkern</code>	–	RadicalExtraAscender	<code>default_rule_thickness</code>
<code>\Umathradicalrule</code>	–	RadicalRuleThickness	$\langle \text{not set} \rangle^2$
<code>\Umathradicalvgap</code>	D, D'	RadicalDisplayStyleVerticalGap	$(\text{default\_rule\_thickness} + (\text{abs}(\text{math\_x\_height})/4))^3$
"	T, T', S, S', SS, SS'	RadicalVerticalGap	$(\text{default\_rule\_thickness} + (\text{abs}(\text{default\_rule\_thickness})/4))^3$
<code>\Umathradicaldegreebefore</code>	–	RadicalKernBeforeDegree	$\langle \text{not set} \rangle^2$
<code>\Umathradicaldegreeafter</code>	–	RadicalKernAfterDegree	$\langle \text{not set} \rangle^2$
<code>\Umathradicaldegreeraise</code>	–	RadicalDegreeBottomRaisePercent	$\langle \text{not set} \rangle^{2,7}$
<code>\Umathspaceafterscript</code>	–	SpaceAfterScript	<code>script_space</code> <sup>4</sup>
<code>\Umathstackdenomdown</code>	D, D'	StackBottomDisplayStyleShiftDown	<code>denom1</code>
"	T, T', S, S', SS, SS'	StackBottomShiftDown	<code>denom2</code>
<code>\Umathstacknumup</code>	D, D'	StackTopDisplayStyleShiftUp	<code>num1</code>
"	T, T', S, S', SS, SS'	StackTopShiftUp	<code>num3</code>
<code>\Umathstackvgap</code>	D, D'	StackDisplayStyleGapMin	$7 * \text{default\_rule\_thickness}$
"	T, T', S, S', SS, SS'	StackGapMin	$3 * \text{default\_rule\_thickness}$
<code>\Umathsubshiftdown</code>	–	SubscriptShiftDown	<code>sub1</code>
<code>\Umathsubshiftdrop</code>	–	SubscriptBaselineDropMin	<code>sub_drop</code>
<code>\Umathsubsupshiftdown</code>	–	SubscriptShiftDownWithSuperscript <sup>8</sup> or SubscriptShiftDown	<code>sub2</code>
<code>\Umathsubtopmax</code>	–	SubscriptTopMax	$(\text{abs}(\text{math\_x\_height} * 4) / 5)$
<code>\Umathsubsupvgap</code>	–	SubSuperscriptGapMin	$4 * \text{default\_rule\_thickness}$
<code>\Umathsupbottommin</code>	–	SuperscriptBottomMin	$(\text{abs}(\text{math\_x\_height}) / 4)$
<code>\Umathsupshiftdrop</code>	–	SuperscriptBaselineDropMax	<code>sup_drop</code>
<code>\Umathsupshiftup</code>	D	SuperscriptShiftUp	<code>sup1</code>
"	T, S, SS,	SuperscriptShiftUp	<code>sup2</code>
"	D', T', S', SS'	SuperscriptShiftUpCramped	<code>sup3</code>
<code>\Umathsupsubbottommax</code>	–	SuperscriptBottomMaxWithSubscript	$(\text{abs}(\text{math\_x\_height} * 4) / 5)$
<code>\Umathunderbarkern</code>	–	UnderbarExtraDescender	<code>default_rule_thickness</code>
<code>\Umathunderbarrule</code>	–	UnderbarRuleThickness	<code>default_rule_thickness</code>
<code>\Umathunderbarvgap</code>	–	UnderbarVerticalGap	$3 * \text{default\_rule\_thickness}$
<code>\Umathconnectoroverlapmin</code>	–	MinConnectorOverlap	$0^5$

Note 1: OPENType fonts set `\Umathlimitabovekern` and `\Umathlimitbelowkern` to zero and set `\Umathquad` to the font size of the used font, because these are not supported in the MATH table,

Note 2: TFM fonts do not set `\Umathradicalrule` because T<sub>E</sub>X82 uses the height of the radical instead. When this parameter is indeed not set when L<sup>A</sup>T<sub>E</sub>X has to typeset a radical, a backward compatibility mode will kick in that assumes that an oldstyle T<sub>E</sub>X font is used. Also, they do not set `\Umathradicaldegreebefore`, `\Umathradicaldegreeafter`, and `\Umathradicaldegreeraise`. These are then automatically initialized to 5/18quad, –10/18quad, and 60.

Note 3: If tfm fonts are used, then the `\Umathradicalvgap` is not set until the first time L<sup>A</sup>T<sub>E</sub>X has to typeset a formula because this needs parameters from both family2 and family3. This provides a partial



backward compatibility with T<sub>E</sub>X82, but that compatibility is only partial: once the `\Umathradicalv-gap` is set, it will not be recalculated any more.

Note 4: (also if tfm fonts are used) A similar situation arises wrt. `\Umathspaceafterscript`: it is not set until the first time L<sup>A</sup>T<sub>E</sub>X has to typeset a formula. This provides some backward compatibility with T<sub>E</sub>X82. But once the `\Umathspaceafterscript` is set, `\scriptspace` will never be looked at again.

Note 5: Tfm fonts set `\Umathconnectoroverlapmin` to zero because T<sub>E</sub>X82 always stacks extensibles without any overlap.

Note 6: The `\Umathoperatorsiz` is only used in `\displaystyle`, and is only set in O<sub>P</sub>E<sub>N</sub>T<sub>E</sub><sub>P</sub>E fonts. In T<sub>F</sub>M font mode, it is artificially set to one scaled point more than the initial attempt's size, so that always the 'first next' will be tried, just like in T<sub>E</sub>X82.

Note 7: The `\Umathradicaldegreeraise` is a special case because it is the only parameter that is expressed in a percentage instead of as a number of scaled points.

Note 8: `SubscriptShiftDownWithSuperscript` does not actually exist in the 'standard' Opentype Math font Cambria, but it is useful enough to be added. New in version 0.38.

Note 9: `FractionDelimiterDisplayStyleSize` and `FractionDelimiterSize` do not actually exist in the 'standard' Opentype Math font Cambria, but were useful enough to be added. New in version 0.47.

## 5.6 Math spacing setting

Besides the parameters mentioned in the previous sections, there are also 64 new primitives to control the math spacing table (as explained in Chapter 18 of the T<sub>E</sub>Xbook). The primitive names are a simple matter of combining two math atom types, but for completeness' sake, here is the whole list:

<code>\Umathordordspacing</code>	<code>\Umathbinopspacing</code>
<code>\Umathordopspacing</code>	<code>\Umathbinbinspacing</code>
<code>\Umathordbinspacing</code>	<code>\Umathbinrelspacing</code>
<code>\Umathordrelspacing</code>	<code>\Umathbinopenspacing</code>
<code>\Umathordopenspacing</code>	<code>\Umathbinclonespacing</code>
<code>\Umathordclonespacing</code>	<code>\Umathbinpunctspacing</code>
<code>\Umathordpunctspacing</code>	<code>\Umathbininnerspacing</code>
<code>\Umathordinnerspacing</code>	<code>\Umathrelordspacing</code>
<code>\Umathopordspacing</code>	<code>\Umathrelopspacing</code>
<code>\Umathopopspacing</code>	<code>\Umathrelbinspacing</code>
<code>\Umathopbinspacing</code>	<code>\Umathrelrelspacing</code>
<code>\Umathoprelspacing</code>	<code>\Umathreloppspacing</code>
<code>\Umathopopenspacing</code>	<code>\Umathrelclonespacing</code>
<code>\Umathopclonespacing</code>	<code>\Umathrelpunctspacing</code>
<code>\Umathoppunctspacing</code>	<code>\Umathrelinnerspacing</code>
<code>\Umathopinnerspacing</code>	<code>\Umathopenordspacing</code>
<code>\Umathbinordspacing</code>	<code>\Umathopenopspacing</code>



<code>\Umathopenbinspacing</code>	<code>\Umathpuncttopspacing</code>
<code>\Umathopenrelspacing</code>	<code>\Umathpunctbinspacing</code>
<code>\Umathopenopenspacing</code>	<code>\Umathpunctrelspacing</code>
<code>\Umathopenclosespacing</code>	<code>\Umathpunctopenspacing</code>
<code>\Umathopenpunctspacing</code>	<code>\Umathpunctclosespacing</code>
<code>\Umathopeninnerspacing</code>	<code>\Umathpunctpunctspacing</code>
<code>\Umathcloseordspacing</code>	<code>\Umathpunctinnerspacing</code>
<code>\Umathcloseopspacing</code>	<code>\Umathinnerordspacing</code>
<code>\Umathclosebinspacing</code>	<code>\Umathinneropspacing</code>
<code>\Umathcloserelspacing</code>	<code>\Umathinnerbinspacing</code>
<code>\Umathcloseopenspacing</code>	<code>\Umathinnerrelspacing</code>
<code>\Umathcloseclosespacing</code>	<code>\Umathinneropenspacing</code>
<code>\Umathclosepunctspacing</code>	<code>\Umathinnerclosespacing</code>
<code>\Umathcloseinnerspacing</code>	<code>\Umathinnerpunctspacing</code>
<code>\Umathpunctordspacing</code>	<code>\Umathinnerinnerspacing</code>

These parameters are of type `\muskip`, so setting a parameter can be done like this:

```
\Umathopordspacing\displaystyle=4mu plus 2mu
```

They are all initialized by `initex` to the values mentioned in the table in Chapter 18 of the `TEXbook`.

Note 1: for ease of use as well as for backward compatibility, `\thinmuskip`, `\medmuskip` and `\thickmuskip` are treated especially. In their case a pointer to the corresponding internal parameter is saved, not the actual `\muskip` value. This means that any later changes to one of these three parameters will be taken into account.

Note 2: Careful readers will realise that there are also primitives for the items marked *\** in the `TEXbook`. These will not actually be used as those combinations of atoms cannot actually happen, but it seemed better not to break orthogonality. They are initialized to zero.

## 5.7 Math accent handling

`LUATEX` supports both top accents and bottom accents in math mode, and math accents stretch automatically (if this is supported by the font the accent comes from, of course). Bottom and combined accents as well as fixed-width math accents are controlled by optional keywords following `\Umathaccent`.

The keyword `bottom` after `\Umathaccent` signals that a bottom accent is needed, and the keyword `both` signals that both a top and a bottom accent are needed (in this case two accents need to be specified, of course).

Then the set of three integers defining the accent is read. This set of integers can be prefixed by the `fixed` keyword to indicate that a non-stretching variant is requested (in case of both accents, this step is repeated).

A simple example:



```
\Umathaccent both fixed 0 0 "20D7 fixed 0 0 "20D7 {example}
```

If a math top accent has to be placed and the accentee is a character and has a non-zero `top_accent` value, then this value will be used to place the accent instead of the `\skewchar` kern used by T<sub>E</sub>X82.

The `top_accent` value represents a vertical line somewhere in the accentee. The accent will be shifted horizontally such that its own `top_accent` line coincides with the one from the accentee. If the `top_accent` value of the accent is zero, then half the width of the accent followed by its italic correction is used instead.

The vertical placement of a top accent depends on the `x_height` of the font of the accentee (as explained in the T<sub>E</sub>Xbook), but if value that turns out to be zero and the font had a `MathConstants` table, then `AccentBaseHeight` is used instead.

If a math bottom accent has to be placed, the `bot_accent` value is checked instead of `top_accent`. Because bottom accents do not exist in T<sub>E</sub>X82, the `\skewchar` kern is ignored.

The vertical placement of a bottom accent is straight below the accentee, no correction takes place.

## 5.8 Math root extension

The new primitive `\Uroot` allows the construction of a radical noad including a degree field. Its syntax is an extension of `\Uradical`:

```
\Uradical <fam integer> <char integer> <radicand>
\Uroot    <fam integer> <char integer> <degree> <radicand>
```

The placement of the degree is controlled by the math parameters `\Umathradicaldegreebefore`, `\Umathradicaldegreeafter`, and `\Umathradicaldegreeraise`. The degree will be typeset in `\scriptscriptstyle`.

## 5.9 Math kerning in super- and subscripts

The character fields in a lua-loaded OpenType math font can have a ‘mathkern’ table. The format of this table is the same as the ‘mathkern’ table that is returned by the `fontloader` library, except that all height and kern values have to be specified in actual scaled points.

When a super- or subscript has to be placed next to a math item, L<sup>A</sup>T<sub>E</sub>X checks whether the super- or subscript and the nucleus are both simple character items. If they are, and if the fonts of both character imtes are OpenType fonts (as opposed to legacy T<sub>E</sub>X fonts), then L<sup>A</sup>T<sub>E</sub>X will use the OpenType MATH algorithm for deciding on the horizontal placement of the super- or subscript.

This works as follows:

- The vertical position of the script is calculated.
- The default horizontal position is flat next to the base character.
- For superscripts, the italic correction of the base character is added.



- For a superscript, two vertical values are calculated: the bottom of the script (after shifting up), and the top of the base. For a subscript, the two values are the top of the (shifted down) script, and the bottom of the base.
- For each of these two locations:
  - find the mathkern value at this height for the base (for a subscript placement, this is the bottom\_right corner, for a superscript placement the top\_right corner)
  - find the mathkern value at this height for the script (for a subscript placement, this is the top\_left corner, for a superscript placement the bottom\_left corner)
  - add the found values together to get a preliminary result.
- The horizontal kern to be applied is the smallest of the two results from previous step.

The mathkern value at a specific height is the kern value that is specified by the next higher height and kern pair, or the highest one in the character (if there is no value high enough in the character), or simply zero (if the character has no mathkern pairs at all).

## 5.10 Scripts on horizontally extensible items like arrows

The new primitives `\Uunderdelimater` and `\Uoverdelimater` (both from 0.35) allow the placement of a subscript or superscript on an automatically extensible item and `\Udelimaterunder` and `\Udelimaterover` (both from 0.37) allow the placement of an automatically extensible item as a subscript or superscript on a nucleus.

The vertical placements are controlled by `\Umathunderdelimaterbgap`, `\Umathunderdelimatervgap`, `\Umathoverdelimaterbgap`, and `\Umathoverdelimatervgap` in a similar way as limit placements on large operators. The superscript in `\Uoverdelimater` is typeset in a suitable scripted style, the subscript in `\Uunderdelimater` is cramped as well.

## 5.11 Extensible delimiters

LUAT<sub>E</sub>X internally uses a structure that supports OPEN<sub>T</sub>YPE ‘MathVariants’ as well as TFM ‘extensible recipes’.

## 5.12 Other Math changes

### 5.12.1 Verbose versions of single-character math commands

LUAT<sub>E</sub>X defines six new primitives that have the same function as `^`, `_`, `$`, and `$$`.

primitive	explanation
<code>\Usuperscript</code>	Duplicates the functionality of <code>^</code>
<code>\Usubscript</code>	Duplicates the functionality of <code>_</code>
<code>\Ustartmath</code>	Duplicates the functionality of <code>\$</code> , when used in non-math mode.



<code>\Ustopmath</code>	Duplicates the functionality of <code>\$</code> , when used in inline math mode.
<code>\Ustartdisplaymath</code>	Duplicates the functionality of <code>\$\$</code> , when used in non-math mode.
<code>\Ustopdisplaymath</code>	Duplicates the functionality of <code>\$\$</code> , when used in display math mode.

All are new in version 0.38. The `\Ustopmath` and `\Ustopdisplaymath` primitives check if the current math mode is the correct one (inline vs. displayed), but you can freely intermix the four `mathon/mathoff` commands with explicit dollar sign(s).

### 5.12.2 Allowed math commands in non-math modes

The commands `\mathchar`, and `\Umathchar` and control sequences that are the result of `\mathchardef` or `\Umathchardef` are also acceptable in the horizontal and vertical modes. In those cases, the `\textfont` from the requested math family is used.

## 5.13 Math todo

The following items are still todo.

- Pre-scripts.
- Multi-story stacks.
- Flattened accents for high characters (?).
- Better control over the spacing around displays and handling of equation numbers.
- Support for multi-line displays using `MATHTML` style alignment points.







## 6 Languages and characters, fonts and glyphs

L<sup>A</sup>T<sub>E</sub>X's internal handling of the characters and glyphs that eventually become typeset is quite different from the way T<sub>E</sub>X82 handles those same objects. The easiest way to explain the difference is to focus on unrestricted horizontal mode (i. e. paragraphs) and hyphenation first. Later on, it will be easy to deal with the differences that occur in horizontal and math modes.

In T<sub>E</sub>X82, the characters you type are converted into `char_node` records when they are encountered by the main control loop. T<sub>E</sub>X attaches and processes the font information while creating those records, so that the resulting 'horizontal list' contains the final forms of ligatures and implicit kerning. This packaging is needed because we may want to get the effective width of for instance a horizontal box.

When it becomes necessary to hyphenate words in a paragraph, T<sub>E</sub>X converts (one word at time) the `char_node` records into a string array by replacing ligatures with their components and ignoring the kerning. Then it runs the hyphenation algorithm on this string, and converts the hyphenated result back into a 'horizontal list' that is consecutively spliced back into the paragraph stream. Keep in mind that the paragraph may contain unboxed horizontal material, which then already contains ligatures and kerns and the words therein are part of the hyphenation process.

The `char_node` records are somewhat misnamed, as they are glyph positions in specific fonts, and therefore not really 'characters' in the linguistic sense. There is no language information inside the `char_node` records. Instead, language information is passed along using `language whatsit` records inside the horizontal list.

In L<sup>A</sup>T<sub>E</sub>X, the situation is quite different. The characters you type are always converted into `glyph_node` records with a special subtype to identify them as being intended as linguistic characters. L<sup>A</sup>T<sub>E</sub>X stores the needed language information in those records, but does not do any font-related processing at the time of node creation. It only stores the index of the font.

When it becomes necessary to typeset a paragraph, L<sup>A</sup>T<sub>E</sub>X first inserts all hyphenation points right into the whole node list. Next, it processes all the font information in the whole list (creating ligatures and adjusting kerning), and finally it adjusts all the subtype identifiers so that the records are 'glyph nodes' from now on.

That was the broad overview. The rest of this chapter will deal with the minutiae of the new process.

### 6.1 Characters and glyphs

T<sub>E</sub>X82 (including P<sub>D</sub>F<sub>T</sub>E<sub>X</sub>) differentiated between `char_nodes` and `lig_nodes`. The former are simple items that contained nothing but a 'character' and a 'font' field, and they lived in the same memory as tokens did. The latter also contained a list of components, and a subtype indicating whether this ligature was the result of a word boundary, and it was stored in the same place as other nodes like boxes and kerns and glues.

In L<sup>A</sup>T<sub>E</sub>X, these two types are merged into one, somewhat larger structure called a `glyph_node`. Besides having the old character, font, and component fields, and the new special fields like 'attr' (see [section 8.1.2.12](#)), these nodes also contain:



- A subtype, split into four main types:
  - `character`, for characters to be hyphenated: the lowest bit (bit 0) is set to 1.
  - `glyph`, for specific font glyphs: the lowest bit (bit 0) is not set.
  - `ligature`, for ligatures (bit 1 is set)
  - `ghost`, for ‘ghost objects’ (bit 2 is set)
 The latter two make further use of two extra fields (bits 3 and 4):
  - `left`, for ligatures created from a left word boundary and for ghosts created from `\leftghost`
  - `right`, for ligatures created from a right word boundary and for ghosts created from `\rightghost`
 For ligatures, both bits can be set at the same time (in case of a single-glyph word).
- `glyph_nodes` of type ‘character’ also contain language data, split into four items that were current when the node was created: the `\setlanguage` (15 bits), `\lefthyphenmin` (8 bits), `\righthyphenmin` (8 bits), and `\uchyph` (1 bit).

Incidentally, L<sup>A</sup>T<sub>E</sub>X allows 16383 separate languages, and words can be 256 characters long.

Because the `\uchyph` value is saved in the actual nodes, its handling is subtly different from T<sub>E</sub>X82: changes to `\uchyph` become effective immediately, not at the end of the current partial paragraph.

Typeset boxes now always have their language information embedded in the nodes themselves, so there is no longer a possible dependency on the surrounding language settings. In T<sub>E</sub>X82, a mid-paragraph statement like `\unhbox0` would process the box using the current paragraph language unless there was a `\setlanguage` issued inside the box. In L<sup>A</sup>T<sub>E</sub>X, all language variables are already frozen.

## 6.2 The main control loop

In L<sup>A</sup>T<sub>E</sub>X’s main loop, almost all input characters that are to be typeset are converted into `glyph_node` records with subtype ‘character’, but there are a few small exceptions.

First, the `\accent` primitive creates nodes with subtype ‘glyph’ instead of ‘character’: one for the actual accent and one for the accented. The primary reason for this is that `\accent` in T<sub>E</sub>X82 is explicitly dependent on the current font encoding, so it would not make much sense to attach a new meaning to the primitive’s name, as that would invalidate many old documents and macro packages. A secondary reason is that in T<sub>E</sub>X82, `\accent` prohibits hyphenation of the current word. Since in L<sup>A</sup>T<sub>E</sub>X hyphenation only takes place on ‘character’ nodes, it is possible to achieve the same effect.

This change of meaning did happen with `\char`, that now generates ‘character’ nodes, consistent with its changed meaning in X<sub>Y</sub>T<sub>E</sub>X. The changed status of `\char` is not yet finalized, but if it stays as it is now, a new primitive `\glyph` should be added to directly insert a font glyph id.

Second, all the results of processing in math mode eventually become nodes with ‘glyph’ subtypes.

Third, the ALEPH-derived commands `\leftghost` and `\rightghost` create nodes of a third subtype: ‘ghost’. These nodes are ignored completely by all further processing until the stage where inter-glyph kerning is added.

Fourth, automatic discretionaries are handled differently. T<sub>E</sub>X82 inserts an empty discretionary after sensing an input character that matches the `\hyphenchar` in the current font. This test is wrong, in our



opinion: whether or not hyphenation takes place should not depend on the current font, it is a language property.

In L<sup>A</sup>T<sub>E</sub>X, it works like this: if L<sup>A</sup>T<sub>E</sub>X senses a string of input characters that matches the value of the new integer parameter `\exhyphenchar`, it will insert an explicit discretionary after that series of nodes. Initex sets the `\exhyphenchar=-1`. Incidentally, this is a global parameter instead of a language-specific one because it may be useful to change the value depending on the document structure instead of the text language.

Note: as of L<sup>A</sup>T<sub>E</sub>X 0.63.0, the insertion of discretionaries after a sequence of explicit hyphens happens at the same time as the other hyphenation processing, *not* inside the main control loop.

The only use L<sup>A</sup>T<sub>E</sub>X has for `\hyphenchar` is at the check whether a word should be considered for hyphenation at all. If the `\hyphenchar` of the font attached to the first character node in a word is negative, then hyphenation of that word is abandoned immediately. **This behavior is added for backward compatibility only, and the use of `\hyphenchar=-1` as a means of preventing hyphenation should not be used in new L<sup>A</sup>T<sub>E</sub>X documents.**

Fifth, `\setlanguage` no longer creates whatsits. The meaning of `\setlanguage` is changed so that it is now an integer parameter like all others. That integer parameter is used in `\glyph_node` creation to add language information to the glyph nodes. In conjunction, the `\language` primitive is extended so that it always also updates the value of `\setlanguage`.

Sixth, the `\noboundary` command (this command prohibits word boundary processing where that would normally take place) now does create whatsits. These whatsits are needed because the exact place of the `\noboundary` command in the input stream has to be retained until after the ligature and font processing stages.

Finally, there is no longer a `main_loop` label in the code. Remember that T<sub>E</sub>X82 did quite a lot of processing while adding `char_nodes` to the horizontal list? For speed reasons, it handled that processing code outside of the ‘main control’ loop, and only the first character of any ‘word’ was handled by that ‘main control’ loop. In L<sup>A</sup>T<sub>E</sub>X, there is no longer a need for that (all hard work is done later), and the (now very small) bits of character-handling code have been moved back inline. When `\tracingcommands` is on, this is visible because the full word is reported, instead of just the initial character.

## 6.3 Loading patterns and exceptions

The hyphenation algorithm in L<sup>A</sup>T<sub>E</sub>X is quite different from the one in T<sub>E</sub>X82, although it uses essentially the same user input.

After expansion, the argument for `\patterns` has to be proper UTF-8 with individual patterns separated by spaces, no `\char` or `\chardef-ed` commands are allowed. (The current implementation is even more strict, and will reject all non-UNICODE characters, but that will be changed in the future. For now, the generated errors are a valuable tool in discovering font-encoding specific pattern files)

Likewise, the expanded argument for `\hyphenation` also has to be proper UTF-8, but here a tiny little bit of extra syntax is provided:



1. three sets of arguments in curly braces (`{-}{-}{-}`) indicates a desired complex discretionary, with arguments as in `\discretionary`'s command in normal document input.
2. `-` indicates a desired simple discretionary, cf. `\-` and `\discretionary{-}{-}{-}` in normal document input.
3. Internal command names are ignored. This rule is provided especially for `\discretionary`, but it also helps to deal with `\relax` commands that may sneak in.
4. `=` indicates a (non-discretionary) hyphen in the document input.

The expanded argument is first converted back to a space-separated string while dropping the internal command names. This string is then converted into a dictionary by a routine that creates key–value pairs by converting the other listed items. It is important to note that the keys in an exception dictionary can always be generated from the values. Here are a few examples:

value	implied key (input)	effect
<code>ta-ble</code>	table	<code>ta\-\ble (= ta\discretionary{-}{-}{-}ble)</code>
<code>ba{k-}{-}{c}ken</code>	backen	<code>ba\discretionary{k-}{-}{c}ken</code>

The resultant patterns and exception dictionary will be stored under the language code that is the present value of `\language`.

In the last line of the table, you see there is no `\discretionary` command in the value: the command is optional in the T<sub>E</sub>X-based input syntax. The underlying reason for that is that it is conceivable that a whole dictionary of words is stored as a plain text file and loaded into L<sup>A</sup>T<sub>E</sub>X using one of the functions in the L<sup>A</sup> `lang` library. This loading method is quite a bit faster than going through the T<sub>E</sub>X language primitives, but some (most?) of that speed gain would be lost if it had to interpret command sequences while doing so.

Starting with L<sup>A</sup>T<sub>E</sub>X 0.63.0, it is possible to specify extra hyphenation points in compound words by using `{-}{-}{-}` for the explicit hyphen character (replace `-` by the actual explicit hyphen character if needed). For example, this matches the word ‘multi-word-boundaries’ and allows an extra break inbetween ‘boun’ and ‘daries’:

```
\hyphenation{multi{-}{-}{-}word{-}{-}{-}boun-daries}
```

The motivation behind the ε-T<sub>E</sub>X extension `\savingshyphcodes` was that hyphenation heavily depended on font encodings. This is no longer true in L<sup>A</sup>T<sub>E</sub>X, and the corresponding primitive is ignored pending complete removal. The future semantics of `\uppercase` and `\lowercase` are still under consideration, no changes have taken place yet.

## 6.4 Applying hyphenation

The internal structures L<sup>A</sup>T<sub>E</sub>X uses for the insertion of discretionaries in words is very different from the ones in T<sub>E</sub>X82, and that means there are some noticeable differences in handling as well.

First and foremost, there is no ‘compressed trie’ involved in hyphenation. The algorithm still reads PATGEN-generated pattern files, but L<sup>A</sup>T<sub>E</sub>X uses a finite state hash to match the patterns against the word to be hyphenated. This algorithm is based on the ‘libhnj’ library used by OpenOffice, which in turn



is inspired by T<sub>E</sub>X. The memory allocation for this new implementation is completely dynamic, so the WEB2C setting for `trie_size` is ignored.

Differences between L<sup>A</sup>T<sub>E</sub>X and T<sub>E</sub>X82 that are a direct result of that:

- L<sup>A</sup>T<sub>E</sub>X happily hyphenates the full UNICODE character range.
- Pattern and exception dictionary size is limited by the available memory only, all allocations are done dynamically. The trie-related settings in `texmf.cnf` are ignored.
- Because there is no ‘trie preparation’ stage, language patterns never become frozen. This means that the primitive `\patterns` (and its LUA counterpart `lang.patterns`) can be used at any time, not only in `initex`.
- Only the string representation of `\patterns` and `\hyphenation` is stored in the format file. At format load time, they are simply re-evaluated. It follows that there is no real reason to preload languages in the format file. In fact, it is usually not a good idea to do so. It is much smarter to load patterns no sooner than the first time they are actually needed.
- L<sup>A</sup>T<sub>E</sub>X uses the language-specific variables `\prehyphenchar` and `\posthyphenchar` in the creation of implicit discretionaries, instead of T<sub>E</sub>X82’s `\hyphenchar`, and the values of the language-specific variables `\preexhyphenchar` and `\postexhyphenchar` for explicit discretionaries (instead of T<sub>E</sub>X82’s empty discretionary).

Inserted characters and ligatures inherit their attributes from the nearest glyph node item (usually the preceding one, but the following one for the items inserted at the left-hand side of a word).

Word boundaries are no longer implied by font switches, but by language switches. One word can have two separate fonts and still be hyphenated correctly (but it can not have two different languages, the `\setlanguage` command forces a word boundary).

All languages start out with `\prehyphenchar=-`, `\posthyphenchar=0`, `\preexhyphenchar=0` and `\postexhyphenchar=0`. When you assign the values of one of these four parameters, you are actually changing the settings for the current `\language`, this behavior is compatible with `\patterns` and `\hyphenation`.

L<sup>A</sup>T<sub>E</sub>X also hyphenates the first word in a paragraph.

Words can be up to 256 characters long (up from 64 in T<sub>E</sub>X82). Longer words generate an error right now, but eventually either the limitation will be removed or perhaps it will become possible to silently ignore the excess characters (this is what happens in T<sub>E</sub>X82, but there the behavior cannot be controlled).

If you are using the LUA function `lang.hyphenate`, you should be aware that this function expects to receive a list of ‘character’ nodes. It will not operate properly in the presence of ‘glyph’, ‘ligature’, or ‘ghost’ nodes, nor does it know how to deal with kerning. In the near future, it will be able to skip over ‘ghost’ nodes, and we may add a less fuzzy function you can call as well.

The hyphenation exception dictionary is maintained as key-value hash, and that is also dynamic, so the `hyph_size` setting is not used either.

A technical paper detailing the new algorithm will be released as a separate document.



## 6.5 Applying ligatures and kerning

After all possible hyphenation points have been inserted in the list, L<sup>A</sup>T<sub>E</sub>X will process the list to convert the ‘character’ nodes into ‘glyph’ and ‘ligature’ nodes. This is actually done in two stages: first all ligatures are processed, then all kerning information is applied to the result list. But those two stages are somewhat dependent on each other: If the used font makes it possible to do so, the ligaturing stage adds virtual ‘character’ nodes to the word boundaries in the list. While doing so, it removes and interprets `noboundary` nodes. The kerning stage deletes those word boundary items after it is done with them, and it does the same for ‘ghost’ nodes. Finally, at the end of the kerning stage, all remaining ‘character’ nodes are converted to ‘glyph’ nodes.

This work separation is worth mentioning because, if you overrule from L<sup>A</sup> only one of the two callbacks related to font handling, then you have to make sure you perform the tasks normally done by L<sup>A</sup>T<sub>E</sub>X itself in order to make sure that the other, non-overruled, routine continues to function properly.

Work in this area is not yet complete, but most of the possible cases are handled by our rewritten ligaturing engine. We are working hard to make sure all of the possible inputs will become supported soon.

For example, take the word `office`, hyphenated `of-fice`, using a ‘normal’ font with all the `f-f` and `f-i` type ligatures:

```
Initial:           {o}{f}{f}{i}{c}{e}
After hyphenation: {o}{f}{f-},{},{f}{f}{i}{c}{e}
First ligature stage: {o}{f-},{f},{f}{f}{i}{c}{e}
Final result:       {o}{f-},{f}{f}{f}{i}{c}{e}
```

That’s bad enough, but let us assume that there is also a hyphenation point between the `f` and the `i`, to create `of-f-ice`. Then the final result should be:

```
{o}{f-},
  {f-},
  {i},
  {f}{f}{f}{i}{c}{e},
  {f}{f}{f}{f}{i}{c}{e},
  {i},
  {f}{f}{f}{f}{f}{i}{c}{e}
```

with discretionaries in the post-break text as well as in the replacement text of the top-level discretionary that resulted from the first hyphenation point.

Here is that nested solution again, in a different representation:

	pre	post	replace
topdisc	<code>f</code> <sup>-1</sup>	sub1	sub2
sub1	<code>f</code> <sup>-2</sup>	<code>i</code> <sup>3</sup>	<code>&lt;fi&gt;</code> <sup>4</sup>
sub2	<code>&lt;ff&gt;</code> <sup>-5</sup>	<code>i</code> <sup>6</sup>	<code>&lt;ffi&gt;</code> <sup>7</sup>



When line breaking is choosing its breakpoints, the following fields will eventually be selected:

```
of-f-ice  f-1
          f-2
          i3
of-fice   f-1
          <fi>4
off-ice   <ff>-5
          i6
office    <ffi>7
```

The current solution in L<sup>A</sup>T<sub>E</sub>X is not able to handle nested discretionaries, but it is in fact smart enough to handle this fictional `of-f-ice` example. It does so by combining two sequential discretionary nodes as if they were a single object (where the second discretionary node is treated as an extension of the first node).

One can observe that the `of-f-ice` and `off-ice` cases both end with the same actual post replacement list (`i`), and that this would be the case even if that `i` was the first item of a potential following ligature like `ic`. This allows L<sup>A</sup>T<sub>E</sub>X to do away with one of the fields, and thus make the whole stuff fit into just two discretionary nodes.

The mapping of the seven list fields to the six fields in this discretionary node pair is as follows:

field	description
<code>disc1.pre</code>	<code>f<sup>-1</sup></code>
<code>disc1.post</code>	<code>&lt;fi&gt;<sup>4</sup></code>
<code>disc1.replace</code>	<code>&lt;ffi&gt;<sup>7</sup></code>
<code>disc2.pre</code>	<code>f<sup>-2</sup></code>
<code>disc2.post</code>	<code>i<sup>3,6</sup></code>
<code>disc2.replace</code>	<code>&lt;ff&gt;-<sup>5</sup></code>

What is actually generated after ligaturing has been applied is therefore:

```
{o}{{f-},
      {<fi>},
      {<ffi>}}
{{f-},
 {i},
 {<ff>-}}{c}{e}
```

The two discretionaries have different subtypes from a discretionary appearing on its own: the first has subtype 4, and the second has subtype 5. The need for these special subtypes stems from the fact that not all of the fields appear in their ‘normal’ location. The second discretionary especially looks odd, with things like the `<ff>-` appearing in `disc2.replace`. The fact that some of the fields have different meanings (and different processing code internally) is what makes it necessary to have different subtypes: this enables L<sup>A</sup>T<sub>E</sub>X to distinguish this sequence of two joined discretionary nodes from the case of two standalone discretionaries appearing in a row.



## 6.6 Breaking paragraphs into lines

This code is still almost unchanged, but because of the above-mentioned changes with respect to discretionary nodes and ligatures, line breaking will potentially be different from traditional  $\text{\TeX}$ . The actual line breaking code is still based on the  $\text{\TeX}$ 82 algorithms, and it does not expect there to be discretionary nodes inside of discretionary nodes.

But that situation is now fairly common in  $\text{\LaTeX}$ , due to the changes to the ligaturing mechanism. And also, the  $\text{\LaTeX}$  discretionary nodes are implemented slightly different from the  $\text{\TeX}$ 82 nodes: the `no_break` text is now embedded inside the disc node, where previously these nodes kept their place in the horizontal list (the discretionary node contained a counter indicating how many nodes to skip).

The combined effect of these two differences is that  $\text{\LaTeX}$  does not always use all of the potential breakpoints in a paragraph, especially when fonts with many ligatures are used.





## 7 Font structure

All T<sub>E</sub>X fonts are represented to Lua code as tables, and internally as C structures. All keys in the table below are saved in the internal font structure if they are present in the table returned by the `define_font` callback, or if they result from the normal TFM/VF reading routines if there is no `define_font` callback defined.

The column ‘from vf’ means that this key will be created by the `font.read_vf()` routine, ‘from TFM’ means that the key will be created by the `font.read_tfm()` routine, and ‘used’ means whether or not the L<sup>A</sup>T<sub>E</sub>X engine itself will do something with the key.

The top-level keys in the table are as follows:

key	from vf	from tfm	used	value type	description
name	yes	yes	yes	string	metric (file) name
area	no	yes	yes	string	(directory) location, typically empty
used	no	yes	yes	boolean	used already? (initial: false)
characters	yes	yes	yes	table	the defined glyphs of this font
checksum	yes	yes	no	number	default: 0
designsize	no	yes	yes	number	expected size (default: 655360 == 10pt)
direction	no	yes	yes	number	default: 0 (TLT)
encodingbytes	no	no	yes	number	default: depends on <code>format</code>
encodingname	no	no	yes	string	encoding name
fonts	yes	no	yes	table	locally used fonts
psname	no	no	yes	string	actual (POSTSCRIPT) name (this is the PS fontname in the incoming font source, also used as fontname identifier in the PDF output, new in 0.43)
fullname	no	no	yes	string	output font name, used as a fallback in the PDF output if the psname is not set
header	yes	no	no	string	header comments, if any
hyphenchar	no	no	yes	number	default: TeX’s <code>\hyphenchar</code>
parameters	no	yes	yes	hash	default: 7 parameters, all zero
size	no	yes	yes	number	loaded (at) size. (default: same as designsize)
skewchar	no	no	yes	number	default: TeX’s <code>\skewchar</code>
type	yes	no	yes	string	basic type of this font
format	no	no	yes	string	disk format type
embedding	no	no	yes	string	PDF inclusion
filename	no	no	yes	string	disk file name
tounicode	no	yes	yes	number	if 1, L <sup>A</sup> T <sub>E</sub> X assumes per-glyph tounicode entries are present in the font
stretch	no	no	yes	number	the ‘stretch’ value from <code>\pdffontexpand</code>



<code>shrink</code>	no	no	yes	number	the ‘shrink’ value from <code>\pdffontexpand</code>
<code>step</code>	no	no	yes	number	the ‘step’ value from <code>\pdffontexpand</code>
<code>auto_expand</code>	no	no	yes	boolean	the ‘autoexpand’ keyword from <code>\pdffontexpand</code>
<code>expansion_factor</code>	no	no	no	number	the actual expansion factor of an expanded font
<code>attributes</code>	no	no	yes	string	the <code>\pdffontattr</code>
<code>cache</code>	no	no	yes	string	this key controls caching of the lua table on the <code>tex</code> end. <b>yes</b> : use a reference to the table that is passed to L <sup>A</sup> T <sub>E</sub> X (this is the default). <b>no</b> : don’t store the table reference, don’t cache any lua data for this font. <b>renew</b> : don’t store the table reference, but save a reference to the table that is created at the first access to one of its fields in <code>font.fonts</code> . (new in 0.40.0, before that caching was always <b>yes</b> ). Note: the saved reference is thread-local, so be careful when you are using coroutines: an error will be thrown if the table has been cached in one thread, but you reference it from another thread ( $\approx$ coroutine)
<code>nomath</code>	no	no	yes	boolean	this key allows a minor speedup for text fonts. if it is present and true, then L <sup>A</sup> T <sub>E</sub> X will not check the character entries for math-specific keys. (0.42.0)
<code>slant</code>	no	no	yes	number	This has the same semantics as the <code>SlantFont</code> operator in font map files. (0.47.0)
<code>extent</code>	no	no	yes	number	This has the same semantics as the <code>ExtendFont</code> operator in font map files. (0.50.0)

The key `name` is always required. The keys `stretch`, `shrink`, `step` and optionally `auto_expand` only have meaning when used together: they can be used to replace a post-loading `\pdffontexpand` command. The `expansion_factor` is value that can be present inside a font in `font.fonts`. It is the actual expansion factor (a value between `-shrink` and `stretch`, with step `step`) of a font that was automatically generated by the font expansion algorithm. The key `attributes` can be used to replace `\pdffontattr`. The key `used` is set by the engine when a font is actively in use, this makes sure that the font’s definition is written to the output file (DVI or PDF). The TFM reader sets it to false. The `direction` is a number signalling the ‘normal’ direction for this font. There are sixteen possibilities:



number	meaning	number	meaning
0	LT	8	TT
1	LL	9	TL
2	LB	10	TB
3	LR	11	TR
4	RT	12	BT
5	RL	13	BL
6	RB	14	BB
7	RR	15	BR

These are OMEGA-style direction abbreviations: the first character indicates the ‘first’ edge of the character glyphs (the edge that is seen first in the writing direction), the second the ‘top’ side.

The `parameters` is a hash with mixed key types. There are seven possible string keys, as well as a number of integer indices (these start from 8 up). The seven strings are actually used instead of the bottom seven indices, because that gives a nicer user interface.

The names and their internal remapping are:

name	internal remapped number
slant	1
space	2
space_stretch	3
space_shrink	4
x_height	5
quad	6
extra_space	7

The keys `type`, `format`, `embedding`, `fullname` and `filename` are used to embed OPENTYPE fonts in the result PDF.

The `characters` table is a list of character hashes indexed by an integer number. The number is the ‘internal code’ T<sub>E</sub>X knows this character by.

Two very special string indexes can be used also: `left_boundary` is a virtual character whose ligatures and kerns are used to handle word boundary processing. `right_boundary` is similar but not actually used for anything (yet!).

Other index keys are ignored.

Each character hash itself is a hash. For example, here is the character ‘f’ (decimal 102) in the font cmr10 at 10 points:

```
[102] = {
  ['width'] = 200250,
  ['height'] = 455111,
  ['depth'] = 0,
  ['italic'] = 50973,
  ['kerns'] = {
```



```

        [63] = 50973,
        [93] = 50973,
        [39] = 50973,
        [33] = 50973,
        [41] = 50973
    },
    ['ligatures'] = {
        [102] = {
            ['char'] = 11,
            ['type'] = 0
        },
        [108] = {
            ['char'] = 13,
            ['type'] = 0
        },
        [105] = {
            ['char'] = 12,
            ['type'] = 0
        }
    }
}

```

The following top-level keys can be present inside a character hash:

key	from vf	from tfm	used	value type	description
width	yes	yes	yes	number	character's width, in sp (default 0)
height	no	yes	yes	number	character's height, in sp (default 0)
depth	no	yes	yes	number	character's depth, in sp (default 0)
italic	no	yes	yes	number	character's italic correction, in sp (default zero)
top_accent	no	no	maybe	number	character's top accent alignment place, in sp (default zero)
bot_accent	no	no	maybe	number	character's bottom accent alignment place, in sp (default zero)
left_protruding	no	no	maybe	number	character's <a href="#">\lpcode</a>
right_protruding	no	no	maybe	number	character's <a href="#">\rpcode</a>
expansion_factor	no	no	maybe	number	character's <a href="#">\efcode</a>
tounicode	no	no	maybe	string	character's Unicode equivalent(s), in UTF-16BE hexadecimal format
next	no	yes	yes	number	the 'next larger' character index
extensible	no	yes	yes	table	the constituent parts of an extensible recipe
vert_variants	no	no	yes	table	constituent parts of a vertical variant set



<code>horiz_variants</code>	no	no	yes	table	constituent parts of a horizontal variant set
<code>kerns</code>	no	yes	yes	table	kerning information
<code>ligatures</code>	no	yes	yes	table	ligaturing information
<code>commands</code>	yes	no	yes	array	virtual font commands
<code>name</code>	no	no	no	string	the character (POSTSCRIPT) name
<code>index</code>	no	no	yes	number	the (OPENTYPE or TRUETYPE) font glyph index
<code>used</code>	no	yes	yes	boolean	typeset already (default: false)?
<code>mathkern</code>	no	no	yes	table	math cut-in specifications

The values of `top_accent`, `bot_accent` and `mathkern` are used only for math accent and superscript placement, see the **math chapter 145** in this manual for details.

The values of `left_protruding` and `right_protruding` are used only when `\pdfprotrudechars` is non-zero.

Whether or not `expansion_factor` is used depends on the font's global expansion settings, as well as on the value of `\pdfadjustspacing`.

The usage of `tounicode` is this: if this font specifies a `tounicode=1` at the top level, then L<sup>A</sup>T<sub>E</sub>X will construct a `/ToUnicode` entry for the PDF font (or font subset) based on the character-level `tounicode` strings, where they are available. If a character does not have a sensible UNICODE equivalent, do not provide a string either (no empty strings).

If the font-level `tounicode` is not set, then L<sup>A</sup>T<sub>E</sub>X will build up `/ToUnicode` based on the T<sub>E</sub>X code points you used, and any character-level `tounicodes` will be ignored. *At the moment, the string format is exactly the format that is expected by Adobe CMAP files (UTF-16BE in hexadecimal encoding), minus the enclosing angle brackets. This may change in the future.* Small example: the `tounicode` for a `fi` ligature would be `00660069`.

The presence of `extensible` will overrule `next`, if that is also present. It in turn can be overruled by `vert_variants`.

The `extensible` table is very simple:

key	type	description
<code>top</code>	number	'top' character index
<code>mid</code>	number	'middle' character index
<code>bot</code>	number	'bottom' character index
<code>rep</code>	number	'repeatable' character index

The `horiz_variants` and `vert_variants` are arrays of components. Each of those components is itself a hash of up to five keys:

key	type	explanation
<code>glyph</code>	number	The character index (note that this is an encoding number, not a name).
<code>extender</code>	number	One (1) if this part is repeatable, zero (0) otherwise.
<code>start</code>	number	Maximum overlap at the starting side (in scaled points).



<b>end</b>	number	Maximum overlap at the ending side (in scaled points).
<b>advance</b>	number	Total advance width of this item (can be zero or missing, then the natural size of the glyph for character <b>component</b> is used).

The **kerns** table is a hash indexed by character index (and ‘character index’ is defined as either a non-negative integer or the string value **right\_boundary**), with the values the kerning to be applied, in scaled points.

The **ligatures** table is a hash indexed by character index (and ‘character index’ is defined as either a non-negative integer or the string value **right\_boundary**), with the values being yet another small hash, with two fields:

key	type	description
<b>type</b>	number	the type of this ligature command, default 0
<b>char</b>	number	the character index of the resultant ligature

The **char** field in a ligature is required.

The **type** field inside a ligature is the numerical or string value of one of the eight possible ligature types supported by T<sub>E</sub>X. When T<sub>E</sub>X inserts a new ligature, it puts the new glyph in the middle of the left and right glyphs. The original left and right glyphs can optionally be retained, and when at least one of them is kept, it is also possible to move the new ‘insertion point’ forward one or two places. The glyph that ends up to the right of the insertion point will become the next ‘left’.

textual (Knuth)	number	string	result
$l + r =: n$	0	<b>=:</b>	$ n$
$l + r =:   n$	1	<b>=:  </b>	$ nr$
$l + r   =: n$	2	<b>  =:</b>	$ ln$
$l + r   =:   n$	3	<b>  =:  </b>	$ lnr$
$l + r =:   > n$	5	<b>=:   &gt;</b>	$n r$
$l + r   =: > n$	6	<b>  =: &gt;</b>	$l n$
$l + r   =:   > n$	7	<b>  =:   &gt;</b>	$l nr$
$l + r   =:   > > n$	11	<b>  =:   &gt; &gt;</b>	$ln r$

The default value is 0, and can be left out. That signifies a ‘normal’ ligature where the ligature replaces both original glyphs. In this table the  $|$  indicates the final insertion point.

The **commands** array is explained below.

## 7.1 Real fonts

Whether or not a T<sub>E</sub>X font is a ‘real’ font that should be written to the PDF document is decided by the **type** value in the top-level font structure. If the value is **real**, then this is a proper font, and the inclusion mechanism will attempt to add the needed font object definitions to the PDF.

Values for **type**:



value	description
real	this is a base font
virtual	this is a virtual font

The actions to be taken depend on a number of different variables:

- Whether the used font fits in an 8-bit encoding scheme or not
- The type of the disk font file
- The level of embedding requested

A font that uses anything other than an 8-bit encoding vector has to be written to the PDF in a different way.

The rule is: if the font table has `encodingbytes` set to 2, then this is a wide font, in all other cases it isn't. The value 2 is the default for `OPENTYPE` and `TRUETYPE` fonts loaded via `LUA`. For `TYPE1` fonts, you have to set `encodingbytes` to 2 explicitly. For `PK` bitmap fonts, wide font encoding is not supported at all.

If no special care is needed, `LUATEX` currently falls back to the mapfile-based solution used by `PDFTEX` and `DVIPS`. This behavior will be removed in the future, when the existing code becomes integrated in the new subsystem.

But if this is a 'wide' font, then the new subsystem kicks in, and some extra fields have to be present in the font structure. In this case, `LUATEX` does not use a map file at all.

The extra fields are: `format`, `embedding`, `fullname`, `cidinfo` (as explained above), `filename`, and the `index` key in the separate characters.

Values for `format` are:

value	description
type1	this is a <code>POSTSCRIPT TYPE1</code> font
type3	this is a bitmapped ( <code>PK</code> ) font
truetype	this is a <code>TRUETYPE</code> or <code>TRUETYPE</code> -based <code>OPENTYPE</code> font
opentype	this is a <code>POSTSCRIPT</code> -based <code>OPENTYPE</code> font

(`type3` fonts are provided for backward compatibility only, and do not support the new wide encoding options.)

Values for `embedding` are:

value	description
no	don't embed the font at all
subset	include and attempt to subset the font
full	include this font in its entirety

It is not possible to artificially modify the transformation matrix for the font at the moment.

The other fields are used as follows: The `fullname` will be the `POSTSCRIPT/PDF` font name. The `cidinfo` will be used as the character set (the `CID /Ordering` and `/Registry` keys). The `filename` points to the actual font file. If you include the full path in the `filename` or if the file is in the local directory,



L<sup>A</sup>T<sub>E</sub>X will run a little bit more efficient because it will not have to re-run the `find_xxx_file` callback in that case.

Be careful: when mixing old and new fonts in one document, it is possible to create PostScript name clashes that can result in printing errors. When this happens, you have to change the `fullname` of the font.

Typeset strings are written out in a wide format using 2 bytes per glyph, using the `index` key in the character information as value. The overall effect is like having an encoding based on numbers instead of traditional (PostScript) name-based reencoding. The way to get the correct `index` numbers for T<sub>E</sub>X fonts is by loading the font via `fontloader.open`; use the table indices as `index` fields.

This type of reencoding means that there is no longer a clear connection between the text in your input file and the strings in the output PDF file. Dealing with this is high on the agenda.

## 7.2 Virtual fonts

You have to take the following steps if you want L<sup>A</sup>T<sub>E</sub>X to treat the returned table from `define_font` as a virtual font:

- Set the top-level key `type` to `virtual`.
- Make sure there is at least one valid entry in `fonts` (see below).
- Give a `commands` array to every character (see below).

The presence of the toplevel `type` key with the specific value `virtual` will trigger handling of the rest of the special virtual font fields in the table, but the mere existence of 'type' is enough to prevent L<sup>A</sup>T<sub>E</sub>X from looking for a virtual font on its own.

Therefore, this also works 'in reverse': if you are absolutely certain that a font is not a virtual font, assigning the value `base` or `real` to `type` will inhibit L<sup>A</sup>T<sub>E</sub>X from looking for a virtual font file, thereby saving you a disk search.

The `fonts` is another LUA array. The values are one- or two-key hashes themselves, each entry indicating one of the base fonts in a virtual font. In case your font is referring to itself, you can use the `font.nextid()` function which returns the index of the next to be defined font which is probably the currently defined one.

An example makes this easy to understand

```
fonts = {
  { name = 'ptmr8a', size = 655360 },
  { name = 'psyr', size = 600000 },
  { id = 38 }
}
```

says that the first referenced font (index 1) in this virtual font is `ptmr8a` loaded at 10pt, and the second is `psyr` loaded at a little over 9pt. The third one is previously defined font that is known to L<sup>A</sup>T<sub>E</sub>X as fontid '38'.





The array index numbers are used by the character command definitions that are part of each character. The `commands` array is a hash where each item is another small array, with the first entry representing a command and the extra items being the parameters to that command. The allowed commands and their arguments are:

command name	arguments	arg type	description
font	1	number	select a new font from the local <code>fonts</code> table
char	1	number	typeset this character number from the current font, and move right by the character's width
node	1	node	output this node (list), and move right by the width of this list
slot	2	number	a shortcut for the combination of a font and char command
push	0		save current position
nop	0		do nothing
pop	0		pop position
rule	2	2 numbers	output a rule $ht * wd$ , and move right.
down	1	number	move down on the page
right	1	number	move right on the page
special	1	string	output a <code>\special</code> command
lua	1	string	execute a LUA script (at <code>\latelua</code> time)
image	1	image	output an image (the argument can be either an <code>&lt;image&gt;</code> variable or an <code>image_spec</code> table)
comment	any	any	the arguments of this command are ignored

Here is a rather elaborate glyph commands example:

```
...
commands = {
  {'push'},           -- remember where we are
  {'right', 5000},    -- move right about 0.08pt
  {'font', 3},        -- select the fonts[3] entry
  {'char', 97},       -- place character 97 (ASCII 'a')
  {'pop'},            -- go all the way back
  {'down', -200000},  -- move upwards by about 3pt
  {'special', 'pdf: 1 0 0 rg'} -- switch to red color
  {'rule', 500000, 20000} -- draw a bar
  {'special', 'pdf: 0 g'} -- back to black
}
...
```

The default value for `font` is always 1 at the start of the `commands` array. Therefore, if the virtual font is essentially only a re-encoding, then you do usually not have create an explicit 'font' command in the array.

Rules inside of `commands` arrays are built up using only two dimensions: they do not have depth. For correct vertical placement, an extra `down` command may be needed.



Regardless of the amount of movement you create within the `commands`, the output pointer will always move by exactly the width that was given in the `width` key of the character hash. Any movements that take place inside the `commands` array are ignored on the upper level.

## 7.2.1 Artificial fonts

Even in a ‘real’ font, there can be virtual characters. When L<sup>A</sup>T<sub>E</sub>X encounters a `commands` field inside a character when it becomes time to typeset the character, it will interpret the commands, just like for a true virtual character. In this case, if you have created no ‘fonts’ array, then the default (and only) ‘base’ font is taken to be the current font itself. In practice, this means that you can create virtual duplicates of existing characters which is useful if you want to create composite characters.

Note: this feature does *not* work the other way around. There can not be ‘real’ characters in a virtual font! You cannot use this technique for font re-encoding either; you need a truly virtual font for that (because characters that are already present cannot be altered).

## 7.2.2 Example virtual font

Finally, here is a plain T<sub>E</sub>X input file with a virtual font demonstration:

```
\directlua {
  callback.register('define_font',
    function (name,size)
      if name == 'cmr10-red' then
        f = font.read_tfm('cmr10',size)
        f.name = 'cmr10-red'
        f.type = 'virtual'
        f.fonts = {{ name = 'cmr10', size = size }}
        for i,v in pairs(f.characters) do
          if (string.char(i)):find('[taco]hanshartmut') then
            v.commands = {
              {'special','pdf: 1 0 0 rg'},
              {'char',i},
              {'special','pdf: 0 g'},
            }
          else
            v.commands = {'char',i}
          end
        end
      else
        f = font.read_tfm(name,size)
      end
      return f
    end
  )
}
```



```
)  
}
```

```
\font\myfont = cmr10-red at 10pt \myfont This is a line of text \par  
\font\myfontx= cmr10 at 10pt \myfontx Here is another line of text \par
```





# 8 Nodes

## 8.1 LUA node representation

TEX's nodes are represented in LUA as userdata object with a variable set of fields. In the following syntax tables, such the type of such a userdata object is represented as `<node>`.

The current return value of `node.types()` is: `hlist` (0), `vlist` (1), `rule` (2), `ins` (3), `mark` (4), `adjust` (5), `disc` (7), `whatsit` (8), `math` (9), `glue` (10), `kern` (11), `penalty` (12), `unset` (13), `style` (14), `choice` (15), `noad` (16), `op` (17), `bin` (18), `rel` (19), `open` (20), `close` (21), `punct` (22), `inner` (23), `radical` (24), `fraction` (25), `under` (26), `over` (27), `accent` (28), `vcenter` (29), `fence` (30), `math_char` (31), `sub_box` (32), `sub_mlist` (33), `math_text_char` (34), `delim` (35), `margin_kern` (36), `glyph` (37), `align_record` (38), `pseudo_file` (39), `pseudo_line` (40), `page_insert` (41), `split_insert` (42), `expr_stack` (43), `nested_list` (44), `span` (45), `attribute` (46), `glue_spec` (47), `attribute_list` (48), `action` (49), `temp` (50), `align_stack` (51), `movement_stack` (52), `if_stack` (53), `unhyphenated` (54), `hyphenated` (55), `delta` (56), `passive` (57), `shape` (58), `fake` (100),.

NOTE: The `\lastnodetype` primitive is  $\epsilon$ -TEX compliant. The valid range is still -1 .. 15 and glyph nodes have number 0 (used to be char node) and ligature nodes are mapped to 7. That way macro packages can use the same symbolic names as in traditional  $\epsilon$ -TEX. Keep in mind that the internal node numbers are different and that there are more node types than 15.

### 8.1.1 Auxiliary items

A few node-typed userdata objects do not occur in the 'normal' list of nodes, but can be pointed to from within that list. They are not quite the same as regular nodes, but it is easier for the library routines to treat them as if they were.

#### 8.1.1.1 glue\_spec items

Skips are about the only type of data objects in traditional TEX that are not a simple value. The structure that represents the glue components of a skip is called a `glue_spec`, and it has the following accessible fields:

key	type	explanation
<code>width</code>	number	
<code>stretch</code>	number	
<code>stretch_order</code>	number	
<code>shrink</code>	number	
<code>shrink_order</code>	number	
<code>writable</code>	boolean	If this is true, you can't assign to this <code>glue_spec</code> because it is one of the preallocated special cases. New in 0.52



These objects are reference counted, so there is actually an extra read-only field named `ref_count` as well. This item type will likely disappear in the future, and the glue fields themselves will become part of the nodes referencing glue items.

### 8.1.1.2 `attribute_list` and attribute items

The newly introduced attribute registers are non-trivial, because the value that is attached to a node is essentially a sparse array of key-value pairs.

It is generally easiest to deal with attribute lists and attributes by using the dedicated functions in the `node` library, but for completeness, here is the low-level interface.

An `attribute_list` item is used as a head pointer for a list of attribute items. It has only one user-visible field:

field	type	explanation
<code>next</code>	<code>&lt;node&gt;</code>	pointer to the first attribute

A normal node's attribute field will point to an item of type `attribute_list`, and the `next` field in that item will point to the first defined 'attribute' item, whose `next` will point to the second 'attribute' item, etc.

Valid fields in `attribute` items:

field	type	explanation
<code>next</code>	<code>&lt;node&gt;</code>	pointer to the next attribute
<code>number</code>	number	the attribute type id
<code>value</code>	number	the attribute value

As mentioned it's better to use the official helpers rather than edit these fields directly. For instance the `prev` field is used for other purposes and there is no double linked list.

### 8.1.1.3 `action` item

Valid fields: `action_type`, `named_id`, `action_id`, `file`, `new_window`, `data`, `ref_count`  
Id: 49

These are a special kind of item that only appears inside pdf start link objects.

field	type	explanation
<code>action_type</code>	number	
<code>action_id</code>	number or string	
<code>named_id</code>	number	
<code>file</code>	string	
<code>new_window</code>	number	
<code>data</code>	string	
<code>ref_count</code>	number	(read-only)



## 8.1.2 Main text nodes

These are the nodes that comprise actual typesetting commands.

A few fields are present in all nodes regardless of their type, these are:

field	type	explanation
next	<node>	The next node in a list, or nil
id	number	The node's type ( <b>id</b> ) number
subtype	number	The node <b>subtype</b> identifier

The **subtype** is sometimes just a stub entry. Not all nodes actually use the **subtype**, but this way you can be sure that all nodes accept it as a valid field name, and that is often handy in node list traversal. In the following tables **next** and **id** are not explicitly mentioned.

Besides these three fields, almost all nodes also have an **attr** field, and there is also a field called **prev**. That last field is always present, but only initialized on explicit request: when the function **node.slide()** is called, it will set up the **prev** fields to be a backwards pointer in the argument node list.

### 8.1.2.1 hlist nodes

Valid fields: **attr**, **width**, **depth**, **height**, **dir**, **shift**, **glue\_order**, **glue\_sign**, **glue\_set**, **head**  
Id: 0

field	type	explanation
subtype	number	0 = unknown origin, 1 = created by linebreaking, 2 = explicit box command. (0.46.0), 3 = paragraph indentation box, 4 = alignment column or row, 5 = alignment cell (0.62.0)
attr	<node>	The head of the associated attribute list
width	number	
height	number	
depth	number	
shift	number	a displacement perpendicular to the character progression direction
glue_order	number	a number in the range 0–4, indicating the glue order
glue_set	number	the calculated glue ratio
glue_sign	number	0 = normal, 1 = stretching, 2 = shrinking
head	<node>	the first node of the body of this list
dir	string	the direction of this box. see 8.1.4.7

A warning: never assign a node list to the **head** field unless you are sure its internal link structure is correct, otherwise an error may result.

Note: the new field name **head** was introduced in 0.65 to replace the old name **list**. Use of the name **list** is now deprecated, but it will stay available until at least version 0.80.



### 8.1.2.2 vlist nodes

Valid fields: As for hlist, except that 'shift' is a displacement perpendicular to the line progression direction, and 'subtype' only has subtypes 0, 4, and 5.

### 8.1.2.3 rule nodes

Valid fields: `attr`, `width`, `depth`, `height`, `dir`

Id: 2

field	type	explanation
subtype	number	unused
attr	<code>&lt;node&gt;</code>	
width	number	the width of the rule; the special value <code>-1073741824</code> is used for 'running' glue dimensions
height	number	the height of the rule (can be negative)
depth	number	the depth of the rule (can be negative)
dir	string	the direction of this rule. see 8.1.4.7

### 8.1.2.4 ins nodes

Valid fields: `attr`, `cost`, `depth`, `height`, `spec`, `head`

Id: 3

field	type	explanation
subtype	number	the insertion class
attr	<code>&lt;node&gt;</code>	
cost	number	the penalty associated with this insert
height	number	
depth	number	
head	<code>&lt;node&gt;</code>	the first node of the body of this insert
spec	<code>&lt;node&gt;</code>	a pointer to the <code>\splittopskip</code> glue spec

A warning: never assign a node list to the `head` field unless you are sure its internal link structure is correct, otherwise an error may be result.

Note: the new field name `head` was introduced in 0.65 to replace the old name `list`. Use of the name `list` is now deprecated, but it will stay available until at least version 0.80.

### 8.1.2.5 mark nodes

Valid fields: `attr`, `class`, `mark`

Id: 4





field	type	explanation
subtype	number	unused
attr	<node>	
class	number	the mark class
mark	table	a table representing a token list

### 8.1.2.6 adjust nodes

Valid fields: [attr](#), [head](#)

Id: 5

field	type	explanation
subtype	number	0 = normal, 1 = 'pre'
attr	<node>	
head	<node>	adjusted material

A warning: never assign a node list to the [head](#) field unless you are sure its internal link structure is correct, otherwise an error may be result.

Note: the new field name [head](#) was introduced in 0.65 to replace the old name [list](#). Use of the name [list](#) is now deprecated, but it will stay available until at least version 0.80.

### 8.1.2.7 disc nodes

Valid fields: [attr](#), [pre](#), [post](#), [replace](#), [penalty](#)

Id: 7

field	type	explanation
subtype	number	indicates the source of a discretionary. 0 = the <a href="#">\discretionary</a> command, 1 = the <a href="#">\-</a> command, 2 = added automatically following a <a href="#">-</a> , 3 = added by the hyphenation algorithm (simple), 4 = added by the hyphenation algorithm (hard, first item), 5 = added by the hyphenation algorithm (hard, second item)
attr	<node>	
pre	<node>	pointer to the pre-break text
post	<node>	pointer to the post-break text
replace	<node>	pointer to the no-break text

The subtype numbers 4 and 5 belong to the 'of-f-ice' explanation given elsewhere.

A warning: never assign a node list to the pre, post or replace field unless you are sure its internal link structure is correct, otherwise an error may be result.

### 8.1.2.8 math nodes

Valid fields: [attr](#), [surround](#)

Id: 9



field	type	explanation
subtype	number	0 = 'on', 1 = 'off'
attr	<code>&lt;node&gt;</code>	
surround	number	width of the <code>\mathsurround</code> kern

### 8.1.2.9 glue nodes

Valid fields: `attr`, `spec`, `leader`

Id: 10

field	type	explanation
subtype	number	0 = <code>\skip</code> , 1–18 = internal glue parameters, 100–103 = 'leader' subtypes
attr	<code>&lt;node&gt;</code>	
spec	<code>&lt;node&gt;</code>	pointer to a <code>glue_spec</code> item
leader	<code>&lt;node&gt;</code>	pointer to a box or rule for leaders

The exact meanings of the subtypes are as follows:

- 1 `\lineskip`
- 2 `\baselineskip`
- 3 `\parskip`
- 4 `\abovedisplayskip`
- 5 `\belowdisplayskip`
- 6 `\abovedisplayshortskip`
- 7 `\belowdisplayshortskip`
- 8 `\leftskip`
- 9 `\rightskip`
- 10 `\topskip`
- 11 `\splittopskip`
- 12 `\tabskip`
- 13 `\spaceskip`
- 14 `\xspaceskip`
- 15 `\parfillskip`
- 16 `\thinmuskip`
- 17 `\medmuskip`
- 18 `\thickmuskip`
- 100 `\leaders`
- 101 `\cleaders`
- 102 `\xleaders`
- 103 `\gleaders`

### 8.1.2.10 kern nodes

Valid fields: `attr`, `kern`, `expansion_factor`

Id: 11



field	type	explanation
subtype	number	0 = from font, 1 = from <code>\kern</code> or <code>\/</code> , 2 = from <code>\accent</code>
attr	<code>&lt;node&gt;</code>	
kern	number	

### 8.1.2.11 penalty nodes

Valid fields: `attr`, `penalty`

Id: 12

field	type	explanation
subtype	number	not used
attr	<code>&lt;node&gt;</code>	
penalty	number	

### 8.1.2.12 glyph nodes

Valid fields: `attr`, `char`, `font`, `lang`, `left`, `right`, `uchyph`, `components`, `xoffset`, `yoffset`, `width`, `height`, `depth`, `expansion_factor`

Id: 37

field	type	explanation
subtype	number	bitfield
attr	<code>&lt;node&gt;</code>	
char	number	
font	number	
lang	number	
left	number	
right	number	
uchyph	boolean	
components	<code>&lt;node&gt;</code>	pointer to ligature components
xoffset	number	
yoffset	number	
width	number	(new in 0.53)
height	number	(new in 0.53)
depth	number	(new in 0.53)
expansion_factor	number	(new in 0.78)

A warning: never assign a node list to the `components` field unless you are sure its internal link structure is correct, otherwise an error may be result.

Valid bits for the `subtype` field are:

bit	meaning
0	character



- 1 ligature
- 2 ghost
- 3 left
- 4 right

See [section 6.1](#) for a detailed description of the [subtype](#) field.

The [expansion\\_factor](#) is relatively new and the result of extensive experiments with a more efficient implementation of expansion. Early versions of L<sup>A</sup>T<sub>E</sub>X already replaced multiple instances of fonts in the backend by scaling but contrary to P<sup>D</sup>F<sub>T</sub>E<sub>X</sub> in L<sup>A</sup>T<sub>E</sub>X we now also got rid of font copies in the frontend and replaced them by expansion factors that travel with glyph nodes. Apart from a cleaner approach this is also a step towards a better separation between front- and backend.

### 8.1.2.13 margin\_kern nodes

Valid fields: [attr](#), [width](#), [glyph](#)

Id: 36

field	type	explanation
subtype	number	0 = left side, 1 = right side
attr	<a href="#">&lt;node&gt;</a>	
width	number	
glyph	<a href="#">&lt;node&gt;</a>	

## 8.1.3 Math nodes

These are the so-called ‘noad’s and the nodes that are specifically associated with math processing. Most of these nodes contain sub-nodes so that the list of possible fields is actually quite small. First, the subnodes:

### 8.1.3.1 Math kernel subnodes

Many object fields in math mode are either simple characters in a specific family or math lists or node lists. There are four associated subnodes that represent these cases (in the following node descriptions these are indicated by the word [<kernel>](#)).

The [next](#) and [prev](#) fields for these subnodes are unused.

#### 8.1.3.1.1 math\_char and math\_text\_char subnodes

Valid fields: [attr](#), [fam](#), [char](#)

Id: 31

field	type	explanation
attr	<a href="#">&lt;node&gt;</a>	



char number  
fam number

The `math_char` is the simplest subnode field, it contains the character and family for a single glyph object. The `math_text_char` is a special case that you will not normally encounter, it arises temporarily during math list conversion (its sole function is to suppress a following italic correction).

### 8.1.3.1.2 sub\_box and sub\_mlist subnodes

Valid fields: `attr`, `head`  
Id: 32

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>head</code>	<code>&lt;node&gt;</code>	

These two subnode types are used for subsidiary list items. For `sub_box`, the `head` points to a 'normal' vbox or hbox. For `sub_mlist`, the `head` points to a math list that is yet to be converted.

A warning: never assign a node list to the `head` field unless you are sure its internal link structure is correct, otherwise an error may be result.

Note: the new field name `head` was introduced in 0.65 to replace the old name `list`. Use of the name `list` is now deprecated, but it will stay available until at least version 0.80.

### 8.1.3.2 Math delimiter subnode

There is a fifth subnode type that is used exclusively for delimiter fields. As before, the `next` and `prev` fields are unused.

#### 8.1.3.2.1 delim subnodes

Valid fields: `attr`, `small_fam`, `small_char`, `large_fam`, `large_char`  
Id: 35

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>small_char</code>	number	
<code>small_fam</code>	number	
<code>large_char</code>	number	
<code>large_fam</code>	number	

The fields `large_char` and `large_fam` can be zero, in that case the font that is used for the `small_fam` is expected to provide the large version as an extension to the `small_char`.



### 8.1.3.3 Math core nodes

First, there are the objects (the T<sub>E</sub>Xbook calls them ‘atoms’) that are associated with the simple math objects: Ord, Op, Bin, Rel, Open, Close, Punct, Inner, Over, Under, Vcent. These all have the same fields, and they are combined into a single node type with separate subtypes for differentiation.

#### 8.1.3.3.1 simple nodes

Valid fields: `attr`, `nucleus`, `sub`, `sup`

Id: 16

field	type	explanation
subtype	number	see below
attr	<code>&lt;node&gt;</code>	
nucleus	<code>&lt;kernel&gt;</code>	
sub	<code>&lt;kernel&gt;</code>	
sup	<code>&lt;kernel&gt;</code>	

Operators are a bit special because they occupy three subtypes. `subtype`.

number	node sub type
0	Ord
1	Op, <code>\displaylimits</code>
2	Op, <code>\limits</code>
3	Op, <code>\nolimits</code>
4	Bin
5	Rel
6	Open
7	Close
8	Punct
9	Inner
10	Under
11	Over
12	Vcent

#### 8.1.3.3.2 accent nodes

Valid fields: `attr`, `nucleus`, `sub`, `sup`, `accent`, `bot_accent`

Id: 28

field	type	explanation
subtype	number	the first bit is used for a fixed top accent flag (if the <code>accent</code> field is present), the second bit for a fixed bottom accent flag (if the <code>bot_accent</code> field is present). Example: the actual value 3 means: do not stretch either accent
attr	<code>&lt;node&gt;</code>	



nucleus	<kernel>
sub	<kernel>
sup	<kernel>
accent	<kernel>
bot_accent	<kernel>

### 8.1.3.3.3 style nodes

Valid fields: `attr`, `style`

Id: 14

field	type	explanation
style	string	contains the style

There are eight possibilities for the string value: one of 'display', 'text', 'script', or 'scriptscript'. Each of these can have a trailing ' ' to signify 'cramped' styles.

### 8.1.3.3.4 choice nodes

Valid fields: `attr`, `display`, `text`, `script`, `scriptscript`

Id: 15

field	type	explanation
attr	<node>	
display	<node>	
text	<node>	
script	<node>	
scriptscript	<node>	

A warning: never assign a node list to the display, text, script, or scriptscript field unless you are sure its internal link structure is correct, otherwise an error may be result.

### 8.1.3.3.5 radical nodes

Valid fields: `attr`, `nucleus`, `sub`, `sup`, `left`, `degree`

Id: 24

field	type	explanation
attr	<node>	
nucleus	<kernel>	
sub	<kernel>	
sup	<kernel>	
left	<delim>	
degree	<kernel>	Only set by <code>\Uroot</code>



A warning: never assign a node list to the nucleus, sub, sup, left, or degree field unless you are sure its internal link structure is correct, otherwise an error may be result.

#### 8.1.3.3.6 fraction nodes

Valid fields: `attr`, `width`, `num`, `denom`, `left`, `right`  
Id: 25

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>width</code>	number	
<code>num</code>	<code>&lt;kernel&gt;</code>	
<code>denom</code>	<code>&lt;kernel&gt;</code>	
<code>left</code>	<code>&lt;delim&gt;</code>	
<code>right</code>	<code>&lt;delim&gt;</code>	

A warning: never assign a node list to the num, or denom field unless you are sure its internal link structure is correct, otherwise an error may be result.

#### 8.1.3.3.7 fence nodes

Valid fields: `attr`, `delim`  
Id: 30

field	type	explanation
<code>subtype</code>	number	1 = <code>\left</code> , 2 = <code>\middle</code> , 3 = <code>\right</code>
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>delim</code>	<code>&lt;delim&gt;</code>	

### 8.1.4 whatsit nodes

Whatsit nodes come in many subtypes that you can ask for by running `node.whatsits()`: `open` (0), `write` (1), `close` (2), `special` (3), `local_par` (6), `dir` (7), `pdf_literal` (8), `pdf_refobj` (10), `pdf_refxform` (12), `pdf_refximage` (14), `pdf_annot` (15), `pdf_start_link` (16), `pdf_end_link` (17), `pdf_dest` (19), `pdf_thread` (20), `pdf_start_thread` (21), `pdf_end_thread` (22), `pdf_save_pos` (23), `pdf_thread_data` (24), `pdf_link_data` (25), `late_lua` (35), `close_lua` (36), `pdf_colorstack` (39), `pdf_setmatrix` (40), `pdf_save` (41), `pdf_restore` (42), `cancel_boundary` (43), `user_defined` (44)

#### 8.1.4.1 open nodes

Valid fields: `attr`, `stream`, `name`, `area`, `ext`  
Id: 8, 0





field	type	explanation
attr	<code>&lt;node&gt;</code>	
stream	number	T <sub>E</sub> X's stream id number
name	string	file name
ext	string	file extension
area	string	file area (this may become obsolete)

#### 8.1.4.2 write nodes

Valid fields: `attr`, `stream`, `data`  
Id: 8, 1

field	type	explanation
attr	<code>&lt;node&gt;</code>	
stream	number	T <sub>E</sub> X's stream id number
data	table	a table representing the token list to be written

#### 8.1.4.3 close nodes

Valid fields: `attr`, `stream`  
Id: 8, 2

field	type	explanation
attr	<code>&lt;node&gt;</code>	
stream	number	T <sub>E</sub> X's stream id number

#### 8.1.4.4 special nodes

Valid fields: `attr`, `data`  
Id: 8, 3

field	type	explanation
attr	<code>&lt;node&gt;</code>	
data	string	the <code>\special</code> information

#### 8.1.4.5 language nodes

L<sup>A</sup>T<sub>E</sub>X does not have language whatsits any more. All language information is already present inside the glyph nodes themselves. This whatsit subtype will be removed in the next release.

#### 8.1.4.6 local\_par nodes

Valid fields: `attr`, `pen_inter`, `pen_broken`, `dir`, `box_left`, `box_left_width`, `box_right`, `box_right_width`  
Id: 8, 6



field	type	explanation
attr	<a href="#">&lt;node&gt;</a>	
pen_inter	number	local interline penalty (from <a href="#">\localinterlinepenalty</a> )
pen_broken	number	local broken penalty (from <a href="#">\localbrokenpenalty</a> )
dir	string	the direction of this par. see <a href="#">8.1.4.7</a>
box_left	<a href="#">&lt;node&gt;</a>	the <a href="#">\localleftbox</a>
box_left_width	number	width of the <a href="#">\localleftbox</a>
box_right	<a href="#">&lt;node&gt;</a>	the <a href="#">\localrightbox</a>
box_right_width	number	width of the <a href="#">\localrightbox</a>

A warning: never assign a node list to the `box_left` or `box_right` field unless you are sure its internal link structure is correct, otherwise an error may be result.

### 8.1.4.7 dir nodes

Valid fields: [attr](#), [dir](#), [level](#), [dvi\\_ptr](#), [dvi\\_h](#)  
Id: [8](#), [7](#)

field	type	explanation
attr	<a href="#">&lt;node&gt;</a>	
dir	string	the direction (but see below)
level	number	nesting level of this direction whatsit
dvi_ptr	number	a saved dvi buffer byte offset
dir_h	number	a saved dvi position

A note on [dir](#) strings. Direction specifiers are three-letter combinations of [T](#), [B](#), [R](#), and [L](#).

These are built up out of three separate items:

- the first is the direction of the ‘top’ of paragraphs.
- the second is the direction of the ‘start’ of lines.
- the third is the direction of the ‘top’ of glyphs.

However, only four combinations are accepted: [TLT](#), [TRT](#), [RTT](#), and [LTL](#).

Inside actual [dir](#) whatsit nodes, the representation of [dir](#) is not a three-letter but a four-letter combination. The first character in this case is always either [+](#) or [-](#), indicating whether the value is pushed or popped from the direction stack.

### 8.1.4.8 pdf\_literal nodes

Valid fields: [attr](#), [mode](#), [data](#)  
Id: [8](#), [8](#)

field	type	explanation
attr	<a href="#">&lt;node&gt;</a>	



mode	number	the 'mode' setting of this literal
data	string	the <code>\pdfliteral</code> information

Mode values:

value	corresponding <code>\pdfTeX</code> keyword
0	setorigin
1	page
2	direct

#### 8.1.4.9 pdf\_refobj nodes

Valid fields: `attr`, `objnum`

Id: 8, 10

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>objnum</code>	number	the referenced PDF object number

#### 8.1.4.10 pdf\_refxform nodes

Valid fields: `attr`, `width`, `depth`, `height`, `objnum`

Id: 8, 12.

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>objnum</code>	number	the referenced PDF object number

Be aware that `pdf_refxform` nodes have dimensions that are used by `LuATEX`.

#### 8.1.4.11 pdf\_refximage nodes

Valid fields: `attr`, `width`, `depth`, `height`, `transform`, `index`

Id: 8, 14

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>objnum</code>	number	the referenced PDF object number



Be aware that `pdf_refximage` nodes have dimensions that are used by L<sup>A</sup>T<sub>E</sub>X.

#### 8.1.4.12 `pdf_annot` nodes

Valid fields: `attr`, `width`, `depth`, `height`, `objnum`, `data`

Id: 8, 15

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>objnum</code>	number	the referenced PDF object number
<code>data</code>	string	the annotation data

#### 8.1.4.13 `pdf_start_link` nodes

Valid fields: `attr`, `width`, `depth`, `height`, `objnum`, `link_attr`, `action`

Id: 8, 16

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>width</code>	number	
<code>height</code>	number	
<code>depth</code>	number	
<code>objnum</code>	number	the referenced PDF object number
<code>link_attr</code>	table	the link attribute token list
<code>action</code>	<code>&lt;node&gt;</code>	the action to perform

#### 8.1.4.14 `pdf_end_link` nodes

Valid fields: `attr`

Id: 8, 17

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	

#### 8.1.4.15 `pdf_dest` nodes

Valid fields: `attr`, `width`, `depth`, `height`, `named_id`, `dest_id`, `dest_type`, `xyz_zoom`, `objnum`

Id: 8, 19

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	



width	number	
height	number	
depth	number	
named_id	number	is the dest_id a string value?
dest_id	number or string	the destination id
dest_type	number	type of destination
xyz_zoom	number	
objnum	number	the PDF object number

#### 8.1.4.16 pdf\_thread nodes

Valid fields: `attr`, `width`, `depth`, `height`, `named_id`, `thread_id`, `thread_attr`  
 Id: 8, 20

field	type	explanation
attr	<node>	
width	number	
height	number	
depth	number	
named_id	number	is the tread_id a string value?
tread_id	number or string	the thread id
thread_attr	number	extra thread information

#### 8.1.4.17 pdf\_start\_thread nodes

Valid fields: `attr`, `width`, `depth`, `height`, `named_id`, `thread_id`, `thread_attr`  
 Id: 8, 21

field	type	explanation
attr	<node>	
width	number	
height	number	
depth	number	
named_id	number	is the tread_id a string value?
tread_id	number or string	the thread id
thread_attr	number	extra thread information

#### 8.1.4.18 pdf\_end\_thread nodes

Valid fields: `attr`  
 Id: 8, 22

field	type	explanation
attr	<node>	



#### 8.1.4.19 pdf\_save\_pos nodes

Valid fields: `attr`

Id: 8, 23

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	

#### 8.1.4.20 late\_lua nodes

Valid fields: `attr`, `reg`, `data`, `name`, `string`

Id: 8, 35

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>data</code>	string	data to execute
<code>string</code>	string	data to execute (0.63)
<code>name</code>	string	the name to use for lua error reporting

The difference between `data` and `string` is that on assignment, the `data` field is converted to a token list, cf. use as `\latelua`. The `string` version is treated as a literal string.

#### 8.1.4.21 pdf\_colorstack nodes

Valid fields: `attr`, `stack`, `cmd`, `data`

Id: 8, 39

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>stack</code>	number	colorstack id number
<code>cmd</code>	number	command to execute
<code>data</code>	string	data

#### 8.1.4.22 pdf\_setmatrix nodes

Valid fields: `attr`, `data`

Id: 8, 40

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>data</code>	string	data



### 8.1.4.23 pdf\_save nodes

Valid fields: `attr`

Id: 8, 41

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	

### 8.1.4.24 pdf\_restore nodes

Valid fields: `attr`

Id: 8, 42

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	

### 8.1.4.25 user\_defined nodes

User-defined whatsit nodes can only be created and handled from LUA code. In effect, they are an extension to the extension mechanism. The L<sup>A</sup>T<sub>E</sub>X engine will simply step over such whatsits without ever looking at the contents.

Valid fields: `attr`, `user_id`, `type`, `value`

Id: 8, 44

field	type	explanation
<code>attr</code>	<code>&lt;node&gt;</code>	
<code>user_id</code>	number	id number
<code>type</code>	number	type of the value
<code>value</code>	number string <code>&lt;node&gt;</code> table	

The `type` can have one of five distinct values:

value	explanation
97	the value is an attribute node list
100	the value is a number
110	the value is a node list
115	the value is a string
116	the value is a token list in LUA table form



## 8.2 Two access models

After doing lots of tests with `LUA $\TeX$`  and `LUAJIT $\TeX$`  with and without just in time compilation enabled, and with and without using `ffi`, we came to the conclusion that userdata prevents a speedup. We also found that the checking of metatables as well as assignment comes with overhead that can't be neglected. This is normally not really a problem but when processing fonts for more complex scripts it could have quite some overhead.

Because the userdata approach has some benefits, this remains the recommended way to access nodes. We did several experiments with faster access using this model, but eventually settled for the 'direct' approach. For code that is proven to be okay, one can use this access model that operates on nodes more directly.

Deep down in  `$\TeX$`  a node has a number which is an entry in a memory table. In fact, this model, where  `$\TeX$`  manages memory is real fast and one of the reasons why plugging in callbacks that operate on nodes is quite fast. No matter what future memory model `LUA $\TeX$`  has, an internal reference will always be a simple data type (like a number or light userdata in `LUA` speak). So, if you use the direct model, even if you know that you currently deal with numbers, you should not depend on that property but treat it an abstraction just like traditional nodes. In fact, the fact that we use a simple basic datatype has the penalty that less checking can be done, but less checking is also the reason why it's somewhat faster. An important aspect is that one cannot mix both methods, but you can cast both models.

So our advice is: use the indexed approach when possible and investigate the direct one when speed might be an issue. For that reason we also provide the `get*` and `set*` functions in the top level node namespace. There is a limited set of getters. When implementing this direct approach the regular index by key variant was also optimized, so direct access only makes sense when we're accessing nodes millions of times (which happens in some font processing for instance).

We're talking mostly of getters because setters are less important. Documents have not that many content related nodes and setting many thousands of properties is hardly a burden contrary to millions of consultations.

Normally you will access nodes like this:

```
local next = current.next
if next then
    -- do something
end
```

Here `next` is not a real field, but a virtual one. Accessing it results in a metatable method being called. In practice it boils down to looking up the node type and based on the node type checking for the field name. In a worst case you have a node type that sits at the end of the lookup list and a field that is last in the lookup chain. However, in successive versions of `LUA $\TeX$`  these lookups have been optimized and the most frequently accessed nodes and fields have a higher priority.

Because in practice the `next` accessor results in a function call, there is some overhead involved. The `next` code does the same and performs a tiny bit faster (but not that much because it is still a function call but one that knows what to look up).





```

local next = node.next(current)
if next then
    -- do something
end

```

There are several such function based accessors now:

getnext	parsing nodelist always involves this one
getprev	used less but is logical companion to getnext
getid	consulted a lot
getsubtype	consulted less but also a topper
getfont	used a lot in offt handling (glyph nodes are consulted a lot)
getchar	idem and also in other places
getlist	we often parse nested lists so this is a convenient one too (only works for hlist and vlist!)
getleader	comparable to list, seldom used in T <sub>E</sub> X (but needs frequent consulting like lists; leaders could have been made a dedicated node type)
getfield	generic getter, sufficient for the rest (other field names are often shared so a specific getter makes no sense then)

It doesn't make sense to add more. Profiling demonstrated that these fields can get accesses way more times than other fields. Even in complex documents, many node and fields types never get seen, or seen only a few times. Most functions in the `node` namespace have a companion in `node.direct`, but of course not the ones that don't deal with nodes themselves. The following table summarized this:

function	node	direct
copy	+	+
copy_list	+	+
count	+	+
current_attr	+	+
dimensions	+	+
do_ligature_n	+	+
end_of_math	+	+
family_font	+	—
fields	+	—
first_character	+	—
first_glyph	+	+
flush_list	+	+
flush_node	+	+
free	+	+
getbox	—	+
getchar	+	+
getfield	+	+
getfont	+	+
getid	+	+
getnext	+	+



getprev	+	+
getlist	+	+
getleader	+	+
getsubtype	+	+
has_glyph	+	+
has_attribute	+	+
has_field	+	+
hpack	+	+
id	+	—
insert_after	+	+
insert_before	+	+
is_direct	—	+
is_node	+	+
kerning	+	—
last_node	+	+
length	+	+
ligaturing	+	—
mlist_to_hlist	+	—
new	+	+
next	+	—
prev	+	—
tostring	+	+
protect_glyphs	+	+
protrusion_skipable	+	+
remove	+	+
set_attribute	+	+
setbox	+	+
setfield	+	+
slide	+	+
subtype	+	—
tail	+	+
todirect	+	+
tonode	+	+
traverse	+	+
traverse_id	+	+
type	+	—
types	+	—
unprotect_glyphs	+	+
unset_attribute	+	+
usedlist	+	+
vpack	+	+
whatsits	+	—
write	+	+



The `node.next` and `node.prev` functions will stay but for consistency there are variants called `get-next` and `getprev`. We had to use `get` because `node.id` and `node.subtype` are already taken for providing meta information about nodes. Note: The getters do only basic checking for valid keys. You should just stick to the keys mentioned in the sections that describe node properties.





200 Nodes

## 9 Modifications

Besides the expected changes caused by new functionality, there are a number of not-so-expected changes. These are sometimes a side-effect of a new (conflicting) feature, or, more often than not, a change necessary to clean up the internal interfaces.

### 9.1 Changes from T<sub>E</sub>X 3.1415926

- The current code base is written in C, not Pascal web (as of L<sup>A</sup>T<sub>E</sub>X 0.42.0).
- See **chapter 6** for many small changes related to paragraph building, language handling, and hyphenation. Most important change: adding a brace group in the middle of a word (like in `of{f}ice`) does not prevent ligature creation.
- There is no pool file, all strings are embedded during compilation.
- `plus 1 filllll` does not generate an error. The extra 'l' is simply typeset.
- The upper limit to `\endlinechar` and `\newlinechar` is 127.

### 9.2 Changes from $\epsilon$ -T<sub>E</sub>X 2.2

- The  $\epsilon$ -T<sub>E</sub>X functionality is always present and enabled (but see below about T<sub>E</sub>X<sub>X</sub>ET), so the prepended asterisk or `-etex` switch for INIT<sub>E</sub>X is not needed.
- T<sub>E</sub>X<sub>X</sub>ET is not present, so the primitives

```
\TeXxETstate
\beginR
\beginL
\endR
\endL
```

are missing.

- Some of the tracing information that is output by  $\epsilon$ -T<sub>E</sub>X's `\tracingassigns` and `\tracingrestores` is not there.
- Register management in L<sup>A</sup>T<sub>E</sub>X uses the ALEPH model, so the maximum value is 65535 and the implementation uses a flat array instead of the mixed flat&sparse model from  $\epsilon$ -T<sub>E</sub>X.
- `savinghyphcodes` is a no-op. See **chapter 6** for details.
- When kpathsea is used to find files, L<sup>A</sup>T<sub>E</sub>X uses the `ofm` file format to search for font metrics. In turn, this means that L<sup>A</sup>T<sub>E</sub>X looks at the `OFMFONTS` configuration variable (like `OMEGA` and `ALEPH`) instead of `TFMFONTS` (like T<sub>E</sub>X and PDF<sub>E</sub>X). Likewise for virtual fonts (L<sup>A</sup>T<sub>E</sub>X uses the variable `OVFONTS` instead of `VFONTS`).

### 9.3 Changes from PDF<sub>T</sub><sub>E</sub>X 1.40



- The (experimental) support for snap nodes has been removed, because it is much more natural to build this functionality on top of node processing and attributes. The associated primitives that are now gone are: `\pdfsnaprefpoint`, `\pdfsnappy`, and `\pdfsnapycomp`.
- The (experimental) support for specialized spacing around nodes has also been removed. The associated primitives that are now gone are: `\pdfadjustinterwordglue`, `\pdfprependkern`, and `\pdfappendkern`, as well as the five supporting primitives `\knbscode`, `\stbscode`, `\shbscode`, `\knbcode`, and `\knacode`.
- A number of ‘utility functions’ is removed:

<code>\pdfelapsedtime</code>	<code>\pdffilemdate</code>	<code>\pdfresettimer</code>
<code>\pdfescapehex</code>	<code>\pdffilesize</code>	<code>\pdfshellescape</code>
<code>\pdfescapename</code>	<code>\pdflastmatch</code>	<code>\pdfstrcmp</code>
<code>\pdfescapestring</code>	<code>\pdfmatch</code>	<code>\pdfunescapehex</code>
<code>\pdffiledump</code>	<code>\pdfmdfivesum</code>	

In 0.81 we also removed:

```

\pdfeachlineheight
\pdfeachlinedepth
\pdffirstlineheight
\pdflastlinedepth

```

- The four primitives that were already marked obsolete in PDF<sub>TEX</sub> 1.40 have been removed since L<sub>A</sub>T<sub>E</sub>X 0.42:

<code>\pdfoptionalwaysusepdfpagebox</code>	<code>\pdfforcepagebox</code>
<code>\pdfoptionpdfinclusionerrorlevel</code>	<code>\pdfmovechars</code>

- A few other experimental primitives are also provided without the extra `pdf` prefix, so they can also be called like this:

<code>\primitive</code>	<code>\ifabsnum</code>
<code>\ifprimitive</code>	<code>\ifabsdim</code>

- The `\pdfTeXversion` was set to 200, but from L<sub>A</sub>T<sub>E</sub>X 0.80.0 returns the luatex version and is declared as obsolete (i.e. it will disappear).
- The PNG transparency fix from 1.40.6 is not applied (high-level support is pending)
- LFS (PDF Files larger than 2GiB) support is not working yet.
- L<sub>A</sub>T<sub>E</sub>X 0.45.0 introduces two extra token lists, `\pdfxformresources` and `\pdfxformattr`, as an alternative to `\pdfxform` keywords.
- As of L<sub>A</sub>T<sub>E</sub>X 0.50.0 is no longer possible for fonts from embedded pdf files to be replaced by / merged with the document fonts of the enveloping pdf document. This regression may be temporary, depending on how the rewritten font backend will look after beta 0.60.

## 9.4 Changes from ALEPH RC4



- Starting with L<sup>A</sup>T<sub>E</sub>X 0.75.0, the extended 16-bit math primitives (`\omathcode` etc. ) have been removed.
- Starting with L<sup>A</sup>T<sub>E</sub>X 0.63.0, OCP processing is no longer supported at all. As a consequence, the following primitives have been removed:

<code>\ocp</code>	<code>\addbeforeocplist</code>
<code>\externalocp</code>	<code>\addafterocplist</code>
<code>\ocplist</code>	<code>\removebeforeocplist</code>
<code>\pushocplist</code>	<code>\removeafterocplist</code>
<code>\popocplist</code>	<code>\ocptracelevel</code>
<code>\clearocplists</code>	

- L<sup>A</sup>T<sub>E</sub>X only understands 4 of the 16 direction specifiers of ALEPH: `TLT` (latin), `TRT` (arabic), `RTT` (cjk), `LTL` (mongolian). All other direction specifiers generate an error (L<sup>A</sup>T<sub>E</sub>X 0.45).
- The input translations from ALEPH are not implemented, the related primitives are not available:

<code>\DefaultInputMode</code>	<code>\DefaultInputTranslation</code>
<code>\noDefaultInputMode</code>	<code>\noDefaultInputTranslation</code>
<code>\noInputMode</code>	<code>\noInputTranslation</code>
<code>\InputMode</code>	<code>\InputTranslation</code>
<code>\DefaultOutputMode</code>	<code>\DefaultOutputTranslation</code>
<code>\noDefaultOutputMode</code>	<code>\noDefaultOutputTranslation</code>
<code>\noOutputMode</code>	<code>\noOutputTranslation</code>
<code>\OutputMode</code>	<code>\OutputTranslation</code>

- The `\hoffset` bug when `\pagedir TRT` is fixed, removing the need for an explicit fix to `\hoffset`
- A bug causing `\fam` to fail for family numbers above 15 is fixed.
- A fair amount of other minor bugs are fixed as well, most of these related to `\tracingcommands` output.
- The internal function `scan_dir()` has been renamed to `scan_direction()` to prevent a naming clash, and it now allows an optional space after the direction is completely parsed.
- The `^^` notation can come in five and six item repetitions also, to insert characters that do not fit in the BMP.
- Glues *immediately after* direction change commands are not legal breakpoints.

## 9.5 Changes from standard WEB2C

- There is no `mltex`
- There is no `enctex`
- The following commandline switches are silently ignored, even in non-LUA mode:

```
-8bit
-translate-file=TCXNAME
-mltex
```



`-enc`  
`-etex`

- `\openout` whatsits are not written to the log file.
- Some of the so-called web2c extensions are hard to set up in non-KPSE mode because `texmf.cnf` is not read: `shell-escape` is off (but that is not a problem because of LUA's `os.execute`), and the paranoia checks on `openin` and `openout` do not happen (however, it is easy for a LUA script to do this itself by overloading `io.open`).
- The 'E' option does not do anything useful.





# 10 Implementation notes

## 10.1 Primitives overlap

The primitives

<code>\pdfpagewidth</code>	<code>\pagewidth</code>
<code>\pdfpageheight</code>	<code>\pageheight</code>
<code>\fontcharwd</code>	<code>\charwd</code>
<code>\fontcharht</code>	<code>\charht</code>
<code>\fontchardp</code>	<code>\chardp</code>
<code>\fontcharic</code>	<code>\charit</code>

are all aliases of each other.

## 10.2 Memory allocation

The single internal memory heap that traditional T<sub>E</sub>X used for tokens and nodes is split into two separate arrays. Each of these will grow dynamically when needed.

The `texmf.cnf` settings related to main memory are no longer used (these are: `main_memory`, `mem_bot`, `extra_mem_top` and `extra_mem_bot`). ‘Out of main memory’ errors can still occur, but the limiting factor is now the amount of RAM in your system, not a predefined limit.

Also, the memory (de)allocation routines for nodes are completely rewritten. The relevant code now lives in the C file `texnode.c`, and basically uses a dozen or so ‘avail’ lists instead of a doubly-linked model. An extra function layer is added so that the code can ask for nodes by type instead of directly requisitioning a certain amount of memory words.

Because of the split into two arrays and the resulting differences in the data structures, some of the macros have been duplicated. For instance, there are now `vlink` and `vinfo` as well as `token_link` and `token_info`. All access to the variable memory array is now hidden behind a macro called `vmem`.

The implementation of the growth of two arrays (via reallocation) introduces a potential pitfall: the memory arrays should never be used as the left hand side of a statement that can modify the array in question.

The input line buffer and pool size are now also reallocated when needed, and the `texmf.cnf` settings `buf_size` and `pool_size` are silently ignored.

## 10.3 Sparse arrays

The `\mathcode`, `\delcode`, `\catcode`, `\sfcode`, `\lccode` and `\uccode` tables are now sparse arrays that are implemented in C. They are no longer part of the T<sub>E</sub>X ‘equivalence table’ and because



each had 1.1 million entries with a few memory words each, this makes a major difference in memory usage.

The `\catcode`, `\sfcode`, `\lccode` and `\uccode` assignments do not yet show up when using the etex tracing routines `\tracingassigns` and `\tracingrestores` (code simply not written yet).

A side-effect of the current implementation is that `\global` is now more expensive in terms of processing than non-global assignments.

See `mathcodes.c` and `textcodes.c` if you are interested in the details.

Also, the glyph ids within a font are now managed by means of a sparse array and glyph ids can go up to index  $2^{21} - 1$ .

## 10.4 Simple single-character csnames

Single-character commands are no longer treated specially in the internals, they are stored in the hash just like the multiletter csnames.

The code that displays control sequences explicitly checks if the length is one when it has to decide whether or not to add a trailing space.

Active characters are internally implemented as a special type of multi-letter control sequences that uses a prefix that is otherwise impossible to obtain.

## 10.5 Compressed format

The format is passed through zlib, allowing it to shrink to roughly half of the size it would have had in uncompressed form. This takes a bit more CPU cycles but much less disk I/O, so it should still be faster.

## 10.6 Binary file reading

All of the internal code is changed in such a way that if one of the `read_xxx_file` callbacks is not set, then the file is read by a C function using basically the same convention as the callback: a single read into a buffer big enough to hold the entire file contents. While this uses more memory than the previous code (that mostly used `getc` calls), it can be quite a bit faster (depending on your I/O subsystem).



# 11 Known bugs and limitations, TODO

There used to be a lists of bugs and planned features below here, but that did not work out too well. There are lists of open bugs and feature requests in the tracker at <http://tracker.luatex.org>.



